

EXAM INFORMATION

RSS feed for this section

This category contains 150 posts

Exam 1z0-060: Upgrade to Oracle Database 12c (update February 4th, 2016)

Identify two ways to rectify the error.

Posted by seenagape on February 22, 2016 5 comments

Your multitenant container (CDB) contains two pluggable databases (PDB), HR_PDB and ACCOUNTS_PDB, both of which use the CDB tablespace. The temp file is called temp01.tmp. A user issues a query on a table on one of the PDBs and receives the following error: Identify two ways to rectify the error.

ERROR at line 1:

ORA-01565: error in identifying file '/u01/app/oracle/oradata/CDB1/temp01.tmp'

ORA-27037: unable to obtain file status

A.

Add a new temp file to the temporary tablespace and drop the temp file that that produced the error.

B.

Take the temporary tablespace offline, recover the missing temp file by applying redo logs, and then bring the temporary tablespace online.

C.

Shutdown the database instance, restore and recover the temp file from the backup, and then open the database with RESETLOGS.

D.

Shut down the database instance and then restart the CDB and PDBs.

E.

Add a new temp file to the temporary tablespace and drop the temp file that that produced the error.

Which two statements are true about redefining the table?

Posted by seenagape on February 22, 2016 1 comment

Examine the following commands for redefining a table with Virtual Private Database (VPD) policies: Which two statements are true about redefining the table?

```

BEGIN
  DBMS_RLS.ADD_POLICY (
    object_schema      => 'hr',
    object_name        => 'employees',
    policy_name        => 'employees_policy',
    function_schema    => 'hr',
    policy_function     => 'auth_emp_dep_100',
    statement_types    => 'select, insert, update',
  );
END;

BEGIN
  DBMS_REDEFINITION.START_REDEF_TABLE (
    uname              => 'hr',
    orig_table         => 'employees',
    int_table          => 'int_employees',
    col_mapping        => NULL,
    options_flag       => DBMS_REDEFINITION.CONSTRAINT_OPTION_OFF,
    orderby_cols       => NULL,
    part_name          => NULL,
    copy_vpd_opt       => DBMS_REDEFINITION.CONS_VPD_COPY,
  );
END;

```

A.

All the triggers for the table are disabled without changing any of the column names or column types in the table.

B.

The primary key constraint on the EMPLOYEES table is disabled during redefinition.

C.

VPD policies are copied from the original table to the new table during online redefinition.

D.

You must copy the VPD policies manually from the original table to the new table during online redefinition.

Which two statements are true about the use of the procedures listed in the

`v$sysaux_occupants.move_procedure` column?

Posted by seenagape on February 22, 2016 1 comment

Which two statements are true about the use of the procedures listed in the

`v$sysaux_occupants.move_procedure` column?

A.

The procedure may be used for some components to relocate component data to the SYSAUX tablespace from its current tablespace.

B.

All the components may be moved from the SYSAUX tablespace.

C.

The procedure may be used for some components to relocate component data from the SYSAUX tablespace to another tablespace.

D.

All the components may be moved into SYSAUX tablespace.

Which statement is true about Oracle Net Listener?

Posted by seenagape on February 22, 2016 2 comments

Which statement is true about Oracle Net Listener?

A.

It acts as the listening endpoint for the Oracle database instance for all local and non-local user connections.

B.

A single listener can service only one database instance and multiple remote client connections.

C.

Service registration with the listener is performed by the process monitor (LREG) process of each database instance.

D.

The listener.ora configuration file must be configured with one or more listening protocol addresses to allow remote users to connect to a database instance.

E.

The listener.ora configuration file must be located in the ORACLE_HOME/network/admin directly.

Explanation:

which three ways can you re-create the lost disk group and restore the data?

Posted by seenagape on February 22, 2016 2 comments

You are administering a database stored in Automatic Storage Management (ASM). You use RMAN to back up the database and the MD_BACKUP command to back up the ASM metadata regularly. You lost an ASM disk group DG1 due to hardware failure. In which three ways can you re-create the lost disk group and restore the data?

A.

Use the MD_RESTORE command to restore both the metadata and data for the failed disk group.

B.

Use the MKDG command to restore the disk group with the same configuration as the backedup disk group name and same set of disks and failure group configuration, and use RMAN to restore the data.

C.

Use the MD_RESTORE command to restore metadata for an existing disk group by passing the existing disk group name as an input parameter and use RMAN to restore the data.

D.

Use the MKDG command to restore the disk group with the same configuration as the backedup disk group and data on the disk group.

E.

Use the MKDG command to add a new disk group DG1 with the same or different specifications for failure group and other attributes and use RMAN to restore the data.

F.

Use the MD_RESTORE command to restore the disk group with the changed disk group specification, failure group specification, name, and other attributes and use RMAN to restore the data.

What should you do before executing the commands to restore and recover the data file in ACCOUNTS_PDB?

Posted by seenagape on February 22, 2016 2 comments

Your multitenant container database, CDB1, is running in ARCHIVELOG mode and has two pluggable databases, HR_PDB and ACCOUNTS_PDB. An RMAN backup exists for the database. You issue the command to open ACCOUNTS_PDB and find that the USERDATA.DBF data file for the default permanent tablespace USERDATA belonging to ACCOUNTS_PDB is corrupted. What should you do before executing the commands to restore and recover the data file in ACCOUNTS_PDB?

A.

Place CDB1 in the mount stage and then the USERDATA tablespace offline in ACCOUNTS_PDB.

B.

Place CDB1 in the mount stage and issue the ALTER PLUGGABLE DATABASE accounts_pdb CLOSE IMMEDIATE command.

C.

Issue the ALTER PLUGGABLE DATABASE accounts_pdb RESTRICTED command.

D.

Take the USERDATA tablespace offline in ACCOUNTS_PDB.

Which Oracle Database component is audited by default if the unified Auditing option is enabled?

Posted by seenagape on February 22, 2016 2 comments

Which Oracle Database component is audited by default if the unified Auditing option is enabled?

A.

Oracle Database Vault

B.

Oracle Real Application Security

C.

Oracle Label Security

D.

Oracle Recovery Manager (RMAN)

E.

Oracle Data Pump

Which option identifies the correct sequence to recover the SYSAUX tablespace?

Posted by seenagape on February 22, 2016 1 comment

Your multitenant container (CDB) containing three pluggable databases (PDBs) is running in ARCHIVELOG mode. You find that the SYSAUX tablespace is corrupted in the root container. The steps to recover the tablespace are as follows: Which option identifies the correct sequence to recover the SYSAUX tablespace?

1. Mount the CDB.
2. Close all the PDBs.
3. Open the database.
4. Apply the archive redo logs.
5. Restore the data file.
6. Take the SYSAUX tablespace offline.
7. Place the SYSAUX tablespace online.
8. Open all the PDBs with RESETLOGS.
9. Open the database with RESETLOGS.
10. Execute the command SHUTDOWN ABORT.

A.

10, 1, 2, 5, 8

B.

10, 1, 2, 5, 4, 9, 8

C.

10, 1, 5, 8, 10

D.

6, 5, 4, 7

Explanation:

Which three are direct benefits of the multiprocess, multithreaded architecture of Oracle Database 12c when it is enabled?

Posted by seenagape on February 22, 2016 2 comments

Which three are direct benefits of the multiprocess, multithreaded architecture of Oracle Database 12c when it is enabled?

A.

Reduced logical I/O

B.

Reduced virtual memory utilization

C.

Improved parallel Execution performance

D.

Improved Serial Execution performance

E.
Reduced physical I/O

F.
Reduced CPU utilization

Which three statements are true about this requirement?

Posted by seenagape on February 22, 20161 comment

In order to exploit some new storage tiers that have been provisioned by a storage administrator, the partitions of a large heap table must be moved to other tablespaces in your Oracle 12c database? Both local and global partitioned B-tree Indexes are defined on the table. A high volume of transactions access the table during the day and a medium volume of transactions access it at night and during weekends. Minimal disruption to availability is required. Which three statements are true about this requirement?

A.
Local indexes must be rebuilt manually after moving the partitions.

B.
The partitions can be compressed in the new tablespaces.

C.
The partitions can be moved online to new tablespaces.

D.
Global indexes must be rebuilt manually after moving the partitions.

E.
The partitions can be compressed in the same tablespaces.

Which three are true about the large pool for an Oracle database instance that supports shared server connections?

Posted by seenagape on February 22, 20162 comments

Which three are true about the large pool for an Oracle database instance that supports shared server connections?

A.
Allocates memory for RMAN backup and restore operations

B.
Allocates memory for shared and private SQL areas

C.
Contains a cursor area for storing runtime information about cursors

D.
Contains stack space

E.
Contains a hash area performing hash joins of tables

What are three purposes of the RMAN “FROM” clause?

Posted by seenagape on February 22, 20163 comments

What are three purposes of the RMAN “FROM” clause?

- A.
to support synchronization of a standby database with the primary database in a Data environment
- B.
To support file recovery over the network in a Data Guard environment
- C.
To support PULL-based active database duplication
- D.
to support PUSH-based active database duplication
- E.
To support file restores over the network in a Data Guard environment

How can you detect the cause of the degraded performance?

Posted by seenagape on February 22, 20161 comment

You notice that the performance of your production 24/7 Oracle database significantly degraded. Sometimes you are not able to connect to the instance because it hangs. You do not want to restart the database instance. How can you detect the cause of the degraded performance?

- A.
Enable Memory Access Mode, which reads performance data from SGA.
- B.
Use emergency monitoring to fetch data directly from SGA analysis.
- C.
Run Automatic Database Diagnostic Monitor (ADDM) to fetch information from the latest Automatic Workload Repository (AWR) snapshots.
- D.
Use Active Session History (ASH) data and hang analysis in regular performance monitoring.
- E.
Run ADDM in diagnostic mode.

Explanation:

Which three storage options support the use of HCC?

Posted by seenagape on February 22, 2016 1 comment

You plan to use the In Database Archiving feature of Oracle Database 12c, and store rows that are inactive for over three months, in Hybrid Columnar Compressed (HCC) format. Which three storage options support the use of HCC?

A.

Database files stored in ZFS and accessed using the Oracle Direct NFS feature

B.

Database files stored in any file system and accessed using the Oracle Direct NFS feature

C.

. ASM disk groups with ASM disks consisting of LUNs on Pillar Axiom Storage arrays

D.

ASM disk groups with ASM disks consisting of any zero padded NFS-mounted files

E.

ASM disk groups with ASM disks consisting of LUNS on any Storage Area Network array

F.

ASM disk groups with ASM disks consisting of Exadata Grid Disks.

G.

Database files stored in ZFS and accessed using conventional NFS mounts.

How does real-time Automatic database Diagnostic Monitor (ADDM) check performance degradation and provide solutions?

Posted by seenagape on February 22, 2016 1 comment

In your multitenant container database (CDB) containing pluggable databases (PDB), users complain about performance degradation. How does real-time Automatic database Diagnostic Monitor (ADDM) check performance degradation and provide solutions?

A.

It collects data from both SGA and PGA, analyzes it, and provides a report.

B.

It collects data from SGA, analyzes it, and provides a report.

C.

It collects data from SGA and compares it with the latest snapshot.

D.

It collects data from SGA and compares it with a preserved snapshot.

What could be the reason for this?

Posted by seenagape on February 22, 2016 4 comments

The tnsnames.ora file has an entry for the service alias ORCL as follows: The TNS ping command executes successfully when tested with ORCL; however, from the same OS user session, you are not able to connect to the database instance with the following command: SQL > CONNECT scott/tiger@orcl What could be the reason for this?

```
ORCL =  
  ( DESCRIPTION =  
    ( ADDRESS = ( PROTOCOL = TCP ) ( HOST = 10.156.24.216 )  
      ( CONNECT_DATA =  
        ( SERVER = DEDICATED )  
        ( SERVICE_NAME = orcl.oracle.com )  
      )  
    )  
  )
```

A.

The listener is running on a different port.

B.

The DEFAULT_DOMAIN parameter is set to the wrong value in the sqlnet.ora file.

C.

The orcl.oracle.com database service is not registered with the listener.

D.

The listener is not running on the database node.

E.

The TNS_ADMIN environment variable is set to the wrong value.

Explanation:

Identify the correct sequence of steps.

Posted by seenagape on February 22, 2016 2 comments

Examine the following steps of privilege analysis for checking and revoking excessive, unused privileges granted to users: Identify the correct sequence of steps.

1. Create a policy to capture the privilege used by a user for privilege analysis.
2. Generate a report with the data captured for a specified privilege capture.
3. Start analyzing the data captured by the policy.
4. Revoke the unused privileges.
5. Compare the used and unused privileges' lists.
6. Stop analyzing the data.

A.

1, 3, 5, 6, 2, 4

B.

1, 3, 5, 2, 6, 4

C.

1, 3, 2, 5, 6, 4

D.

1, 3, 6, 2, 5, 4

E.

1, 3, 2, 5, 6, 4

Which statement is true about the archived redo log files?

Posted by seenagape on February 22, 2016 1 comment

Your database is running in ARCHIVELOG mode. The following parameters are set in your database instance: Which statement is true about the archived redo log files?

```
LOG_ARCHIVE_FORMAT = arch+_%t_%.arc
```

```
LOG_ARCHIVE_DEST_1 = 'LOCATION = /disk1/archive'
```

```
DB_RECOVERY_FILE_DEST_SIZE = 50G
```

```
DB_RECOVERY_FILE = '/u01/oradata'
```

A.

They are created only in the location specified by the LOG_ARCHIVE_DEST_1 parameter.

B.

They are created only in the Fast Recovery Area.

C.

They are created in the location specified by the LOG_ARCHIVE_DEST_1 parameter and in the default location \$ORACLE_HOME/dbs/arch.

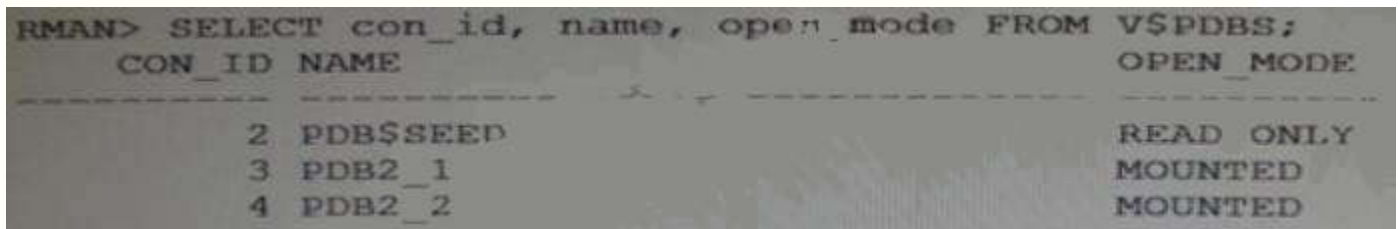
D.

They are created in the location specified by the LOG_ARCHIVE_DEST_1 parameter and the location specified by the DB_RECOVERY_FILE_DEST parameter.

Which data files will be backed up?

Posted by seenagape on February 22, 2016 2 comments

Your multitenant container database (CDB) is running in ARCHIVELOG mode. You connect to the CDB RMAN. Examine the following command and its output: You execute the following command: RMAN > BACKUP DATABASE PLUS ARCHIVELOG; Which data files will be backed up?



```
RMAN> SELECT con_id, name, open_mode FROM V$PDBS;
```

CON_ID	NAME	OPEN_MODE
2	PDB\$SEED	READ ONLY
3	PDB2_1	MOUNTED
4	PDB2_2	MOUNTED

A.

Data files that belong to only the root container

B.

Data files that belong to the root container and all the pluggable databases (PDBs)

C.

Data files that belong to only the root container and PDB\$SEED

D.

Data files that belong to the root container and all the PDBs excluding PDB\$SEED

What is the result?

Posted by seenagape on February 22, 2016 1 comment

You are administering a database stored in Automatic Storage management (ASM). The files are stored in the DATA disk group. You execute the following command: What is the result?

```
SQL > ALTER DISKGROUP data ADD ALIAS '+data/prod/myfile.dbf' FOR '+data/prod/myfile.dbf'
```

A.

The file '+data.231.54769' remains in the same location and a synonym 'myfile.dbf' is created.

B.

The file 'myfile.dbf' is created in '+data/prod' and the reference to '+data.231.54769' in the data dictionary removed.

C.

The file '+data.231.54769' is physically relocated to '+data/prod' and renamed as 'myfile.dbf'.

D.

The file '+data.231.54769' is renamed as 'myfile.dbf', and copied to '+data/prod'.

Explanation:

Which three functions are performed by the SQL Tuning Advisor?

Posted by seenagape on February 22, 2016 1 comment

Which three functions are performed by the SQL Tuning Advisor?

A.

Building and implementing SQL profiles

B.

Recommending the optimization of materialized views

C.

Checking query objects for missing and stale statistics

D.

Recommending bitmap, function-based, and B-tree indexes

E.

Recommending the restructuring of SQL queries that are using bad plans

Which statement is true?

Posted by seenagape on February 22, 2016 1 comment

Examine the following command: ALTER SYSTEM SET enable_ddl_logging=FALSE;
Which statement is true?

A.

Only DDL commands that resulted in the creation of new database files are logged.

B.

None of the data definition language (DDL) statements are logged in the trace file.

C.

Only DDL commands that resulted in errors are logged in the alert log file.

D.

A new log.xml file that contains the DDL statements is created, and the DDL command details are removed from the alert log file.

Which three steps should you perform to recover the control file and make the database fully operational?

Posted by seenagape on February 22, 2016 3 comments

Your multitenant container database (CDB) contains three pluggable database (PDBs). You find that the control file is damaged. You plan to use RMAN to recover the control file. There are no startup triggers associated with the PDBs. Which three steps should you perform to recover the control file and make the database fully operational?

A.

Start the database instance in the nomount stage and restore the control file from control file auto backup.

B.

Recover each pluggable database.

C.

Open all the pluggable databases.

D.

Mount the CDB and then recover and open the database, with the RESETLOGS option.

E.

Mount the container database (CDB) and restore the control file from the control file auto backup.

F.

Recover and open the CDB in NORMAL mode.

What should you do to accomplish this task?

Posted by seenagape on February 22, 2016 2 comments

A new report process containing a complex query is written, with high impact on the database. You want to collect basic statistics about query, such as the level of parallelism,

total database time, and the number of I/O requests. For the database instance STATISTICS_LEVEL, the initialization parameter is set to TYPICAL and the CONTROL_MANAGEMENT_PACK_ACCESS parameter is set to DIAGNOSTIC+TUNING. What should you do to accomplish this task?

A.

Execute the query and view Active Session History (ASH) for information about the query.

B.

Enable SQL trace for the query.

C.

Create a database operation, execute the query, and use the DBMS_SQL_MONITOR.REPORT_SQL_MONITOR function to view the report.

D.

Use the DBMS_APPLICATION_INFO.SET_SESSION_LONGOPS procedure to monitor query execution and view the information from the V\$SESSION_LONGOPS view.

Identify two valid options for adding a pluggable database (PDB) to an existing multitenant container database (CDB).

Posted by seenagape on February 22, 20163 comments

Identify two valid options for adding a pluggable database (PDB) to an existing multitenant container database (CDB).

A.

Use the CREATE DATABASE . . . ENABLE PLUGGABLE DATABASE statement to provision a PDB by copying file from the SEED.

B.

Use the DBMS_PDB package to plug an Oracle 11 g Release 2 (11.2.0.3.0) non-CDB database into an existing CDB.

C.

Use the CREATE PLUGGABLE DATABASE statement to create a PDB using the files from the SEED.

D.

Use the DBMS_PDB package to clone an existing PDB.

E.

Use the DBMS_PDB package to plug an Oracle 12c non-CDB database into an existing CDB.

What must you do to receive recommendations about the efficient use of indexes and materialized views to improve query performance?

Posted by seenagape on February 22, 20161 comment

Your database supports a DSS workload that involves the execution of complex queries: Currently, the library cache contains the ideal workload for analysis. You want to analyze some of the queries for an application that are cached in the library cache. What must you do to receive recommendations about the efficient use of indexes and materialized views to improve query performance?

A.

Create a SQL Tuning Set (STS) that contains the queries cached in the library cache and run the SQL Tuning Advisor (STA) on the workload captured in the STS.

B.

Run the Automatic Workload Repository Monitor (AWR).

C.

Create an STS that contains the queries cached in the library cache and run the SQL Performance Analyzer (SPA) on the workload captured in the STS.

D.

Create an STS that contains the queries cached in the library cache and run the SQL Access Advisor on the workload captured in the STS.

Identify the correct sequence of steps:

Posted by seenagape on February 22, 2016 1 comment

The following parameters are set for your Oracle 12c database instance: You want to manage the SQL plan evolution task manually. Examine the following steps: 1. Set the evolve task parameters. 2. Create the evolve task by using the DBMS_SPM.CREATE_EVOLVE_TASK function. 3. Implement the recommendations in the task by using the DBMS_SPM.IMPLEMENT_EVOLVE_TASK function. 4. Execute the evolve task by using the DBMS_SPM.EXECUTE_EVOLVE_TASK function. 5. Report the task outcome by using the DBMS_SPM.REPORT_EVOLVE_TASK function. Identify the correct sequence of steps:

```
OPTIMIZER_CAPTURE_SQL_PLAN_BASELINES=FALSE
```

```
OPTIMIZER_USE_SQL_PLAN_BASELINES=TRUE
```

A.

2, 4, 5

B.

2, 1, 4, 3, 5

C.

1, 2, 3, 4, 5

D.

1, 2, 4, 5

Which option would you consider first to decrease the wait event immediately?

Posted by seenagape on February 22, 2016 1 comment

In a recent Automatic Workload Repository (AWR) report for your database, you notice a high number of buffer busy waits. The database consists of locally managed tablespaces with free list managed segments. On further investigation, you find that buffer busy waits are caused by contention on data blocks. Which option would you consider first to decrease the wait event immediately?

- A.
Decreasing PCTUSED
- B.
Increasing db_buffer_cache based on the V\$DB_CACHE_ADVICE recommendation
- C.
Using Automatic Segment Space Management (ASSM)
- D.
Decreasing PCTFREE
- E.
Increasing the number of DBWN process

Which three statements are true about the effect of this command?

Posted by seenagape on February 22, 2016 2 comments

Examine this command: SQL > exec DBMS_STATS.SET_TABLE_PREFS ('SH', 'CUSTOMERS', 'PUBLISH', 'false'); Which three statements are true about the effect of this command?

- A.
Statistics collection is not done for the CUSTOMERS table when schema stats are gathered.
- B.
Statistics collection is not done for the CUSTOMERS table when database stats are gathered.
- C.
Any existing statistics for the CUSTOMERS table are still available to the optimizer at parse time.
- D.
Statistics gathered on the CUSTOMERS table when schema stats are gathered are stored as pending statistics.
- E.
Statistics gathered on the CUSTOMERS table when database stats are gathered are stored as pending statistics.

Explanation:

Which three are prerequisites for successful execution of the command?

Posted by seenagape on February 22, 2016 2 comments

Examine the following impdp command to import a database over the network from a pre-12c Oracle database (source): Which three are prerequisites for successful execution of the command?

```
$> impdp <user_name> full=Y network_link=hrdb_test transport_datafiles=
'/u01/app/oracle/oradata/hrdb/sales01.dbf',
'/u01/app/oracle/oradata/hrdb/cust01.dbf',
'/u01/app/oracle/oradata/hrdb/emp01.dbf'
version=12 logfile=import.log
```

- A.
The import operation must be performed by a user on the target database with the DATAPUMP_IMP_FULL_DATABASE role, and the database link must connect to a user on the source database with the DATAPUMP_EXD_FULL_DATABASE role.
- B.
The impdp operation must be performed by the same user that performed the expdp operation.
- C.
The path of data files on the target database must be the same as that on the source database.
- D.
All the user-defined tablespaces must be in read-only mode on the source database.
- E.
The export dump file must be created before starting the import on the target database.
- F.
The source and target database must be running on the same platform with the same endianness.

Which two are true concerning a multitenant container database with three pluggable database?

Posted by seenagape on February 22, 2016 3 comments

Which two are true concerning a multitenant container database with three pluggable database?

- A.
All administration tasks must be done to a specific pluggable database.
- B.
The pluggable databases increase patching time.
- C.
The pluggable databases reduce administration effort.

D.

The pluggable databases are patched together.

E.

Pluggable databases are only used for database consolidation.

Which statement is true?

Posted by seenagape on February 22, 2016 2 comments

Examine the current value for the following parameters in your database instance: You issue the following command to increase the value of DB_8K_CACHE_SIZE: Which statement is true?

```
SGA_MAX_SIZE = 1024M
SGA_TARGET = 700M
DB_8K_CACHE_SIZE = 124M
LOG_BUFFER = 200M
SQL> ALTER SYSTEM SET DB_8K_CACHE_SIZE=140M;
```

A.

It succeeds only if memory is available from the autotuned components if SGA.

B.

It fails because an increase in DB_8K_CACHE_SIZE cannot be accommodated within

SGA_TARGET.

C.

It fails because the DB_8K_CACHE_SIZE parameter cannot be changed dynamically.

D.

It fails because an increase in DB_8K_CACHE_SIZE cannot be accommodated within SGA_MAX_SIZE.

Which three statements are true concerning unplugging a pluggable database (PDB)?

Posted by seenagape on February 22, 2016 4 comments

Which three statements are true concerning unplugging a pluggable database (PDB)?

A.

The PDB must be open in read only mode.

B.

The PDB must be closed.

C.

The unplugged PDB becomes a non-CDB.

D.

The unplugged PDB can be plugged into the same multitenant container database (CDB)

E.

The unplugged PDB can be plugged into another CDB.

F.

The PDB data files are automatically removed from disk.

Which three statements are true about using an invisible column in the PRODUCTS table?

Posted by seenagape on February 22, 2016 1 comment

Examine the following command: Which three statements are true about using an invisible column in the PRODUCTS table?

```
* SELECT * FROM statements in SQL
* DESCRIBE commands in SQL*Plus
* %ROWTYPE attribute declarations in PL/SQL
* Describes in Oracle Call Interface (OCI)
```

A.

Referential integrity constraint cannot be set on the invisible column.

B.

The invisible column cannot be made visible and can only be marked as unused.

C.

The %ROWTYPE attribute declarations in PL/SQL to access a row will not display the invisible column in the output.

D.

The DESCRIBE commands in SQL *Plus will not display the invisible column in the output.

E.

A primary key constraint can be added on the invisible column.

which database users is the audit policy now active?

Posted by seenagape on February 22, 2016 1 comment

You wish to enable an audit policy for all database users, except SYS, SYSTEM, and SCOTT. You issue the following statements: For which database users is the audit policy now active?

```
SQL> AUDIT POLICY ORA_DATABASE_PARAMETER EXCEPT SYS;  
SQL> AUDIT POLICY ORA_DATABASE_PARAMETER EXCEPT SYSTEM;  
SQL> AUDIT POLICY ORA_DATABASE_PARAMETER EXCEPT SCOTT;
```

- A.
All users except SYS
- B.
All users except SCOTT
- C.
All users except sys and SCOTT
- D.
All users except sys, system, and SCOTT

Which two statements are true regarding the command?

Posted by seenagape on February 22, 2016 2 comments

On your Oracle 12c database, you invoked SQL *Loader to load data into the EMPLOYEES table in the HR schema by issuing the following command: \$> sqlldr hr/hr@pdb table=employees Which two statements are true regarding the command?

- A.
It fails because no SQL *Loader control file location is specified.
- B.
It succeeds with default settings if the EMPLOYEES table belonging to HR is already defined in the database.
- C.
It fails because no SQL *Loader data file location is specified.
- D.
It fails if the HR user does not have the CREATE ANY DIRECTORY privilege.

Explanation:

What must you do to activate the new default value for numeric full redaction?

Posted by seenagape on February 22, 2016 1 comment

After implementing full Oracle Data Redaction, you change the default value for the NUMBER data type as follows: After changing the value, you notice that FULL redaction continues to redact numeric data with zero. What must you do to activate the new default value for numeric full redaction?

```

SQL> SELECT NUMBER_VALUE FROM REDACTION_VALUES_FOR_TYPE_
NUMBER_VALUE
-----
0

SQL> EXEC DBMS_REDACT.UPDATE_FULL_REDACTION_VALUES(-1)
PL/SQL procedure successfully completed.

SQL> select number_value from redaction_values_for_type_
NUMBER_VALUE
-----
-1

```

- A.
Restart the database instance.
- B.
Flush the shared pool.
- C.
Re-connect the sessions that access objects with redaction policies defined on them.
- D.
Re-enable redaction policies that use FULL data redaction.
- E.
Re-create redaction policies that use FULL data redaction.

Which two must you do to track the transactions?

Posted by seenagape on February 22, 2016 1 comment

You must track all transactions that modify certain tables in the sales schema for at least three years. Automatic undo management is enabled for the database with a retention of one day. Which two must you do to track the transactions?

- A.
Enable supplemental logging for the database.
- B.
Specify undo retention guarantee for the database.
- C.
Create a Flashback Data Archive in the tablespace where the tables are stored.

D.

Create a Flashback Data Archive in any suitable tablespace.

E.

Enable Flashback Data Archiving for the tables that require tracking.

Which technique will move the table and indexes while maintaining the highest level of availability to the application?

Posted by seenagape on February 22, 2016 2 comments

You are the DBA supporting an Oracle 11g Release 2 database and wish to move a table containing several DATE, CHAR, VARCHAR2, and NUMBER data types, and the table's indexes, to another tablespace. The table does not have a primary key and is used by an OLTP application. Which technique will move the table and indexes while maintaining the highest level of availability to the application?

A.

Edition-Based Table Redefinition.

B.

Online Table Redefinition.

C.

An ALTER TABLE MOVE to move the table and ALTER INDEX REBUILD to move the indexes.

D.

Oracle Data Pump.

E.

An ALTER TABLE MOVE to move the table and ALTER INDEX REBUILD ONLINE to move the indexes.

Identify the reason the instance failed to start.

Posted by seenagape on February 22, 2016 1 comment

To implement Automatic Management (AMM), you set the following parameters: When you try to start the database instance with these parameter settings, you receive the following error message: SQL > startup ORA-00824: cannot set SGA_TARGET or MEMORY_TARGET due to existing internal settings, see alert log for more information. Identify the reason the instance failed to start.

```
MEMORY_MAX_TARGET=600M
SGA_MAX_SIZE=500M
MEMORY_TARGET=600M
OPEN_CURSORS=300
SGA_TARGET=300M
PROCESSES=150
STATISTICS_LEVEL=BASIC
PGA_AGGREGATE_TARGET=0
```

- A.
The PGA_AGGREGATE_TARGET parameter is set to zero.
- B.
The STATISTICS_LEVEL parameter is set to BASIC.
- C.
Both the SGA_TARGET and MEMORY_TARGET parameters are set.
- D.
The SGA_MAX_SIZE and SGA_TARGET parameter values are not equal.

Explanation:

What are two benefits of installing Grid Infrastructure software for a stand-alone server before installing and creating an Oracle database?

Posted by seenagape on February 22, 20163 comments

What are two benefits of installing Grid Infrastructure software for a stand-alone server before installing and creating an Oracle database?

- A.
Effectively implements role separation
- B.
Enables you to take advantage of Oracle Managed Files.
- C.
Automatically registers the database with Oracle Restart.
- D.
Helps you to easily upgrade the database from a prior release.
- E.
Enables the Installation of Grid Infrastructure files on block or raw devices.

Identify two correct statements about multitenant architectures.

Posted by seenagape on February 22, 20162 comments

Identify two correct statements about multitenant architectures.

A.

Multitenant architecture can be deployed only in a Real Application Clusters (RAC) configuration.

B.

A PDB can have a private undo tablespace.

C.

Patches are always applied at the CDB level.

D.

Multiple non-RAC CDB instances can mount the same PDB as long as they are on the same server.

E.

Multiple pluggable databases (PDBs) share certain multitenant container database (CDB) resources.

F.

Multiple CDBs share certain PDB resources.

Which two actions does the script perform?

Posted by seenagape on February 22, 2016 1 comment

You upgrade your Oracle database in a multiprocessor environment. As a recommended you execute the following script: SQL > @utlrp.sql Which two actions does the script perform?

A.

Sequential recompilation of Java code

B.

Parallel compilation of only the stored PL/SQL code

C.

Sequential recompilation of only the stored PL/SQL code

D.

Parallel recompilation of any stored PL/SQL code

E.

Sequential recompilation of any stored PL/SQL code

F.

Parallel recompilation of Java code

Which two statement is true concerning dropping a pluggable database (PDB)?

Posted by seenagape on February 22, 20166 comments

Which two statement is true concerning dropping a pluggable database (PDB)?

A.

A dropped PDB can never be plugged back into a multitenant container database (CDB).

B.

The PDB data files are always removed from disk.

C.

The PDB must be open in read-only mode.

D.

The PDB must be in mount state.

E.

The PDB must be unplugged.

Identify three possible reasons for this.

Posted by seenagape on February 22, 20161 comment

You notice a high number of waits for the db file scattered read and db file sequential read events in the recent Automatic Database Diagnostic Monitor (ADDM) report. After further investigation, you find that queries are performing too many full table scans and indexes are not being used even though the filter columns are indexed. Identify three possible reasons for this.

A.

Missing or stale histogram statistics

B.

Undersized shared pool

C.

High clustering factor for the indexes

D.

High value for the DB_FILE_MULTIBLOCK_READ_COUNT parameter

E.

Oversized buffer cache

Explanation:

Which three features work together, to allow a SQL statement to have different cursors for the same statement based on different selectivity ranges?

Posted by seenagape on February 22, 2016 4 comments

Which three features work together, to allow a SQL statement to have different cursors for the same statement based on different selectivity ranges?

A.

Adaptive Cursor Sharing

B.

Bind variable used in a SQL statement

C.

SQL Plan Baselines

D.

Bind Variable Peeking

E.

Literals in a SQL statement

You want to know which change caused this performance difference.

Posted by seenagape on February 22, 2016 2 comments

You notice a performance change in your production Oracle 12c database. You want to know which change caused this performance difference.

A.

Compare Period ADDM report

B.

AWR Compare Period report

C.

Active Session History (ASH) report

D.

Taking a new snapshot and comparing it with a preserved snapshot

Explanation:

Identify the correct sequence of steps.

Posted by seenagape on February 22, 2016 1 comment

You want to capture column group usage and gather extended statistics for better cardinality estimates for the CUSTOMERS table in the SH schema. Examine the following steps: 1. Issue the SELECT DBMS_STATS.CREATE_EXTENDED_STATS ('SH', 'CUSTOMERS') FROM dual statement. 2. Execute the DBMS_STATS.SEED_COL_USAGE (null, 'SH', 500) procedure. 3. Execute the required queries on the CUSTOMERS table. 4. Issue the SELECT

DBMS_STATS.REPORT_COL_USAGE ('SH', 'CUSTOMERS') FROM dual statement.
Identify the correct sequence of steps.

- A.
3, 2, 4, 1
- B.
3, 2, 1, 4
- C.
2, 3, 4, 1
- D.
4, 1, 3, 2

Which three statements are true about Automatic Workload Repository (AWR)?

Posted by seenagape on February 22, 20164 comments

Which three statements are true about Automatic Workload Repository (AWR)?

- A.
AWR contains system wide tracing and logging information.
- B.
AWR computes time model statistics based on time usage for activities, which are displayed in the v\$SYS time model and V\$SESS_TIME_MODEL views.
- C.
All AWR tables belong to the SYSTEM schema.
- D.
The snapshots collected by AWR are used by the self-tuning components in the database
- E.
The AWR data is stored in memory and in the database.

Which two tasks must you perform to add users with SYSBACKUP, SYSDG, and SYSKM privilege to the password file?

Posted by seenagape on February 22, 20161 comment

You upgraded your database from pre-12c to a multitenant container database (CDB) containing pluggable databases (PDBs). Examine the query and its output: Which two tasks must you perform to add users with SYSBACKUP, SYSDG, and SYSKM privilege to the password file?

```
SQL> SELECT * FROM v$PWFILE_USERS;
```

USERNAME	SYSDBA	SYSOP	SYSAS	SYSBA	S
SYS	TRUE	TRUE	FALSE	FALSE	F

- A.
Assign the appropriate operating system groups to SYSBACKUP, SYSDG, SYSKM.
- B.
Grant SYSBACKUP, SYSDG, and SYSKM privileges to the intended users.
- C.
Re-create the password file with SYSBACKUP, SYSDG, and SYSKM privilege and the FORCE argument set to No.
- D.
Re-create the password file with SYSBACKUP, SYSDG, and SYSKM privilege, and FORCE arguments set to Yes.
- E.
Re-create the password file in the Oracle Database 12c format.

How would you guarantee that the blocks for the table never age out?

Posted by seenagape on February 22, 2016 1 comment

An application accesses a small lookup table frequently. You notice that the required data blocks are getting aged out of the default buffer cache. How would you guarantee that the blocks for the table never age out?

- A.
Configure Automata Shared Memory Management.
- B.
Configure Automatic Memory Management
- C. Configure the RECYCLE buffer pool and alter the table with the corresponding storage clause.
- D.
Increase the database buffer cache size.
- E.
Configure the KEEP buffer pool and alter the table with the corresponding storage clause.

What happens after issuing the SHUTDOWN TRANSACTIONAL statement?

Posted by seenagape on February 22, 2016 1 comment

You conned using SQL Plus to the root container of a multitenant container database (CDB) with SYSDBA privilege. The CDB has several pluggable databases (PDBs) open in the read/write mode. There are ongoing transactions in both the CDB and PDBs. What happens after issuing the SHUTDOWN TRANSACTIONAL statement?

A.

The shutdown proceeds as soon as all transactions in the CDB are either committed or rolled back.

B.

The shutdown proceeds as soon as all transactions in both the CDB and PDBs are either committed or rolled back.

C.

The statement results in an error because there are open PDBs.

D.

The shutdown proceeds immediately. The shutdown proceeds as soon as all transactions in the PDBs are either committed or rolled back.

Which three techniques can you use to achieve this?

Posted by seenagape on February 22, 2016 1 comment

You are planning the creation of a new multitenant container database (CDB) and want to store the ROOT and SEED container data files in separate directories. You plan to create the database using SQL statements. Which three techniques can you use to achieve this?

A.

Specify all files in the CREATE DATABASE statement without using Oracle managed Files (OMF).

B.

Use Oracle Managed Files (OMF).

C.

Specify the SEED_FILE_NAME_CONVERT clause.

D.

Specify the PDB_FILE_NAME_CONVERT initialization parameter.

E.

Specify the DB_FILE_NAME_CONVERT initialization parameter.

Which technique should you use to minimize down time while plugging this non-CDB into the CDB?

Posted by seenagape on February 22, 2016 3 comments

You are about to plug a multi-terabyte non-CDB into an existing multitenant container database (CDB). The characteristics of the non-CDB are as follows: – Version: Oracle Database 11g Release 2 (11.2.0.2.0) 64-bit – Character set: AL32UTF8 – National character set: AL16UTF16 – O/S: Oracle Linux 6 64-bit The characteristics of the CDB are as follows: – Version: Oracle Database 12c Release 1 64-bit – Character Set: AL32UTF8 – National character set: AL16UTF16 – O/S: Oracle Linux 6 64-bit Which technique should you use to minimize down time while plugging this non-CDB into the CDB?

- A.
RMAN
- B.
The DBMS_PDB package
- C.
Transportable database
- D.
Data Pump full export/import
- E.
Transportable tablespace

What should you use to achieve this?

Posted by seenagape on February 22, 2016 5 comments

Your database supports an online transaction processing (OLTP) application. The application is undergoing some major schema changes, such as addition of new indexes and materialized views. You want to check the impact of these changes on workload performance. What should you use to achieve this?

- A.
Database replay
- B.
SQL Tuning Advisor
- C.
SQL Access Advisor
- D.
SQL Performance Analyzer

E.

Automatic Workload Repository compare reports

Which four statements are true about this administrator establishing connections to root in a CDB that has been opened in read only mode?

Posted by seenagape on February 22, 2016 1 comment

An administrator account is granted the CREATE SESSION and SET CONTAINER system privileges. A multitenant container database (CDB) instant has the following parameter set: THREADED_EXECUTION = FALSE Which four statements are true about this administrator establishing connections to root in a CDB that has been opened in read only mode?

A.

You can connect as a local user by using the connect statement.

B.

You can connect as a local user by using the SET CONTAINER statement.

C.

You can connect by using a Net Service name.

D.

You can connect by using easy connect.

E.

You can conned as a common user by using the connect statement.

F.

You can connect by using OS authentication.

Which statement is true?

Posted by seenagape on February 22, 2016 1 comment

Examine the following query output: You issue the following command to import tables into the hr schema: Which statement is true?

```
SQL> SELECT name, force_logging FROM v$database;
NAME          FORCE_LOGGING
-----
PROD          NO
```

```
$ > impdp hr/hr directory = dumpdir dumpfile = hr_new.dmp schemas=hr
TRANSFORM=DISABLE_ARCHIVE_LOGGING: Y
```

A.

All database operations performed by the impdp command are logged.

B.

Only CREATE INDEX and CREATE TABLE statements generated by the import are logged.

C.

Only CREATE TABLE and ALTER TABLE statements generated by the import are logged.

D.

None of the operations against the master table used by Oracle Data Pump to coordinate its activities are logged.

Which three findings would you get from the report?

Posted by seenagape on February 22, 2016 1 comment

You notice a performance change in your production Oracle database and you want to know which change has made this performance difference. You generate the Compare Period Automatic Database Diagnostic Monitor (ADDM) report to further investigation. Which three findings would you get from the report?

A.

It detects the top wait events causing performance degradation.

B.

It gives information about statistics collection in both time periods.

C.

It shows the resource usage for CPU, memory, and I/O in both time periods.

D.

It detects any configuration change that caused a performance difference in both time periods.

E.

It identifies any workload change that caused a performance difference in both time periods.

F.

It shows the difference in the size of memory pools in both time periods.

Identify the reason why the optimizer chose different execution plans.

Posted by seenagape on February 22, 2016 4 comments

Examine the parameter for your database instance: You generated the execution plan for the following query in the plan table and noticed that the nested loop join was done. After actual execution of the query, you notice that the hash join was done in the execution plan: Identify the reason why the optimizer chose different execution plans.

NAME	TYPE	VALUE
optimizer_adaptive_reporting_only	boolean	FALSE
optimizer_capture_sql_plan_baselines	boolean	FALSE
optimizer_dynamic_sampling	integer	2
optimizer_features_enable	string	12.1.0

```

SQL> SELECT product_name
FROM    order_items o, product_information p
WHERE   o.unit_price < 10
AND     quantity > 1
AND     p.product_id = o.product_id;

30 rows selected.

```

- A.
The optimizer used a dynamic plan for the query.
- B.
The optimizer chose different plans because automatic dynamic sampling was enabled.
- C.
The optimizer used re-optimization cardinality feedback for the query.
- D.
The optimizer chose different plan because extended statistics were created for the columns used.

Which three statements are true about adaptive SQL plan management?

Posted by seenagape on February 22, 2016 3 comments

Which three statements are true about adaptive SQL plan management?

- A.
It automatically performs verification or evolves non-accepted plans, in COMPREHENSIVE mode when they perform better than existing accepted plans.
- B.
The optimizer always uses the fixed plan, if the fixed plan exists in the plan baseline.
- C.
It adds new, better plans automatically as fixed plans to the baseline.
- D.
The non-accepted plans are automatically accepted and become usable by the optimizer if they perform better than the existing accepted plans.
- E.
The non-accepted plans in a SQL plan baseline are automatically evolved, in

COMPREHENSIVE mode, during the nightly maintenance window and a persistent verification report is generated.

Which three tablespaces are created by default in HR_PDB?

Posted by seenagape on February 22, 2016 1 comment

You create a new pluggable database, HR_PDB, from the seed database. Which three tablespaces are created by default in HR_PDB?

- A.
SYSTEM
- B.
USERS
- C.
TEMP
- D.
UNDO
- E.
SYSAUX
- F.
EXAMPLE

Which two statements are true about variable extent size support for large ASM files?

Posted by seenagape on February 22, 2016 1 comment

Which two statements are true about variable extent size support for large ASM files?

- A.
The metadata used to track extents in SGA is reduced.
- B.
Rebalance operations are completed faster than with a fixed extent size
- C.
An ASM Instance automatically allocates an appropriate extent size.
- D.
Resync operations are completed faster when a disk comes online after being taken offline.
- E.
Performance improves in a stretch cluster configuration by reading from a local copy of an extent.

Explanation:

What is the quickest way to recover the contents of the OCA.EXAM_RESULTS table to the OCP schema?

Posted by seenagape on February 22, 2016 2 comments

You executed a DROP USER CASCADE on an Oracle 11g release 1 database and immediately realized that you forgot to copy the OCA.EXAM_RESULTS table to the OCP schema. The RECYCLE_BIN enabled before the DROP USER was executed and the OCP user has been granted the FLASHBACK ANY TABLE system privilege. What is the quickest way to recover the contents of the OCA.EXAM_RESULTS table to the OCP schema?

- A.
Execute FLASHBACK TABLE OCA.EXAM_RESULTS TO BEFORE DROP RENAME TO OCP.EXAM_RESULTS; connected as SYSTEM.
- B.
Execute FLASHBACK TABLE OCA.EXAM_RESULTS TO BEFORE DROP RENAME TO EXAM_RESULTS; connected as the OCP user.
- C.
Recovery the table using Database Point In Time Recovery.
- D.
Recover the table using traditional Tablespace Point In Time Recovery.
- E.
Recover the table using Automated Tablespace Point In Time Recovery.

which they do not have any privileges?

Posted by seenagape on February 22, 2016 1 comment

In your multitenant container database (CDB) containing pluggable database (PDBs), the HR user executes the following commands to create and grant privileges on a procedure:
CREATE OR REPLACE PROCEDURE create_test_v (v_emp_id NUMBER, v_ename VARCHAR2, v_salary NUMBER, v_dept_id NUMBER) BEGIN INSERT INTO hr.test VALUES (V_emp_id, V_ename, V_salary, V_dept_id); END; / GRANT EXECUTE ON CREATE_TEST TO john, jim, smith, king; How can you prevent users having the EXECUTE privilege on the CREATE_TEST procedure from inserting values into tables on which they do not have any privileges?

- A.
Create the CREATE_TEST procedure with definer's rights.
- B.
Grant the EXECUTE privilege to users with GRANT OPTION on the CREATE_TEST procedure.

C.

Create the CREATE_TEST procedure with invoker's rights.

D.

Create the CREATE_TEST procedure as part of a package and grant users the EXECUTE privilege the package.

What are two effects of not using the "ENABLE PLUGGABLE database" clause?

Posted by seenagape on February 22, 2016 1 comment

You created a new database using the "create database" statement without specifying the "ENABLE PLUGGABLE" clause. What are two effects of not using the "ENABLE PLUGGABLE database" clause?

A.

The database is created as a non-CDB but will become a CDB whenever the first PDB is plugged in.

B.

The database is created as a non-CDB and can never contain a PDB.

C.

The database is treated as a PDB and must be plugged into an existing multitenant container database (CDB).

D.

The database is created as a non-CDB and can never be plugged into a CDB.

E.

The database is created as a non-CDB but can be plugged into an existing CDB.

Explanation:

What is the effect of specifying the "ENABLE PLUGGABLE DATABASE" clause in a "CREATE DATABASE" statement?

Posted by seenagape on February 22, 2016 3 comments

What is the effect of specifying the "ENABLE PLUGGABLE DATABASE" clause in a "CREATE DATABASE" statement?

A.

It will create a CDB with root opened and seed mounted.

B.

It will create a CDB that must be plugged into an existing CDB.

C.

It will create a CDB with root and seed opened and one PDB mounted.

D.

It will create a multitenant container database (CDB) with only the root opened.

E.

It will create a CDB with root opened and seed read only.

How should the DB_FLASH_CACHE_SIZE be configured to use both devices?

Posted by seenagape on February 22, 2016 1 comment

You have installed two 64G flash devices to support the Database Smart Flash Cache feature on your database server that is running on Oracle Linux. You have set the DB_SMART_FLASH_FILE parameter: DB_FLASH_CACHE_FILE= '/dev/flash_device_1 ','/dev/flash_device_2' How should the DB_FLASH_CACHE_SIZE be configured to use both devices?

A.

Set DB_FLASH_CACHE_SIZE = 64G.

B.

Set DB_FLASH_CACHE_SIZE = 64G, 64G

C.

Set DB_FLASH_CACHE_SIZE = 128G.

D.

DB_FLASH_CACHE_SIZE is automatically configured by the instance at startup.

Which three initialization parameters are not controlled by Automatic Shared Memory Management (ASMM)?

Posted by seenagape on February 22, 2016 1 comment

Examine the following parameters for a database instance: Which three initialization parameters are not controlled by Automatic Shared Memory Management (ASMM)?

```
MEMORY_MAX_TARGET=0
```

```
MEMORY_TARGET=0
```

```
SGA_TARGET=0
```

```
PGA_AGGREGATE_TARGET=500m
```

A.

LOG_BUFFER

B.
DB_KEEP_CACHE_SIZE

C.
DB_16K_CACHE_SIZE

D.
SORT_AREA_SIZE

E.
JAVA_POOL_SIZE

F.
STREAMS_POOL_SIZE

Explanation:

Which three statements are true regarding the SQL* Loader operation performed using the control file?

Posted by seenagape on February 22, 2016 2 comments

Examine the contents of SQL loader control file: Which three statements are true regarding the SQL* Loader operation performed using the control file?

A.
An EMP table is created if a table does not exist. Otherwise, if the EMP table is appended with the loaded data.

B.
The SQL* Loader data file myfile1.dat has the column names for the EMP table.

C.
The SQL* Loader operation fails because no record terminators are specified.

D.
Field names should be the first line in the both the SQL* Loader data files.

E.
The SQL* Loader operation assumes that the file must be a stream record format file with the normal carriage return string as the record terminator.

What is the result?

Posted by seenagape on February 22, 2016 3 comments

In your multitenant container database (CDB) containing pluggable database (PDBs), you granted the CREATE TABLE privilege to the common user C ## A_ADMIN in root and all PDBs. You execute the following command from the root container: SQL > REVOKE create table FROM C ## A_ADMIN; What is the result?

A.

It fails and reports an error because the CONTAINER=ALL clause is not used.

B.

It executes successfully and the CREATE TABLE privilege is revoked from C ##

A_ADMIN in all PDBs.

C.

It excludes successfully and the CREATE TABLE privilege is revoked from C ##

A_ADMIN in root and all PDBs.

D.

It executes successfully and the CREATE TABLE privilege is revoked from C ##

A_ADMIN in root only.

E.

It fails and reports an error because the CONTAINER=CURRENT clause is not used.

Which two statements are true concerning the Resource Manager plans for individual pluggable databases (PDB plans) in a multitenant container database (CDB)?

Posted by seenagape on February 22, 2016 3 comments

Which two statements are true concerning the Resource Manager plans for individual pluggable databases (PDB plans) in a multitenant container database (CDB)?

A.

If no PDB plan is enabled for a pluggable database, then all sessions for that PDB are treated to an equal degree of the resource share of that PDB.

B.

In a PDB plan, subplans may be used with up to eight consumer groups.

C.

If a PDB plan is enabled for a pluggable database, then resources are allocated to consumer groups across all PDBs in the CDB.

D.

If no PDB plan is enabled for a pluggable database, then the PDB share in the CDB plan is dynamically calculated.

E.

If a PDB plan is enabled for a pluggable database, then resources are allocated to consumer groups based on the shares provided to the PDB in the CDB plan and the shares provided to the consumer groups in the PDB plan.

Explanation:

Which two statements are true?

Posted by seenagape on February 22, 2016 2 comments

You use a recovery catalog for maintaining your database backups. You execute the following command: Which two statements are true?

You execute the following command:

```
$rman TARGET / CATALOG rman / cat@catdb
```

```
RMAN > BACKUP VALIDATE DATABASE ARCHIVELOG ALL;
```

- A.
Checks are performed to confirm whether all database files exist in correct locations
- B.
Backup sets containing both data files and archive logs are created.
- C.
Checks are performed for physical corruptions.
- D.
Corrupted blocks, if any, are repaired.
- E.
Checks are performed for logical corruptions.

Which three statements are true concerning the multitenant architecture?

Posted by seenagape on February 22, 2016 1 comment

Which three statements are true concerning the multitenant architecture?

- A.
A PDB can have a private temp tablespace.
- B.
PDBs can share the sysaux tablespace.
- C.
Log switches occur only at the multitenant container database (CDB) level.
- D.
Different PDBs can have different default block sizes.
- E.
PDBs share a common system tablespace.
- F.
Instance recovery is always performed at the CDB level.

G.

Each pluggable database (PDB) has its own set of background processes.

Which two actions would reduce the job's elapsed time?

Posted by seenagape on February 22, 2016 4 comments

You notice that the elapsed time for an important database scheduler Job is unacceptably long. The job belongs to a scheduler job class and window. Which two actions would reduce the job's elapsed time?

A.

Increasing the job's relative priority within the Job class to which it belongs

B.

Increasing the priority of the scheduler window to which the job belongs

C.

Increasing the value of the JOB_QUEUE_PROCESSES parameter

D.

Increasing the resource allocation for the consumer group mapped to the scheduler job's job class within the plan mapped to the scheduler window

E.

Increasing the priority of the job class to which the job belongs

F.

Moving the job to an existing higher priority scheduler window with the same schedule and duration

Explanation:

Which two methods or commands would you use to accomplish this task?

Posted by seenagape on February 22, 2016 1 comment

You plan to migrate your database from a File system to Automata Storage Management (ASM) on same platform. Which two methods or commands would you use to accomplish this task?

A.

RMAN CONVERT command

B.

Data Pump Export and import

C.

Conventional Export and Import

D.

The BACKUP AS COPY DATABASE . . . command of RMAN

E.

DBMS_FILE_TRANSFER with transportable tablespace

Which two statements are true about the outcome after running the script?

Posted by seenagape on February 22, 2016 1 comment

You run a script that completes successfully using SQL*Plus that performs these actions: 1. Creates a multitenant container database (CDB) 2. Plugs in three pluggable databases (PDBs) 3. Shuts down the CDB instance 4. Starts up the CDB instance using STARTUP OPEN READ WRITE Which two statements are true about the outcome after running the script?

A.

The PDBs will be opened read/write.

B.

The seed will be in mount state.

C.

The seed will be opened read-only.

D.

The seed will be opened read/write.

E.

The other PDBs will be opened read-only.

F.

The other PDBs will be in mount state.

Which two statements are true when a session logged in as SCOTT queries the SAL column in the view and the table?

Posted by seenagape on February 21, 2016 1 comment

You execute the following piece of code with appropriate privileges: User SCOTT has been granted the CREATE SESSION privilege and the MGR role. Which two statements are true when a session logged in as SCOTT queries the SAL column in the view and the table?

```

BEGIN
  DBMS_REDACT.ADD_POLICY(
    OBJECT_SCHEMA => 'SCOTT',
    OBJECT_NAME   => 'EMP',
    POLICY_NAME   => 'SCOTT_EMP',
    COLUMN_NAME   => 'SAL',
    FUNCTION_TYPE => DBMS_REDACT.FULL,
    EXPRESSION    => 'SYS_CONTEXT(''SYS_SESSION_ROLES'',
END;
/

CREATE VIEW SCOTT.EMP_V AS SELECT * FROM SCOTT.EMP;

BEGIN
  DBMS_REDACT.ADD_POLICY(
    OBJECT_SCHEMA => 'SCOTT',
    OBJECT_NAME   => 'EMP_V',
    POLICY_NAME   => 'SCOTT_EMP_V',
    COLUMN_NAME   => 'SAL',
    FUNCTION_TYPE => DBMS_REDACT.NONE,
    EXPRESSION    => 'SYS_CONTEXT(''SYS_SESSION_ROLES'',
END;
/

```

A.

Data is redacted for the EMP.SAL column only if the SCOTT session does not have the MGR role set.

B.

Data is redacted for EMP.SAL column only if the SCOTT session has the MGR role set.

C.

Data is never redacted for the EMP_V.SAL column.

D.

Data is redacted for the EMP_V.SAL column only if the SCOTT session has the MGR role set.

E.

Data is redacted for the EMP_V.SAL column only if the SCOTT session does not have the MGR role set.

What happens to the sessions that are presently connected to the database Instance?

Posted by seenagape on February 21, 2016 1 comment

Your database is open and the LISTENER listener running. You stopped the wrong listener LISTENER by issuing the following command: lsnrctl > STOP What happens to the sessions that are presently connected to the database Instance?

- A.
They are not allowed to perform any operations until the listener LISTENER is started.
- B.
They are able to perform only queries.
- C.
They are not affected and continue to function normally.
- D.
They are terminated and the active transactions are rolled back.

Which three statements are true about using flashback database in a multitenant container database (CDB)?

Posted by seenagape on February 21, 20163 comments

Which three statements are true about using flashback database in a multitenant container database (CDB)?

- A.
A CDB can be flashed back specifying the desired target point in time or an SCN, but not a restore point.
- B.
The DB_FLASHBACK RETENTION_TARGET parameter must be set to enable flashback of the CDB.
- C.
Individual PDBs can be flashed back without flashing back the entire CDB.
- D.
The root container can be flashed back without flashing back the pluggable databases (PDBs).
- E.
To enable flashback database, the CDB must be mounted.

Which two statements are true?

Posted by seenagape on February 21, 20163 comments

You execute the following PL/SQL: Which two statements are true?

```

BEGIN
DBMS_FGA.add_policy(
object_schema => 'JIM',
object_name   > 'PRODUCTS',
policy_name   > 'PROD_AUDIT',
audit_condition => 'PRICE > 10000',
audit_column => 'PRICE');
END;
/

```

A.

Fine-Grained Auditing (FGA) is enabled for the PRICE column in the PRODUCTS table for SELECT statements only when a row with PRICE > 10000 is accessed.

B.

FGA is enabled for the PRODUCTS.PRICE column and an audit record is written whenever a row with PRICE > 10000 is accessed.

C.

FGA is enabled for all DML operations by JIM on the PRODUCTS.PRICE column.

D.

FGA is enabled for the PRICE column of the PRODUCTS table and the SQL statements is captured in the FGA audit trail.

Which statement is true about the audit record that generated when auditing after instance restarts?

Posted by seenagape on February 21, 20166 comments

You execute the following commands to audit database activities: Which statement is true about the audit record that generated when auditing after instance restarts?

```
SQL > ALTER SYSTEM SET AUDIT_TRIAL=DB, EXTENDED SCOPE=SPFILE;
```

```
SQL > AUDIT SELECT TABLE, INSERT TABLE, DELETE TABLE BY JOHN By SESSION
WHENEVER SUCCESSFUL;
```

A.

One audit record is created for the whole session if john successfully executes a SELECT, INSERT, or DELETE command on a table, and contains the execution plan, SQL text, and bind variables used.

B.

One audit record is created for the whole session if JOHN successfully executes a select command, and contains the SQL text and bind variables used.

C.

One audit record is created for every successful execution of a SELECT, INSERT OR DELETE command, and contains the execution plan for the SQL statements.

D.

One audit record is created for every successful execution of a SELECT, INSERT OR DELETE command on a table, and contains the SQL text for the SQL Statements.

E.

One audit record is created for the whole session if john successfully executes a SELECT, INSERT, or DELETE command, and contains the execution plan for the SQL statements.

Explanation:

Which three statements are true about the ASM disk group compatibility attributes that are set for a disk group?

Posted by seenagape on February 21, 2016 1 comment

You support Oracle Database 12c Oracle Database 11g, and Oracle Database log on the same server. All databases of all versions use Automatic Storage Management (ASM).

Which three statements are true about the ASM disk group compatibility attributes that are set for a disk group?

A.

RDBMS compatibility together with the database version determines whether a database Instance can mount the ASM disk group.

B.

The RDBMS compatibility setting allows only databases set to the same version as the compatibility value, to mount the ASM disk group.

C.

The ASM compatibility attribute determines some of the ASM features that may be used by the Oracle disk group.

D.

The ADVM compatibility attribute determines the ACFS features that may be used by the Oracle 10 g database.

E.

The ASM compatibility attribute controls the format of the disk group metadata.

What is the result when you start up the database instance?

Posted by seenagape on February 21, 2016 1 comment

To enable the Database Smart Flash Cache, you configure the following parameters: What is the result when you start up the database instance?

```
DB_FLASH_CACHE_FILE = '/dev/flash_device_1', '/dev/flash_device_2'  
DB_FLASH_CACHE_SIZE=64G
```

A.

It results in an error because these parameter settings are invalid.

B.

One 64G flash cache file will be used.

C.

Two 64G flash cache files will be used.

D.

Two 32G flash cache files will be used.

Which two statements are true about the password file?

Posted by seenagape on February 21, 2016 1 comment

You executed this command to create a password file: \$ orapwd file = orapworcl entries = 10 ignorecase = N Which two statements are true about the password file?

A.

It will permit the use of uppercase passwords for database users who have been granted the SYSOPER role.

B.

It will not permit the use of mixed case passwords for the database users who have been granted the SYSDBA role.

C.

It will permit the use of lowercase passwords for database users who have granted the SYSDBA role.

D.

It contains username and passwords of database users who are members of the OSOPER operating system group.

E.

It contains usernames and passwords of database users who are members of the OSDBA operating system group.

Explanation:

Identify three valid methods of opening, pluggable databases (PDBs).

Posted by seenagape on February 21, 2016 1 comment

Identify three valid methods of opening, pluggable databases (PDBs).

A.

ALTER PLUGGABLE DATABASE OPEN ALL ISSUED from the root

- B.
ALTER PLUGGABLE DATABASE OPEN ALL ISSUED from a PDB
- C.
ALTER PLUGGABLE DATABASE PDB OPEN issued from the seed
- D.
ALTER DATABASE PDB OPEN issued from the root
- E.
ALTER DATABASE OPEN issued from that PDB
- F.
ALTER PLUGGABLE DATABASE PDB OPEN issued from another PDB
- G.
ALTER PLUGGABLE DATABASE OPEN issued from that PDB

Which two recommendations should you make to speed up the rebalance operation if this type of failure happens again?

Posted by seenagape on February 21, 2016 1 comment

You administer an online transaction processing (OLTP) system whose database is stored in Automatic Storage Management (ASM) and whose disk group use normal redundancy. One of the ASM disks goes offline, and is then dropped because it was not brought online before DISK_REPAIR_TIME elapsed. When the disk is replaced and added back to the disk group, the ensuing rebalance operation is too slow. Which two recommendations should you make to speed up the rebalance operation if this type of failure happens again?

- A.
Set the DISK_REPAIR_TIME disk attribute to a lower value.
- B.
Increase the number of DBWR_IO_SLAVES in the ASM instance.
- C.
Increase the number of ASMB processes.
- D.
Increase the value of the ASM_POWER_LIMIT parameter.
- E.
Specify the statement that adds the disk back to the disk group.

How would you accomplish these requirements?

Posted by seenagape on February 21, 2016 1 comment

You are administering a database and you receive a requirement to apply the following restrictions: 1. A connection must be terminated after four unsuccessful login attempts by user. 2. A user should not be able to create more than four simultaneous sessions. 3. User session must be terminated after 15 minutes of inactivity. 4. Users must be prompted to change their passwords every 15 days. How would you accomplish these requirements?

- A.
by granting a secure application role to the users
- B.
by creating and assigning a profile to the users and setting the REMOTE_OS_AUTHENT parameter to FALSE
- C.
By creating and assigning a profile to the users and setting the SEC_MAX_FAILED_LOGIN_ATTEMPTS parameter to 4
- D.
By Implementing Fine-Grained Auditing (FGA) and setting the REMOTE_LOGIN_PASSWORD_FILE parameter to NONE.
- E.
By implementing the database resource Manager plan and setting the SEC_MAX_FAILED_LOGIN_ATTEMPTS parameters to 4.

Explanation:

What could be a reason for this recommendation?

Posted by seenagape on February 21, 2016 1 comment

A senior DBA asked you to execute the following command to improve performance: SQL> ALTER TABLE subscribe log STORAGE (BUFFER_POOL recycle); You checked the data in the SUBSCRIBE_LOG table and found that it is a large table containing one million rows. What could be a reason for this recommendation?

- A.
Automatic Shared Memory Management is not enabled.
- B.
The data blocks in the SUBSCRIBE_LOG table are rarely accessed.
- C.
Automatic Workarea Management is not configured.
- D.
The keep pool is not configured.

E.

All the queries on the SUBSCRIBE_LOG table are rewritten to a materialized view.

Which three tasks can be automatically performed by the Automatic Data Optimization feature of Information lifecycle Management (ILM)?

Posted by seenagape on February 21, 2016 2 comments

Which three tasks can be automatically performed by the Automatic Data Optimization feature of Information lifecycle Management (ILM)?

A.

Tracking the most recent read time for a table segment in a user tablespace

B.

Tracking the most recent write time for a table segment in a user tablespace

C.

Tracking insert time by row for table rows

D.

Tracking the most recent write time for a table block

E.

Tracking the most recent read time for a table segment in the SYSAUX tablespace

F.

Tracking the most recent write time for a table segment in the SYSAUX tablespace

Which two partitioned table maintenance operations support asynchronous Global Index Maintenance in Oracle database 12c?

Posted by seenagape on February 21, 2016 1 comment

Which two partitioned table maintenance operations support asynchronous Global Index Maintenance in Oracle database 12c?

A.

ALTER TABLE DROP PARTITION

B.

ALTER TABLE ADD PARTITION

C.

ALTER TABLE MOVE PARTITION

D.

ALTER TABLE MERGE PARTITION

E.

ALTER TABLE SPLIT PARTITION

F.

ALTER TABLE TRUNCATE PARTITION

Which two memory areas that are part of PGA are stored in SGA instead, for shared server connection?

Posted by seenagape on February 21, 2016 1 comment

You configure your database Instance to support shared server connections. Which two memory areas that are part of PGA are stored in SGA instead, for shared server connection?

A.

Location of a part of the runtime area for SELECT statements

B.

User session data

C.

Stack space

D.

Private SQL area

E.

Location of the runtime area for DML and DDL Statements

Which two statements are true about Oracle Managed Files (OMF)?

Posted by seenagape on February 21, 2016 1 comment

Which two statements are true about Oracle Managed Files (OMF)?

A.

OMF automatically creates unique file names for table spaces and control files.

B.

OMF can be used with ASM disk groups, as well as with raw devices, for better file management.

C.

OMF cannot be used in a database that already has data files created with user-specified directions.

D.

OMF may affect the location of the redo log files and archived log files.

E.

The file system directions that are specified by OMF parameters are created automatically.

Explanation:

Which three actions are possible during an Online Data file Move operation?

Posted by seenagape on February 21, 2016 1 comment

Which three actions are possible during an Online Data file Move operation?

A.

Performing Block Media Recovery for a data block in the data file being moved

B.

Flashing back the database

C.

Executing DML statements on objects stored in the data file being moved

D.

Querying tables in the data file being moved

E.

Creating and dropping tables in the data file being moved

F.

Performing file shrink of the data file being moved

Which task should you perform before issuing the command?

Posted by seenagape on February 21, 2016 1 comment

Your multitenant container database (CDB) contains a pluggable database, HR_PDB. The default permanent tablespace in HR_PDB is USERDATA. The container database (CDB) is open and you connect RMAN. You want to issue the following RMAN command: RMAN > BACKUP TABLESPACE hr_pdb:userdata; Which task should you perform before issuing the command?

A.

Place the root container in ARCHIVELOG mode.

B.

Take the user data tablespace offline.

C.

Place the root container in the nomount stage.

D.

Ensure that HR_PDB is open.

Explanation:

Identify three scenarios in which you would recommend the use of SQL Performance Analyzer to analyze impact on the performance of SQL statements.

Posted by seenagape on February 21, 2016 1 comment

Identify three scenarios in which you would recommend the use of SQL Performance Analyzer to analyze impact on the performance of SQL statements.

A.

Change in the Oracle Database version

B.

Database and operating system upgrade

C.

Migration of database storage from non-ASM to ASM storage

D.

Change in your network infrastructure

E.

Change in the hardware configuration of the database server

Which two statements are true about the RMAN validate database command?

Posted by seenagape on February 21, 2016 2 comments

Which two statements are true about the RMAN validate database command?

A.

It checks the database for intrablock corruptions.

B.

It can detect corrupt pfiles.

C.

It can detect corrupt spfiles.

D.

It checks the database for interblock corruptions.

E.

It can detect corrupt block change tracking files.

Which statement is true?

Posted by seenagape on February 21, 2016 1 comment

You install a non-RAC Oracle Database. During Installation, the Oracle Universal Installer (OUI) prompts you to enter the path of the Inventory directory and also to specify an operating system group name. Which statement is true?

A.

The operating system group that is specified must have permission to write to the inventory directory.

B.

The ORACLE_BASE base parameter is not set.

C.

The installation is being performed by the root user.

D.

The operating system group that is specified should have the root user as its member.

Identify the correct order of the required steps.

Posted by seenagape on February 21, 2016 1 comment

You are required to migrate your 11.2.0.3 database as a pluggable database (PDB) to a multitenant container database (CDB). The following are the possible steps to accomplish this task: 1. Place all the user-defined tablespaces in read-only mode on the source database. 2. Upgrade the source database to a 12c version. 3. Create a new PDB in the target container database. 4. Perform a full transportable export on the source database with the VERSION parameter set to 12 using the expdp utility. 5. Copy the associated data

files and export the dump file to the desired location in the target database. 6. Invoke the Data Pump import utility on the new PDB database as a user with the DATAPUMP_IMP_FULL_DATABASE role and specify the full transportable import options. 7. Synchronize the PDB on the target container database by using the DBMS_PDS.SYNC_ODB function. Identify the correct order of the required steps.

A.

1, 5, 6, 4, 3, 2

B.

2, 1, 3, 4, 5, 6, 7

C.

1, 4, 3, 5, 6, 7

D.

2, 1, 3, 4, 5, 6

E.
1, 3, 4, 5, 6, 7

Which statement is true?

Posted by seenagape on February 21, 2016 2 comments

In your multitenant container database (CDB) with two pluggable database (PDBs). You want to create a new PDB by using SQL Developer. Which statement is true?

- A.
The CDB must be open.
- B.
The CDB must be in the mount stage.
- C.
The CDB must be in the nomount stage.
- D.
All existing PDBs must be closed.

Explanation:

Which two statements are true about the Oracle Direct Network File system (DNFS)?

Posted by seenagape on February 21, 2016 1 comment

Which two statements are true about the Oracle Direct Network File system (DNFS)?

- A.
Direct NFS can load-balance I/O traffic across multiple network adapters.
- B.
Direct NFS is available only in UNIX platforms.
- C.
A traditional NFS mount is not required when using Direct NFS.
- D.
It utilizes the OS file system cache.
- E.
Oracle Disk Manager can manage NFS on its own, without using the operating kernel NFS driver.

Which three statements are true about the process of automatic optimization by using cardinality feedback?

Posted by seenagape on February 21, 2016 2 comments

Examine the parameters for your database instance: Which three statements are true about the process of automatic optimization by using cardinality feedback?

NAME	TYPE	VALUE
optimizer_adaptive_reporting_only	boolean	FALSE
optimizer_capture_sql_plan_baselines	boolean	FALSE
optimizer_dynamic_sampling	integer	2
optimizer_features_enable	string	12.1.0

- A.
After the optimizer identifies a query as a re-optimization candidate, statistics collected by the collectors are submitted to the optimizer.
- B.
The optimizer automatically changes a plan during subsequent execution of a SQL statement if there is a huge difference in optimizer estimates and execution statistics.
- C.
The optimizer can re optimize a query only once using cardinality feedback.
- D.
The optimizer enables monitoring for cardinality feedback after the first execution of a query.
- E.
The optimizer does not monitor cardinality feedback if dynamic sampling and multicolumn statistics are enabled.

Which three statements are true when the listener handles connection requests to an Oracle 12c database instance with multithreaded architecture enabled In UNIX?

Posted by seenagape on February 21, 2016 1 comment

Which three statements are true when the listener handles connection requests to an Oracle 12c database instance with multithreaded architecture enabled In UNIX?

- A.
Each multithreaded Oracle process has an SCM thread.
- B.
Each Oracle process runs an SCM thread.
- C.
Thread creation must be routed through a dispatcher process

D.

The local listener may pass the request to an existing process which in turn will create a thread.

E.

The local listener may spawn a new process and have that new process create a thread

Which three operations can be performed as multipartition operations in Oracle?

Posted by seenagape on February 21, 2016 1 comment

Which three operations can be performed as multipartition operations in Oracle?

A.

Move partitions of a range-partitioned table

B.

Rename partitions of a range partitioned table

C.

Merge partitions of a reference partitioned index

D.

Coalesce partitions of a hash-partitioned global index.

E.

Merge partitions of a list partitioned table

F.

Drop partitions of a list partitioned table

What is the result of the last SET CONTAINER statement and why is it so?

Posted by seenagape on February 21, 2016 2 comments

You are connected using SQL*Plus to a multitenant container database (CDB) with SYSDBA privileges and execute the following sequence statements: What is the result of the last SET CONTAINER statement and why is it so?

```

SQL> CREATE PLUGGABLE DATABASE NEW_PDB ADMIN USER PDB_ADMIN;
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE NEW_PDB OPEN;
Pluggable database altered.
SQL> ALTER SESSION SET CONTAINER = NEW_PDB;
Session altered.
SQL> GRANT CONNECT TO PDB_ADMIN;
Grant succeeded.
SQL> CONNECT PDB_ADMIN/SECRET@LOCALHOST/NEW_PDB
Connected.
SQL> SELECT * FROM SESSION_PRIVS;

PRIVILEGE
-----
CREATE SESSION
SET CONTAINER

SQL> ALTER SESSION SET CONTAINER = PDB$SEED;

```

A.

It succeeds because the PDB_ADMIN user has the required privileges.

B.

It fails because common users are unable to use the SET CONTAINER statement.

C.

It fails because local users are unable to use the SET CONTAINER statement.

D.

If fails because the SET CONTAINER statement cannot be used with PDB\$SEED as the target pluggable database (PDB).

What are three possible causes for the latch-related wait events?

Posted by seenagape on February 21, 2016 3 comments

Examine the details of the Top 5 Timed Events in the following Automatic Workloads Repository (AWR) report: What are three possible causes for the latch-related wait events?

Top 5 Timed Foreground Events					
Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
DB CPU		67		98.21	
db file sequential read	8,371	0	0	0.52	User I/O
latch: row cache objects	10	0	8	0.19	Concurrency
latch: shared pool	956	0	0	0.15	Concurrency
log file sync	25	0	2	0.06	Commit

- A.
The size of the shared pool is too small.
- B.
The buffers are being read into the buffer cache, but some other session is changing the buffers.
- C.
There are frequent logons and logoffs.
- D.
Cursors are not being shared.
- E.
A large number COMMITS are being performed.

Explanation:

which executions is the audit policy now active?

Posted by seenagape on February 21, 2016 2 comments

You enabled an audit policy by issuing the following statements: For which database users and for which executions is the audit policy now active? Select two.

```
SQL> AUDIT POLICY ORA_DATABASE_PARAMETER BY SCOTT;
```

```
SQL> AUDIT POLICY ORA_DATABASE_PARAMETER BY SYS, SYSTEM;
```

- A.
SYS, SYSTEM
- B.
SCOTT
- C.
Only for successful executions
- D.
Only for failed executions

E.

Both successful and failed executions

which three situations will data not be redacted?

Posted by seenagape on February 21, 20162 comments

A redaction policy was added to the SAL column of the SCOTT.EMP table: All users have their default set of system privileges. For which three situations will data not be redacted?

```
BEGIN
  DBMS_REDACT.ADD_POLICY(
    OBJECT_SCHEMA => 'SCOTT',
    OBJECT_NAME   => 'EMP',
    POLICY_NAME   => 'SCOTT_EMP',
    COLUMN_NAME   => 'SAL',
    EXPRESSION    => 'SYS_CONTEXT(''SYS_SESSION_ROLES'', ''MGR'') = ''MGR''',
  );
END;
```

A.

SCOTT sessions, because he is the owner of the table

B.

SYSTEM session, only if the MGR role is set in the session

C.

SCOTT sessions, only if the MGR role is granted to SCOTT

D.

SCOTT sessions, only if the MGR role is set in the session

E.

SYSTEM sessions, regardless of the roles that are set in the session

F.

SYS sessions, regardless of the roles that are set in the session

What is the result of executing a TRUNCATE TABLE command on a table that has Flashback Archiving enabled?

Posted by seenagape on February 21, 20162 comments

What is the result of executing a TRUNCATE TABLE command on a table that has Flashback Archiving enabled?

A.

The rows in the table are truncated without being archived.

B.

The rows in the table are archived, and then truncated.

C.

The rows in both the table and the archive are truncated.

D.

It fails with the ORA-665610 Invalid DDL statement on history-tracked message

Explanation:

Which three activities are supported by the Data Recovery Advisor?

Posted by seenagape on February 21, 20162 comments

Which three activities are supported by the Data Recovery Advisor?

A.

Advising on invalid block header field values

B.

Advising on block checksum failures

C.

Advising on inaccessible control files

D.

Advising on inaccessible block change tracking files

E.

Advising on empty password files

Which three statements are true concerning the use of the Valid Time Temporal feature for the EMPLOYEES table?

Posted by seenagape on February 21, 20163 comments

You create a table with the PERIOD FOR clause to enable the use of the Temporal Validity feature of Oracle Database 12c. Examine the table definition: Which three statements are true concerning the use of the Valid Time Temporal feature for the EMPLOYEES table?

```
create table employees
(empno number, salary number,
deptid number, name varchar2(100),
period for employee_time);
```

A.

Setting the session valid time using

DBMS_FLASHBACK_ARCHIVE.ENABLE_AT_VALID_TIME sets the visibility for data manipulation language (DML), data definition language (DDL), and queries performed by the session.

B.

The valid time columns are visible by default when the table is described.

C.

The valid time columns employee_time_start and employee_time_end are automatically created.

D.

The same statement may filter on both transaction time and valid temporal time by using the AS OF TIMESTAMP and PERIOD FOR clauses.

E.

The valid time columns are not populated by the Oracle Server automatically.

Which three statements are true regarding the use of the Database Migration Assistant for Unicode (DMU)?

Posted by seenagape on February 21, 2016 2 comments

Which three statements are true regarding the use of the Database Migration Assistant for Unicode (DMU)?

A.

A DBA can check specific tables with the DMU

B.

The database to be migrated must be opened read-only.

C.

The release of the database to be converted can be any release since 9.2.0.8.

D.

The DMU can report columns that are too long in the converted character set.

E.

The DMU can report columns that are not represented in the converted character set.

What does this imply?

Posted by seenagape on February 21, 2016 1 comment

Oracle Grid Infrastructure for a stand-alone server is installed on your production host before installing the Oracle Database server. The database and listener are configured by using Oracle Restart. Examine the following command and its output: \$ crsctl config has CRS-4622: Oracle High Availability Services auto start is enabled. What does this imply?

A.

When a database instance is started by using the SRVCTL utility and listener startup fails, the instance is still started.

B.

When a database is created by using SQL* Plus, it is automatically added to the Oracle Restart configuration.

C.

When you start an instance on a high with SQL *Plus dependent listeners and ASM disk groups are automatically started.

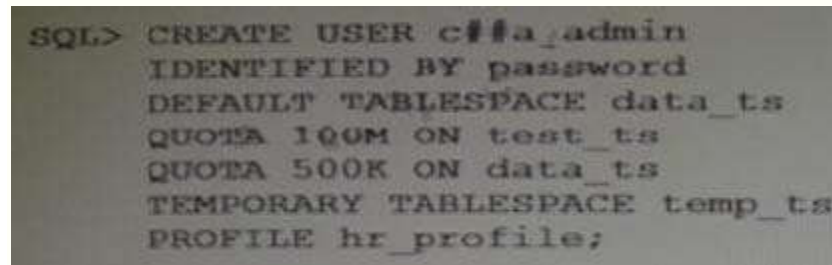
D.

When you create a database service by modifying the SERVICE_NAMES initialization parameter, it is automatically added to the Oracle Restart configuration.

Which two statements are true?

Posted by seenagape on February 21, 2016 2 comments

Your multitenant container database (CDB) contains some pluggable databases (PDBs), you execute the following command in the root container: Which two statements are true?



```
SQL> CREATE USER c##a_admin
IDENTIFIED BY password
DEFAULT TABLESPACE data_ts
QUOTA 100M ON test_ts
QUOTA 500K ON data_ts
TEMPORARY TABLESPACE temp_ts
PROFILE hr_profile;
```

A.

Schema objects owned by the C##A_ADMIN common user can be shared across all PDBs.

B.

The C##A_ADMIN user will be able to use the TEMP_TS temporary tablespace only in root.

C.

The command will, create a common user whose description is contained in the root and each PDB.

D.

The schema for the common user C##A_ADMIN can be different in each container.

E.

The command will create a user in the root container only because the container clause is not used.

Which three statements are true?

Posted by seenagape on February 21, 2016 3 comments

You performed an incremental level 0 backup of a database: To enable block change tracking after the incremental level 0 backup, you issued this command: To perform an incremental level 1 cumulative backup, you issued this command: Which three statements are true?


```
RMAN > BACKUP INCREMENTAL LEVEL 0 DATABASE;  
SQL > ALTER DATABASE ENABLE BLOCK CHANGE TRACKING USING FILE  
' /mydir/rman_change_track.f';  
RMAN> BACKUP INCREMENTAL LEVEL 1 CUMULATIVE DATABASE;
```

A.

Backup change tracking will sometimes reduce I/O performed during cumulative incremental backups.

B.

The incremental level 1 backup that immediately follows the enabling of block change tracking will not read the change tracking file to discover changed blocks.

C.

More than one database block may be read by an incremental backup for a change made to a single block.

D.

The change tracking file must always be backed up when you perform a full database backup.

E.

Block change tracking will always reduce I/O performed during cumulative incremental backups.

Which method is used by the optimizer to limit the rows being returned?

Posted by seenagape on February 21, 2016 3 comments

You find this query being used in your Oracle 12c database: Which method is used by the optimizer to limit the rows being returned?

```
select employee_id, first_name, salary  
from hr.employees  
order by employee_id  
fetch first 20 percent rows only;
```

A.

A filter is added to the table query dynamically using ROWNUM to limit the rows to 20 percent of the total rows

B.

All the rows are returned to the client or middle tier but only the first 20 percent are returned to the screen or the application.

C.

A view is created during execution and a filter on the view limits the rows to 20 percent of the total rows.

D.

A TOP-N query is created to limit the rows to 20 percent of the total rows

Which three resources might be prioritized between competing pluggable databases...?

Posted by seenagape on February 21, 2016 1 comment

Which three resources might be prioritized between competing pluggable databases when creating a multitenant container database plan (CDB plan) using Oracle Database Resource Manager?

A.

Local file system I/O

B.

Maximum Undo per consumer group

C.

Maximum Idle time

D.

Parallel server limit

E.

CPU

F.

Exadata I/O

Explanation:

Which is true about the result?

Posted by seenagape on February 21, 2016 2 comments

You created an encrypted tablespace: You then closed the encryption wallet because you were advised that this is secure. Later in the day, you attempt to create the EMPLOYEES table in the SECURESPACE table space with the SALT option on the EMPLOYEE column. Which is true about the result?

```
SQL> CREATE TABLESPACE securespace  
      DATAFILE '/home/user/oradata/secure01.dbf'  
      SIZE 150M  
      ENCRYPTION USING '3DES168'  
      DEFAULT STORAGE (ENCRYPT) ;
```

A.

It generates error when creating the table, because the salt option cannot be used with encrypted tablespaces.

B.

It creates the table successfully, and encrypts any inserted data in the EMPNAME column because the wallet needs to be open only for tablespace creation.

C.

It creates the table successfully but does not encrypt any inserted data in the EMPNAME column because the wallet must be opened to encrypt columns with SALT.

D.

It generates an error when creating the table because the wallet is closed.

Which two statements are true?

Posted by seenagape on February 21, 2016 3 comments

On your Oracle Database, you issue the following commands to create indexes: Which two statements are true?

```
SQL > CREATE INDEX oe.ord_customer_ix1 ON oe.orders (customer_id, sales_rep_id)
INVISIBLE;
```

```
SQL> CREATE BITMAP INDEX oe.ord_customer_ix2 ON oe.orders (customer_id, sales_rep_id);
```

A.

Only the ORD_CUSTOMER_IX1 index created.

B.

Both the indexes are updated when a row is inserted, updated, or deleted in the ORDERS table.

C.

Both the indexes are created: however, only ORD_CUSTOMERS_IX1 is used by the optimizer for queries on the ORDERS table.

D.

The ORD_CUSTOMER_IX1 index is not used by the optimizer even when the OPTIMIZER_USE_INVISIBLE_INDEXES parameters is set to true.

E.

Both the indexes are created and used by the optimizer for queries on the ORDERS table.

F.

Both the indexes are created: however, only ORD_CUSTOMERS_IX2 is used by the optimizer for queries on the ORDERS table.

Which two statements are true when row archival management is enabled?

Posted by seenagape on February 21, 2016 2 comments

Which two statements are true when row archival management is enabled?

A.

The ORA_ARCHIVE_STATE column is visible if referenced in the select list of a query.

B.

The ORA_ARCHIVE_STATE column is updated automatically by the Oracle Server based on activity tracking columns, to indicate that a row is no longer considered active.

C.

The ORA_ARCHIVE_STATE column visibility is controlled by the ROW ARCHIVAL VISIBILITY session parameter.

D.

The ROW ARCHIVAL VISIBILITY session parameter defaults to active rows only.

E.

The ORA_ARCHIVE_STATE column is updated manually or by a program that could reference activity tracking columns, to indicate that a row is no longer considered active.

Which three methods could transparently help to achieve this result?

Posted by seenagape on February 21, 2016 3 comments

A warehouse fact table in your Oracle 12c Database is range-partitioned by month and accessed frequently with queries that span multiple partitions. The table has a local prefixed, range-partitioned index. Some of these queries access very few rows in some partitions and all the rows in other partitions, but these queries still perform a full scan for all accessed partitions. This commonly occurs when the range of dates begins at the end of a month or ends close to the start of a month. You want an execution plan to be generated that uses indexed access when only a few rows are accessed from a segment, while still allowing full scans for segments where many rows are returned. Which three methods could transparently help to achieve this result?

A.

Using a partial local Index on the warehouse fact table month column with indexing disabled to the table partitions that return most of their rows to the queries.

B.

Using a partial local Index on the warehouse fact table month column with indexing disabled for the table partitions that return a few rows to the queries.

C.

Using a partitioned view that does a UNION ALL query on the partitions of the warehouse fact table, which retains the existing local partitioned column.

D.

Converting the partitioned table to a partitioned view that does a UNION ALL query on the monthly tables, which retains the existing local partitioned column.

E.

Using a partial global index on the warehouse fact table month column with indexing disabling for the table partitions that return most of their rows to the queries.

F.

Using a partial global index on the warehouse fact table month column with indexing disabled for the table partitions that return a few rows to the queries.

Which three statements are true about the advisor given by the segment advisor?

Posted by seenagape on February 21, 2016 2 comments

You use the segment advisor to help determine objects for which space may be reclaimed.

Which three statements are true about the advisor given by the segment advisor?

A.

It may advise the use of online table redefinition for tables in dictionary managed tablespace.

B.

It may advise the use of segment shrink for free list managed tables.

C.

It will detect and advise about chained rows.

D.

It may advise the use of segment shrink for tables in dictionary managed tablespaces if there are no chained rows.

E.

It may advise the use of online table redefinition for tables in locally managed tablespaces

which affect the invisible index columns?

Posted by seenagape on February 21, 2016 1 comment

You have altered a non-unique index to be invisible to determine if queries execute within an acceptable response time without using this index. Which two are possible if table updates are performed which affect the invisible index columns?

A.

The index remains invisible.

B.

The index is not updated by the DML statements on the indexed table.

C.

The index automatically becomes visible in order to have it updated by DML on the table.

D.

The index becomes unusable but the table is updated by the DML.

E.

The index is updated by the DML on the table.

Which two statements are true?

Posted by seenagape on February 21, 2016 1 comment

In your multitenant container database (CDB) containing same pluggable databases (PDBs), you execute the following commands in the root container: Which two statements are true?

```
SQL> CREATE ROLE c##role1;  
SQL> GRANT create view, create procedure to c##role1;  
SQL> GRANT c##role1 to c##a_admin;
```

A.

The statement for granting a role to a user fails because the CONTAINER clause is not used.

B.

Privileges are granted to the C##A_ADMIN user in the root database and all PDBs.

C.

The C ## ROLE1 role is created only in the root database because the container clause is not used.

D.

The C ## ROLE1 role is created in the root database and all the PDBs.

E.

Privileges are granted to the C##A_ADMIN user only in the root database.

Identify four RMAN commands that produce a multi-section backup.

Posted by seenagape on February 21, 2016 1 comment

The persistent configuration settings for RMAN have default for all parameters. Identify four RMAN commands that produce a multi-section backup.

A.

BACKUP TABLESPACE SYSTEM SECTION SIZE 100M;

- B.
BACKUP AS COPY TABLESPACE SYSTEM SECTION SIZE 100M;
- C.
BACKUP ARCHIVELOG ALL SECTION SIZE 25M;
- D.
BACKUP TABLESPACE “TEMP” SECTION SIZE 10M;
- E.
BACKUP TABLESPACE “UNDO” INCLUDE CURRENT CONTROLFILE SECTION SIZE 100M;
- F.
BACKUP SPFILE SECTION SIZE 1M;
- G.
BACKUP INCREMENTAL LEVEL 0 TABLESPACE SYSAUX SECTION SIZE 100M;

Which command or commands should you execute next to allow updates to the flashback back schema?

Posted by seenagape on February 21, 2016 1 comment

Flashback is enabled for your multitenant container database (CDB), which contains two pluggable database (PDBs). A local user was accidentally dropped from one of the PDBs. You want to flash back the PDB to the time before the local user was dropped. You connect to the CDB and execute the following commands: Examine following commands: 1. ALTER PLUGGABLE DATABASE ALL OPEN; 2. ALTER DATABASE OPEN; 3. ALTER DATABASE OPEN RESETLOGS; Which command or commands should you execute next to allow updates to the flashback back schema?

```
SQL > SHUTDOWN IMMEDIATE
```

```
SQL > STARTUP MOUNT
```

```
SQL > FLASHBACK DATABASE to TIME "TO_DATE ('08/20/12' , 'MM/DD/YY')";
```

- A.
Only 1
- B.
1 and 2
- C.
3 and 1
- D.
Only 2

E.
Only 3

Which two statements are true?

Posted by seenagape on February 21, 2016 2 comments

Examine the commands executed to monitor database operations: Which two statements are true?

```
$> conn sys oracle/oracle@prod as sysdba  
  
SQL > VAR eid NUMBER  
  
SQL > EXEC: eid := DBMS_SQL_MONITOR.BEGIN_OPERATION ('batch_job' ,  
FORCED_TRACKING => 'Y');
```

- A.
Database operations will be monitored only when they consume a significant amount of resource.
- B.
Database operations for all sessions will be monitored.
- C.
Database operations will be monitored only if the STATISTICS_LEVEL parameter is set to TYPICAL and CONTROL_MANAGEMENT_PACK_ACCESS is set DIAGNOSTIC + TUNING.
- D.
Only DML and DDL statements will be monitored for the session.
- E.
All subsequent statements in the session will be treated as one database operation and will be monitored.

Which three statements are true about the working of system privileges in a multitenant control database (CDB) that has pluggable databases (PDBs)?

Posted by seenagape on February 21, 2016 2 comments

Which three statements are true about the working of system privileges in a multitenant control database (CDB) that has pluggable databases (PDBs)?

- A.
Common users connected to a PDB can exercise privileges across other PDBs.
- B.
System privileges with the with grant option container all clause must be granted to a common user before the common user can grant privileges to other users.

C.

Local users cannot use local system privileges on the schema of a common user.

D.

System privileges apply only to the PDB in which they are used.

E.

The granter of system privileges must possess the set container privilege.

Which technique should you use to minimize down time while plugging this non-CDB into the CDB?

Posted by seenagape on February 21, 2016 3 comments

You are about to plug a multi-terabyte non-CDB into an existing multitenant container database (CDB) as a pluggable database (PDB). – The characteristics of the non-CDB are as follows: – Version: Oracle Database 12c Release 1 64-bit – Character set: WE8ISO8859P15 – National character set: AL16UTF16 – O/S: Oracle Linux 6 64-bit The characteristics of the CDB are as follows: – Version: Oracle Database 12c Release 1 64-bit – Character set: AL32UTF8 – O/S: Oracle Linux 6 64-bit Which technique should you use to minimize down time while plugging this non-CDB into the CDB?

A.

Transportable tablespace

B.

Data Pump full export / import

C.

The DBMS_PDB package

D.

RMAN

E.

Transportable database

Identify the correct outcome and the step to aggregate by using tkprof utility?

Posted by seenagape on February 21, 2016 2 comments

Your database has the SRV1 service configured for an application that runs on middle-tier application server. The application has multiple modules. You enable tracing at the service level by executing the following command: SQL > exec DBMS_MONITOR.SERV_MOD_ACT_TRACE_ENABLE ('SRV1'); The possible outcome and actions to aggregate the trace files are as follows: 1. The command fails because a module name is not specified. 2. A trace file is created for each session that is running the SRV1 service. 3. An aggregated trace file is created for all the sessions that are running the SRV1 service. 4. The trace files may be aggregated by using the trcess utility. 5. The trace

files be aggregated by using the tkprof utility. Identify the correct outcome and the step to aggregate by using tkprof utility?

- A.
3 and 5
- B.
1
- C.
2 and 4
- D.
2 and 5
- E.
3 and 4

What is the result?

Posted by seenagape on February 21, 2016 1 comment

Your multitenant container database (CDB) contains pluggable databases (PDBs), you are connected to the HR_PDB. You execute the following command: SQL > CREATE UNDO TABLESPACE undotb01 DATAFILE 'u01/oracle/rddb1/undotbs01.dbf' SIZE 60M AUTOEXTEND ON; What is the result?

- A.
It executes successfully but neither tablespace nor the data file is created.
- B.
It fails and reports an error because the CONTAINER=CURRENT clause is not specified in the command.
- C.
It fails and reports an error because the CONTAINER=ALL clause is not specified in the command.
- D.
It executes successfully and creates an UNDO tablespace in HR_PDB.
- E.
It fails and reports an error because there can be only one undo tablespace in a CDB.

Which three statements are true about SQL plan directives?

Posted by seenagape on February 21, 2016 1 comment

Which three statements are true about SQL plan directives?

A.

They are tied to a specific statement or SQL ID.

B.

They instruct the maintenance job to collect missing statistics or perform dynamic sampling to generate a more optimal plan.

C.

They are used to gather only missing statistics.

D.

They are created for a query expression where statistics are missing or the cardinality estimates by the optimizer are incorrect.

E.

They instruct the optimizer to create only column group statistics.

F.

Improve plan accuracy by persisting both compilation and execution statistics in the SYSAUX tablespace.

Which two statements are true about this flashback scenario?

Posted by seenagape on February 21, 2016 2 comments

You want to flash back a test database by five hours. You issue this command: SQL > FLASHBACK DATABASE TO TIMESTAMP (SYSDATE - 5/24); Which two statements are true about this flashback scenario?

A.

The database must have multiplexed redo logs for the flashback to succeed.

B.

The database must be opened in read-only mode to check if the database has been flashed back to the correct SCN.

C.

The database must be opened with the RESETLOGS option after the flashback is complete.

D.

The database must be MOUNTED for the flashback to succeed.

E.

The database must use block change tracking for the flashback to succeed.

F.

The database must be opened in restricted mode for the flashback to succeed.

Which three are true about the MRKT tablespace?

Posted by seenagape on February 21, 2016 1 comment

Examine these two statements: Which three are true about the MRKT tablespace?

```
SQL> CREATE BIGFILE TABLESPACE MRKT
  2  DATAFILE '/u01/app/oracle/oradata/orcl/mrkt.dbf' SIZE 1024M
  3  EXTENT MANAGEMENT LOCAL SEGMENT SPACE MANAGEMENT AUTO
Tablespace created.

SQL> ALTER DATABASE DEFAULT TABLESPACE MRKT;
Database altered.
```

A.

The MRKT tablespace is created as a small file tablespace, because the file size is less than the minimum required for big file files.

B.

The MRKT tablespace may be dropped if it has no contents.

C.

Users who were using the old default tablespace will have their default tablespaces changed to the MRKT tablespace.

D.

No more data files can be added to the tablespace.

E.

The relative file number of the tablespace is not stored in rowids for the table rows that are stored in the MRKT tablespace.

How would you accomplish this?

Posted by seenagape on February 21, 2016 2 comments

In your database, you want to ensure that idle sessions that are blocking active are automatically terminated after a specified period of time. How would you accomplish this?

A.

Decreasing the value of the IDLE_TIME resource limit in the default profile

B.

Setting a metric threshold

C.

Implementing Database Resource Manager

D.
Enabling resumable timeout for user sessions

Which two statements are true about the password file?

Posted by seenagape on February 21, 2016 1 comment

You Execute the Following command to create a password file in the database server: \$
orapwd file = '+DATA/PROD/orapwprod entries = 5 ignorecase = N format = 12' Which two

statements are true about the password file?

- A.
It supports the SYSBACKUP, SYSDG, and SYSKM system privileges.
- B.
It records the usernames and passwords of all users when they are added to the OSDBA or OSOPER operating system groups.
- C.
Is used by Oracle to authenticate users for remote database administration.
- D.
It records the usernames and passwords of users when granted the DBA role.
- E.
It contains the usernames and passwords of users for whom auditing is enabled.

Identify two situations in which the alert log file is updated.

Posted by seenagape on February 21, 2016 2 comments

Identify two situations in which the alert log file is updated.

- A.
Rebuilding an index using ALTER INDEX . . . REBUILD fails with an ORA-01578: ORACLE data block corrupted (file # 14, block # 50) error.
- B.
Running a query on a table returns ORA-600: Internal Error.
- C.
Inserting a value into a table returns ORA-01722: invalid number.
- D.
Creating a table returns ORA-00955: name us already in used by an existing objects.
- E.
Inserting a value into a table returns ORA-00001: unique constraint (SYS.OK_TECHP) violated.

Explanation:

Which three statements are true about Oracle Data Pump export and import operations?

Posted by seenagape on February 21, 2016 1 comment

Which three statements are true about Oracle Data Pump export and import operations?

A.

You can detach from a data pump import job and reattach later.

B.

The master table is the last object to be exported by the data pump.

C.

You can detach from a data pump export job and reattach later.

D.

Data pump uses parallel execution server processes to implement parallel import.

E.

Data pump import requires the import file to be in a directory owned by the oracle owner.

Which three statements are true about the users (other than sys) in the output?

Posted by seenagape on February 21, 2016 2 comments

Examine the query and its output executed in an RDBMS Instance: Which three statements are true about the users (other than sys) in the output?

```
SQL> SELECT * FROM v$pwfile_users;
```

USERNAME	SYSDB	SYSOP	SYSAS	SYSBA	SYSDG	SYSK
SYS	TRUE	TRUE	FALSE	FALSE	FALSE	FALSE
C##B_ADMIN	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE
C##C_ADMIN	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE
C##A_ADMIN	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
C##D_ADMIN	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE

A.

The C ## B_ADMIN user can perform all backup and recovery operations using RMAN only.

B.

The C ## C_ADMIN user can perform the data guard operation with Data Guard Broker.

C.

The C ## A_ADMIN user can perform wallet operations.

D.

The C ## D_ADMIN user can perform backup and recovery operations for Automatic Storage Management (ASM).

E.

The C ## B_ADMIN user can perform all backup and recovery operations using RMAN or SQL* Plus.

Which two storage-tiering actions might be automated when using information Lifecycle Management (ILM) to automate data movement?

Posted by seenagape on February 21, 2016 1 comment

In your Database, the TBS PERCENT USED parameter is set to 60 and the TBS PERCENT FREE parameter is set to 20. Which two storage-tiering actions might be automated when using information Lifecycle Management (ILM) to automate data movement?

A.

Setting the target tablespace to read-only

B.

The movement of some blocks to a target tablespace with a lower degree of compression, on a different storage tier, when the source tablespace exceeds TBS PERCENT USED

C.

The movement of some segments to a target tablespace with a higher degree of compression, on a different storage tier, when the source tablespace exceeds TBS PERCENT USED

D.

The movement of all segments to a target tablespace with a higher degree of compression, on a different storage tier, when the source tablespace exceeds TBS PERCENT USED

E.

Setting the target tablespace offline

Explanation:

Which three statements are true about Flashback Database?

Posted by seenagape on February 21, 2016 1 comment

Which three statements are true about Flashback Database?

A.

Flashback logs are written sequentially, and are archived.

B.

Flashback Database uses a restored control file to recover a database.

C.

The Oracle database automatically creates, deletes, and resides flashback logs in the Fast Recovery Area.

D.

Flashback Database can recover a database to the state that it was in before a reset logs operation.

E.

Flashback Database can recover a data file that was dropped during the span of time of the flashback.

F.

Flashback logs are used to restore to the blocks' before images, and then the redo data may be used to roll forward to the desired flashback time.

Which statement is true about Enterprise Manager (EM) express in Oracle Database 12c?

Posted by seenagape on February 21, 2016 3 comments

Which statement is true about Enterprise Manager (EM) express in Oracle Database 12c?

A.

You can use EM express to manage multiple databases running on the same server.

B.

You can create and configure pluggable databases by using EM express.

C.

By default, EM express is available for a database after database creation.

D.

You can perform basic administrative tasks for pluggable databases by using the EM express interface.

E.

You cannot start up or shut down a database Instance by using EM express.

Which statement is true?

Posted by seenagape on February 21, 2016 2 comments

Examine the following command; ALTER SYSTEM SET enable_ddl_logging = TRUE;

Which statement is true?

A.

Only the data definition language (DDL) commands that resulted in errors are logged in the alert log file.

B.

All DDL commands are logged in the alert log file.

C.

All DDL commands are logged in a different log file that contains DDL statements and their execution dates.

D.

Only DDL commands that resulted in the creation of new segments are logged.

E.

All DDL commands are logged in XML format in the alert directory under the Automatic Diagnostic Repository (ADR) home.

which two scenarios do you use SQL* Loader to load data?

Posted by seenagape on February 21, 2016 1 comment

In which two scenarios do you use SQL* Loader to load data?

A.

Generate unique sequential key values in specified columns.

B.

Transform the data while it is being loaded into the database.

C.

Use transparent parallel processing without having to split the external data first.

D.

Load data into multiple tables during the same load statement.

Which is true about the result of this command?

Posted by seenagape on February 21, 2016 1 comment

You are connected to a pluggable database (PDB) as a common user with DBA privileges.

The STATISTICS_LEVEL parameter is PDB_MODIFIABLE. You execute the following: SQL > ALTER SYSTEM SET STATISTICS_LEVEL = ALL SID = '*' SCOPE = SPFILE; Which is true about the result of this command?

A.

Nothing happens; because there is no SPFILE for each PDB, the statement is ignored.

B.

The STATISTICS_LEVEL parameter is set to all whenever the multitenant container database (CDB) is restarted.

C.

The STATISTICS_LEVEL parameter is set to all whenever this PDB is re-opened.

D.

The STATISTICS_LEVEL parameter is set to ALL whenever any PDB is reopened.

Which two are prerequisites for performing a flashback transaction?

Posted by seenagape on February 21, 2016 2 comments

Which two are prerequisites for performing a flashback transaction?

A.

Flashback Database must be enabled.

B.

Undo retention guarantee for the database must be configured.

C.

EXECUTE privilege on the DBMS_FLASHBACK package must be granted to the user flashing back transaction.

D.

Supplemental logging must be enabled.

E.

Recycle bin must be enabled for the database.

F.

Block change tracking must be enabled for the database.

What happens if the CONTROLLER1 failure group becomes unavailable due to error of for maintenance?

Posted by seenagape on February 21, 2016 3 comments

A database is stored in an Automatic Storage Management (ASM) disk group, disk group, DGROUP1 with SQL: There is enough free space in the disk group for mirroring to be done. What happens if the CONTROLLER1 failure group becomes unavailable due to error of for maintenance?

```
SQL> CREATE DISKGROUP dgroup1 NORMAL REDUNDANCY  
        FAILGROUP controller1 DISK '/devices/disk1', '/devic  
        FAILGROUP controller2 DISK '/devices/disk1', '/devic
```

A.

Transactions and queries accessing database objects contained in any tablespace stored in DGROUP1 will fail.

B.

Transactions accessing database objects contained in any tablespace stored in DGROUP1 will fail but queries will succeed.

C.

ASM does not mirror any data until the controller failure group is brought back online, and newly allocated primary allocation units (AU) are stored in the controller2 failure group, without mirroring.

D.

The data in the CONTROLLER1 failure group is copied to the controller2 failure group and rebalancing is initiated.

E.

Mirroring of allocation units will be done to ASM disks in the CONTROLLER2 failure group until the CONTROLLER1 for failure group is brought back online.

Which two statement are correct?

Posted by seenagape on February 21, 2016 1 comment

On your Oracle 12c database, you Issue the following commands to create indexes Which two statement are correct?

```
SQL > CREATE INDEX oe.ord_customer_ix1 ON oe.orders (customers_id, sales_rep_id)
INVISIBLE;
```

```
SQL> CREATE BITMAP INDEX oe.ord_customer_ix2 ON oe.orders (customers_id, sales_rep_id);
```

A.

Both the indexes are created; however, only the ORD_COSTOMER index is visible.

B.

The optimizer evaluates index access from both the Indexes before deciding on which index to use for query execution plan.

C.

Only the ORD_CUSTOMER_IX1 index is created.

D.

Only the ORD_CUSTOMER_IX2 index is created.

E.

Both the indexes are updated when a new row is inserted, updated, or deleted In the orders table.

Explanation:

Which two RMAN commands may be; used to back up only the PDB1 pluggable database?

Posted by seenagape on February 21, 20161 comment

Your multitenant container database has three pluggable databases (PDBs): PDB1, PDB2, and PDB3. Which two RMAN commands may be; used to back up only the PDB1 pluggable database?

- A.
BACKUP DATABASE while connected to the PDB1 container
- B.
BACKUP DATABASE while connected to the boot container
- C.
BACKUP PLUGGABLE DATABASE PDB1 while connected to the PDB1 container
- D.
BACKUP PLUGGABLE DATABASE PDB1 while connected to the root container
- E.
BACKUP PLUGGABLE database PDB1 while connected to PDB2

Identify three benefits of Unified Auditing.

Posted by seenagape on February 21, 20162 comments

Identify three benefits of Unified Auditing.

- A.
Decreased use of storage to store audit trail rows in the database.
- B.
It improves overall auditing performance.
- C.
It guarantees zero-loss auditing.
- D.
The audit trail cannot be easily modified because it is read-only.
- E.
It automatically audits Recovery Manager (RMAN) events.

Explanation:

How do you accomplish this?

Posted by seenagape on February 21, 20162 comments

You upgraded from a previous Oracle database version to Oracle Database version to Oracle Database 12c. Your database supports a mixed workload. During the day, lots of insert, update, and delete operations are performed. At night, Extract, Transform, Load

(ETL) and batch reporting jobs are run. The ETL jobs perform certain database operations using two or more concurrent sessions. After the upgrade, you notice that the performance of ETL jobs has degraded. To ascertain the cause of performance degradation, you want to collect basic statistics such as the level of parallelism, total database time, and the number of I/O requests for the ETL jobs. How do you accomplish this?

A.

Enable real-time database operation monitoring using the DBMS_SQL_MONITOR.BEGIN_OPERATION function, and then use the DBMS_SQL_MONITOR.REPORT_SQL_MONITOR function to view the required information.

B.

Examine the Active Session History (ASH) reports for the time period of the ETL or batch reporting runs.

C.

Enable SQL tracing for the queries in the ETL and batch reporting queries and gather diagnostic data from the trace file.

D.

Enable real-time SQL monitoring for ETL jobs and gather diagnostic data from the V\$SQL_MONITOR view.