

Oracle Database 11g: Performance Tuning

Volume III • Student Guide

D50317GC20

Edition 2.0

June 2010

D67629

ORACLE®

Author

James Spiller

**Technical Contributors
and Reviewers**

John Beresniewicz

Yanti Chang

Kurt Engleiter

Gerlinde Frenzen

Joel Goodman

Martin Jensen

Pete Jones

Sean Kim

Roderick Manalac

Wayne Reeser

Branislav Valny

Sergiusz Wolicki

Editors

Aju Kumar

Raj Kumar

Vijayalakshmi Narasimhan

Graphic Designer

Satish Bettgowda

Publishers

Syed Ali

Shaik Mahaboob Basha

Michael Sebastian

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Disclaimer

This document contains proprietary information and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except where your use constitutes "fair use" under copyright law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

Restricted Rights Notice

If this documentation is delivered to the United States Government or anyone using the documentation on behalf of the United States Government, the following notice is applicable:

U.S. GOVERNMENT RIGHTS

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

Trademark Notice

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Contents

1 Introduction

- Course Objectives 1-2
- Organization 1-3
- Agenda 1-4
- What Is Not Included 1-6
- Who Tunes? 1-7
- What Does the DBA Tune? 1-8
- How to Tune 1-10
- Tuning Methodology 1-11
- Effective Tuning Goals 1-13
- General Tuning Session 1-15
- Quiz 1-17
- Summary 1-18

2 Basic Tuning Diagnostics

- Objectives 2-2
- Performance Tuning Diagnostics 2-3
- Performance Tuning Tools 2-4
- Tuning Objectives 2-6
- Top Timed Events 2-7
- DB Time 2-8
- CPU and Wait Time Tuning Dimensions 2-9
- Time Model: Overview 2-10
- Time Model Statistics Hierarchy 2-11
- Time Model Example 2-13
- Quiz 2-14
- Dynamic Performance Views 2-15
- Dynamic Performance Views: Usage Examples 2-16
- Dynamic Performance Views: Considerations 2-17
- Statistic Levels 2-18
- Instance Activity and Wait Event Statistics 2-20
- System Statistic Classes 2-22
- Displaying Statistics 2-23
- Displaying SGA Statistics 2-25
- Wait Events 2-26
- Using the V\$EVENT_NAME View 2-27
- Wait Classes 2-28

- Displaying Wait Event Statistics 2-29
- Commonly Observed Wait Events 2-31
- Using the V\$SESSION_WAIT View 2-32
- Precision of System Statistics 2-34
- Quiz 2-35
- Using Features of the Packs 2-36
- Accessing the Database Home Page 2-38
- Performance Information on the Database Home Page 2-39
- Viewing the Alert Log 2-41
- Using Alert Log Information as an Aid in Tuning 2-42
- User Trace Files 2-44
- Background Processes Trace Files 2-45
- Quiz 2-46
- Summary 2-47
- Practice 2 Overview: Using Basic Tools 2-48

3 Using Automatic Workload Repository

- Objectives 3-2
- Automatic Workload Repository: Overview 3-3
- Automatic Workload Repository Data 3-4
- Workload Repository 3-5
- Database Control and AWR 3-6
- AWR Snapshot Purging Policy 3-7
- AWR Snapshot Settings 3-8
- Manual AWR Snapshots 3-9
- Managing Snapshots with PL/SQL 3-10
- Generating AWR Reports in EM 3-11
- Generating AWR Reports in SQL*Plus 3-12
- Reading the AWR Report 3-13
- Snapshots and Periods Comparisons 3-14
- Compare Periods: Benefits 3-15
- Compare Periods: Results 3-16
- Compare Periods: Report 3-17
- Compare Periods: Load Profile 3-18
- Compare Periods: Top Events 3-19
- Quiz 3-20
- Summary 3-21
- Practice 3 Overview: Using AWR-Based Tools 3-22

4 Defining Problems

Objectives	4-2
Defining the Problem	4-3
Limit the Scope	4-4
Setting the Priority	4-5
Top 5 Timed Events	4-6
Setting the Priority: Example	4-7
Top SQL Reports	4-8
Common Tuning Problems	4-9
Tuning Life Cycle Phases	4-11
Tuning During the Life Cycle	4-12
Application Design and Development	4-13
Testing: Database Configuration	4-14
Deployment	4-15
Production	4-16
Migration, Upgrade, and Environment Changes	4-17
ADDM Tuning Session	4-18
Performance Versus Business Requirements	4-19
Performance Tuning Resources	4-20
Filing a Performance Service Request	4-21
RDA Report	4-22
Monitoring and Tuning Tool: Overview	4-23
Quiz	4-25
Summary	4-26
Practice 4 Overview: Identifying the Problem	4-27
5 Using Metrics and Alerts	
Objectives	5-2
Metrics and Alerts	5-3
Limitation of Base Statistics	5-4
Typical Delta Tools	5-5
Oracle Database 11g Solution: Metrics	5-6
Benefits of Metrics	5-7
Viewing Metric History Information	5-8
Using EM to View Metric Details	5-9
Statistic Histograms	5-10
Histogram Views	5-11
Server-Generated Alerts	5-12
Alert Usage Model	5-13
Setting Thresholds	5-14
Creating and Testing an Alert	5-15
Metric and Alert Views	5-16

Quiz 5-17
Summary 5-18
Practice Overview 5: Working with Metrics 5-19

6 Using Baselines

Objectives 6-2
Comparative Performance Analysis with AWR Baselines 6-3
Automatic Workload Repository Baselines 6-4
Moving Window Baseline 6-5
Baselines in Performance Page Settings 6-6
Baseline Templates 6-7
AWR Baselines 6-8
Creating AWR Baselines 6-9
Single AWR Baseline 6-10
Creating a Repeating Baseline Template 6-11
Managing Baselines with PL/SQL 6-12
Generating a Baseline Template for a Single Time Period 6-13
Creating a Repeating Baseline Template 6-14
Baseline Views 6-15
Performance Monitoring and Baselines 6-16
Defining Alert Thresholds Using a Static Baseline 6-18
Using EM to Quickly Configure Adaptive Thresholds 6-19
Changing Adaptive Threshold Settings 6-21
Quiz 6-22
Summary 6-23
Practice 6: Overview Using AWR Baselines 6-24

7 Using AWR-Based Tools

Objectives 7-2
Automatic Maintenance Tasks 7-3
Maintenance Windows 7-4
Default Maintenance Plan 7-5
Automated Maintenance Task Priorities 7-6
Tuning Automatic Maintenance Tasks 7-7
ADDM Performance Monitoring 7-8
ADDM and Database Time 7-9
DBTime-Graph and ADDM Methodology 7-10
Top Performance Issues Detected 7-12
Database Control and ADDM Findings 7-13
ADDM Analysis Results 7-14
ADDM Recommendations 7-15

Database Control and ADDM Task	7-16
Changing ADDM Attributes	7-17
Retrieving ADDM Reports by Using SQL	7-18
Quiz	7-19
Active Session History: Overview	7-20
Active Session History: Mechanics	7-21
ASH Sampling: Example	7-22
Accessing ASH Data	7-23
Analyzing the ASH Data	7-24
Generating ASH Reports	7-25
ASH Report Script	7-26
ASH Report: General Section	7-27
ASH Report Structure	7-28
ASH Report: Activity Over Time	7-29
New or Enhanced Automatic Workload Repository Views	7-30
Quiz	7-31
Summary	7-32
Practice 7 Overview: Using AWR-Based Tools	7-33
8 Monitoring Applications	
Objectives	8-2
What Is a Service?	8-3
Service Attributes	8-4
Service Types	8-5
Creating Services	8-6
Managing Services in a Single-Instance Environment	8-7
Where Are Services Used?	8-8
Using Services with Client Applications	8-9
Using Services with the Resource Manager	8-10
Services and Resource Manager with EM	8-11
Services and the Resource Manager: Example	8-12
Services and the Scheduler with EM	8-13
Services and the Scheduler: Example	8-15
Using Services with Metric Thresholds	8-16
Changing Service Thresholds by Using EM	8-17
Services and Metric Thresholds: Example	8-18
Service Aggregation and Tracing	8-19
Top Services Performance Page	8-20
Service Aggregation Configuration	8-21
Service Aggregation: Example	8-22
Client Identifier Aggregation and Tracing	8-23

trcsess Utility 8-24
 Service Performance Views 8-25
 Quiz 8-27
 Summary 8-28
 Practice 8 Overview: Using Services 8-29

9 Identifying Problem SQL Statements

Objectives 9-2
 SQL Statement Processing Phases 9-3
 Parse Phase 9-4
 SQL Cursor Storage 9-5
 Cursor Usage and Parsing 9-6
 SQL Statement Processing Phases: Bind 9-8
 SQL Statement Processing Phases: Execute and Fetch 9-9
 Processing a DML Statement 9-10
 COMMIT Processing 9-12
 Role of the Oracle Optimizer 9-13
 Quiz 9-15
 Identifying Bad SQL 9-16
 TOP SQL Reports 9-17
 SQL Monitoring 9-18
 Monitored SQL Executions 9-19
 SQL Monitoring List 9-20
 Monitored SQL Execution Details 9-21
 The SQL Monitoring Report 9-22
 Quiz 9-23
 What Is an Execution Plan? 9-24
 Methods for Viewing Execution Plans 9-25
 Uses of Execution Plans 9-27
 DBMS_XPLAN Package: Overview 9-28
 EXPLAIN PLAN Command 9-30
 EXPLAIN PLAN Command: Example 9-31
 EXPLAIN PLAN Command: Output 9-32
 Reading an Execution Plan 9-33
 Using the V\$SQL_PLAN View 9-34
 V\$SQL_PLAN Columns 9-35
 Querying V\$SQL_PLAN 9-36
 V\$SQL_PLAN_STATISTICS View 9-37
 Querying the AWR 9-38
 SQL*Plus AUTOTRACE 9-40

Using SQL*Plus AUTOTRACE	9-41
SQL*Plus AUTOTRACE: Statistics	9-42
SQL Trace Facility	9-43
How to Use the SQL Trace Facility	9-45
Initialization Parameters	9-46
Enabling SQL Trace	9-48
Disabling SQL Trace	9-49
Formatting Your Trace Files	9-50
TKPROF Command Options	9-51
Output of the TKPROF Command	9-53
TKPROF Output with No Index: Example	9-58
TKPROF Output with Index: Example	9-59
Generate an Optimizer Trace	9-60
Quiz	9-61
Summary	9-62
Practice Overview 9: Using Execution Plan Utilities	9-63
10 Influencing the Optimizer	
Objectives	10-2
Functions of the Query Optimizer	10-3
Selectivity	10-5
Cardinality and Cost	10-6
Changing Optimizer Behavior	10-7
Optimizer Statistics	10-9
Extended Statistics	10-10
Optimizer Parameters	10-11
Controlling the Behavior of the Optimizer with Parameters	10-12
Enabling Query Optimizer Features	10-14
Using Hints	10-15
Influencing the Optimizer Approach	10-16
Optimizing SQL Statements	10-17
Quiz	10-18
Access Paths	10-19
Choosing an Access Path	10-20
Full Table Scans	10-21
Row ID Scans	10-23
Index Operations	10-24
B*Tree Index Operations	10-25
Bitmap Indexes	10-26
Bitmap Index Access	10-27

Combining Bitmaps 10-28
Bitmap Operations 10-29
Join Operations 10-30
Join Methods 10-31
Nested Loop Joins 10-32
Hash Joins 10-34
Sort-Merge Joins 10-35
Join Performance 10-37
How the Query Optimizer Chooses Execution Plans for Joins 10-38
Sort Operations 10-40
Tuning Sort Performance 10-41
Quiz 10-42
Summary 10-43
Practice 10 Overview: Influencing the Optimizer 10-44

11 Reducing the Cost

Objectives 11-2
Reducing the Cost 11-3
Index Maintenance 11-4
Dropping Indexes 11-6
Creating Indexes 11-7
Other Index Options 11-8
SQL Access Advisor 11-9
Quiz 11-10
Table Maintenance for Performance 11-11
Table Reorganization Methods 11-12
Space Management 11-14
Extent Management 11-15
Locally Managed Extents 11-16
Large Extents: Considerations 11-17
How Table Data Is Stored 11-19
Anatomy of a Database Block 11-20
Minimize Block Visits 11-21
The DB_BLOCK_SIZE Parameter 11-22
Small Block Size: Considerations 11-23
Large Block Size: Considerations 11-24
Block Allocation 11-25
Free Lists 11-26
Block Space Management 11-27
Block Space Management with Free Lists 11-28
Automatic Segment Space Management 11-30

Automatic Segment Space Management at Work	11-31
Block Space Management with ASSM	11-33
Creating an Automatic Segment Space Management Segment	11-34
Quiz	11-35
Migration and Chaining	11-36
Guidelines for PCTFREE and PCTUSED	11-38
Detecting Migration and Chaining	11-39
Selecting Migrated Rows	11-40
Eliminating Migrated Rows	11-41
Shrinking Segments: Overview	11-43
Shrinking Segments: Considerations	11-44
Shrinking Segments by Using SQL	11-45
Segment Shrink: Basic Execution	11-46
Segment Shrink: Execution Considerations	11-47
Using EM to Shrink Segments	11-48
Table Compression: Overview	11-49
Table Compression Concepts	11-50
Compressing Table Data	11-51
Using Table Compression	11-53
Using the Compression Advisor	11-54
Viewing Table Compression Information	11-55
Quiz	11-56
Summary	11-57
Practice 11 Overview: Reducing the Cost	11-58

12 Using SQL Performance Analyzer

Objectives	12-2
Real Application Testing: Overview	12-3
Real Application Testing: Use Cases	12-4
SQL Performance Analyzer: Process	12-5
Capturing the SQL Workload	12-7
Creating a SQL Performance Analyzer Task	12-8
SQL Performance Analyzer: Tasks	12-9
Parameter Change	12-10
SQL Performance Analyzer Task Page	12-12
Comparison Report	12-13
Comparison Report SQL Detail	12-14
Tuning Regressing Statements	12-15
Tuning Regressed SQL Statements	12-17
Preventing Regressions	12-18
Parameter Change Analysis	12-19

Guided Workflow Analysis 12-20
SQL Performance Analyzer: PL/SQL Example 12-21
SQL Performance Analyzer: Data Dictionary Views 12-23
Quiz 12-24
Summary 12-25
Practice 12: Overview 12-26

13 SQL Performance Management

Objectives 13-2
Maintaining SQL Performance 13-3
Maintaining Optimizer Statistics 13-4
Automated Maintenance Tasks 13-5
Statistic Gathering Options 13-6
Setting Statistic Preferences 13-7
Restore Statistics 13-9
Deferred Statistics Publishing: Overview 13-10
Deferred Statistics Publishing: Example 13-12
Automatic SQL Tuning: Overview 13-13
SQL Statement Profiling 13-14
Plan Tuning Flow and SQL Profile Creation 13-15
SQL Tuning Loop 13-16
Using SQL Profiles 13-17
SQL Tuning Advisor: Overview 13-18
Using the SQL Tuning Advisor 13-19
SQL Tuning Advisor Options 13-20
SQL Tuning Advisor Recommendations 13-21
Using the SQL Tuning Advisor: Example 13-22
Alternative Execution Plans 13-23
Quiz 13-25
Using the SQL Access Advisor 13-26
View Recommendations 13-28
View Recommendation Details 13-29
SQL Plan Management: Overview 13-30
SQL Plan Baseline: Architecture 13-31
Loading SQL Plan Baselines 13-33
Evolving SQL Plan Baselines 13-34
Important Baseline SQL Plan Attributes 13-35
SQL Plan Selection 13-37
Possible SQL Plan Manageability Scenarios 13-39
SQL Performance Analyzer and SQL Plan Baseline Scenario 13-40
Loading a SQL Plan Baseline Automatically 13-41

Purging SQL Management Base Policy	13-42
Enterprise Manager and SQL Plan Baselines	13-43
Quiz	13-44
Summary	13-45
Practice 13: Overview Using SQL Plan Management	13-46

14 Using Database Replay

Objectives	14-2
Using Database Replay	14-3
The Big Picture	14-4
System Architecture: Capture	14-5
System Architecture: Processing the Workload	14-7
System Architecture: Replay	14-8
Capture Considerations	14-9
Replay Considerations: Preparation	14-10
Replay Considerations	14-11
Replay Options	14-12
Replay Analysis	14-13
Database Replay Workflow in Enterprise Manager	14-15
Capturing Workload with Enterprise Manager	14-16
Capture Wizard: Plan Environment	14-17
Capture Wizard: Options	14-18
Capture Wizard: Parameters	14-19
Viewing Capture Progress	14-20
Viewing Capture Report	14-21
Export Capture AWR Data	14-22
Quiz	14-23
Viewing Workload Capture History	14-24
Processing Captured Workload	14-25
Using the Preprocess Captured Workload Wizard	14-26
Using the Replay Workload Wizard	14-27
Replay Workload: Prerequisites	14-28
Replay Workload: Choose Initial Options	14-29
Replay Workload: Customize Options	14-30
Replay Workload: Prepare Replay Clients	14-31
Replay Workload: Wait for Client Connections	14-32
Replay Workload: Replay Started	14-33
Viewing Workload Replay Progress	14-34
Viewing Workload Replay Statistics	14-35
Packages and Procedures	14-37
Data Dictionary Views: Database Replay	14-38

Database Replay: PL/SQL Example 14-39
 Calibrating Replay Clients 14-41
 Quiz 14-42
 Summary 14-43
 Practice 14: Overview 14-44

15 Tuning the Shared Pool

Objectives 15-2
 Shared Pool Architecture 15-3
 Shared Pool Operation 15-4
 The Library Cache 15-5
 Latch and Mutex 15-7
 Latch and Mutex: Views and Statistics 15-9
 Diagnostic Tools for Tuning the Shared Pool 15-11
 AWR/Statspack Indicators 15-13
 Load Profile 15-14
 Instance Efficiencies 15-15
 Top Timed Events 15-16
 Time Model 15-17
 Library Cache Activity 15-19
 Avoid Hard Parses 15-20
 Are Cursors Being Shared? 15-21
 Sharing Cursors 15-23
 Adaptive Cursor Sharing: Example 15-25
 Adaptive Cursor Sharing Views 15-27
 Interacting with Adaptive Cursor Sharing 15-28
 Reduce the Cost of Soft Parses 15-29
 Quiz 15-30
 Sizing the Shared Pool 15-31
 Shared Pool Advisory 15-32
 Shared Pool Advisor 15-34
 Avoiding Fragmentation 15-35
 Large Memory Requirements 15-36
 Tuning the Shared Pool Reserved Space 15-38
 Keeping Large Objects 15-40
 Data Dictionary Cache 15-42
 Dictionary Cache Misses 15-43
 SQL Query Result Cache: Overview 15-44
 Managing the SQL Query Result Cache 15-45
 Using the `RESULT_CACHE` Hint 15-47
 Using Table Annotation to Control Result Caching 15-48

Using the DBMS_RESULT_CACHE Package	15-49
Viewing SQL Result Cache Dictionary Information	15-50
SQL Query Result Cache: Considerations	15-51
Quiz	15-52
Summary	15-53
Practice Overview 15: Tuning the Shared Pool	15-54
16 Tuning the Buffer Cache	
Objectives	16-2
Oracle Database Architecture	16-3
Buffer Cache: Highlights	16-4
Database Buffers	16-5
Buffer Hash Table for Lookups	16-6
Working Sets	16-7
Tuning Goals and Techniques	16-9
Symptoms	16-11
Cache Buffer Chains Latch Contention	16-12
Finding Hot Segments	16-13
Buffer Busy Waits	16-14
Buffer Cache Hit Ratio	16-15
Buffer Cache Hit Ratio Is Not Everything	16-16
Interpreting Buffer Cache Hit Ratio	16-17
Read Waits	16-19
Free Buffer Waits	16-21
Solutions	16-22
Sizing the Buffer Cache	16-23
Buffer Cache Size Parameters	16-24
Quiz	16-25
Dynamic Buffer Cache Advisory Parameter	16-26
Buffer Cache Advisory View	16-27
Using the V\$DB_CACHE_ADVICE View	16-28
Using the Buffer Cache Advisory with EM	16-29
Caching Tables	16-30
Multiple Buffer Pools	16-31
Enabling Multiple Buffer Pools	16-33
Calculating the Hit Ratio for Multiple Pools	16-34
Multiple Block Sizes	16-36
Multiple Database Writers	16-37
Multiple I/O Slaves	16-38
Use Multiple Writers or I/O Slaves	16-39
Private Pool for I/O Intensive Operations	16-40

Automatically Tuned Multiblock Reads	16-41
DB Smart Flash Cache Overview	16-42
Using DB Smart Flash Cache	16-43
DB Smart Flash Cache Architecture Overview	16-44
Configuring DB Smart Flash Cache	16-45
Sizing DB Smart Flash Cache	16-46
Specifying DB Smart Flash Cache for a Table	16-47
Flushing the Buffer Cache (for Testing Only)	16-48
Quiz	16-49
Summary	16-50
Practice 16: Overview Tuning the Buffer Cache	16-51

17 Tuning PGA and Temporary Space

Objectives	17-2
SQL Memory Usage	17-3
Performance Impact	17-4
Automatic PGA Memory	17-5
SQL Memory Manager	17-6
Configuring Automatic PGA Memory	17-8
Setting <code>PGA_AGGREGATE_TARGET</code> Initially	17-9
Monitoring SQL Memory Usage	17-10
Monitoring SQL Memory Usage: Examples	17-12
Tuning SQL Memory Usage	17-13
PGA Target Advice Statistics	17-14
PGA Target Advice Histograms	17-15
Automatic PGA and Enterprise Manager	17-16
Automatic PGA and AWR Reports	17-17
Temporary Tablespace Management: Overview	17-18
Temporary Tablespace: Best Practice	17-19
Configuring Temporary Tablespace	17-20
Temporary Tablespace Group: Overview	17-22
Temporary Tablespace Group: Benefits	17-23
Creating Temporary Tablespace Groups	17-24
Maintaining Temporary Tablespace Groups	17-25
View Tablespace Groups	17-26
Monitoring Temporary Tablespace	17-27
Temporary Tablespace Shrink	17-28
Tablespace Option for Creating Temporary Table	17-29
Quiz	17-30
Summary	17-31

Practice Overview 17: Tuning PGA Memory 17-32

18 Automatic Memory Management

Objectives 18-2

Oracle Database Architecture 18-3

Dynamic SGA 18-4

Granule 18-5

Memory Advisories 18-6

Manually Adding Granules to Components 18-7

Increasing the Size of an SGA Component 18-8

Automatic Shared Memory Management: Overview 18-9

SGA Sizing Parameters: Overview 18-10

Dynamic SGA Transfer Modes 18-11

Memory Broker Architecture 18-12

Manually Resizing Dynamic SGA Parameters 18-13

Behavior of Auto-Tuned SGA Parameters 18-14

Behavior of Manually Tuned SGA Parameters 18-15

Using the `V$PARAMETER` View 18-16

Resizing `SGA_TARGET` 18-17

Disabling Automatic Shared Memory Management 18-18

Configuring ASMM 18-19

SGA Advisor 18-20

Monitoring ASMM 18-21

Automatic Memory Management: Overview 18-22

Oracle Database Memory Parameters 18-24

Automatic Memory Parameter Dependency 18-25

Enabling Automatic Memory Management 18-26

Monitoring Automatic Memory Management 18-27

DBCA and Automatic Memory Management 18-29

Quiz 18-30

Summary 18-31

Practice 18: Overview Using Automatic Memory Tuning 18-32

19 Tuning I/O

Objectives 19-2

I/O Architecture 19-3

File System Characteristics 19-4

I/O Modes 19-5

Direct I/O 19-6

Bandwidth Versus Size 19-7

Important I/O Metrics for Oracle Databases 19-8

I/O Calibration and Enterprise Manager	19-10
I/O Calibration and the PL/SQL Interface	19-11
I/O Statistics: Overview	19-13
I/O Statistics and Enterprise Manager	19-14
Stripe and Mirror Everything	19-16
Using RAID	19-17
RAID Cost Versus Benefits	19-18
Should I Use RAID 1 or RAID 5?	19-20
Diagnostics	19-21
Database I/O Tuning	19-22
Quiz	19-23
What Is Automatic Storage Management?	19-24
Tuning ASM	19-25
How Many Disk Groups per Database	19-26
Which RAID Configuration for Best Availability?	19-27
ASM Mirroring Guidelines	19-28
ASM Striping Granularity	19-29
What Type of Striping Works Best?	19-30
ASM Striping Only	19-31
Hardware RAID Striped LUNs	19-32
ASM Guidelines	19-33
ASM Instance Initialization Parameters	19-34
Dynamic Performance Views	19-35
Monitoring Long-Running Operations by Using V\$ASM_OPERATION	19-37
ASM Instance Performance Diagnostics	19-38
ASM Performance Page	19-39
Database Instance Parameter Changes	19-40
ASM Scalability	19-41
Quiz	19-42
Summary	19-43

20 Performance Tuning Summary

Objectives	20-2
Necessary Initialization Parameters with Little Performance Impact	20-3
Important Initialization Parameters with Performance Impact	20-4
Sizing Memory Initially	20-6
UGA and Oracle Shared Server	20-7
Large Pool	20-8
Tuning the Large Pool	20-9
Database High Availability: Best Practices	20-10
Undo Tablespace: Best Practices	20-11

Temporary Tablespace: Best Practices	20-12
General Tablespace: Best Practices	20-14
Internal Fragmentation Considerations	20-15
Block Size: Advantages and Disadvantages	20-16
Automatic Checkpoint Tuning	20-17
Sizing the Redo Log Buffer	20-18
Sizing Redo Log Files	20-19
Increasing the Performance of Archiving	20-20
Automatic Statistics Gathering	20-22
Automatic Statistics Collection: Considerations	20-23
Commonly Observed Wait Events	20-24
Additional Statistics	20-25
Top 10 Mistakes Found in Customer Systems	20-26
Quiz	20-28
Summary	20-29

Appendix A: Practices and Solutions

B Using Statspack

Objectives	B-2
Introduction to Statspack	B-3
Statspack Scripts	B-4
Installing Statspack	B-6
Capturing Statspack Snapshots	B-7
Configuring Snapshot Data Capture	B-8
Statspack Snapshot Levels	B-9
Statspack Baselines and Purging	B-11
Reporting with Statspack	B-13
Statspack Considerations	B-14
Statspack and AWR Reports	B-16
Reading a Statspack or AWR Report	B-17
Statspack/AWR Report Drilldown Sections	B-18
Report Drilldown Examples	B-20
Load Profile Section	B-21
Time Model Section	B-22
Statspack and AWR	B-23
Summary	B-24
Practice Overview: Use Statspack	B-25

Preface

Profile

Before You Begin This Course

Before you begin this course, you should be able to perform the duties of an entry level DBA, have taken the Oracle Database 11g: Administration Workshop 1 and 2 courses or equivalent experience. It is recommended that you also have taken the Oracle Database 11g: SQL and PL/SQL Fundamentals course or have equivalent experience.

How This Course Is Organized

Oracle Database 11g: Performance Tuning is an instructor-led course featuring lectures and hands-on exercises. Online demonstrations and written practice sessions reinforce the concepts and skills that are introduced.

Related Publications

Oracle Publications

Title	Part Number
Oracle Database Administrator's Guide 11g (11.2)	E10595-06
Oracle Database 2 Day + Performance Tuning Guide 11g Release 2 (11.2)	E10822-02
Oracle Database Performance Tuning Guide 11g Release 2 (11.2)	E10821-04
Oracle Database Reference 11g Release 2 (11.2)	E10820-03

Additional Publications

- System release bulletins
- Installation and user's guides
- *read.me* files
- International Oracle User's Group (IOUG) articles
- *Oracle Magazine*

Typographic Conventions

The following two lists explain Oracle University typographical conventions for words that appear within regular text or within code samples.

1. Typographic Conventions for Words Within Regular Text

Convention	Object or Term	Example
Courier New	User input; commands; column, table, and schema names; functions; PL/SQL objects; paths	Use the <code>SELECT</code> command to view information stored in the <code>LAST_NAME</code> column of the <code>EMPLOYEES</code> table. Enter <code>300</code> . Log in as <code>scott</code>
Initial cap	Triggers; user interface object names, such as button names	Assign a When-Validate-Item trigger to the ORD block. Click the Cancel button.
Italic	Titles of courses and manuals; emphasized words or phrases; placeholders or variables	For more information on the subject see <i>Oracle SQL Reference</i> <i>Manual</i> Do <i>not</i> save changes to the database. Enter <i>hostname</i> , where <i>hostname</i> is the host on which the password is to be changed.
Quotation marks	Lesson or module titles referenced within a course	This subject is covered in Lesson 3, “Working with Objects.”

Typographic Conventions (continued)

2. Typographic Conventions for Words Within Code Samples

Convention	Object or Term	Example
Uppercase	Commands, functions	<code>SELECT employee_id FROM employees;</code>
Lowercase, italic	Syntax variables	<code>CREATE ROLE <i>role</i>;</code>
Initial cap	Forms triggers	Form module: ORD Trigger level: S_ITEM.QUANTITY item Trigger name: When-Validate-Item . . .
Lowercase	Column names, table names, filenames, PL/SQL objects	. . . <code>OG_ACTIVATE_LAYER (OG_GET_LAYER ('prod_pie_layer'))</code> . . . <code>SELECT last_name FROM employees;</code>
Bold	Text that must be entered by a user	<code>CREATE USER scott IDENTIFIED BY tiger;</code>

19

Tuning I/O

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Objectives

After completing this lesson, you should be able to do the following:

- Diagnose database I/O issues
- Describe the Stripe and Mirror Everything (SAME) concept
- Set `filesystemio_options`
- Choose appropriate I/O solutions
- Tune Automatic Storage Management (ASM)

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

I/O Architecture

Oracle Database 11g includes three standard storage options:

- File system
 - Network attached storage (NAS)
 - Storage area network (SAN)
 - Direct attached storage
- Automatic Storage Management (ASM)

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

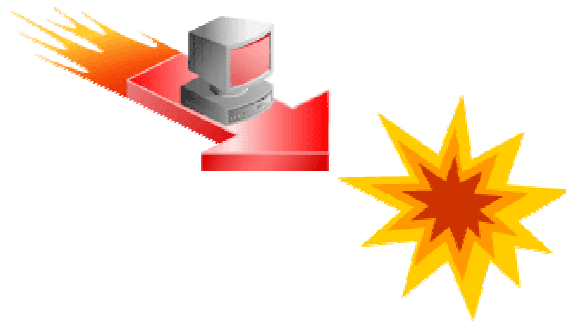
I/O Architecture

- The Oracle database uses a logical storage container called a tablespace to store all permanent data. The tablespace is associated with one or more data files. Data files are the physical storage entities. A data file is mapped to an operating system (OS) file, or an ASM file. OS files can be on direct attached storage, network attached storage (NAS), or storage area network (SAN) devices that the OS can represent as files. Each of these methods has certain advantages.
- OS files provide easy naming and administration, but standard file system access requires the use of the OS buffer cache, which has a longer code path and requires additional copies of the data blocks to be made. The `FILESYSTEMIO_OPTIONS` initialization parameter allows the Oracle instance to use direct I/O or asynchronous I/O to OS files.
- The Automatic Storage Management (ASM) feature allows the convenience of OS files with the speed of direct I/O. ASM uses disk partitions, with no intervening filesystem code to provide high performance data access.

File System Characteristics

Certain characteristics are better for database use:

- Write-through-cache ability
- Write acknowledgement
- Security
- Journaling
- High performance



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

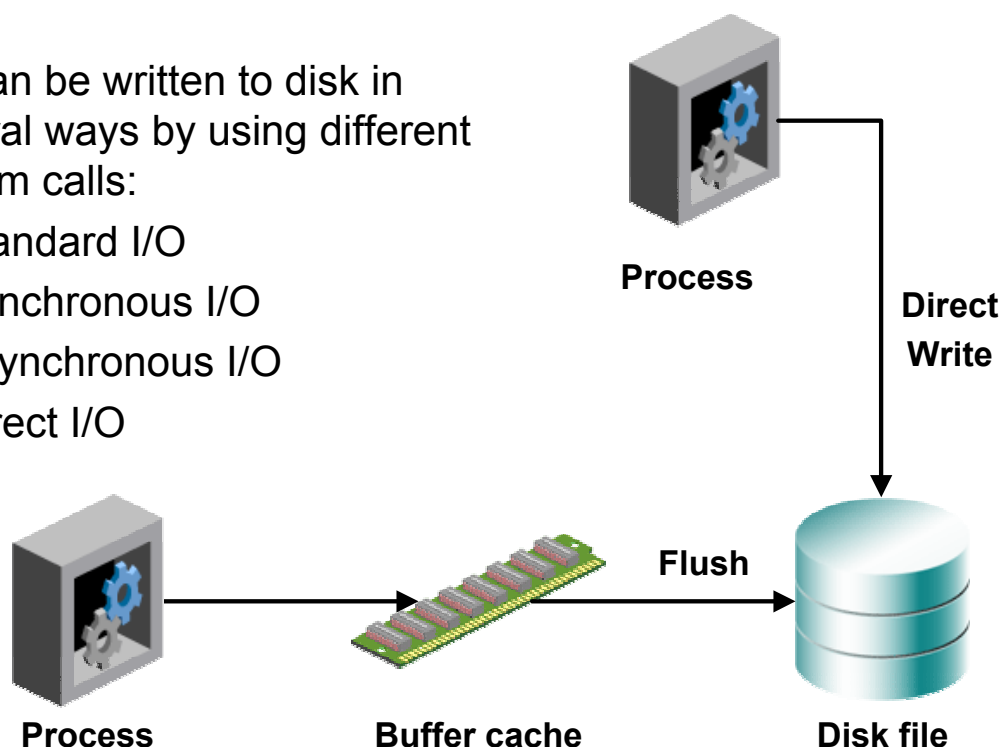
File System Characteristics

- In choosing a file system, performance is important, but not the primary consideration. It does not matter how fast the file system is if the data is corrupted, lost, or compromised. The Oracle database does not support database files on file systems that do not have a write-through-cache capability. Supported file systems must acknowledge writes. The Oracle database has security requirements as well. Database files require file permissions for proper database security. Data files should be accessible only to the database owner, and all other access should be controlled by the database itself.
- Journaling records the changes to the file system in a journal file similarly to how database changes are recorded in redo logs. If the server crashes or shuts down without synchronizing the disks, the journal file can be applied to the file system, thus restoring the file system integrity very quickly at OS startup.
- File systems and OS combinations that support direct I/O and asynchronous I/O generally have better performance than other combinations that do not support these features.

I/O Modes

I/O can be written to disk in several ways by using different system calls:

- Standard I/O
- Synchronous I/O
- Asynchronous I/O
- Direct I/O



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

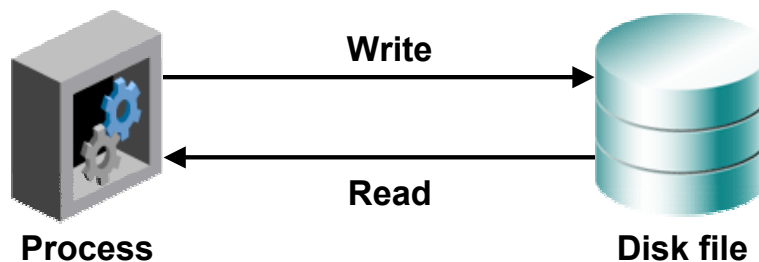
I/O Modes

- In any OS, there are multiple methods of writing to disk.
- The standard method writes the data to the operating system buffer cache, and then the buffer is written to disk by a kernel process at a later time. If the machine crashes or loses power before the buffer is written to disk, the data is lost. This method is fast, but not acceptable for Oracle database file I/O because of the potential for data loss.
- Synchronous I/O, also known as write-thru-cache, writes the data to the buffer but also forces an immediate write to the disk. The process that makes the write request is suspended until the write is completed. The completion is signaled by the disk controller when the buffer is written to disk. On some disk subsystems with large caches, the controller signals a completion when the buffer is written into the subsystem cache. Although synchronous writes are slower, they are the default method for DBWR and LGWR processes to perform their writes because this I/O method is reliable and supported on all operating systems.
- Asynchronous writes allow instance background processes to make multiple write requests without being suspended after each request. The completion must still be signaled back to the process. This permits a much larger throughput because the background process can continue work without waiting for the disk I/O to complete. The ability to perform asynchronous writes varies by operating system.

Direct I/O

Direct I/O is considered to be the high-performance solution.

- Direct reads and writes do not use the OS buffer cache.
- Direct reads and writes can move larger buffers than file system I/Os.
- Direct I/O can be synchronous or asynchronous.



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

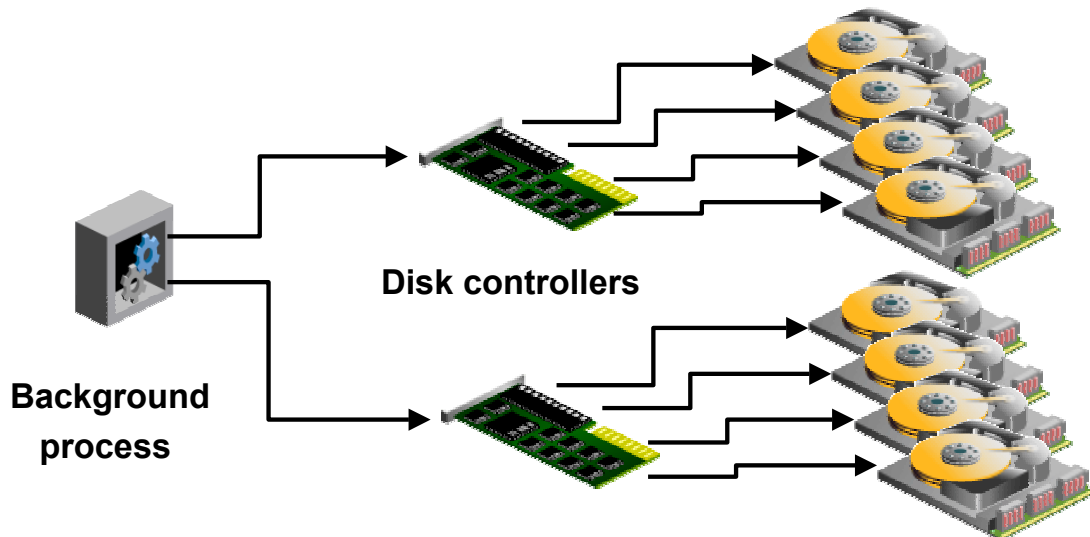
Direct I/O

- The main benefit of direct I/O is the lack of caching. Oracle server caches its own database blocks. For the normal reads in an online transaction processing (OLTP) environment, the simple read-ahead that is done by the file system is not effective. By avoiding the OS cache, the direct reads and writes are not copied in memory, and the application performing the reads and writes can choose much larger buffers for the data transfers than are allowed by the file system.
- This also frees OS memory used for caching I/O that can be applied to the SGA, where the database can make more effective use of the space.
- Oracle recommends the use of direct I/O on any OS file system that supports it.
- Oracle Cluster File System (OCFS and OCFS2), ASM, NFS, various Storage Area Networks, and Clustered File Systems are used for Oracle Real Application Clusters.
- I/O can be set to be direct and asynchronous or direct and synchronous. Direct and asynchronous is the recommended setting for high performance systems.

Bandwidth Versus Size

I/O performance depends on bandwidth.

- Number of disks, not size
- Number of controllers



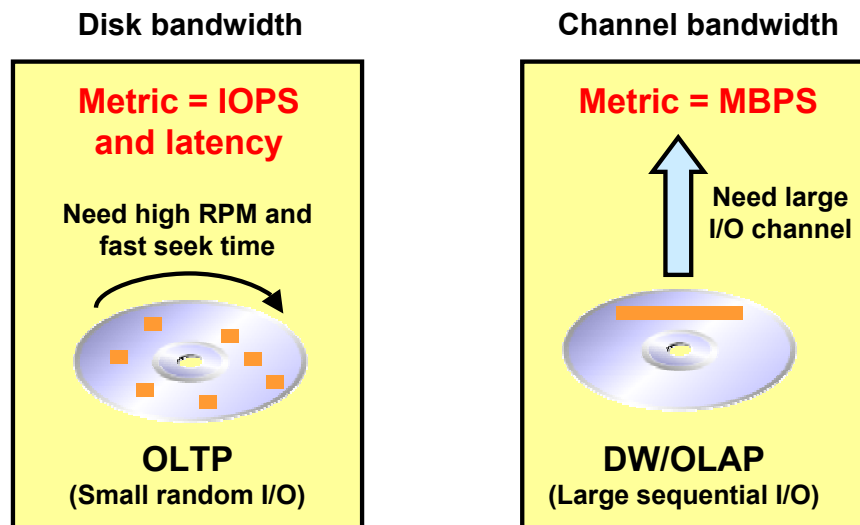
ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Bandwidth Versus Size

Many times, database disk subsystems are sized only by storage capacity. The main consideration for performance in the I/O subsystem is I/Os per second or bandwidth. More bandwidth becomes available as you increase the number of disks you spread the I/Os across. For example, each disk may have a transfer rate of 160 MB/sec; therefore, four disks will have a combined bandwidth of 640 MB/sec. The disk controllers play an important role as well, because each disk controller has a limited bandwidth. Continuing the example, a SCSI disk controller with a transfer rate of 600 MB/sec will become a bottleneck on disk transfers if there are more than 4–5 disks, even though the specifications say that the controller can manage up to 16 disks.

Important I/O Metrics for Oracle Databases



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Important I/O Metrics for Oracle Databases

The database I/O workload typically consists of small random I/O and large sequential I/O. Small random I/O is more prevalent in an OLTP application environment in which each foreground process reads a data block into the buffer cache for updates and the changed blocks are written in batches by the dbwr process. Large sequential I/O is common in an OLAP application environment. The OLTP application performance depends on how fast a small I/O is serviced, which depends on how fast the disk can spin and find the data. Large I/O performance depends on the capacity of the I/O channel that connects the server to the storage array; large I/O throughput is better when the capacity of the channel is larger.

IOPS (I/O per second): This metric represents the number of small random I/O that can be serviced in a second. The IOPS rate mainly depends on how fast the disk media can spin. The IOPS rate from a storage array can be increased either by adding more disk drives or by using disk drives with a higher RPM (revolutions per minute) rate.

MBPS (megabytes per second): The rate at which data can be transferred between the computing server node and the storage array depends on the capacity of the I/O channel that is used to transfer data. More data can be transferred through a wider pipe.

Important I/O Metrics for Oracle Databases (continued)

The throughput of a streaming data application depends on how fast this data can be transferred. Throughput is measured by using the MBPS metric.

Even though the disks themselves have an upper limit on the amount of sequential data that they can transfer, it is often channel capacity that limits the overall throughput of the system. For example, a host connected to an NAS server through a GigE switch is limited by a transfer capacity of 128 MBPS. There are usually multiple disks in the NAS that can together provide more than 128 MBPS. The channel resource becomes the limiting resource.

I/O latency: Latency is another important metric that is used in measuring the performance of an I/O subsystem. Traditionally, latency is simply the time for the disk to access a particular sector on the disk. But from the database point of view, *latency* represents all the time it takes for a submitted I/O request to be serviced by the storage. In other words, it represents the overhead before the first byte of a transfer arrives after an I/O request has been submitted. Latency values are more commonly used for small random I/O when tuning a system. If there are too many I/O requests queued up against a disk, the latency increases as the wait time in the queue increases. To improve the latency of I/O requests, data is usually striped across multiple spindles so that all I/O requests to a file do not go to the same disk. A higher latency usually indicates an overloaded system.

Other than the main resources mentioned above, there are also other storage array components that can affect I/O performance. Array vendors provide caching mechanisms to improve read throughput, but their real benefit is questionable in a database environment because Oracle Database uses caches and read-ahead mechanisms so that data is available directly from RAM rather than from the disks.

I/O Calibration and Enterprise Manager

I/O Calibration

Running the I/O calibration tool will place an extremely high I/O load on your storage system and will negatively effect the performance of any active sessions. For the best run the I/O calibration when there are no other active sessions. The calibration should complete in approximately 10 minutes.

Existing Calibration Results

Calibration Date **None**

Inputs

Max I/O per Second **None**

Max MB per Second **None**

Measurements

Maximum Tolerable Latency (ms) for Single-Block I/O Requests **None**

Status **NOT A**

Inputs for a New Calibration

Approximate Number of Physical Disks in the Database's Storage System

Maximum Tolerable Latency in Milliseconds for Single-Block I/O Requests

Schedule

☒ Immediately

I/O Calibration

Running the I/O calibration tool will place an extremely high I/O load on your storage system and will negatively effect the performance of any active sessions. For the best run the I/O calibration when there are no other active sessions. The calibration should complete in approximately 10 minutes.

Existing Calibration Results

Calibration Date **Mar 25, 2010 12:56:47 PM**

Inputs

Max I/O per Second **91**

Max MB per Second **27**

Measurements

Maximum Tolerable Latency (ms) for Single-Block I/O Requests **21**

Status **READY**

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

I/O Calibration and Enterprise Manager

- To determine the previously discussed I/O metrics, you can use the I/O Calibration tool from the Enterprise Manager Performance page or PL/SQL. I/O Calibration is a modified version of the ORION tool and is based on the asynchronous I/O library. Because calibration requires issuing enough I/O to saturate the storage system, any performance-critical sessions will be negatively impacted. Thus, you should run I/O calibration only when there is little activity on your system.
- I/O Calibration takes approximately 10 minutes to run. You can launch an I/O Calibration task directly from Enterprise Manager (as shown in the slide). You do this by clicking the Performance tab. On the Performance page, you can click the I/O tab and then the I/O Calibration button.
- On the I/O Calibration page, specify the number of physical disks attached to your database storage system as well as the maximum tolerable latency for a single-block I/O request. Then, in the Schedule section of the I/O Calibration page, specify when to execute the calibration operation. Click the Submit button to create a corresponding Scheduler job.
- On the Scheduler Jobs page, you can see the amount of time that it takes for the calibration task to run. When the job is finished, you can view the results of the calibration operation on the I/O Calibration page: maximum I/O per second, maximum megabytes per second, and average actual latency metrics.

I/O Calibration and the PL/SQL Interface

```
SQL> exec dbms_resource_manager.calibrate_io(
                                num_disks=>1,
                                max_latency=>10,
                                max_iops=>:max_iops,
                                max_mbps=>:max_mbps,
                                actual_latency=>:actual_latency);
```

PL/SQL procedure successfully completed.

```
SQL>
  2  FROM DBA_RSRC_IO_CALIBRATE;
```

MAX_IOPS	MAX_MBPS	MAX_PMBPS	LATENCY
137	12	6	64

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

I/O Calibration and the PL/SQL Interface

You can run the I/O Calibration task by using the PL/SQL interface. Execute the `CALIBRATE_IO` procedure from the `DBMS_RESOURCE_MANAGER` package. This procedure calibrates the I/O capabilities of the storage. The calibration status and results are available from the `V$IO_CALIBRATION_STATUS` and `DBA_RSRC_IO_CALIBRATE` views.

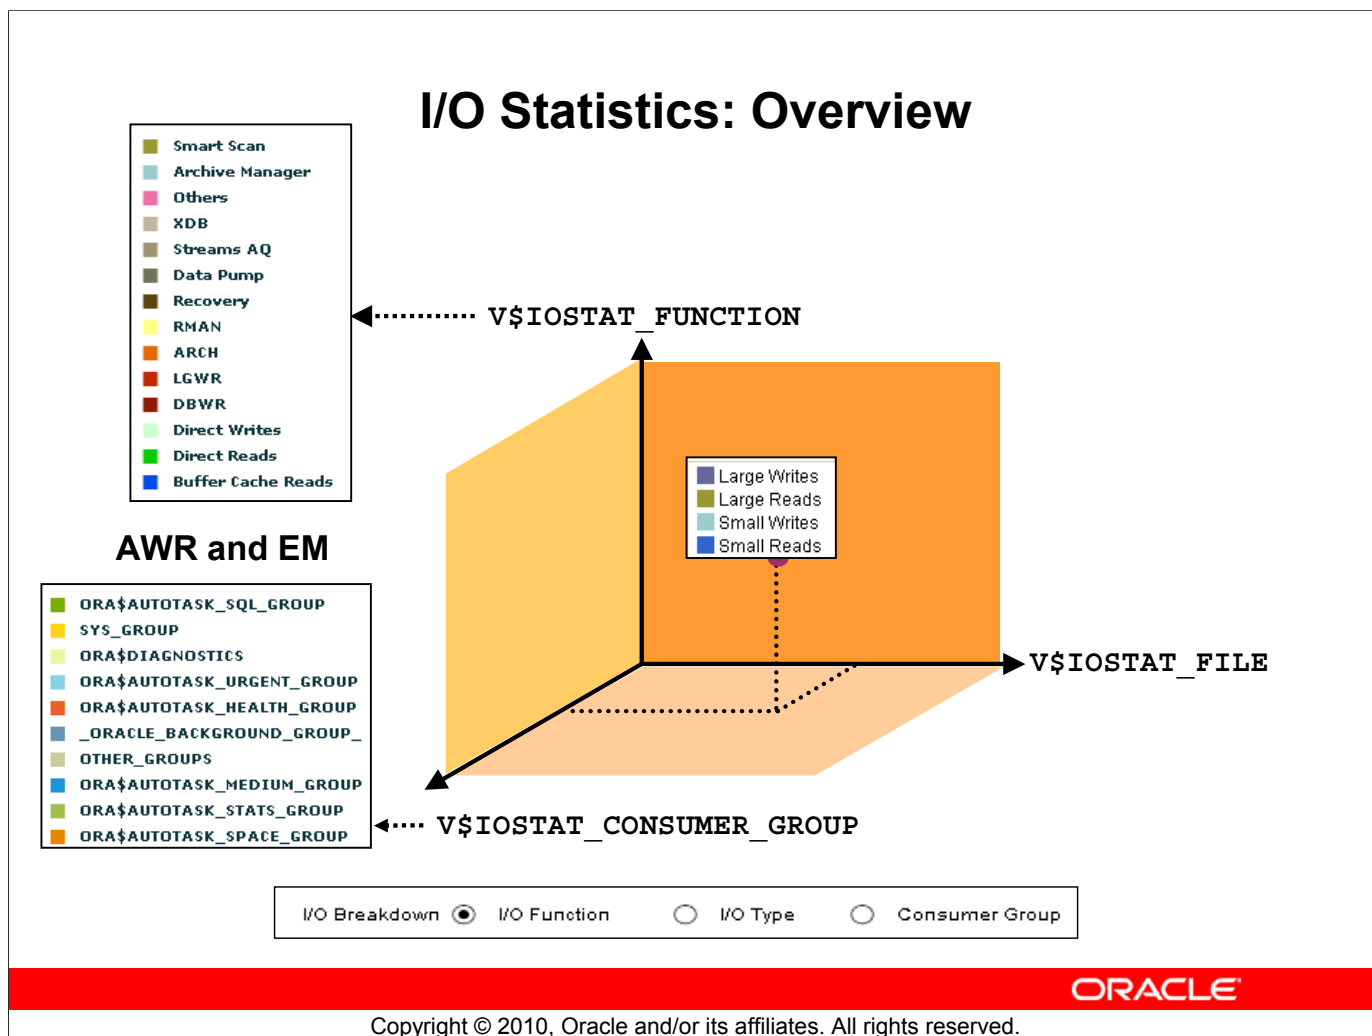
Here is a brief description of the parameters you can specify for the `CALIBRATE_IO` procedure:

- `num_disks`: Approximate number of physical disks in the database storage
- `max_latency`: Maximum tolerable latency (in milliseconds) for database-block-sized I/O requests
- `max_iops`: Maximum number of I/O requests per second that can be sustained. The I/O requests are randomly distributed, database-block-sized reads.
- `max_bps`: Maximum throughput of I/O that can be sustained (in megabytes per second). The I/O requests are randomly distributed 1 MB reads.
- `actual_latency`: Average latency of database-block-sized I/O requests at `max_iops` rate (in milliseconds)

I/O Calibration and the PL/SQL Interface (continued)

Usage notes:

- Only users with the SYSDBA privilege can run this procedure.
- Only one calibration can run at a time. If another calibration is initiated at the same time, it fails.
- For a RAC database, the workload is simultaneously generated from all instances.
- The latency time is computed only when the TIMED_STATISTICS initialization parameter is set to TRUE (which is set when STATISTICS_LEVEL is set to TYPICAL or ALL).
- The disks to be tested must be configured to use asynchronous I/O both at the OS level and in the database. The database requires the FILESYSTEMIO_OPTIONS parameter be set to ASYNCH or BOTH.



I/O Statistics: Overview

The Oracle Database provides a consistent set of statistics for all I/O issued from an Oracle instance with a set of virtual views that collect I/O statistics in three dimensions:

- **RDBMS components:** Components are grouped by functional categories, `V$IOSTAT_FUNCTION`.
- When Resource Management is enabled, I/O statistics are collected for all consumer groups that are part of the currently enabled resource plan.
- I/O statistics are also collected for the individual files that have been opened.

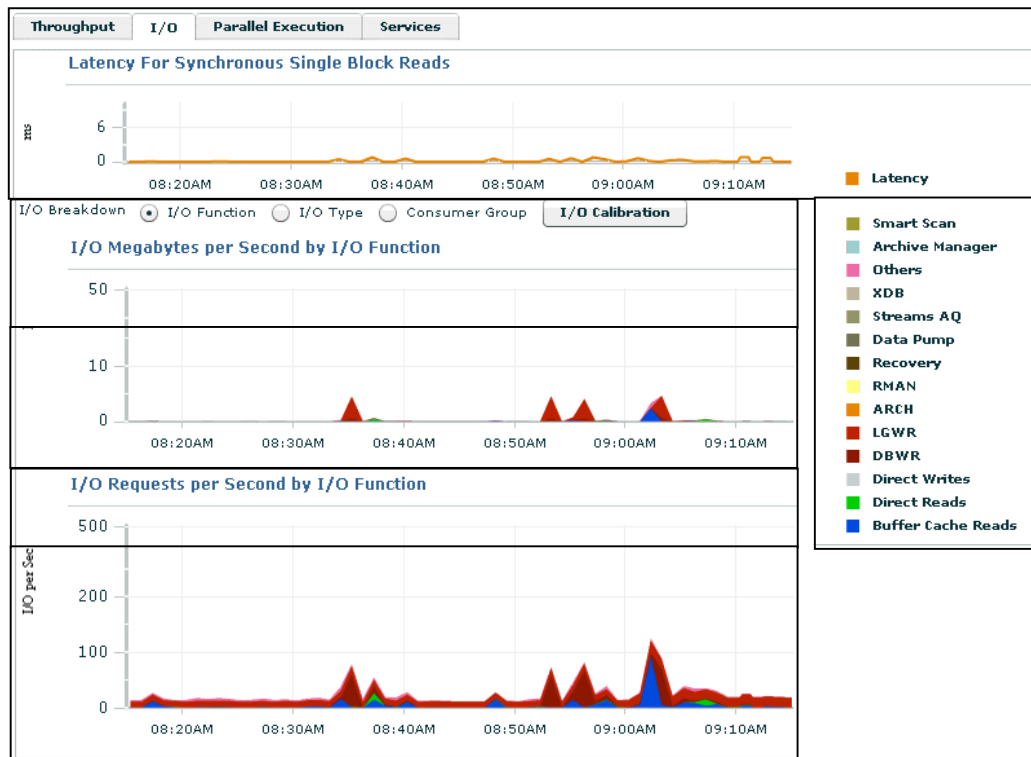
Each dimension includes statistics for read and write operations. Because read/write can occur in single block or multiblock operations, they are separated into four different operations (as shown in the slide). For each operation type, the number of corresponding requests and megabytes are cumulated. In addition to these, the total I/O wait time in milliseconds and the total waits are also cumulated for both component and consumer group statistics.

Total service time, and statistics for single block reads are collected for file statistics.

Virtual views show cumulative values for statistics. Component and consumer group statistics are transformed into AWR metrics that are sampled regularly and stored in the AWR repository. You can retrieve those metrics across a timeline directly on the Performance page of Enterprise Manager.

Note: `V$IOSTAT_NETWORK` collects network I/O statistics that are related to accessing files on a remote database instance.

I/O Statistics and Enterprise Manager



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

I/O Statistics and Enterprise Manager

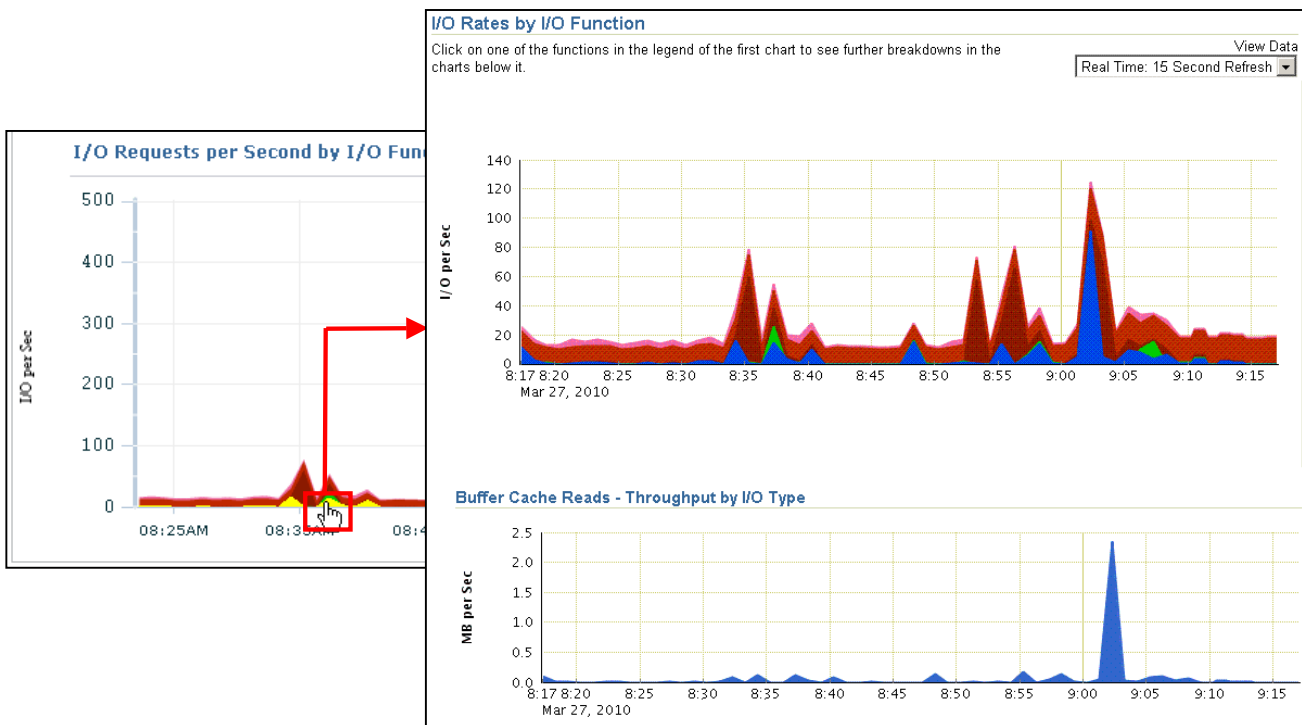
You can retrieve I/O statistics directly on the Performance page in Enterprise Manager. On the Performance page, click the I/O subtab located under the Average Active Session graph.

On the I/O subpage, you see a breakdown of I/O statistics on three possible dimensions: I/O Function, I/O Type, and Consumer Group. Select one of the options to look at the corresponding dimension graphs. The slide shows the two graphs corresponding to the I/O function dimension:

- I/O Megabytes per Second (by RDBMS component)
- I/O Requests per Second (by RDBMS component)

Note: The “Others” RDBMS component category corresponds to everything that is not directly issued from SQL (for example, PL/SQL and Java).

I/O Statistics and Enterprise Manager



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ORACLE

I/O Statistics and Enterprise Manager (continued)

From the I/O Function statistic graphs, you can drill down to a specific component by clicking that component. In the example in the slide, you drill down to the Buffer Cache Reads component. This takes you to the “I/O Rates by I/O Function” page, where you can see three important graphs for that particular component: MBPS, IOPS, and wait time.

These statistics do not affect the optimizer. The major benefit of the I/O calibration and the I/O statistics is being able to compare the current I/O performance against a measured performance level. When the current performance level deviates from the measured (expected) level, You can drill down by component, function, or file to find the source of the problem.

Stripe and Mirror Everything

- All data files to access all available bandwidth
- All database files to be on the same logical devices
- Highest performance configuration

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Stripe and Mirror Everything

Database performance is directly related to I/O performance and a major part of I/O performance is bandwidth. The problem is how to make use of all of the bandwidth all of the time. With file systems data files are placed on particular mount points that map to a single disk or a logical volume. Manual striping by placing active tables and indexes in tablespaces that have multiple data files spread across several disks or volumes is time consuming and requires constant rebalancing. Oracle Corporation recommends the Stripe and Mirror Everything (SAME) approach. You form all the disks available for the database files into a logical volume set, and stripe all the data files across this set. This method means that redo log files, archive log files, data files, and control files have equal access to all the bandwidth that is available. If high availability is desired, use two volume sets and mirror them at the OS level.

There are three general methods available to implement the SAME methodology: use a logical volume manager (LVM), use hardware-based RAID technology, or use Automatic Storage Management (ASM). LVM can be expensive, incompatible with certain Oracle database features such as Real Application Clusters, or limited in functionality.

Note: For best performance, redo log files should be created on the same disk group as the data files. Test results show SAME is a better method of utilizing the I/O system where all database files including the redo logs were striped across all the available disks, see the Oracle white paper *Optimal Storage Configuration Made Easy* at

http://technet.oracle.com/deploy/availability/pdf/oow2000_sane.pdf

Using RAID

Redundant Array of Inexpensive Devices (RAID) levels:

- Level 0
 - Striped for performance
 - No redundancy
- Level 1
 - Mirrored for safety
 - Little performance benefit
- Level 5
 - Block level redundancy (rebuild algorithm)
 - Improved read performance
 - Additional write cost

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

RAID Levels

- RAID is a Redundant Array of Inexpensive Devices (Disks). The goal of RAID is to group together several inexpensive disks in such a way as to provide fast, reliable, and inexpensive storage. The goal is met in a variety of ways, each having its own associated cost.
- RAID 0 is a set of disks with the data striped across all the members of the array. Typically, a RAID 0 array has 2–10 devices. The more devices in the array, the better the performance. The stripe size affects performance and should be relatively small (64 KB–2 MB) so that the data for a single read or write is spread across several disks.
- RAID 1 is a very reliable method that creates “mirrors” of the data on each disk in the array. The data is duplicated on each member or “mirror.” Typically, these arrays have two or three members in each array. RAID 1 arrays have better performance in high-read environments than the same array in a mixed read/write environment.
- RAID 5 is a striped set of disks with one stripe segment out of every n devoted to parity information. The parity segment is distributed across all the disks. This parity information allows “Hot Swap” and continued operation with the loss of any one disk of the array. RAID 6 is an alternate form of RAID 5 that uses two parity disks and can tolerate the loss of two disks in the set.

RAID Cost Versus Benefits

RAID cost is measured in performance and reliability.

- RAID 0:
 - Fast
 - The loss of any device damages the array.
- RAID 1:
 - Safe and expensive
 - Slight benefit in high-read environments
- RAID 5
 - Fast
 - Safe with loss of any one device
 - Possible high write cost

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

RAID Cost Versus Benefits

RAID 0 is fast. With a 4-disk array, the read and write times are approximately 80% of the time to read and write to a single disk. There is no penalty on the storage. The cost comes in reliability. The mean time between failure (MTBF) is $1/n$, where n is the number of disks in the array. The failure of any member of the array destroys the array, and the entire array must be recovered.

RAID 1: The controllers are generally “smart” enough to know that the same data is on all the mirrors, so multiple read requests are sent to different disks parallelizing the requests. The reliability increases with the number of mirrors. A failure of the array means all the mirrors in the array would have to fail at one time. The cost is measured in space, requiring each byte of data to be stored on all the mirrors.

RAID 5: The read cost of a normal operation is similar to RAID 0. The reliability is quite good; the failure of the array can be caused by a failure of two or more disks at the same time. The storage cost translates to the loss of one disk out of the array for storing parity data. This pushes the cost conscious to add more disks to the array, but this also lowers reliability. The major cost of RAID 5 is in the write speed. The write algorithm requires four I/Os for each write. Some vendors have optimized the controllers for these operations but the cost of calculating and writing the parity stripe is still there. The write cost can vary between vendors from a little more than writing to a single disk to more than two times the standard disk write.

RAID Cost Versus Benefits (continued)

RAID 1+0, 0+1, 10

These designations are vendor specific and must be looked at carefully because the designation and the configuration may not match expectations. There are two cases: RAID 0 over 1 and RAID 1 over 0. The cost and reliability measures are the same. The difference is in mean time to recover (MTTR). Vendors are not consistent with the nomenclature and RAID 10 could be either of these methods. Each of these configurations has the speed of RAID 0 and the reliability and cost of RAID 1.

The first method, RAID 1 over 0 or 0+1, combines two or more RAID 0 arrays in a RAID 1 configuration. RAID 1 over 0 has a high MTTR. If one disk in the set fails, the entire striped set must be rebuilt. During the rebuild period, the entire array will have a lowered performance while the one striped set is rebuilt. The protection of two sets is lost for the duration of the rebuild.

The other method, RAID 0 over 1 or 1+0, involves creating RAID 1 arrays of individual disks and then striping them together in a RAID 0 array. RAID 0 over 1 has a lower MTTR. If one disk fails, only that disk is lost. The rebuild time is only for that one disk. This configuration has a lowered performance and loss of protection for one disk while the disk is being rebuilt, and the time to rebuild is reduced.

RAID 50

RAID 50, as with other two digit RAID designations, refers to combined RAID types, in this case RAID 0 striping over a set of RAID 5 arrays. This configuration has a higher write performance than a RAID 5 array. The vulnerability comes when any disk fails, the RAID 5 array with the failed disk will continue with degraded service. If a second disk fails in the same RAID 5 array the entire RAID 50 array must be rebuilt.

RAID 6

RAID 6 is very similar to RAID 5, but it has two parity blocks in each stripe so the loss of two disks at the same time will not cause a loss of data. This configuration will cost more data space, 2 out of n , but also safely allows more disks to be included in a single array. Like RAID 5 there is a write penalty for the parity blocks.

Should I Use RAID 1 or RAID 5?

RAID 1 (Mirroring)

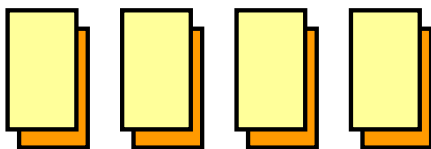
- Recommended by Oracle
- Most demanding applications

Advantages

- Best redundancy
- Best performance
- Low recovery overhead

Disadvantages

- Requires higher capacity



RAID 5 (Parity)

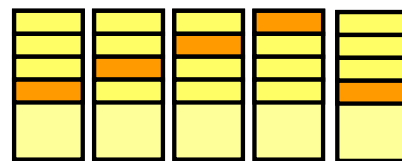
- DSS and moderate OLTP

Advantages

- Requires less capacity

Disadvantages

- Less redundancy
- Less performance
- High recovery overhead



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Should I Use RAID 1 or RAID 5?

The slide lists the advantages and disadvantages of the techniques. Although Oracle recommends using RAID 1, you need to take into account the additional disks required. The general rule is to deploy RAID 5 where the cost of storage is critical, performance is not the primary goal, and applications are primarily read operations such as data warehouse applications. The flash recovery area can be a good use of RAID 5, where the storage capacity requirement is the highest and predominantly sequential I/O.

Note: Although RAID 5 implementations have come a long way, the performance is vastly different from one storage array product to another and caution should be exercised when choosing a storage platform. The ORION tools (<http://www.oracle.com/technology/software/index.html#util>) can be used to help determine the advantages and disadvantages of arrays for your application.

Diagnostics

Indicators of I/O issues:

- Top waits are reads and writes plus:
 - Buffer busy waits
 - Write complete waits
 - DB file parallel writes
 - Enqueue waits
- The file I/O Statistics section shows high waits and AVG Buffer Wait time higher than average on certain files.

Note: On a well-performing system, the top events are likely to be CPU time, DB file scattered read, and DB file sequential read.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Diagnostics

I/O issues can become apparent in several ways. Checkpointing and redo issues often involve I/O issues as well. When the primary issue is I/O, the top wait events show read and write events taking the most time. Even these indicators require more research. If the I/O is saturated, buffer busy waits, write complete waits, DB file parallel writes, and enqueues may all appear in the top wait events.

The File I/O Statistics section of the Statspack or AWR report shows whether the waits are concentrated in one or a few data files. This could indicate a hot spot or a slow device. Hot spots can also be caused by poorly designed SQL or data structures. Tuning the SQL to reduce I/O is the first step. Tuning the access methods is the second. Consider moving tables and indexes to different tablespaces on different devices. This last step is unnecessary if the SAME principle has been applied.

Database I/O Tuning

- Configuring storage for a database depends on many variables:
 - Which data to put on which disk; complicated by vendor-configured logical units (LUNs)
 - DB application workloads: OLTP, DSS, batch versus online
 - Trade-offs between available options
 - Ongoing tuning: changes in workloads
 - Expanding or contracting your database

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Database I/O Tuning

When planning your database design storage layout, you need to determine:

- How many volumes and database files are needed?
- Where should each database file reside to reduce contention?
- What is the mapping of physical disks to logical unit (LUN)?
- What are the I/O characteristics of the primary workload?
- How will changes in the workload change the layout?
- Can the layout be changed without shutting down the database?

Consider the types of database files needed for the Oracle database. File types include control files, redo log files, data files, backup files, temp files, and flashback logs.

Choose a naming convention for your files that allows clear distinctions for file type and the database it belongs to. Managing these files can be problematic if you manage multiple databases, especially during recovery operations. Is there spare space for recovery operations? Where should it be located and managed?

One solution is to use Automatic Storage Management (ASM). ASM implements the Stripe and Mirror Everything (SAME) technique to optimize storage I/O utilization in terms of bandwidth, sequential and random access, and workload types.

Quiz

Which of the following advantages is provided by the Stripe and Mirror everything approach?

- a. Administrators must continually monitor the data file growth.
- b. Data is protected by mirroring.
- c. The redo log files should be placed on a different storage for best performance.
- d. All database files share all the available I/O bandwidth.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

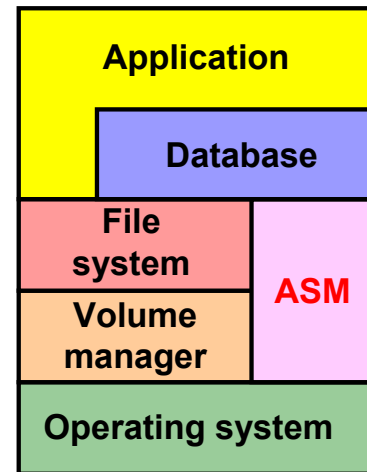
Answer: b, d

Stripe and mirror everything allows all the data files to share all the available I/O bandwidth when all database files including redo log file are placed in the same volume. By using the common volume or two mirrored volumes the growth of the data files does not require monitoring.

What Is Automatic Storage Management?

ASM:

- Is a portable and high-performance cluster file system
- Manages Oracle database files
- Distributes data across disks to balance load
- Provides integrated mirroring across disks
- Solves many storage management challenges
- Encapsulates the SAME methodology



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

What Is Automatic Storage Management?

Automatic Storage Management (ASM):

- Is an Oracle Database 10g feature
- Provides a file system and volume manager that is specifically built for database files
- Can provide management for single SMP machines or across multiple nodes of a cluster for Oracle Real Application Clusters support
- Distributes I/O load across all available resources to optimize performance while removing the need for manual I/O tuning
- Helps you to manage a dynamic database environment by allowing you to increase the size of the database without having to shut down the database to adjust the storage allocation
- Can maintain redundant copies of data to provide fault tolerance, or it can be built on top of vendor-supplied reliable storage mechanisms.

Data management is done by selecting the desired reliability and performance characteristics for classes of data rather than with human interaction on a per-file basis.

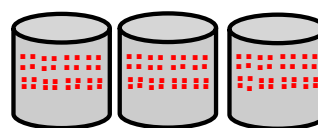
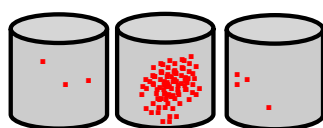
ASM capabilities save you time by automating manual storage and thereby increasing your ability to manage larger databases (and more of them) with increased efficiency.

Note: ASM encapsulates the Stripe and Mirror Everything (SAME) methodology.

Tuning ASM

- Adjust rebalancing speed
- Set redundancy on a file basis
- Set `DISK_ASYNCH_IO` to TRUE

Select	Disk Δ	Failure Group	Path	Read/Write Errors	State	Mode	Size (GB)	Used (GB)	Used (%)
<input type="checkbox"/>	DATA_0000	DATA_0000	/dev/xvdc	0	NORMAL	ONLINE	2.00	0.46	23.14
<input type="checkbox"/>	DATA_0001	DATA_0001	/dev/xvdd	0	NORMAL	ONLINE	2.00	0.45	22.71
<input type="checkbox"/>	DATA_0002	DATA_0002	/dev/xvde	0	NORMAL	ONLINE	2.00	0.46	22.80
<input type="checkbox"/>	DATA_0003	DATA_0003	/dev/xvdf	0	NORMAL	ONLINE	2.00	0.46	22.75



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

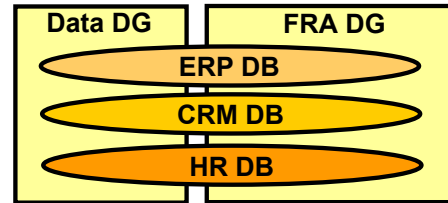
Tuning ASM

- ASM divides files into chunks called data extents made of allocation units (AUs) and spreads the AUs for each file evenly across all the disks. These data extents are variable in size within the same disk group and file.
- When your storage capacity changes, ASM does not re-stripe all of the data, but moves only the amount of data needed to evenly redistribute the files and maintain a balanced I/O load across the disks. The result is that all disks have the about the same amount of data. Rebalancing is done while the database is active and you can adjust the speed of a rebalance operation to increase its speed or to lower the impact on the I/O subsystem with the `POWER` clause. The variable size extents rarely cause fragmentation and a rebalance operation also defragments.
- ASM includes mirroring protection without the need to purchase a third-party logical volume manager. One unique advantage of ASM is that the mirroring is applied on a file basis, rather than on a volume basis. Therefore, the same disk group can contain a combination of files protected by mirroring, or not protected at all. Redundancy has an overhead. Writes for `NORMAL` redundancy are doubled, and writes for `HIGH` redundancy are tripled. Therefore, only mirror what is required.
- ASM uses direct I/O to access the disk storage, ignoring the `FILESYSTEMIO_OPTIONS` initialization parameter. The `DISK_ASYNCH_IO` parameter controls the ASM behavior. `DISK_ASYNCH_IO` should be set to `TRUE` on the database instance using ASM storage.

How Many Disk Groups per Database

- Two disk groups are recommended:

- Leverage maximum of LUNs
- Backup for each other
- Lower performance may be used for FRA (or inner tracks)



- Exceptions:

- Additional disk groups for different capacity or performance characteristics
- Different ILM storage tiers

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

How Many Disk Groups per Database

- Most of the time two disk groups are enough in ASM to share the storage between multiple databases.
- You can create an ASM disk group of LUNs on the outer edge of your disks, which gives you the best performance for active database files.
- Using a second disk group allows you to have a backup of your data by using it as your common flash recovery area (FRA). You can create a disk group of LUNs on the inner edge of your disks with lower performance for the FRA.
- The two noticeable exceptions to this rule are whenever you are using disks with different capacity or performance characteristics, or when you want to archive your data on lower-end disks for Information Lifecycle Management (ILM) purposes.

Note: Disk partitioning utilities start with cylinder 1, which is the outermost cylinder of the disk, and continue to cylinder n , which is the innermost. The first half of the disk has the highest performance, the inner half the lowest. This is due to how the sectors are arranged on the disk platter. The outer cylinder can have a maximum transfer rate that is twice the rate of the inner cylinder.

Which RAID Configuration for Best Availability?

- A. ASM mirroring
- B. Hardware RAID 1 (mirroring)
- C. Hardware RAID 5 (Parity Protection)
- ~~D. Both ASM mirroring and hardware RAID~~

Answer: **Depends on business requirement and budget (cost, availability, performance, and utilization)**

ASM leverages hardware RAID.

ORACLE

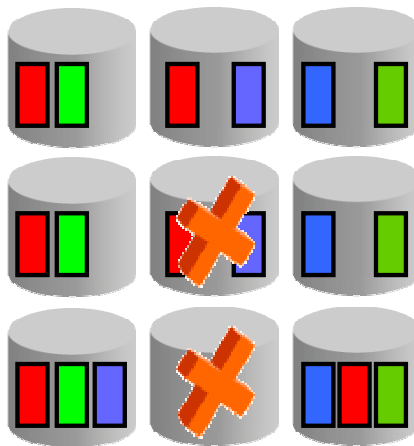
Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Which RAID Configuration for Best Availability?

- To favor availability, you have multiple choices as shown in the slide.
- You could just use ASM mirroring capabilities, or hardware RAID 1 (which is a hardware mirroring technique), or hardware RAID 5. Another possibility, although not recommended, is to use both ASM mirroring and hardware mirroring. Oracle recommends that you use external redundancy disk groups when using hardware mirroring techniques to avoid an unnecessary overhead.
- Choose between A, B, and C, depending on your business requirements and budget.
- ASM mirroring works with the standard disks without specialized hardware.
- RAID 1 has the best performance but requires twice the storage capacity.
- RAID 5 is a much more economical solution, but with a performance penalty for write-intensive workloads.

ASM Mirroring Guidelines

- Best choice for low-cost storage
- Enables Extended Clustering solutions
- No hardware mirroring



ORACLE

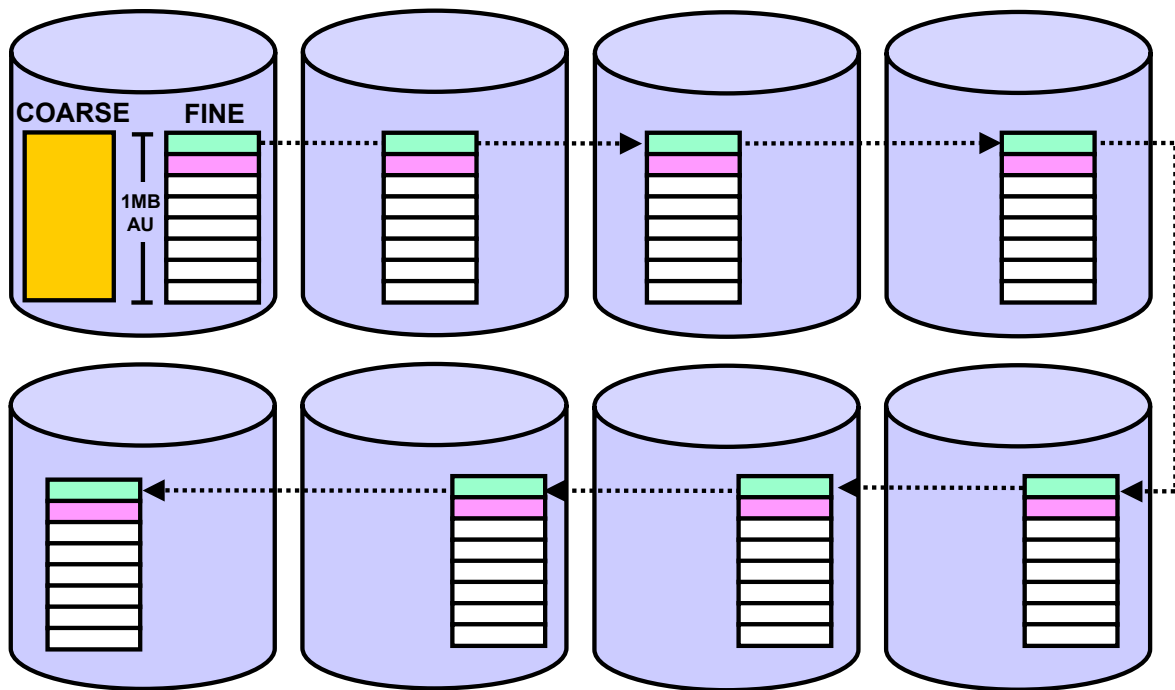
Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Mirroring Guidelines

Leverage the storage array hardware RAID 1 mirroring protection when possible to offload the mirroring overhead from the server. Use ASM mirroring in the absence of a hardware RAID capability.

Because the storage cost can grow very rapidly whenever you want to achieve extended clustering solutions, ASM mirroring should be used as an alternative to hardware mirroring for low-cost storage solutions.

ASM Striping Granularity



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Striping Granularity

- The difference between ASM COARSE and FINE striping is illustrated in this slide.
- For a file with a coarse stripe, striping is done at the allocation unit (AU) size across all the disks in the disk group. This does not account for any mirroring that would require additional blocks. Data files have the coarse-striping property. A diskgroup with a 1MB AU would be striped as shown above.
- When a block is allocated for a file with a fine stripe, a set of allocation units are set aside for this purpose in a round-robin allocation across all the available disks. The write is then made to these eight allocation units in 128 KB chunks with each successive write going to the next allocation unit in the set. Therefore, with a 1 MB AU, 8 MB can be written before the next allocation is made. Redo log files have the fine-striping property.
- Basically, one redo log write of 1 MB is striped in 128 KB chunks across eight allocation units. This is done automatically by the system when the redo log file is created.
- Coarse striping is done with allocation unit size. The AU is set at disk group creation and can be 1, 2, 4, 8, 16, 32, or 64MB in size.

What Type of Striping Works Best?

- A. ASM striping only (no RAID 0)
- B. RAID 0 and ASM striping
- ~~C. Use LVM~~
- ~~D. No striping~~

Answer: **A and B**

ASM and RAID striping are complementary.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

What Type of Striping Works Best?

- As shown in the slide, you can use ASM striping only, or you can use ASM striping in combination with RAID 0.
- With RAID 0, multiple disks are configured together as a set or a bank. Data from any one data file is spread, or striped, across all the disks in the bank.
- Combining both ASM striping and RAID striping is called stripe-on-stripe. This combination offers good performance.
- It is recommended that you use some type of striping: either ASM or RAID.
- Because ASM provides striping, there is no longer a need to use a logical volume manager for your database files.

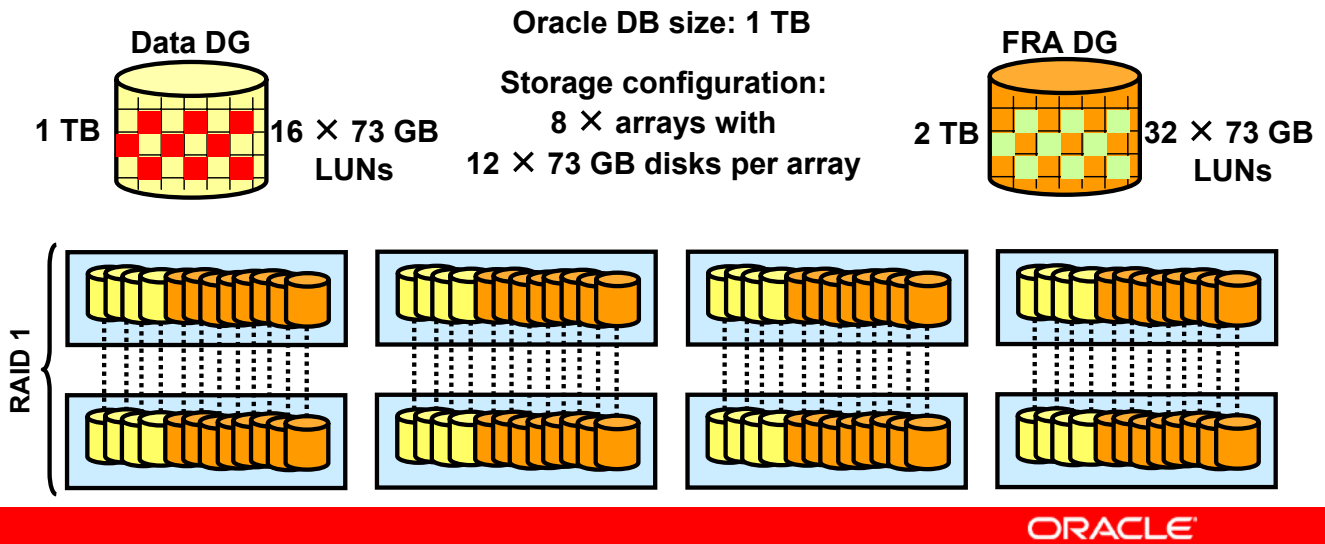
ASM Striping Only

Advantages

- Drives evenly distributed for Data & FRA
- Higher bandwidth
- Allows small incremental growth (73 GB)
- No drive contention

Disadvantages

- Not well balanced across *all* disks
- LUN size limited to disk size



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Striping Only

- In the case shown in the slide, you want to store a one-terabyte database with a corresponding two-terabytes flash recovery area. You have eight arrays of twelve disks, with each disk being 73 GB. You use RAID 1 to mirror each disk creating 48 logical units (LUNs) of 73 GB each. ASM mirroring and hardware RAID 0 are not used.
- Using ASM, the Data disk group is allocated 16 LUNs of 73 GB each. The Flash Recovery Area disk group is assigned 32 LUNs of 73 GB each.
- This configuration allows you to evenly distribute disks for your data and backups, achieving good performance and allowing you to manage your storage in small incremental chunks.
- However, using a limited number of disks in your Data disk group does not balance your data well across all your disks, and you still have many LUNs to manage at the storage level.
- If you want to place all the data on fast regions of the disk, partition the disks and create LUNs over the fast partitions and slow partitions. This is possible; it just multiplies the number of LUNs to manage by the number of partitions on each disk.

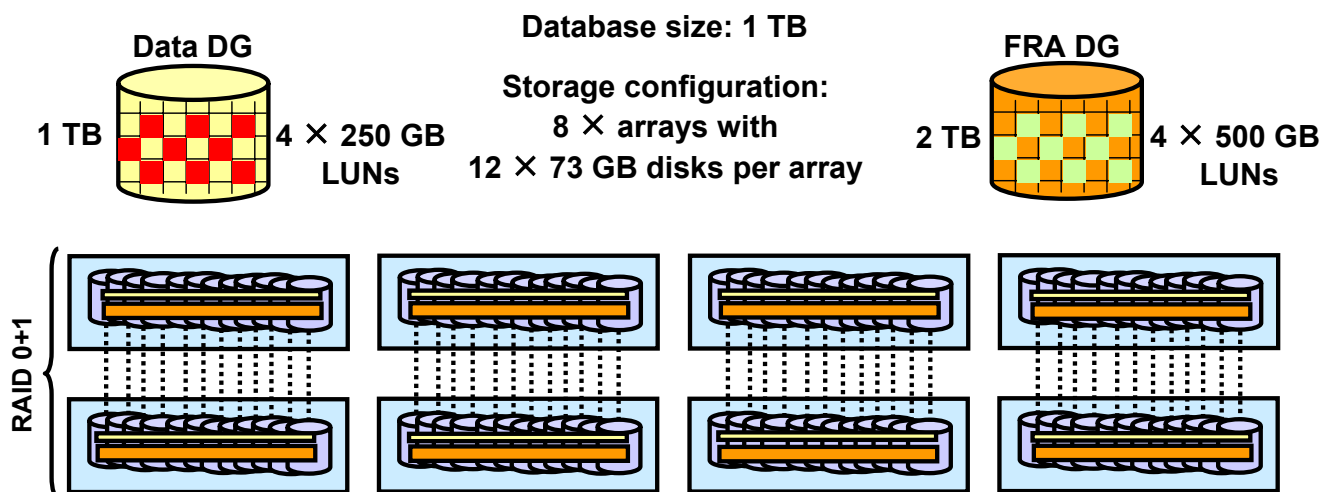
Hardware RAID Striped LUNs

Advantages

- Fastest region for Data DG
- Balanced data distribution
- Fewer LUNs to manage while max spindles

Disadvantages

- Large incremental growth
- Data and FRA “contention”



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Hardware RAID Striped LUNs

In the case shown in the slide, you have a one-terabyte database with a corresponding two-terabytes flash recovery area. You have eight arrays of twelve disks, with each disk being 73 GB. You use RAID 0+1, a combination of hardware striping and mirroring, to mirror and stripe each partition, creating four logical units of 876 GB each. These can be defined as an ASM single disk group. ASM mirroring is not used.

Here, you define bigger LUNs, not restricted to the size of one of your disks. ASM striping spreads the allocation units across the LUNs. The RAID-level striping further spreads I/Os across the physical disks. By doing this, you achieve a better data distribution across all your disks, and you end up managing significantly fewer LUNs.

However, you must manipulate your storage in much larger chunks than in the previous configuration. To force the use of the faster region for data, an additional LUN would have to be created across a set of partitions of the disks resulting in two LUNs.

Note: The hardware stripe size you choose is very important. You want 1 MB alignment as much as possible, to match the ASM AUs. Selecting “power of two” stripe sizes (64 KB, 128 KB, or 256 KB) is better than selecting other numbers. The choices offered by storage vendors for stripe size and partitioning depend on their storage array technology. These can create unnecessary I/O bottlenecks if not carefully considered.

ASM Guidelines

- Use external RAID protection when possible.
- Create logical units (LUNs) using:
 - Outside half of disk drives for highest performance
 - Small disk, high RPM (for example, 73 GB/15k RPM)
- Use LUNs with the same performance characteristics.
- Use LUNs with the same capacity.
- Maximize the number of spindles in your disk group.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Guidelines

Use Oracle Database ASM for volume and file management to equalize the workload across disks of logical units and eliminate hot spots.

Follow these simple guidelines and best practices when configuring ASM disk groups:

- Use external RAID protection when possible.
- Create LUNs using:
 - Outside half of disk drives for highest performance
 - Small disks with high RPM (for example, 73 GB with 15k RPM) for best performance. Spindle (platter) speed directly impacts both positioning time and data transfer. Faster spindle speed drives have improved performance regardless of whether they are used for many small, random accesses, or for streaming large contiguous blocks from the disk.
- Maximize the number of spindles in your disk group for greater bandwidth.
- Disks or LUNs assigned to ASM disk groups should have the same storage performance and availability characteristics. Configuring mixed speed drives will default to the lowest common denominator.
- ASM data distribution policy is capacity based. Therefore, LUNs provided to ASM should have the same capacity for each disk group to avoid imbalance and hot spots.

ASM Instance Initialization Parameters

- ASM instances have static memory needs.
- Use Automatic Memory Management.
- Using default SGA sizing parameters should be enough for most configurations:

```
INSTANCE_TYPE = ASM
DB_UNIQUE_NAME = +ASM
ASM_POWER_LIMIT = 1
ASM_DISKSTRING = '/dev/rdisk/*s2', '/dev/rdisk/c1*'
ASM_DISKGROUPS = dgroupA, dgroupB
MEMORY_TARGET = 256M /* Default value */
```

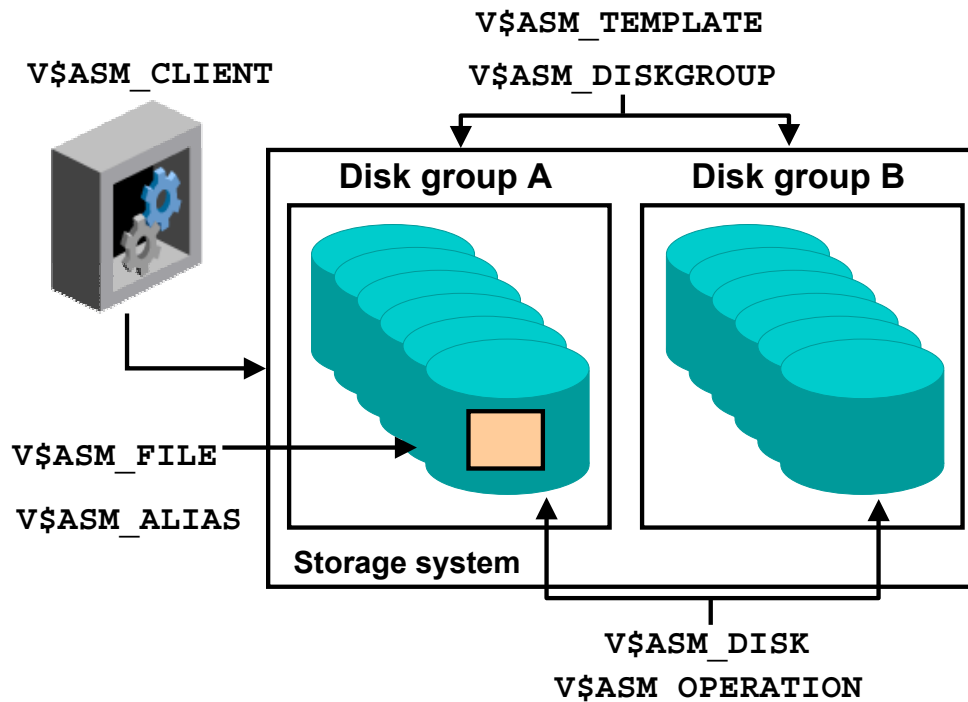
ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Instance Initialization Parameters

- `INSTANCE_TYPE` must be set to `ASM` for ASM instances.
- `DB_UNIQUE_NAME` specifies the service name. Use default value of `+ASM`.
- `ASM_POWER_LIMIT` controls the speed of a rebalance operation. Possible values range from the default low value of 1 through a high value of 11. The rebalance speed may be set per operation by the (`POWER`) setting of the rebalance command.
- `ASM_DISKSTRING` is an OS-dependent value that limits the set of disks considered for discovery. If not specified, it is assumed to be `NULL` and ASM disk discovery finds all disks to which an ASM instance has read and write access.
- `ASM_DISKGROUPS` is the list of names of disk groups to be mounted by an ASM instance at startup. ASM automatically adds a disk group to this parameter when a disk group is successfully mounted, and automatically removes the disk group when it is dismounted except for dismounts at instance shutdown.
- `PROCESSES` is the maximum number of processes started by the instance. The default value is sufficient for most instances.
- `MEMORY_TARGET` is the memory size allocated to the SGA and PGA of the ASM instance. 256M is the default value, and if `MEMORY_TARGET` is set below 100M, `MEMORY_TARGET` will be reset to 256M.

Dynamic Performance Views



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Dynamic Performance Views

The dynamic performance views that display the ASM information behave differently depending on the type of instance they are accessed from. The following gives the view name and the behavior on an ASM instance and a database instance.

V\$ASM_CLIENT

- ASM—One row for every database instance using a disk group in the ASM instance
- Database—One row for each disk group with the database and ASM instance name

V\$ASM_DISKGROUP

- ASM—One row for every discovered disk group
- Database—A row for all disk groups mounted or dismounted

V\$ASM_TEMPLATE

- ASM—One row for every template present in every mounted disk group
- Database—Rows for all templates in mounted disk groups

V\$ASM_DISK

- ASM—One row for every discovered disk, including disks that are not part of any disk group
- Database—Rows for disks in the disk groups in use by the database instance

Dynamic Performance Views (continued)

V\$ASM_OPERATION

- ASM—One row for every active ASM long-running operation executing in the ASM instance
- Database—Contains no rows

V\$ASM_FILE

- ASM—One row for every ASM file in every disk group mounted by the ASM instance
- Database—Contains no rows

V\$ASM_ALIAS

- ASM—One row for every alias present in every disk group mounted by the ASM instance
- Database—Contains no rows

The V\$ASM_DISK_STAT and V\$ASM_DISKGROUP_STAT views return the same information as their counterparts without _STAT. They return cached results, rather than doing a new disk discovery. Use V\$ASM_DISK_STAT instead of V\$ASM_DISK, and V\$ASM_DISKGROUP_STAT instead of V\$ASM_DISKGROUP, to query performance statistics. The V\$ASM_DISK and V\$ASM_DISKGROUP views require a discovery of disks during every execution, which makes it a relatively expensive operation, and repeats of this query at high frequency are not advisable. For example:

```
SELECT path, reads, writes, read_time, write_time,
       read_time/decode(reads,0,1,reads) "avgrdtime",
       write_time/decode(writes,0,1,writes) "avgwrtime"
FROM V$ASM_DISK_STAT;
```

Note: For more information about ASM data dictionary views, refer to the *Oracle Database Reference* guide.

Monitoring Long-Running Operations by Using V\$ASM_OPERATION

Column	Description
GROUP_NUMBER	Disk group
OPERATION	Type of operation: REBAL
STATE	State of operation: WAIT or RUN
POWER	Power requested for this operation
ACTUAL	Power allocated to this operation
SOFAR	Number of allocation units moved so far
EST_WORK	Estimated number of remaining allocation units
EST_RATE	Estimated number of allocation units moved per minute
EST_MINUTES	Estimated amount of time (in minutes) for operation termination

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Monitoring Long-Running Operations by Using V\$ASM_OPERATION

The ALTER DISKGROUP DROP, RESIZE, and REBALANCE commands return before the operation is complete. To monitor the progress of these long-running operations, you can query the V\$ASM_OPERATION fixed view. This view is described in the table in this slide.

Note: In addition to the REBAL operation, you may see other operations relating to recovery and ASM metadata.

ASM Instance Performance Diagnostics

```
SELECT  event, total_waits t_wait,
        total_timeouts t_timeout,
        time_waited t_waittm,
        average_wait a_waittm, wait_class
FROM    V$SYSTEM_EVENT
WHERE    wait_class <> 'Idle' and time_waited > 0
ORDER BY 4 DESC;
```

EVENT	WAIT	TOUT	WAITT	AVG	CLASS
ASM mount : wait for heartbeat	1	1	439	438.85	Admin...
kfk: async disk IO	578	0	377	.65	SystI/O
log write(odd)	7	3	296	42.33	Other
rdbms ipc reply	37	1	259	7.01	Other
log write(even)	8	2	197	24.58	Other
SQL*Net message to client	139249	0	103	0	Network
os thread startup	9	0	79	8.77	Conc...
buffer write wait	1	0	60	60.31	Other
DBFG waiting for reply	16	0	1	.04	Other

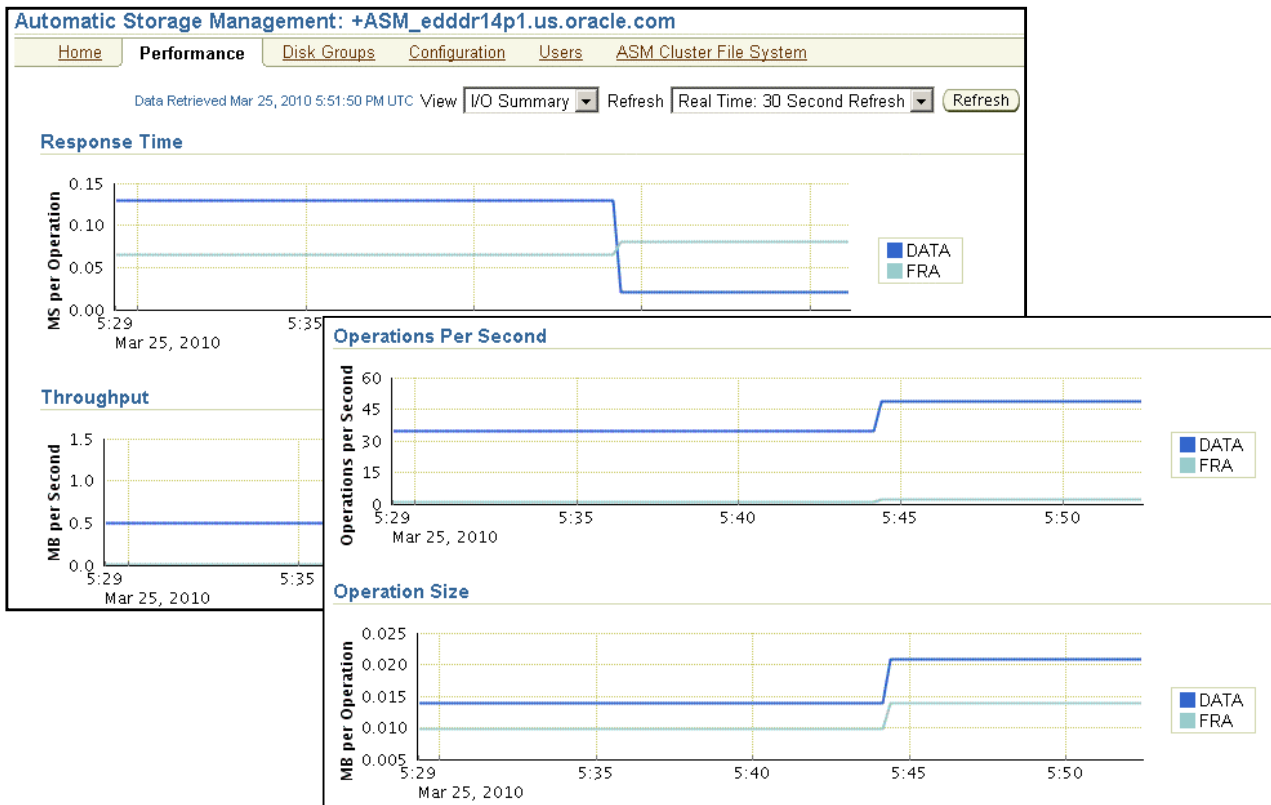
ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Instance Performance Diagnostics

The ASM instance architecture does not allow the creation of tables, so it is not possible to generate AWR or Statspack reports. However, ASM instances support querying all the dynamic performance (V\$) views, so you need to base performance diagnostics on SQL scripts. One possible example is illustrated in the slide.

ASM Performance Page



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Performance Page

The Performance tab of the Automatic Storage Management page shows the I/O response time and throughput for each disk group. You can further drill down to view disk-level performance metrics.

Database Instance Parameter Changes

- Add the following to SHARED_POOL_SIZE:

```
(DB_SPACE/100)*#_External_Redundancy +2 OR  
(DB_SPACE/50)*#_Normal_Redundancy +4 OR  
(DB_SPACE/33)*#_High_Redundancy +6
```

```
SELECT d+l+t DB_SPACE  
FROM (SELECT SUM(bytes)/(1024*1024*1024) d  
      FROM V$DATAFILE),  
      (SELECT SUM(bytes)/(1024*1024*1024) l  
      FROM V$LOGFILE a, V$LOG b WHERE a.group#=b.group#),  
      (SELECT SUM(bytes)/(1024*1024*1024) t  
      FROM V$TEMPFILE WHERE status='ONLINE');
```

- Add at least 16 to PROCESSES.
- Add 600 KB to LARGE_POOL_SIZE.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Database Instance Parameter Changes

If the database instance is not using Automatic Shared Memory Management, the SGA parameters for a database instance need slight modification to support AU maps and other ASM information. The following are guidelines for SGA sizing on the database instance:

- Additional memory is required to store AU maps in the shared pool. Use the result of the query shown in the slide to obtain the current database storage size (DB_SPACE) that either is already on ASM or will be stored in ASM. Then, determine the redundancy type that is used (or will be used), and add to the shared pool size one of the following values:
 - For disk groups using external redundancy: Every 100 GB of space needs 1 MB of extra shared pool plus a fixed amount of 2 MB of shared pool.
 - For disk groups using normal redundancy: Every 50 GB of space needs 1 MB of extra shared pool plus a fixed amount of 4 MB of shared pool.
 - For disk groups using high redundancy: Every 33 GB of space needs 1 MB of extra shared pool plus a fixed amount of 6 MB of shared pool.
- Add at least 16 to the value of the PROCESSES initialization parameter.
- Add 600 KB to the LARGE_POOL_SIZE parameter

ASM Scalability

ASM imposes the following limits:

- 63 disk groups
- 10,000 ASM disks
- 4 petabytes (PB) per ASM disk
- 40 exabytes (EB) of storage
- 1 million files per disk group
- Maximum file size:
 - External redundancy: 140 PB
 - Normal redundancy: 42 PB
 - High redundancy: 15 PB

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

ASM Scalability

ASM imposes the following limits:

- 63 disk groups in a storage system
- 10,000 ASM disks in a storage system
- 4 petabyte maximum storage for each ASM disk
- 40 exabyte maximum storage for each storage system
- 1 million files for each disk group
- Maximum file sizes depending on the redundancy type of the disk groups used:
140 PB for external redundancy, 42 PB for normal redundancy, and 15 PB for high redundancy

Oracle Database supports data file sizes up to 128 TB. ASM supports file sizes greater than 128 TB in any redundancy mode. This provides near unlimited capacity for future growth.

Quiz

Automatic Storage Management (ASM) is a replacement for a logical volume manager, which enables you to administer a group of disks as a logical unit. ASM uses the SAME methodology.

- a. True
- b. False

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Answer: a

Summary

In this lesson, you should have learned how to:

- Diagnose database I/O issues
- Describe the Stripe and Mirror Everything (SAME) concept
- Explain the benefits of asynchronous I/O
- Choose appropriate I/O solutions
- Tune I/O using Automatic Storage Management (ASM)

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

20

Performance Tuning Summary

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Objectives

After completing this lesson, you should be able to do the following:

- List best practices identified throughout the course
- Summarize the performance tuning methodology

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Necessary Initialization Parameters with Little Performance Impact

Parameter	Description
DB_NAME	Name of the database. This should match the ORACLE_SID environment variable.
DB_DOMAIN	Location of the database in Internet dot notation
OPEN_CURSORS	Limit on the maximum number of cursors for each session. The setting is application-dependent; 500 is recommended.
CONTROL_FILES	Set to contain at least two files on different disk drives to prevent failures from control file loss
DB_FILES	Set to the maximum number of files that can be assigned to the database
STATISTICS_LEVEL	Set to TYPICAL

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Necessary Initialization Parameters with Little Performance Impact

A running Oracle instance is configured using initialization parameters, which are set in the initialization parameter file. These parameters influence the behavior of the running instance, including influencing performance. In general, a very simple initialization file with a few relevant settings covers most situations. The initialization file should not be the first place you expect to do performance tuning.

The table in the slide describes the parameters necessary in a minimal initialization parameter file. Although these parameters are necessary, they have little performance impact.

STATISTICS_LEVEL can have a large impact on the performance of your database. Timed statistics are turned off when STATISTICS_LEVEL is set to BASIC, removing the effective information from the Statspack and AWR reports. STATISTICS_LEVEL set to TYPICAL has a small performance overhead. The entire AWR suite of tools is supposed to have no more than a 5% performance overhead.

Important Initialization Parameters with Performance Impact

Parameter	Description
COMPATIBLE	To take advantage of the latest improvements of a new release
DB_BLOCK_SIZE	8192 for OLTP and higher for DSS
MEMORY_TARGET	Automatically sized memory components
SGA_TARGET	Automatic SGA component management
PGA_AGGREGATE_TARGET	Automatic PGA management
PROCESSES	Maximum number of processes that can be started by that instance
SESSIONS	To be used essentially with shared server
UNDO_MANAGEMENT	AUTO mode recommended
UNDO_TABLESPACE	Undo tablespace to be used by instance

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Important Initialization Parameters with Performance Impact

The table in the slide includes the most important parameters to set with performance implications:

- **COMPATIBLE:** Specifies the release with which the Oracle server must maintain compatibility. It lets you take advantage of the maintenance improvements of a new release immediately in your production systems without testing the new functionality in your environment. If your application is designed for a specific release, and you are actually installing a later release, then you might want to set this parameter to the version of the previous release. This parameter does not guarantee that the behavior of the new release will be identical to the release specified by COMPATIBLE. COMPATIBLE allows you to use a new release, while at the same time guaranteeing backward compatibility with an earlier release. This is helpful if it becomes necessary to revert to the earlier release. Some features of the release may be restricted.
- **DB_BLOCK_SIZE:** Sets the size of the database blocks stored in the database files and cached in the SGA. The range of values depends on the operating system, but it is typically 8192 for transaction processing systems and higher values for database warehouse systems.

Important Initialization Parameters with Performance Impact (continued)

- **MEMORY_TARGET**: Specifies the total memory allocated to the instance, SGA and PGA.
- **SGA_TARGET**: Specifies the total size of all automatically tuned SGA components
- **PGA_AGGREGATE_TARGET**: Specifies the target aggregate PGA memory available to all server processes attached to the instance
- **PROCESSES**: Sets the maximum number of processes that can be started by the instance. This is the most important primary parameter to set, because many other parameter values are derived from this.
- **SESSIONS**: This is set by default from the value of **PROCESSES**. However, if you are using the shared server, the derived value is likely to be insufficient.

Sizing Memory Initially

As an initial guess for memory allocation:

- Leave 20% of available memory to other applications.
- Give 80% of memory to the Oracle instance.

```
MEMORY_MAX_TARGET=(total_mem*80%)
```

- By default the memory distribution starts at 60% SGA and 40% PGA.
- Memory distribution will adjust to workload.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

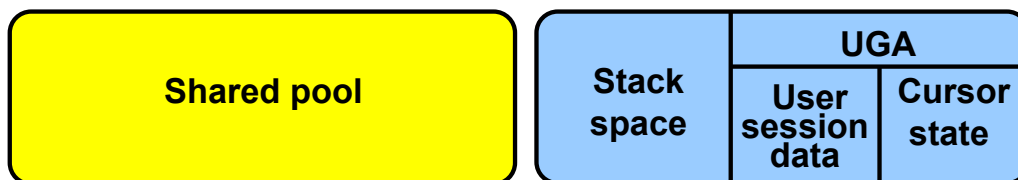
Sizing Memory Initially

Oracle strongly recommends that you use Automatic Memory Management. AMM allows you to set the total amount of memory to be used by the instance. Especially for a instance where there is no history, this method allows the instance to adjust the memory distribution to the workload. You may also set `MEMORY_MAX_TARGET` parameter to some value higher than `MEMORY_TARGET`. This will allow you to adjust the `MEMORY_TARGET` value without having to restart the instance.

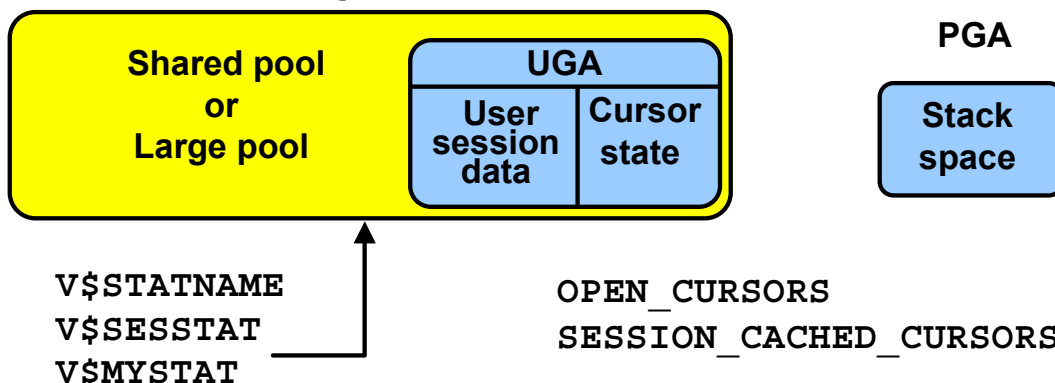
Example: Your Oracle instance will run on a system with 4 GB of physical memory. Part of that memory will be used for the operating system and other applications running on the same hardware system. You decide to dedicate only 80% (3.2 GB) of the available memory to the Oracle instance. OLTP and DSS systems have different patterns of memory usage. For OLTP systems, the PGA memory typically accounts for a small fraction of the total memory available (for example, 20% of the instance memory), leaving 80% for the SGA. For DSS systems running large, memory-intensive queries, PGA memory can typically use up to 70% of the instance memory. With AMM the memory will initially be allocated 1280 MB for `PGA_AGGREGATE_TARGET` and 1920 MB for `SGA_TARGET`, but as the instance monitors the memory usage the distribution will be adjusted. These new values will be recorded in the SPFILE, and applied at each instance startup.

UGA and Oracle Shared Server

Dedicated server configuration



Shared server configuration



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

The User Global Area

In a dedicated server configuration, the user global area (UGA) does not use any memory within the SGA. If you use Oracle Shared Server, then the UGA (which includes the user session and cursor state information) is stored in the SGA instead of in a private user memory. Sort areas and private SQL areas are included in the session information. This is because shared servers work on a per-call basis, so any server may need access to any user's information.

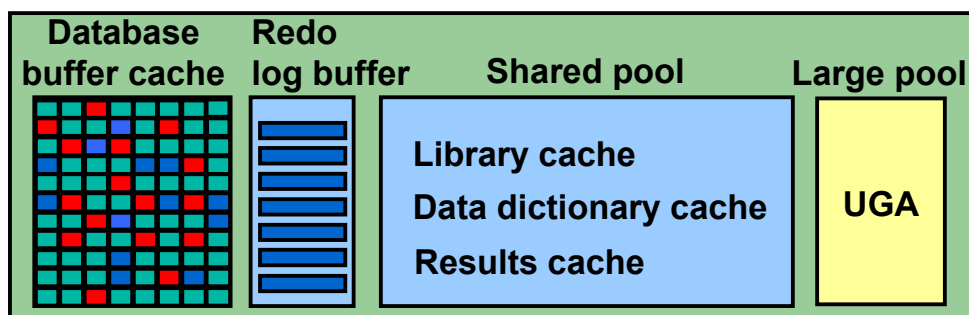
The memory requirement for the UGA with Oracle Shared Server is no larger than if you use dedicated servers, but it is allocated from the SGA instead of process memory. You may need to increase the SGA size, but your private user memory is lower.

In a shared server environment, it is recommended to configure the large pool. If a large pool has not been configured, the UGA is stored in the shared pool. Increase the `SHARED_POOL_SIZE` parameter to compensate. Allocate approximately 250 KB for each concurrent session as a starting point.

If you set the large pool, the UGA for shared servers is allocated from the large pool. By allocating session memory from the large pool, the shared pool is used primarily for caching shared SQL and avoids additional contention, with the dependent performance reduction, due to the reduced space used by the UGA.

Large Pool

- Can be configured as a separate memory area in the SGA, used for memory with:
 - I/O server processes: DBWR_IO_SLAVES
 - Backup and restore operations
 - Session memory for the shared servers
 - Parallel execution messaging
- Is used to avoid performance overhead caused by shrinking the shared SQL cache
- Is sized by the `LARGE_POOL_SIZE` parameter



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Large Pool

Oracle Support recommends that you configure the large pool anytime you configure shared servers, parallel query server, DBWR_IO_SLAVES, or ASM. The large pool participates in Automatic Shared Memory Management (ASMM). ASMM is the recommended method for managing the memory requirements of the main portions of the SGA.

Existence of the Large Pool

The memory of the large pool is part of the SGA, and adds to the amount of shared memory the Oracle server needs for an instance at startup.

Advantages of the Large Pool

The large pool is used to provide large allocations of session memory for:

- I/O server processes
- Backup and restore operations (RMAN Slave Buffers)
- Block change tracking buffers

The memory for backup and restore operations, and for I/O server processes is allocated in buffers of a few hundred kilobytes. The large pool is better able to satisfy such requests than the shared pool.

Tuning the Large Pool

- The large pool has one parameter, `LARGE_POOL_SIZE`.
- `V$SGASTAT` shows used and free memory.

```
SELECT * FROM V$SGASTAT
WHERE pool = 'large pool';
```

POOL	NAME	BYTES
large pool	free memory	2814112
large pool	session heap	1380192

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Tuning the Large Pool

The large pool is not governed by a least recently used algorithm. Memory allocations are released by the processes that allocated them or by PMON as part of a cleanup operation.

The only tuning that is available is to monitor the free memory and adjust the `LARGE_POOL_SIZE` parameter as needed.

The large pool area shown as `session heap` is the User Global Area used by Oracle Shared Servers.

Database High Availability: Best Practices

- Use `SPFILE`.
- Multiplex redo logs.
- Use resumable space allocation.
- Create at least two control files.
- Enable Flashback Database.
- Use Automatic Undo Management.
- Use Automatic Segment Space Management.
- Use locally managed tablespaces.
- Use locally managed temporary tablespaces.
- Enable `ARCHIVELOG` mode and use a flash recovery area.
- Set time long enough for `CONTROL_FILE_RECORD_KEEP_TIME`.
- Designate a default permanent tablespace other than `SYSTEM` and `SYSAUX`.
- Enable block checking.
- Use auto-tune checkpointing.
- Log checkpoints to the alert log.
- Use database resource manager.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Database High Availability: Best Practices

The list in the slide gives you a short summary of the recommended practices that apply to single-instance databases, RAC databases, and Data Guard standby databases.

These practices affect the performance, availability, and mean time to recover (MTTR) of your system. Some of these practices may reduce performance, but they are necessary to reduce or avoid outages. The minimal performance impact is outweighed by the reduced risk of corruption or the performance improvement for recovery.

Note: For more information about how to configure the features listed above, refer to the following documents:

- *Oracle Database Administrator's Guide*
- *Oracle Data Guard Concepts and Administration*
- *Oracle Database Net Services Administrator's Guide*

Undo Tablespace: Best Practices

- Use Automatic Undo Management.
- UNDO_RETENTION:
 - Is automatically adjusted when the UNDO tablespace has AUTOEXTEND enabled.
- Undo tablespace size:
 - Initial size: Small with AUTOEXTEND enabled
 - Steady state: Fix size using the Undo Advisor and add a 20% safe margin.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Undo Tablespace: Best Practices

Automatic Undo Management (AUM) is the preferred option because it greatly simplifies undo management and also provides an easy way to avoid errors such as ORA-1555.

UNDO_RETENTION: This parameter sets the minimum time that undo will be retained. The main purpose of this parameter is to support flashback requirements because the read-consistency aspect of undo retention no longer needs manual configuration. Undo retention is automatically tuned to the maximum allowable by the tablespace when the tablespace is a fixed size, even if it is greater than the longest running query in the system. This gives the best possible retention and in effect makes the setting of the UNDO_RETENTION parameter unnecessary. Therefore, the only consideration in undo management is now the proper sizing of the undo tablespace.

The most accurate and recommended way of sizing the undo tablespace is by using the Undo Advisor. However, a newly created database does not have the kind of historical workload information that is needed by the advisor to make sound recommendations. For such new systems, the best way to size the undo tablespace is to start with a small size, say 100 MB, but to set the AUTOEXTEND data file attribute to ON. This allows the tablespace to grow automatically as needed. After the database has reached a steady state, it is recommended to adjust the size of the tablespace. To determine the appropriate size of the undo tablespace, you should use the Undo Advisor. Add 20 percent to the size provided by the Undo Advisor for safety purposes and disable the AUTOEXTEND attribute.

Temporary Tablespace: Best Practices

Locally managed temporary tablespaces use a uniform extent. Extent size should be:

- 1 MB to 10 MB extent size:
 - For DSS, OLAP applications involving huge work areas
 - Where large temporary LOBs are predominant.
- 64 KB or multiples, less than 1 MB:
 - Small global temporary tables are predominant.
 - OLTP

Temporary tablespace group increases addressability from terabytes to petabytes.



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Temporary Tablespace: Best Practices

It is important to understand the types of workloads on the temporary tablespace:

- **Temporary LOBs:** Temporary LOBs are created explicitly by the user as well as implicitly by the database to store transient unstructured data. If the temporary space usage is dominated by large temporary LOBs, then larger extent sizes should be used. Oracle Corporation recommends using 1 MB to 10 MB as the extent size.
- **DSS:** The DSS workload is typically dominated by complex queries performing intensive sort and hash join operations. The general rule for good performance and efficient space usage is to set an extent size of 1 MB.
- **OLTP operations** needing temporary space tend to be small. Extent sizes that are small multiples of 64 KB up to 1 MB give better performance and use less space.
- **Global temporary tables:** These are transient tables created by the user as well as the system in the temporary tablespace. Each global temporary table requires at least one extent to be per instantiation. If the volume of data loaded into the temporary tables is pretty small, choosing smaller multiples of 64 KB for the extent size should avoid wasting space.

Temporary Tablespace: Best Practices (continued)

- **General:** 1 MB extent size is a good initial choice until workload patterns are established

Temporary tablespace groups: These are several temporary tablespaces assigned to a temporary tablespace group. Different sort operations can be assigned to different temporary tablespaces in the same group. Each parallel query server assigned to a parallel operation can use a different temporary tablespace in the same temporary tablespace group. This increases the available temporary space and reduces contention on a single temporary tablespace. A single sort operation (non-parallel) uses only one temporary tablespace in a temporary tablespace group.

Note: The primary purpose of a temporary tablespace group is to increase the addressability of the temporary space. A single temporary tablespace can have a maximum of four billion blocks. This is 8 TB for a 2 KB block size and 64 TB for a 16 KB block size. With temporary tablespace groups, the addressability is increased to several petabytes.

General Tablespace: Best Practices

- Use locally managed tablespaces with auto-allocate extents policy.
- Use Automatic Segment Space Management (ASSM).
- Use online segment shrink to eliminate internal fragmentation.
- Periodically review the results of the Automatic Segment Advisor.
- Monitor tablespace space usage using server-generated alerts.
- Size of extents matter more than the number of extents in the segments.



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

General Tablespace: Best Practices

- Use locally managed tablespaces with auto-allocate to avoid fragmentation of tablespace.
Note: The largest auto-allocate extent size possible is 64 MB. This should not impact segments that are less than a terabyte in size.
- Use ASSM for better segment space management performance.
- The Automatic Segment Advisor examines segments residing in tablespaces with outstanding tablespace-full alerts and those with heavy user access for signs of fragmentation. If the segments are fragmented, a recommendation is made to rectify.
- Use server-generated alerts to monitor tablespaces for impending out-of-space conditions.
- Use properly sized extents, when assigning the extent value manually. If the extents are too small, then multiblock reads are prevented. If the extents are too large, space is wasted in the last extent. There are only a few cases where the number of extents should be a cause for concern:
 - Tablespace is dictionary-managed and DDL commands such TRUNCATE and DROP are common.
 - Parallel queries are common and extents are noncontiguous in file.

Internal Fragmentation Considerations

- Watch for:
 - Bad choices of PCTFREE and PCTUSED for heap segments
 - Bad choices of PCTVERSION and RETENTION for LOB segments
 - Low density of data in segment
 - Direct loads followed by deletes (no inserts)
 - Indexes on tables with random updates and deletes with no further inserts
- Remedy:
 - Online segment shrink
 - Online redefinition
 - MOVE operations

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Internal Fragmentation Considerations

Segment fragmentation is a serious issue for two main reasons:

- Poor space utilization, because the free space under the HWM of one segment cannot be reused by other segments without performing a segment shrink first
- Poor performance, certain data access patterns cause more I/O than is necessary to read the data. For example, a full table scan reads all the blocks below the HWM, whether or not there are rows in those blocks. In the case of LOBs, data in such segments is always accessed through an index. Therefore, there is no impact in access performance.

Internal fragmentation can impact heap, index, as well as LOB segments. The slide shows you the main reasons leading to internal fragmentation.

Fragmentation is an issue that database administrators (DBAs) have been battling for a long time. Use of locally managed tablespaces in conjunction with Automatic Segment Space Management (ASSM) offers the best way to minimize internal fragmentation.

However, after internal fragmentation has occurred, the traditional methods that you can use to solve the problem include the following:

- Alter table MOVE, CTAS (create table as select), Import/Export.
- Reorganize the object using online redefinition set of procedures.

These techniques are effective in removing fragmentation, but require additional disk space and involve considerable manual labor. The online segment shrink feature is more efficient.

Block Size: Advantages and Disadvantages

Block Size	Advantages	Disadvantages
Smaller	Small rows with lots of random access	Relatively high overhead (block header)
	Reduce block contention (OLTP)	For large rows (chaining)
Larger	Lower overhead: more room for data	Waste cache space with random access of small rows
	Good for sequential access	Index leaf block contention (OLTP)

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Block Size: Advantages and Disadvantages

A block size of 8 KB is optimal for most systems. However, OLTP systems occasionally use smaller block sizes and DSS systems occasionally use larger block sizes. For read-intensive operations, and regardless of the size of the data, the goal is to minimize the number of reads required to retrieve the desired data.

- If the rows are small and access is predominantly random, choose a smaller block size.
- If the rows are small and access is predominantly sequential, choose a larger block size.
- If the rows are small and access is both random and sequential, it might be effective to choose a larger block size.
- If the rows are large, such as rows containing large object (LOB) data, choose a larger block size.

Note: The use of multiple block sizes in a single database instance is not encouraged because of manageability issues.

Automatic Checkpoint Tuning

- There is no longer a continuous manual tuning effort.
- Automatic checkpoint tuning is the best-effort checkpointing, without much overhead.
- It reduces average recovery time by making use of unused bandwidth.
- Automatic checkpoint tuning is enabled when `FAST_START_MTTR_TARGET` is not explicitly set to zero.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Automatic Checkpoint Tuning

The automatic checkpoint tuning included with Oracle Database 11g greatly reduces the need to set any of the checkpoint parameters. These parameters can still be used to set an upper bound on recovery time. In earlier versions, if you had large log files and large buffer caches so that very few buffers were written for aging purposes, you would set parameters such as `FAST_START_MTTR_TARGET` to reduce crash recovery times.

By default, Oracle Database 11g supports automatic checkpoint tuning by making the best effort to write out dirty buffers without an adverse impact on the throughput. Thus, a reasonable crash recovery time can be achieved even if you do not set `FAST_START_MTTR_TARGET` or if you set it to a very large value.

You enable automatic checkpoint tuning by setting the `FAST_START_MTTR_TARGET` parameter to a nonzero value. You can also enable this feature by not setting the `FAST_START_MTTR_TARGET` parameter. However, you disable it by setting the `FAST_START_MTTR_TARGET` parameter to zero.

Sizing the Redo Log Buffer

The size of the redo log buffer is determined by the:

- LOG_BUFFER parameter
- Remaining space in the fixed area granule

Default value: Either 2 MB or 128 KB × the value of CPU_COUNT, whichever is greater

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Sizing the Redo Log Buffer

If transactions are long or numerous, then larger values of LOG_BUFFER reduce log file I/O. The smaller the buffer is, the more often it will get to be one-third full before a COMMIT operation clears the buffer.

The default value of LOG_BUFFER is 512 KB. But the actual allocation will include the remainder of the granule allocated to the fixed area of the SGA. The log buffer size may be different than the LOG_BUFFER parameter value. The V\$SGASTAT view displays the actual log buffer size.

Example:

```
SELECT 'v$parameter' "View name", name,
       to_number (value,'999999999') "Value"
FROM V$PARAMETER WHERE name = 'log_buffer'
UNION
SELECT 'v$sgastat' "View name", name, bytes
FROM V$SGASTAT WHERE name = 'log_buffer';
```

View name	NAME	Value
-----	-----	-----
v\$parameter	log_buffer	23576576
v\$sgastat	log_buffer	23945216

Sizing Redo Log Files

- Size of redo log files can influence performance.
- Larger redo log files provide better performance.
- Generally, redo log files should range between 100 MB and a few gigabytes.
- Switch redo log file at most once every twenty minutes.
- Use the Redo Logfile Size Advisor to correctly size your redo logs.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Sizing Redo Log Files

- The size of the redo log files can influence performance, because the behavior of the database writer and archiver processes depend on the redo log sizes. Generally, larger redo log files provide better performance. Undersized log files increase checkpoint activity and reduce performance.
- If the `FAST_START_MTTR_TARGET` parameter is set to limit the instance recovery time, the system automatically tries to checkpoint as frequently as necessary. Under this condition, the size of the log files should be large enough to avoid additional checkpointing. The optimal size can be obtained by querying the `OPTIMAL_LOGFILE_SIZE` column from the `V$INSTANCE_RECOVERY` view. You can also obtain sizing advice on the Redo Log Groups page of Oracle Enterprise Manager Database Control.
- It is not possible to provide a specific size recommendation for redo log files, but redo log files in the range of a hundred megabytes to a few gigabytes are considered reasonable. Size your online redo log files according to the amount of redo your system generates.
A rough guide is to switch logs at most once every twenty minutes.

Increasing the Performance of Archiving

- Share the archiving work during a temporary increase in workload:

```
ALTER SYSTEM ARCHIVE LOG ALL  
TO <log_archive_dest>
```

- Increase the number of archiver processes with LOG_ARCHIVE_MAX_PROCESSES.
- Multiplex the redo log files, and add more members.
- Change the number of archive destinations:
 - LOG_ARCHIVE_DEST_n

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Increasing the Performance of Archiving

Share the Archiving Workload Temporarily

If you anticipate a heavy workload for archiving, you can add another process to share the work by executing the `ALTER SYSTEM ARCHIVE LOG ALL TO 'directory_name'` command. This command causes the server process to behave as an archiver until all redo log files have been archived. These processes are not affected by LOG_ARCHIVE_MAX_PROCESSES.

Increase the Number of Archiver Processes

Occasionally, on busy databases, a single ARCn process cannot keep up with the volume of information written to the redo logs. With the Oracle database, you can define multiple archiver processes by using the LOG_ARCHIVE_MAX_PROCESSES parameter. This parameter sets the maximum number of ARCn processes that can be started.

The LGWR process starts a new ARCn process whenever the current number of ARCn processes is insufficient to handle the workload.

Increasing the Performance of Archiving (continued)

Multiplex Redo Log Files

If you are archiving, it is even more important to have more than two redo log groups with multiple members. When logging switches to another group, the DBW_n process must perform checkpointing as usual, and one file must be archived. You must allow time for both of these operations before the LGWR process needs to overwrite the file again by having sufficient groups. Take advantage of the ARC_n behavior of reading a block from each member of the redo log group to spread the I/O across multiple disks by adding more members to the redo log groups.

Decrease the Number of Archive Destinations

This is not often a possible choice. More archive destinations increase the overhead of archiving. An alternative to several archive destinations could be to use a fetch archive log server to make additional archive log file copies on another machine.

Automatic Statistics Gathering

- `STATISTICS_LEVEL = TYPICAL | ALL`
- Statistics are gathered by the Optimizer Statistics Gathering automatic maintenance task.
- This task implicitly determines the following:
 - Database objects with missing or stale statistics
 - Appropriate sampling percentage necessary to gather good statistics on those objects
 - Appropriate columns that require histograms and the size of those histograms
 - Degree of parallelism for statistics gathering
 - Prioritization of objects on which to collect statistics

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Automatic Statistics Gathering

- Optimizer statistics are automatically gathered by the Optimizer Statistics Gathering task part of the Automatic Maintenance tasks. This task gathers statistics on all objects in the database that have missing or stale statistics.
- This task is created automatically at database creation time and is managed by the Scheduler. This task gathers optimizer statistics by calling the `DBMS_STATS.GATHER_DATABASE_STATS_JOB_PROC` procedure. This procedure operates in a similar fashion to the `DBMS_STATS.GATHER_DATABASE_STATS` procedure by using the `GATHER AUTO` option. The primary difference is that the `GATHER_DATABASE_STATS_JOB_PROC` procedure prioritizes the database objects that require statistics, so that those objects that need updated statistics the most are processed first. This ensures that the most needed statistics are gathered before the maintenance window closes.
- You should make sure that automatic maintenance tasks are enabled, and that `STATISTICS_LEVEL` is set to at least `TYPICAL`.

Automatic Statistics Collection: Considerations

- You should still manually gather statistics in the following cases:
 - After bulk operations
 - When using external tables
 - To collect system statistics
 - To collect statistics on fixed objects
- Prevent automatic gathering for volatile tables:
 - Lock with statistics for representative values
 - Lock without statistics implies dynamic sampling.
- Set Optimizer Statistic Preferences for objects that need special handling.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Automatic Statistics Collection: Considerations

- For tables that are being bulk-loaded, the statistics-gathering procedures should be run immediately following the load process.
- Automatic statistics collection is not supported for external tables, system statistics, and statistics on fixed objects (such as the dynamic performance tables). These statistics should be gathered. External tables should have statistics gathered after a change in size. Statistics should be gathered on fixed objects after a change in workload.
- With the Oracle Database, you can lock statistics on a specified table using the `LOCK_TABLE_STATS` procedure of the `DBMS_STATS` package.
- You can lock statistics on a volatile table at a point when it is fully populated so that the table statistics are representative of the table population. This is good for interim tables that are volatile, but have a reasonable consistent size.
- You can lock statistics on a table without statistics to prevent automatic statistics collection so that you can use dynamic sampling on a volatile table with no statistics. This option works best on tables that have a widely ranging size.
- Optimizer statistics preferences can be set to direct the automatic task to use specific parameters when gathering statistics on specific objects or schemas.

Commonly Observed Wait Events

Wait Event	Area	Possible cause	Examine
buffer busy waits	Buffer cache, DBWR	Depends on buffer type. PK index and seq.	V\$SESSION_WAIT (block) while issue is occurring
free buffer waits	Buffer cache, DBWR, I/O	Slow DBWR	Write time using OS stats. Buffer cache stats
db file scattered read, db file sequential read	I/O, SQL Tuning	Poorly tuned SQL, Slow I/O system	V\$SQLAREA disk reads. V\$FILESTAT read time
Enqueue waits (enq:)	Locks	Depends on enq type	V\$ENQUEUE_STAT
Library cache waits	Latches	SQL parsing/sharing	V\$SQLAREA parse calls, child cursors
log buffer space	Log buffer I/O	Small buffer, slow I/O	V\$SYSSTAT redo buffer allocation retries, disk
Log file sync	Over-commit, I/O	Slow I/O, un-batched commits	commits + rollbacks from V\$SYSSTAT, Disks.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Commonly Observed Wait Events

This slide links wait events to possible causes and gives an overview of the diagnostics that could be most useful to review next.

You may also want to review the following Oracle MetaLink notices on `buffer busy waits` (34405.1) and `free buffer waits` (62172.1):

- http://metalink.oracle.com/metalink/plsql/ml2_documents.showDocument?p_database_id=NOT&p_id=34405.1
- http://metalink.oracle.com/metalink/plsql/ml2_documents.showDocument?p_database_id=NOT&p_id=62172.1

Additional Statistics

Statistic name	Description	Recommended action
Redo Log Space Requests	How many times that a server process had to wait for space in the online redo log	Checkpoints, DBWR, or archiver activity should be tuned, larger log files
Consistent changes	How many rollbacks done for consistent read purposes	Use automatic undo, tune the workload
Consistent gets	How many blocks are read in Consistent Read mode	Use automatic undo, tune the workload
Table Fetch by Continued Row	Migrated or chained rows	Reorganize

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Additional Statistics

- Wait events and statistics generally should not drive tuning . The statistics in the slide can help identify areas that need tuning or confirm symptoms that appear in the top timed events.
- The redo log space requests statistic indicates how many times a server process had to wait for space in the online redo log, not for space in the redo log buffer. A significant value for this statistic and the wait events should be used as an indication that checkpoints, DBWR, or archiver activity should be tuned, not LGWR. Increasing the size of the redo log buffer does not help. Increase the log file size, reducing the amount of time spent switching log files, or increase the speed of the archiver processes.
- A large number of consistent gets and consistent changes can indicate that the workload has long-running queries and many short transactions against the same tables. Other scenarios that produce these statistics are the result of rollback segment contention. Consistent changes and consistent gets may contribute to "CPU used by this session" and "DB CPU time." They could also lead to cache buffers chains latch activity and even buffer busy waits. Some consistent gets also trigger physical reads, thus leading to some of the LRU chains activity.
- Table fetch by continued row will also increase consistent gets with all its side effects mentioned above. You may also see it reflected as db file sequential reads peppered between db file scattered reads as full table scans encounter chained rows.

Top 10 Mistakes Found in Customer Systems

1. Bad connection management
2. Bad use of cursors and shared pool
3. Bad SQL
4. Use of nonstandard initialization parameters
5. Getting database I/O wrong
6. Redo log setup problems
7. Serialization of data blocks in the buffer cache
8. Long full-table scans
9. High amount of recursive SQL
10. Deployment and migration errors



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Top 10 Mistakes Found in Customer Systems

This slide lists the most common mistakes found in Customer systems:

1. The application connects and disconnects for each database interaction. This problem is common with stateless middleware in application servers. It has more than two orders of magnitude impact on performance, and is not scalable.
2. Not using shared cursors results in repeated parses. If bind variables are not used, then there is hard parsing of all SQL statements. This has an order of magnitude impact in performance, and it is not scalable. Use cursors with bind variables that open the cursor and execute it many times. Be suspicious of applications generating dynamic SQL.
3. Bad SQL is SQL that uses more resources than appropriate for the application requirement. This can be a decision support system (DSS) query that runs for more than 24 hours or a query from an online application that takes more than a minute. Any SQL that consumes significant system resources should be investigated for potential improvement. ADDM identifies high-load SQL and the SQL Tuning Advisor can be used to provide recommendations for improvement.

Top 10 Mistakes Found in Customer Systems (continued)

4. These might have been implemented on the basis of poor advice or incorrect assumptions. Most systems give acceptable performance using only the set of basic parameters. In particular, parameters associated undocumented optimizer features can cause a great deal of problems that can require considerable investigation. Likewise, optimizer parameters set in the initialization parameter file can override proven optimal execution plans. For these reasons, schemas, schema statistics, and optimizer settings should be managed together as a group to ensure consistency of performance.
5. Many sites lay out their databases poorly over the available disks. Other sites specify the number of disks incorrectly, because they configure the I/O subsystem by disk space and not I/O bandwidth.
6. Many sites run with too few redo logs that are too small. Small redo logs cause system checkpoints to continuously put a high load on the buffer cache and I/O system. If there are too few redo logs, the archive cannot keep up, and the database waits for the archive process to catch up.
7. Serialization of data blocks in the buffer cache due to lack of free lists, free list groups, transaction slots (INITTRANS), or shortage of rollback segments. This is particularly common on insert-heavy applications, in applications that have raised the block size above 8 KB, or in applications with large numbers of active users and few rollback segments. Use Automatic Segment Space Management (ASSM) and Automatic Undo Management to solve this problem.
8. Long full-table scans for high-volume or interactive online operations could indicate poor transaction design, missing indexes, or poor SQL optimization. Long table scans, by nature, are I/O intensive and not scalable.
9. Large amounts of recursive SQL executed by SYS could indicate space management activities, such as extent allocations, taking place. This is not scalable and impacts user response time. Use locally managed tablespaces to reduce recursive SQL due to extent allocation. Recursive SQL executed under another user ID is probably SQL and PL/SQL, and this is not a problem.
10. In many cases, an application uses too many resources because the schema owning the tables has not been successfully migrated from the development environment or from an older implementation. Examples of this are missing indexes or incorrect statistics. These errors can lead to suboptimal execution plans and poor interactive user performance. When migrating applications of known performance, export the schema statistics to maintain plan stability by using the DBMS_STATS package. SQL Plan Baselines provide a method of migrating the execution plans and preventing plan regression.

Quiz

Automatic Statistics gathering collects optimizer statistics:

- a. During the Automatic Maintenance tasks
- b. On fixed objects
- c. On external tables
- d. On all tables, including the data dictionary

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Answer: a, d

Statistics are not gathered on fixed objects, such as dynamic performance views or external tables.

The statistics are gathered on all tables in priority order as the statistics become stale or if the table is missing statistics.

Summary

In this lesson, you should have learned how to:

- Create your initial database following the best practices identified throughout the course
- Summarize the performance tuning methodology

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

Appendix A

Practices and Solutions

Table of Contents

Practices for Lesson 1	4
Practices for Lesson 2	5
Practice 2-1: Viewing Diagnostic Information	6
Practice 2-2: Using the Alert Log	10
Practice 2-3: Viewing System Statistics and Wait Events	13
Practices for Lesson 3	18
Practice 3-1: Using Automatic Workload Repository	19
Practices for Lesson 4	29
Practice 4-1: Using Enterprise Manager to Identify OS Issues	30
Practices for Lesson 5	32
Practice 5-1: Using Metrics	33
Practices for Lesson 6	43
Practice 6-1: Using Baselines	44
Practices for Lesson 7	51
Practice 7-1: Using AWR-Based Tools	52
Practices for Lesson 8	63
Practice 8-1: Using Services with a Single-Instance Oracle Database	64
Practice 8-2: Tracing Services in a Single-Instance Oracle Environment	76
Practices for Lesson 9	83
Practice 9-1: Using AUTOTRACE and DBMS_XPLAN	84
Practice 9-2: Using the SQL TRACE Facility	89
Practices for Lesson 10	93
Practice 10-1: Capturing Extended Statistics	94
Practice 10-2: Stale Statistics	99
Practices for Lesson 11	104
Practice 11-1: Excess Blocks	105
Practice 11-2: Coalescing an Index	110
Practices for Lesson 12	117
Practice 12-1: Using the SQL Performance Analyzer	118
Practices for Lesson 13	128
Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor	129
Practice 13-2: SQL Plan Management (SPM)	136
Practices for Lesson 14	166
Practice 14-1: Using Database Replay	167
Practices for Lesson 15	176
Practice 15-1: Sizing the Shared Pool	177
Practice 15-2: Tuning a Hard-Parse Workload	189
Practice 15-3: Keeping Objects in the Shared Pool	198
Practices for Lesson 16	201
Practice 16-1: Using the DB Cache Advisor	202
Practice 16-2: Using the Keep Pool	208
Practices for Lesson 17	216
Practice 17-1: Tuning PGA_AGGREGATE_TARGET	217
Practices for Lesson 18	229

Practice 18-1: Enabling Automatic Shared Memory	230
Practice 18-2: Adjusting Memory as Workloads Change.....	237
Practices for Lesson 19	241
Practices for Lesson 20	242
Practices for Appendix B	243
Practice for Appendix B-1: Installing Statspack.....	244
Practice for Appendix B-2: Creating Snapshots	245
Practice for Appendix B-3: Generating Statspack Reports.....	248

Practices for Lesson 1

There are no practices for this lesson.

Practices for Lesson 2

In this practice, you look at various diagnostics to view the types of information available in the database.

Practice 2-1: Viewing Diagnostic Information

In this practice, you use the alert log to gather information about the performance of the instance while running a workload. You review the alert log for information that is relevant to performance tuning.

1. Set up the orcl database for a workload. Execute the prepare script:

```
./prepare 2 1
```

```
$ cd /home/oracle/workshops
$ ./prepare 2 1

Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area  836976640 bytes
Fixed Size                  1339740 bytes
Variable Size               805310116 bytes
Database Buffers            25165824 bytes
Redo Buffers                 5160960 bytes
Database mounted.
Database opened.
```

2. Run a workload by using the workload generator: `./workgen 2 1`.
Check whether the workload has started by using the `ps` command in the same terminal window. You see a set of processes as illustrated in the following screenshot. The order and number of each vary. In most cases, you see several `.sql` scripts, `sqlplus` and `sleep` processes. Not all these processes may be active at the time that the `ps` command is issued. This workload continues until it is stopped. Continue to the next step in the practice.

```
$ cd /home/oracle/workshops
$ ./workgen 2 1

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

$ ps
  PID TTY          TIME CMD
 25946 pts/2    00:00:00 bash
 30124 pts/2    00:00:00 insert_orders.s
 30125 pts/2    00:00:00 update_orders.s
 30126 pts/2    00:00:00 delete_orders.s
 30130 pts/2    00:00:00 insert_orders.s
 30131 pts/2    00:00:00 update_orders.s
 30134 pts/2    00:00:00 delete_orders.s
 30135 pts/2    00:00:00 insert_orders.s
```

Practice 2-1: Viewing Diagnostic Information (continued)

```
...  
$
```

3. Start the Enterprise Manager Database Control by starting the browser with the URL: <https://<hostname>:1158/em>.

The <hostname> is the name of your PC. You can find that name with the hostname command.

```
$ hostname
```

Log in as the SYS user with the password, oracle_4U, and connect as SYSDBA.

Note: Passwords are case-sensitive by default in Oracle Database 11g.

On the Database home page in Enterprise Manager, set View Data to Automatically (15 Sec). On the Database home page, what diagnostic information is available?

_____ (Host CPU and Active Sessions graphs are available after a few refreshes.)

4. Observe the characteristics of the workload on the server. This is shown in the Host CPU graph.
 - a) What is the load shown in the Host CPU graph? _____
(This number varies at each refresh.)
 - b) What percentage of the CPU is the orcl processes using? _____
5. The characteristics of the workload on the database instance are shown in the Active Sessions graph. When there are a large number of waiting sessions relative to the number of sessions that use the CPU, you must check other diagnostics.
 - a) What is the average number of sessions that use the CPU?

 - b) What is the average number of sessions that are waiting? _____, number waiting on I/O _____
6. How many CPU cores are in your system? _____

The Core Count below the Active Session Graph shows the number of CPU cores in your system. Another way to determine the number of CPUs in your system: Click the Database tab. Click the Load link in the Host CPU section. Click CPU Usage in the CPU Utilization section. The CPU Usage page shows one row for each CPU in your system.

7. Navigate to the Performance page. Click the Database tab, and then the Performance tab at the top of the page.
8. On the Performance page, load average (Average Runnable Processes) and average active sessions are graphed. The number of CPU cores can be shown in the graph by selecting the Show CPU Cores check box.

Practice 2-1: Viewing Diagnostic Information (continued)

- a) What is the value of the load average in relation to the maximum CPU?
 _____ (The load average is greater than the maximum CPU.) When the load average is greater than the maximum CPU, processes need to wait on the run queue.
9. In the Average Active Sessions graph, the various colors of the graph show various components of the waits and service times that the sessions are experiencing. Each component shows the number of active sessions in that state. For example, the component at the bottom of the graph shows the number of active sessions that use the CPU.
 - a) Which component of the graph, other than the CPU, consistently has the highest number of active sessions? _____ (This can vary. CPU Wait, Concurrency, Configuration, User I/O, and Commit are valid answers.)
10. Click Concurrency in the Average Active Sessions graph legend to view the waits that are included in this category. From a visual scan of the Active Sessions Waiting: Concurrency page, what are the highest waits?
 _____ (Buffer busy waits, Enq TX: index contention, Library Cache Mutex X, and Latch: Cache Buffer chains are all valid answers.) **Note:** In 11g Release 2, the internal algorithms have changed making buffer busy waits less likely.
11. Navigate to the Server page. Click the Database tab at the top of the page, and then click the Server tab.
12. On the Server page, click Automatic Workload Repository in the Statistics Management section.
13. On the Automatic Workload Repository page, click the number beside Snapshots in the Manage Snapshots and Baselines section.
14. On the Snapshots page, note the last snapshot number shown: _____ (This number may be over 100.)
 - a) Create a new snapshot. Click Create.
 - b) On the Confirmation page, click Yes.
15. Create an AWR report.
 - a) On the Snapshots page, select the snapshot that you recorded in step 14. Select View Report from the Actions drop-down menu, and then click Go beside the Action menu.
 - b) On the View Report page, select the next snapshot after the one chosen in step 14 a. Click OK.
16. When the report appears, scroll down to Top 5 Timed Foreground Wait Events.
 - a) What is the top timed event? _____
 - b) What is the second timed event? _____

Practice 2-1: Viewing Diagnostic Information (continued)

- c) Are any of the other timed events found in step 10 listed here? If so, list the timed events that are found.
-

17. Scroll down to the Wait Events Statistics section. Click the Time Model Statistics link.

- a) Which category collects the most %DB Time? _____
(sql execute elapsed time is expected but this item might be second.)
- b) Can you guess what performance issue is slowing this instance?
_____ (Poor SQL execution plans that produce index contention may be contributing to the sql elapsed time. In addition, log file waits may be holding up commit processing also contributing to SQL elapsed time)
- c) Stop the workload. In the terminal window, issue the following command:

```
$ rm $HOME/workshops/runload
```

NOTE: Most workload scripts in this course run until stopped. They repeat as long as the \$HOME/workshops/runload file exists.

18. Click the Database tab at the top of the page to return to the Database Instance home page.

Practice 2-2: Using the Alert Log

In this practice, you use the alert log to gather information about instance configuration and performance.

1. View the alert log through Enterprise Manager (EM) Database Control. Find the last 500 entries.
 - a) Invoke the browser (Mozilla) and enter the URL for Enterprise Manager Database Control for the `orcl` database. Use the following URL:
`https://<hostname>:1158/em`
 - b) Log in to EM for the `orcl` database.
 - c) From the Database home page, navigate to the Alert Log Contents page. Scroll to the bottom of the Database home page and click the Alert Log Contents link in the Related Links section.
 - d) In the View Entries box, select Last 500 and click Go.
 - e) Scroll through the entries. Note that the entries are in reverse chronological order. The most recent entries are at the top of the page.
2. Find the current Redo Log number and how many redo logs have been generated since yesterday. Each time the log file is full, a log switch occurs and a message is sent to the alert log giving the current log number.
 - a) Click Search. On the Search Alert Log Contents page, enter Most Recent 25 Hours. In the Message Text box, enter `. *Current log. *` and select the Regular Expression check box. Then click Go.
 Current sequence number (seq#): _____ (140, varies)
 Number of log switches: _____
 - b) It is recommended that log switches take place no more than every 20 minutes. Is the number of log switches excessive? _____
 (Yes. To reduce the number of log switches, increase the size of the redo log files.)
3. Find the name of the current SPFILE and the names of any other SPFILES that have been used in the last 10 days. Use the Search feature of the View Alert Log Contents page.
 - a) Use the Search feature again. Set the date range to Most Recent 10 Days, and then change the message text to `. *spfile. *`. Ensure that the Regular Expression check box is selected.

 (/u01/app/oracle/product/11.2.0/dbhome_1/dbs/spfileo
 rcl.ora)

Practice 2-2: Using the Alert Log (continued)

- b) Why is the SPFILE that is used to start the database important?

(Changes in the initialization parameters can have a large effect on database performance.)

- c) How else can initialization parameters change?
_____(by the use of ALTER SYSTEM commands)
- d) Have any initialization parameters been changed using the ALTER SYSTEM command? (Hint: Use . *ALTER SYSTEM. * in the message text field.)

Note: The prepare script issues ALTER SYSTEM commands.

4. From the command line, find and view the alert log file for the orcl database. Use the operating system login, and locate and view the log for the orcl database. The text version of the alert log in Oracle Database 11g is located in the directory specified by the Diag Trace entry in the V\$DIAG_INFO view.
- a) Open a terminal window. Start sqlplus and connect as sysdba. Execute the following command to find the Diag Trace directory. Exit sqlplus.

```
$ sqlplus / as sysdba

/* version and option banner omitted */

SQL> select * from V$DIAG_INFO
  2  where name = 'Diag Trace';

INST_ID NAME
-----
VALUE
-----
          1 Diag Trace
/u01/app/oracle/diag/rdbms/orcl/orcl/trace

SQL> exit
```

- b) In the terminal window, change directories to the value of the Diag Trace directory. Then, view the alert_orcl.log file. (gedit, a WYSIWYG editor, and text file viewers—less and more—are available.)
- c) Find the time of the last startup. Start at the end of the file. (Search from the most recent entries to the “Starting ORACLE instance” string.) The time stamp for this event is the time stamp immediately above the entry.

Time of the last startup _____

5. Determine how often the log files have switched in the last 10 minutes, while the workload is running.

Practice 2-2: Using the Alert Log (continued)

- a) The log file may have switched several times. The `tail` command lists the last 10 lines of a file. The `-n 100` option causes the `tail` command to display 100 lines; you may need to use `tail` with more lines to see a full 10 minutes.
6. While viewing the alert log, find the settings for the `db_cache_size`, `log_buffer`, `sga_max_size`, and `compatible` parameters.
Hint: All nondefault parameter values are listed in the alert log on every instance startup.
- a) Scroll through the alert log entries. Start from the most recent entries in the file to find the current parameter setting. Note that with this method, you see only those parameters that have been set to a nondefault value.
 (log_buffer is not visible because it has a default value.)
7. Use SQL*Plus to find the values for the `db_cache_size`, `log_buffer`, `sga_max_size`, and `compatible` parameters. Use the `SHOW PARAMETERS` command.

```
$ sqlplus / as sysdba

/* version and option banner omitted */

SQL> show parameter db_cache_size
```

NAME	TYPE	VALUE
db_cache_size	big integer	24M

```
SQL> show parameter log_buffer
```

NAME	TYPE	VALUE
log_buffer	integer	4980736

```
SQL> show parameter sga_max_size
```

NAME	TYPE	VALUE
sga_max_size	big integer	800M

```
SQL> show parameter COMPATIBLE
```

NAME	TYPE	VALUE
compatible	string	11.2.0.0.0

```
SQL> exit
```


Practice 2-3: Viewing System Statistics and Wait Events

In this practice, you use dynamic performance views to view tuning information.

1. Start the workload generator and continue. The workload continues until it is stopped.

```
$ cd /home/oracle/workshops
$ ./workgen 2 3

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
$
```

2. View the system statistics in SQL*Plus as the SYS user. Use the V\$SYSSTAT view and select NAME, CLASS, and VALUE. Order the output by CLASS. These statistics are cumulative from the time that the instance was started.
 - a) What can you deduce about the instance workload and performance from a listing of the system statistics?

Answer: There is very little that you can determine from a listing of the system statistics by itself. Two listings that are used to create a difference in values over a known period of time can yield rate and load information.

```
$ sqlplus / as sysdba
SQL> COLUMN name FORMAT A40 WORD_WRAPPED

SQL> SELECT name, class, value
2 FROM V$SYSSTAT
3 ORDER BY class;
```

3. View system wait events. Use the V\$SYSTEM_EVENT view to see wait event statistics. Find the total number of waits, the time waited, and the average time waited for the log file sync and enq: TX - index contention events, and the time waited relative to the other wait events. The number of rows returned will vary; events that have zero counts will not be listed.

Answer: The log file sync and enq: TX - index contention events have collected some waits and the average wait time is displayed. There are several wait events that are not important, such as rdbms ipc message and smon timer. These are internal waits for a command to be issued, or for a timer to expire to initiate a task. They do not affect performance. Your output from this query will vary from the example shown.

```
SQL> COLUMN event FORMAT A30 WORD_WRAPPED
SQL> COLUMN average_wait HEADING AVERAGE|WAIT
SQL> COLUMN TIME_WAITED HEADING TIME|WAITED
SQL> COLUMN TOTAL_WAITS HEADING TOTAL|WAITS
SQL> SET PAGESIZE 60
```

Practice 2-3: Viewing System Statistics and Wait Events (continued)

```
SQL> SELECT event, total_waits, time_waited, average_wait
2 FROM V$SYSTEM_EVENT
3 ORDER BY time_waited DESC;
```

EVENT	TOTAL WAITS	TIME WAITED	AVERAGE WAIT
-----	-----	-----	-----
rdbms ipc message	127563	1866712	14.63
SQL*Net message from client	66606	1351490	20.29
DIAG idle wait	3258	347413	106.63
jobq slave wait	3782	189368	50.07
pmon timer	999	173713	173.89
wait for unread message on broadcast channel	1735	172055	99.17
 Streams AQ: qmn slave idle wait	 63	 171667	 2724.88
 shared server idle wait	 58	 171575	 2958.19
Space Manager: slave idle wait	413	170408	412.61
Streams AQ: qmn coordinator idle wait	68	170226	2503.32
 dispatcher timer	 28	 168073	 6002.59
Streams AQ: waiting for messages in the queue	336	168025	500.07
smon timer	63	161556	2564.38
log file parallel write	331932	38612	.12
class slave wait	19	29965	1577.11
log file sync	7952	10467	1.32
Streams AQ: waiting for time management or cleanup tasks	4	6977	1744.37
db file sequential read	25281	5859	.23
db file async I/O submit	871	4727	5.43
Data file init write	440	1116	2.54
enq: HW - contention	26	1036	39.83
db file scattered read	2852	759	.27
buffer busy waits	63	751	11.92
os thread startup	99	583	5.89
Streams AQ: qmn coordinator waiting for slave to start	3	500	166.67
 control file parallel write	 1119	 428	 .38
control file heartbeat	1	400	399.99
library cache: mutex X	11	357	32.47
LGWR wait for redo copy	1223	275	.22
SQL*Net break/reset to client	138	217	1.58
db file parallel read	88	209	2.37
latch: cache buffers chains	25	182	7.27

Practice 2-3: Viewing System Statistics and Wait Events (continued)

buffer deadlock	2	153	76.59
cursor: pin S	2	144	72.06
log file switch completion	5	142	28.4
Disk file operations I/O	9510	129	.01
latch free	105	129	1.22
JS coord start wait	2	100	50.07
ADR block file read	135	79	.58
latch: redo allocation	1126	75	.07
rdbms ipc reply	42	69	1.64
latch: In memory undo latch	22	56	2.53
latch: redo copy	1	56	56.03
latch: cache buffers lru chain	9	45	4.96
EVENT	TOTAL WAITS	TIME WAITED	AVERAGE WAIT
-----	-----	-----	-----
asynch descriptor resize	6982	44	.01
direct path read	398	40	.1
log file single write	56	40	.71
enq: TX - row lock contention	11	36	3.25
latch: shared pool	20	27	1.36
enq: TX - allocate ITL entry	1	23	23.38
db file single write	242	19	.08
latch: call allocation	4	16	3.93
SQL*Net message to client	62550	15	0
latch: enqueue hash chains	11	14	1.29
control file sequential read	8236	14	0
local write wait	6	11	1.81
ADR block file write	6	11	1.78
row cache lock	1	9	8.52
direct path sync	5	7	1.44
enq: TX - index contention	1	7	6.8
latch: messages	150	6	.04
enq: CN - race with init	1	4	4.12
read by other session	2	3	1.34
log file switch (private strand flush incomplete)	3	2	.79
direct path write	56	2	.04
reliable message	6	1	.14
SQL*Net more data to client	165	1	.01
CRS call completion	1	1	1.17
log file sequential read	56	0	0
enq: KO - fast object checkpoint	1	0	.3
direct path write temp	1	0	.01
enq: US - contention	2	0	.03
latch: checkpoint queue latch	1	0	0
instance state change	1	0	.06
ADR file lock	8	0	0

Practice 2-3: Viewing System Statistics and Wait Events (continued)

latch: session allocation	1	0	.04
SQL*Net more data from client	82	0	0
latch: row cache objects	3	0	.13

78 rows selected.

4. Check the session wait events. Find the sessions that are currently waiting on the log file sync event. Use the V\$SESSION_WAIT view. Why are very few or no sessions listed?

Answer: The V\$SESSION_WAIT view shows waits only for currently connected sessions. The waits for all sessions, past and current, are summed in the V\$SYSTEM_EVENTS view. Your output from this query will vary from the example shown.

```
SQL> COLUMN time_waited HEADING TIME|WAITED
SQL> COLUMN AVERAGE_WAIT HEADING AVERAGE|WAIT
SQL> COLUMN seconds_in_wait HEADING SECONDS|IN_WAIT
SQL> COLUMN EVENT FORMAT A20
SQL> COLUMN total_waits HEADING TOTAL|WAITS
SQL>
SQL> select sid, wait_time, seconds_in_wait, state
  2      from v$session_wait
  3      where event = 'log file sync';
```

no rows selected

```
SQL>
SQL>
SQL> SELECT sid, event, total_waits,
  2      time_waited, average_wait
  3      FROM V$SESSION_EVENT
  4      WHERE event = 'log file sync';
```

SID	EVENT	TOTAL WAITS	TIME WAITED	AVERAGE WAIT
1	log file sync	1	0	.11
20	log file sync	262	106	.4
25	log file sync	3	7	2.37
26	log file sync	24	13	.56
32	log file sync	4	2	.39
33	log file sync	160	114	.71
36	log file sync	1	0	.03
37	log file sync	1	0	.07
38	log file sync	1	0	.05

9 rows selected.

SQL> exit

Practice 2-3: Viewing System Statistics and Wait Events (continued)

5. Stop the workload by executing the `rm runload` command. Then, clean up the practice environment by running the `./cleanup 2 3` script. (Cleanup may not be needed.)

```
$ cd /home/oracle/workshops
$ rm /home/oracle/workshops/runload
$ ./cleanup 2 3
System altered.
$
```

Practices for Lesson 3

During this practice, you use AWR snapshots and reports to identify and fix the issues that you discover in a running workload.

Practice 3-1: Using Automatic Workload Repository

In this practice, you run a workload on the `orcl` database, capture AWR snapshots, and diagnose and fix a performance issue. Note that during this practice, you explicitly capture AWR snapshots. However, by default, the system automatically captures them every hour.

1. Execute the `./prepare 3 1` script to set up this practice. You find this script in your `$HOME/workshops` directory. Log out of Enterprise Manager.

```
$ cd $HOME/workshops
$ ./prepare 3 1

drop tablespace tbsspc including contents and datafiles
*
ERROR at line 1:
ORA-00959: tablespace 'TBSSPC' does not exist

Tablespace created.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

User dropped.

User created.

Grant succeeded.

drop table spct purge
*
ERROR at line 1:
ORA-00942: table or view does not exist

Table created.

PL/SQL procedure successfully completed.

Database closed.
Database dismounted.
ORACLE instance shut down.
```

Practice 3-1: Using Automatic Workload Repository (continued)

```

ORACLE instance started.

Total System Global Area  836976640 bytes
Fixed Size                  1339740 bytes
Variable Size              805310116 bytes
Database Buffers           25165824 bytes
Redo Buffers                5160960 bytes
Database mounted.
Database opened.
$

```

2. Take your first AWR snapshot. You can use either Enterprise Manager or the DBMS_WORKLOAD_REPOSITORY package. The PL/SQL solution is shown in the sol_03_01_02.sh script located in your \$HOME/solutions directory. The script creates an AWR snapshot by using the DBMS_WORKLOAD_REPOSITORY package. The solution in this document uses Enterprise Manager.
 - a) First, open your browser by using the URL:
https://<hostname>:1158/em. This takes you to the EM Login page. Enter the credentials for the SYS user and click Login.
 - b) On the Database Instance page, click the Server tab.
 - c) On the Server page, click the Automatic Workload Repository link in the Statistics Management section.
 - d) On the Automatic Workload Repository page, click the number corresponding to the Snapshots field. The number of snapshots that you see will vary.

Manage Snapshots and Baselines

Snapshots 2

Baselines 1

Latest Snapshot Time **Feb 13, 2008 10:23:23 PM**

Earliest Snapshot Time **Feb 13, 2008 11:30:07 AM**

- e) On the Snapshots page, click Create.
 - f) On the Confirmation page, click Yes.
 - g) The processing page appears. Wait until the snapshot is created.
 - h) The Confirmation box appears on the Snapshots page. You can see that a new snapshot is created. Record the snapshot number: _____
3. Execute the ./workgen 3 1 script to generate the workload for this lab. The script is located in your \$HOME/workshops directory. Wait until the “Load is finished” message appears, and then continue. This script takes about two minutes. Proceed to step 4 immediately after the script finishes.

Note: The last “\$” prompt does not appear until you press Enter one more time.

Practice 3-1: Using Automatic Workload Repository (continued)

```
$ ./workgen 3 1

PL/SQL procedure successfully completed.
...
PL/SQL procedure successfully completed.

Load is finished
$
```

4. Take a second AWR snapshot right after workload generation is completed.
 - a) Navigate to the Snapshots page, and then click Create.
 - b) On the Confirmation page, click Yes.
 - c) The processing page appears. Wait until the snapshot is created.
 - d) The Confirmation box appears on the Snapshots page, where you can see that a new snapshot has been created. Record the snapshot number: _____
5. Generate an AWR report. You can use the `awrrpt.sql` script located in the `$ORACLE_HOME/rdbms/admin` directory to create an AWR report. The solution in this document uses Enterprise Manager.
 - a) Still on the Snapshots page, select View Report from the Actions drop-down list. Make sure that your first AWR snapshot is selected. This is the snapshot number that you recorded in step 2h. Click Go.
 - b) On the View Report page, make sure that your second AWR snapshot is selected. This is the snapshot that you recorded in step 4d. Click OK.
 - c) The processing page appears. Wait until the report is created.
6. Review the report. What can you determine from this report about the performance of this workload?
 - a) After the report is created, you are taken to the Snapshot Details page where you can see the report.
 - b) Scroll down to Top 5 Timed Events. What wait event accounts for most of the %DB Time? _____ (buffer busy waits)

Top 5 Timed Foreground Events

Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
buffer busy waits	1,225	268	219	42.11	Concurrency
DB CPU		59		9.28	
enq: HW - contention	47	54	1153	8.52	Configuration
log file sync	252	34	133	5.27	Commit
free buffer waits	748	21	28	3.25	Configuration

Practice 3-1: Using Automatic Workload Repository (continued)

- c) Scroll down to Time Model Statistics. Which category collects the most %DB Time? _____(sql execute elapsed time)

Note: Elapsed time includes wait time.

From these first-level diagnostics, you know that the buffer busy waits occurring during the SQL execution are the cause of the performance issue.

Time Model Statistics		
<ul style="list-style-type: none">• Total time in database user-calls (DB Time): 635.9s• Statistics including the word "background" measure background• Ordered by % or DB time desc, Statistic name		
Statistic Name	Time (s)	% of DB Time
sql execute elapsed time	564.37	88.75
parse time elapsed	74.92	11.78
hard parse elapsed time	74.07	11.65
DB CPU	59.02	9.28
connection management call elapsed time	18.46	2.90
PL/SQL execution elapsed time	9.49	1.49
PL/SQL compilation elapsed time	5.45	0.86
hard parse (sharing criteria) elapsed time	0.23	0.04
repeated bind elapsed time	0.07	0.01
hard parse (bind mismatch) elapsed time	0.01	0.00
sequence load elapsed time	0.00	0.00
DB time	635.91	
background elapsed time	93.76	
background cpu time	0.76	

- d) The SQL statements that have high elapsed time will be the most likely causes of the waits. You can scroll down to the "SQL ordered by Elapsed Time" section or click the following links under Main Report: SQL Statistics, and then SQL ordered by Elapsed Time. What is the SQL ID of the statements that have the highest percentage of DB time?

Practice 3-1: Using Automatic Workload Repository (continued)

SQL ordered by Elapsed Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- % Total DB Time is the Elapsed Time of the SQL statement divided into the Total Database Time multiplied by 100
- %Total - Elapsed Time as a percentage of Total DB time
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Captured SQL account for 82.2% of Total DB Time (s): 906
- Captured PL/SQL account for 69.7% of Total DB Time (s): 906

Elapsed Time (s)	Executions	Elapsed Time per Exec (s)	%Total	%CPU	%IO	SQL Id	SQL Module	SQL Text
603.99	23	26.26	66.70	3.59	0.03	fcu83t10rr413	SQL*Plus	declare t number; begin for t ...
538.15	480,000	0.00	59.43	2.11	0.03	3csh3g3mjhmzh	SQL*Plus	INSERT INTO SPCT VALUES (NULL,...
109.21	28	3.90	12.06	0.02	0.27	4vs91dcv7u1p6	OMS	insert into sys.aud\$(sessioni...
43.29	27	1.60	4.78	0.04	7.88	f711myt0q6cma	perl@edrsr10p1.us.oracle.com (TNS V1-V3)	insert into sys.aud\$(sessioni...
11.32	4	2.83	1.25	1.18	1.28	6gvch1xu9ca3g		DECLARE job BINARY_INTEGER := ...

- e) Click the SQL ID value to see the full SQL text for each of the SQL statements that you identified in step 6d. What do you observe?

Answer: The first is a PL/SQL statement with an INSERT statement. The second is the INSERT statement from the same PL/SQL block. The statement inserts records into the SPCT table.

- f) What types of blocks have the buffer busy waits? _____
 Scroll down to the Buffer Wait Statistics report in the Wait Statistics section, or click the following links in succession: Back to Top, Wait Statistics, and Buffer Wait Statistics. (Data blocks; these could be table or index blocks.)

Buffer Wait Statistics

- ordered by wait time desc, waits desc

Class	Waits	Total Wait Time (s)	Avg Time (ms)
data block	1,764	264	150
segment header	22	3	139
undo header	6	0	80

Practice 3-1: Using Automatic Workload Repository (continued)

- g) Are the buffer busy waits against one or a few tables and indexes? _____
 If so, which segments are affected? _____
 In what tablespace do these segments exist? _____
 Scroll down to “Segments by Buffer Busy Waits,” or click the following links in succession: Back to Top, Segment Statistics, and Segments by Buffer Busy Waits. (Yes, one table is experiencing most of the waits. The SPCT table is experiencing the waits. The SPCT table is in the TBSSPC tablespace.)

Segments by Buffer Busy Waits

- % of Capture shows % of Buffer Busy Waits for each top segment compared
- with total Buffer Busy Waits for all segments captured by the Snapshot

Owner	Tablespace Name	Object Name	Subobject Name	Obj. Type	Buffer Busy Waits	% of Capture
SPC	TBSSPC	SPCT		TABLE	1,779	99.61
SYS	SYSTEM	AUD\$		TABLE	7	0.39

Diagnostics: This workload is causing buffer busy waits and block contention by performing inserts from multiple processes into the same table.

Additional Information: If the tablespace has segment space management set to manual, all the processes attempt to insert rows into the same small set of blocks. If the segment space management is set to automatic, the processes will insert rows into different blocks, choosing blocks based on a hash of the process_id.

- h) Check the properties of the TBSSPC tablespace. On the Database home page, click Server, and then Tablespaces. Select the TBSSPC tablespace, and then click View.

View Tablespace: TBSSPC

Actions

Name **TBSSPC**
 Bigfile tablespace **No**
 Status **ReadWrite**
 Type **Permanent**
 Extent Management **local**
 Encryption **NO**

Storage

Allocation Type **Automatic**
Segment Space Management **Manual**
 Enable logging **Yes**
 Compression **No Compression**
 Block Size (B) **8192**

Conclusion: The TBSSPC tablespace is set to Manual Segment Space Management. Moving the SPC table to a tablespace with Automatic Segment Space Management could make a significant difference in performance.

Practice 3-1: Using Automatic Workload Repository (continued)

7. To implement the changes to the SPCT table, execute the `wksh_03_01_07.sql` script. This script drops the SPCT table, drops the TBSSPC tablespace, re-creates the tablespace with the ASSM property, and creates the SPCT table. You can locate the `wksh_03_01_07.sql` script in your `$HOME/workshops` directory.

```
$ sqlplus / as sysdba

SQL> @wksh_03_01_07.sql
SQL>
SQL> drop tablespace tbsspc
      2 including contents and datafiles;

Tablespace dropped.

SQL>
SQL> CREATE SMALLFILE TABLESPACE "TBSSPC"
      2 DATAFILE 'tbsspc1.dbf' SIZE 50M
      3 AUTOEXTEND ON NEXT 10M MAXSIZE 200M
      4 LOGGING
      5 EXTENT MANAGEMENT LOCAL
      6 SEGMENT SPACE MANAGEMENT AUTO;

Tablespace created.

SQL>
SQL> connect spc/spc
Connected.
SQL>
SQL> drop table spct purge;
drop table spct purge
      *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> create table spct(id number, name varchar2(2000))
      2 tablespace tbsspc;

Table created.

SQL>
SQL> exec DBMS_STATS.GATHER_TABLE_STATS(-
> ownname=>'SPC', tabname=>'SPCT',-
> estimate_percent=>DBMS_STATS.AUTO_SAMPLE_SIZE);

PL/SQL procedure successfully completed.

SQL>
SQL> exit;
$
```

Practice 3-1: Using Automatic Workload Repository (continued)

8. Execute the `./workgen 3 1` script again to start generating the workload for this practice. Wait until “Load is finished” appears.

```
$ ./workgen 3 1
...
Load is finished
$
```

9. Take a third AWR snapshot.
 - a) Navigate to the Snapshots page and click Create.
 - b) On the Confirmation page, click Yes.
 - c) The processing page appears. Wait until the snapshot is created.
 - d) The Confirmation box appears on the Snapshots page. You can clearly see that a third snapshot has been created.
10. Generate and review a Compare Periods report comparing the periods between the three snapshots. You can use the `awrrdrpt.sql` script located in your `$ORACLE_HOME/rdbms/admin` directory to generate the Compare Periods report. The solution in this document uses Enterprise Manager.
 - a) On the Snapshots page, select Compare Periods from the Actions drop-down list. Make sure that the first AWR snapshot is selected. Click Go.
 - b) On the Compare Periods: First Period End page, select your second AWR snapshot and click Next.
 - c) On the Compare Periods: Second Period Start page, select your second AWR snapshot and click Next.
 - d) On the Compare Periods: Second Period End page, select your last AWR snapshot and click Next.
 - e) On the Compare Periods: Review page, click Finish.
 - f) On the Compare Periods: Results page, select Per Transaction from the View Data drop-down list on the General tabbed page.
 - g) You can see the performance differences between the two analyzed periods. Click Report.
 - h) The processing page appears. Wait until the report is generated.
 - i) On the Report tabbed page, you can now see Workload Repository Compare Period Report.
 - j) In the second section from the top of the report, notice the difference between the Average Active Sessions, and the DB Time between the first and second snapshot sets. Because the workload was identical, the reduction in Average Active Sessions indicates the waits were lower, and the DB Time reduction indicates that less time was spent in database calls. Both indicate that the instance is performing more efficiently.

Practice 3-1: Using Automatic Workload Repository (continued)

- k) Scroll down to the Top 5 Timed Events section. You can see the difference between the two periods for the buffer busy waits event. (Your results will vary from the following example.) In the second run, the order of the waits will likely change.

Top Timed Events

- Events with a "-" did not make the Top list in this set of snapshots, but are displayed for comparison purposes

1st						2nd					
Event	Wait Class	Waits	Time(s)	Avg Time(ms)	%DB time	Event	Wait Class	Waits	Time(s)	Avg Time(ms)	%DB time
buffer busy waits	Concurrency	1,225	267.76	218.58	42.11	buffer busy waits	Concurrency	180	85.51	475.07	16.89
CPU time			59.02		9.28	log file switch (private strand flush incomplete)	Configuration	22	74.17	3,371.43	14.65
enq: HW - contention	Configuration	47	54.21	1,153.43	8.52	free buffer waits	Configuration	90	68.18	757.53	13.47
db file async I/O submit	System I/O	29	35.02	1,207.66	5.51	log file sync	Commit	169	62.81	371.67	12.41
log file sync	Commit	252	33.52	133.00	5.27	CPU time			49.14		9.71
-free buffer waits	Configuration	748	20.65	27.61	3.25	-db file async I/O submit	System I/O	148	25.94	175.28	5.12
-log file switch (private strand flush incomplete)	Configuration	7	12.71	1,816.05	2.00	-					

- l) Scroll down to the Buffer Wait Statistics in the Wait Stats section. Here also, you can clearly see the difference.

Buffer Wait Statistics

- Ordered by absolute value of 'Diff' column of 'Wait Time % of DB time' descending

Class	Wait Time % of DB time			Total Wait Time (s)		# Waits		Avg Wait Time (ms)		
	1st	2nd	Diff	1st	2nd	1st	2nd	1st	2nd	%Diff
data block	0.42	0.16	-0.25	264.37	82.65	1,764	300	1.50	2.76	84.00
undo header	0.00	0.00	0.00	0.48	2.09	6	5	0.80	4.18	422.50
segment header	0.00	0.00	-0.00	3.05	0.74	22	6	1.39	1.23	-11.51

Practice 3-1: Using Automatic Workload Repository (continued)

- m) Scroll down to the Top Segments by Buffer Busy Waits in the Segment Statistics section. The number of buffer busy waits has changed. Possibly, the classes of buffer busy waits have changed as well. (This can be seen in some reports where one period or the other is reporting 0 waits.) The bitmap block (BMB) class of wait cannot occur in the first period because the BMB block exists only in the ASSM tablespaces.

Top Segments by Buffer Busy Waits

- Ordered by absolute value of 'Diff' column of '% of Total Buffer Busy Waits'
- Top 5 Segments from each period are compared
- Total Buffer Busy Waits First: 1,225, Second: 180
- Buffer Busy Wait Time First: 267.76 seconds, Second: 85.51 seconds
- Buffer Busy Wait Time as % of DB time First: 42.11%, Second: 16.89%
- Captured Buffer Busy Waits First: 1,786, Second: 306

Owner	Tablespace	Object Name	Subobject Name	Type	% of Total Buffer Busy Waits					Buffer Busy Waits			% of Captured Buffer Busy Waits		
					1st	1st Total	2nd	2nd Total	Diff	1st	2nd	%Diff	1st	2nd	Diff
SPC	TBSSPC	SPCT		TABLE	145.22	145.22	0.00	0.00	-145.22	1,779	0	-100.00	99.61	0.00	-99.61
SYS	SYSTEM	AUD\$		TABLE	0.57	145.80	24.44	24.44	23.87	7	44	528.57	0.39	14.38	13.99
SPC	TBSSPC	SPCT		TABLE	145.22	291.02	145.56	170.00	0.33	1,779	262	-85.27	99.61	85.62	-13.99

11. Run the `./cleanup 3 1` cleanup script.

```
$ ./cleanup 3 1
```


Practices for Lesson 4

The goal of this practice is to identify the sources of performance issues.

Practice 4-1: Using Enterprise Manager to Identify OS Issues

In this practice, the workload generator starts an application load on the database and an OS load.

1. In a terminal window, use the `date` command to determine the start time of the workload.

```
$ date
Tue Mar 25 06:19:57 GMT-7 2008
```

2. In a terminal window, change directory to `/home/oracle/workshops`. Start the workload generator with the command: `./workgen 4 1`. Verify that the workload generator has started the workload processes by executing the `ps` command.

```
$ cd /home/oracle/workshops
$ ./workgen 4 1
$ ps
  PID TTY          TIME CMD
 2365 pts/2        00:00:00 bash
 15312 pts/2        00:00:00 wkld_shell.sh
 15313 pts/2        00:00:00 insert_orders.s
 15314 pts/2        00:00:00 update_orders.s
 15317 pts/2        00:00:00 delete_orders.s
 15318 pts/2        00:00:00 sqlplus
 15319 pts/2        00:00:00 insert_orders.s
 15320 pts/2        00:00:00 sqlplus
 15321 pts/2        00:00:00 update_orders.s
 15322 pts/2        00:00:00 delete_orders.s
 15323 pts/2        00:00:00 sqlplus
 15329 pts/2        00:00:00 insert_orders.s
 15330 pts/2        00:00:00 sqlplus
 15445 pts/2        00:00:00 sqlplus
 15452 pts/2        00:00:00 sqlplus
 15459 pts/2        00:00:00 sleep
 15460 pts/2        00:00:00 sleep
 15461 pts/2        00:00:00 sleep
 15462 pts/2        00:00:00 sleep
 15463 pts/2        00:00:00 sleep
 15464 pts/2        00:00:00 sleep
...
```

3. Use Enterprise Manager to determine the source of the loads. In the browser's address bar, enter the URL, <https://<hostname>:1158/em>, to view the `orcl` database information. On the Login page, enter `SYS` as the username and `oracle` as the password, select `SYSDBA` from the Connect As drop-down list, and click Login.
4. Wait for a while until the Active Sessions graph shows some waits. Find the OS process that is consuming the most CPU.

Practice 4-1: Using Enterprise Manager to Identify OS Issues (continued)

- The Database home page for the `orcl.us.oracle.com` database appears. Observe the Host CPU measurement.
 - A significant percentage of the Host CPU is dedicated to **Other** processes. Click the Other link to view the Performance Summary page.
 - Scroll down to view Top 10 Processes by specifying View By CPU Utilization (%).
- Note:** Refresh this view a few times until an OS process (not owned by `oracle`) appears. The `dd` process should appear. It may not be the top process. What is the process ID of the `dd` process?

Processes
Processes [244](#)

Top 10 Processes
View By: CPU Utilization (%)

Process ID	Command	CPU Utilization (%)	CPU Total (seconds)	Resident Size (KB)	Virtual Size (KB)	Owner
16321	oracleorcl (DESCRIPTION=(LOCAL=YES)(ADDRESS=(PROTOCOL=beq)))	1.75	0	21,396	949,720	oracle
16324	oracleorcl (DESCRIPTION=(LOCAL=YES)(ADDRESS=(PROTOCOL=beq)))	1.35	0	21,996	949,720	oracle
14705	dd if /dev/sda5 of /dev/null conv ebodic bs 102400 count 23520	0.85	5	624	4,188	root
16311	oracleorcl (DESCRIPTION=(LOCAL=YES)(ADDRESS=(PROTOCOL=beq)))	0.75	0	22,652	949,720	oracle

- Stop the `dd` process using the `kill` command. Because this process is running as `root`, use the following command:

```
$ sudo kill <process ID>
```

- Stop the workload by executing the `rm runload` command.

```
$ rm runload
```

In this practice, you found and stopped an OS process that was consuming CPU.

Practices for Lesson 5

The goal of this practice is to create and monitor a server-generated alert.

Practice 5-1: Using Metrics

This practice shows you how to set up server-generated alerts and how to retrieve alert information from Enterprise Manager. During this practice, you see how alerts are triggered. However, this practice does not involve analysis of the cause of the generated alerts. This analysis is performed during subsequent practices by using AWR snapshots and ADDM.

1. Change to the `/home/oracle/workshops` directory. Execute the `./prepare 5 1` script to set up the necessary objects before starting the workload. Log out of Enterprise Manager.

```
$ cd /home/oracle/workshops
$ ./prepare 5 1
Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area  836976640 bytes
Fixed Size                  1339740 bytes
Variable Size              805310116 bytes
Database Buffers           25165824 bytes
Redo Buffers                5160960 bytes
Database mounted.
Database opened.
drop tablespace tbsjfv including contents and datafiles
*
ERROR at line 1:
ORA-00959: tablespace 'TBSJFV' does not exist

Tablespace created.

Table created.

PL/SQL procedure successfully completed.

User dropped.

User created.

Grant succeeded.

drop table jfvvt purge
*
```

Practice 5-1: Using Metrics (continued)

```
ERROR at line 1:
ORA-00942: table or view does not exist

Table created.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

$
```

2. Using EM Database Control for the `orcl` instance, determine the current threshold values set for the `DB_TIME_WAITING` metric. Pay special attention to the Concurrency wait class. Use both Enterprise Manager and SQL*Plus to find those values.
 - a) On the Database home page, scroll down to the Related Links section, and click All Metrics.

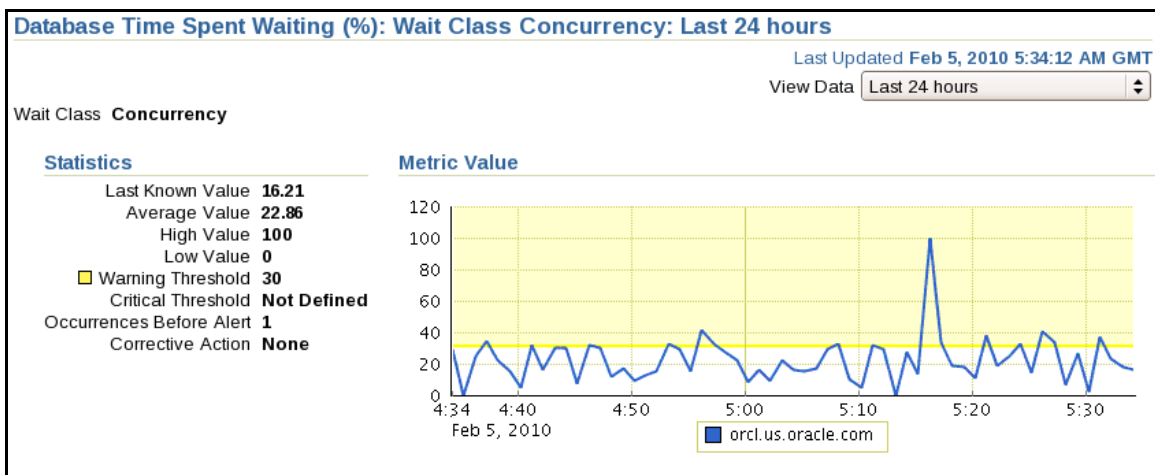
Related Links		
Access	Add Exadata Cell Targets	Advisor Central
Alert History	Alert Log Contents	All Metrics
Baseline Metric Thresholds	Blackouts	EM SQL History
Jobs	Metric and Policy Settings	Metric Collection Errors
Monitoring Configuration	Monitor in Memory Access Mode	Policy Groups
Scheduler Central	SQL Worksheet	Target Properties
User-Defined Metrics		

- b) On the All Metrics page, scroll to Waits by Wait Class and expand it.
 - c) Click Database Time Spent Waiting (%).
 - d) On the Database Time Spent Waiting (%) page, you can see all thresholds set for each class. Click Concurrency.

Practice 5-1: Using Metrics (continued)

Database Time Spent Waiting (%)					
				Latest Data Collected From Target Feb 5, 2010 5:16:52 AM GMT	
				View Data	Last 24 hours
Wait Class	Average Value	High Value	Low Value	Last Known Value	Collection Timestamp
Cluster	0	0	0	0	Feb 5, 2010 5:30:12 AM
Application	1.41	43.92	0.02	0.02	Feb 5, 2010 5:30:12 AM
Idle	0	0	0	0	Feb 5, 2010 5:30:12 AM
System I/O	15.87	100	1.27	4.18	Feb 5, 2010 5:30:12 AM
User I/O	2.02	33.23	0	33.23	Feb 5, 2010 5:30:12 AM
Network	0.14	0.81	0.03	0.1	Feb 5, 2010 5:30:12 AM
Scheduler	0	0	0	0	Feb 5, 2010 5:30:12 AM
Queueing	0	0	0	0	Feb 5, 2010 5:30:12 AM
Configuration	0	0	0	0	Feb 5, 2010 5:30:12 AM
Commit	7.25	39.47	0.12	3.75	Feb 5, 2010 5:30:12 AM
Administrative	0	0	0	0	Feb 5, 2010 5:30:12 AM
Other	0.86	11.9	0.03	0.03	Feb 5, 2010 5:30:12 AM
Concurrency	22.44	100	0	2.22	Feb 5, 2010 5:30:12 AM

- e) View the Wait Class Concurrency page.
 What is the default Warning threshold? _____
 Have there been any alerts generated? _____
 Do the alerts correspond to the graph? _____ Check the alert history below the graph. The graph does not display information before the last database restart.



- f) From a SQL*Plus session, determine the threshold values for the wait classes shown in the preceding screenshot. Change the directory to /home/oracle/labs. Connect as a sysdba user and execute the lab_05_01_02.sql script.

```
$ cd /home/oracle/labs
$ sqlplus / as sysdba

SQL> @lab_05_01_02.sql
SQL>
SQL> variable wo number
SQL> variable wv varchar2(10)
```

Practice 5-1: Using Metrics (continued)

```

SQL> variable co number
SQL> variable cv varchar2(10)
SQL> variable op number
SQL> variable ct number
SQL>
SQL> BEGIN
  2  DBMS_SERVER_ALERT.get_threshold(
  3      metrics_id => dbms_server_alert.DB_TIME_WAITING,
  4      warning_operator      => :wo,
  5      warning_value         => :wv,
  6      critical_operator     => :co,
  7      critical_value        => :cv,
  8      observation_period    => :op,
  9      consecutive_occurrences => :ct,
 10      instance_name         => 'orcl',
 11      object_type =>
 12          dbms_server_alert.OBJECT_TYPE_EVENT_CLASS,
 13      object_name           => 'Concurrency');
 14  END;
 15  /

```

PL/SQL procedure successfully completed.

```

SQL>
SQL> print  wv

```

```

WV
-----
30

```

```

SQL>
SQL> col object_name format a20
SQL> col metrics_name format a25
SQL> col warning_value format a10
SQL> col critical_value format a10
SQL>
SQL> select object_name,warning_value,critical_value
  2  from dba_thresholds
  3  where metrics_name='Database Time Spent Waiting (%)';

```

OBJECT_NAME	WARNING_VA	CRITICAL_V
Concurrency	30	
Administrative	30	
Application	30	
Configuration	30	
Network	30	
Other	30	
Cluster	50	
Commit	50	

8 rows selected.

Practice 5-1: Using Metrics (continued)

```
SQL>
SQL> exit;
$
```

3. Using a PL/SQL command or the lab_05_01_03.sql script, change the thresholds for the Concurrency class to 60% for the warning level and 90% for the critical level. Using both SQL*Plus and Enterprise Manager, verify that your change took effect.
 - a) Use the DBMS_SERVER_ALERT package to set the new threshold values.

```
$ sqlplus / as sysdba
```

```
SQL> @lab_05_01_03.sql
```

```
SQL>
```

```
SQL> BEGIN
```

```
 2     DBMS_SERVER_ALERT.set_threshold(
 3     metrics_id => dbms_server_alert.DB_TIME_WAITING,
 4     warning_operator => dbms_server_alert.OPERATOR_GE,
 5     warning_value    => 60,
 6     critical_operator=> dbms_server_alert.OPERATOR_GE,
 7     critical_value   => 90,
 8     observation_period=> 1,
 9     consecutive_occurrences => 1,
10     instance_name    => 'orcl',
11     object_type       =>
12         dbms_server_alert.OBJECT_TYPE_EVENT_CLASS,
13     object_name       => 'Concurrency');
14     END;
15 /
```

PL/SQL procedure successfully completed.

```
SQL>SQL> col object_name format a20
```

```
col metrics_name format a25
```

```
col warning_value format a10
```

```
col critical_value format a10
```

```
SQL> SQL> SQL> SQL>
```

```
SQL> select object_name,warning_value,critical_value
 2     from dba_thresholds
 3     where metrics_name=
 4         'Database Time Spent Waiting (%)';
```

OBJECT_NAME	WARNING_VA	CRITICAL_V
Concurrency	60	90
Administrative	30	
Application	30	
Configuration	30	
Network	30	
Other	30	

Practice 5-1: Using Metrics (continued)

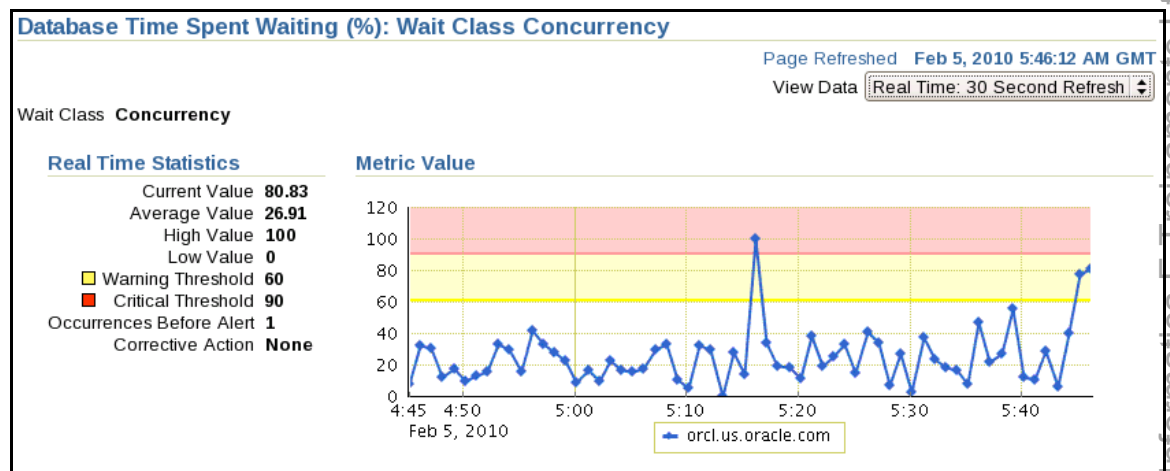
```
Cluster          50
Commit           50
```

8 rows selected.

```
SQL> exit;
```

```
$
```

- b) Still on the Wait Class Concurrency page, select Real Time: 30 Second Refresh from the View Data drop-down list. Note that the Warning and Critical values have been changed.



4. Change the directory to \$HOME/workshops. Execute the ./workgen 5 1 script to start the workload for this practice. While the script is executing, observe the variation of the curve on the Wait Class Concurrency page. After the workload is finished, find the alert history from a SQL*Plus session. Check both the outstanding alerts and the alert history. The workload script takes up to six minutes to complete on some machines.

```
$ cd $HOME/workshops
$ ./workgen 5 1
$
... The following messages will appear for 1 to 6 minutes ...
PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
```

Practice 5-1: Using Metrics (continued)

```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

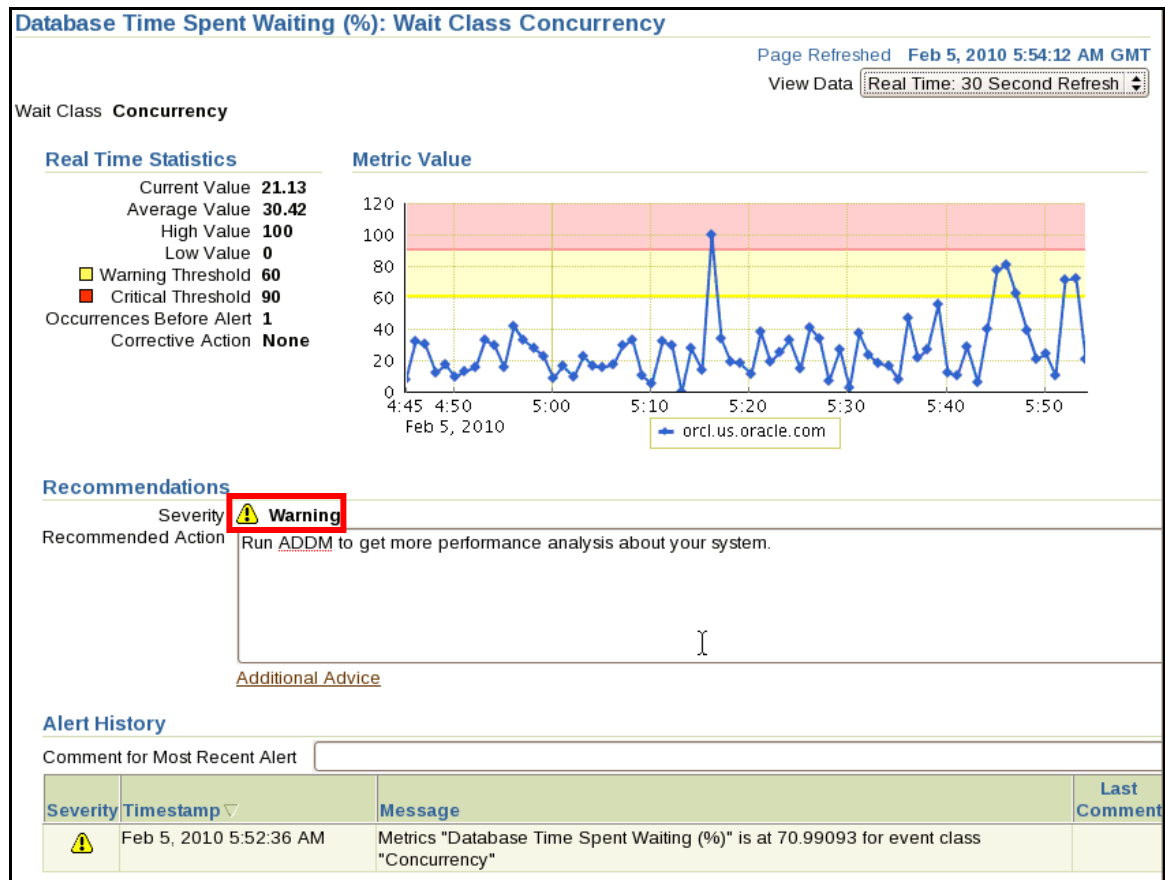
```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

```
$
```

Practice 5-1: Using Metrics (continued)

- While the script executes, look at the evolution of the curve on the Wait Class Concurrency graph. After a while, you should see a warning or critical alert.



- After the script finishes, you can retrieve the history of your alerts by using SQL*Plus. Depending on your timing, the alert may appear in the DBA_OUTSTANDING_ALERTS or DBA_ALERT_HISTORY view. If the alert is not yet cleared, it shows in outstanding Alerts view. After it is cleared, it is visible in the Alert History view. The following example shows the output when the alert has been cleared (**Hint:** lab_05_01_06.sql in the \$HOME/labs directory).

```
$ sqlplus / as sysdba

SQL> @../labs/lab_05_01_06.sql
SQL>
SQL> col reason format a75
SQL>
SQL> SELECT reason FROM DBA_OUTSTANDING_ALERTS;

REASON
-----
Metrics "Database Time Spent Waiting (%)" is at 63.59094 for
event class "Concurrency"
```

Practice 5-1: Using Metrics (continued)

```

SQL>
SQL> SELECT TO_CHAR(begin_time,'hh24:mi:ss'),
2         dbtime_in_wait,average_waiter_count
3 FROM V$WAITCLASSMETRIC_HISTORY
4 WHERE wait_class#=4
5 AND wait_class_id=3875070507
6 AND begin_time > SYSDATE-(10/1440);

TO_CHAR( DBTIME_IN_WAIT AVERAGE_WAITER_COUNT
-----
05:55:12      63.5909417          .001635947
05:54:12      24.6778322          .004029849
05:53:11      21.1284229          .055115698
05:52:11      71.9247151         13.4691388
05:51:12       70.990928          7.82854524
05:50:12      10.3871741          .001729537
05:49:12      24.3581563          .003170473
05:48:11      20.7383737          .001333433
05:47:11      38.9326013          .001199334

9 rows selected.

SQL>
SQL>
SQL> -- Cleared!
SQL>
SQL> SELECT reason,resolution
2 FROM DBA_ALERT_HISTORY
3 WHERE reason like
4      '%Database Time Spent Waiting (%)%Concurrency%'
5 AND TO_DATE(SUBSTR(TO_CHAR(creation_time),1,18)||
6      SUBSTR(TO_CHAR(creation_time),26,3) ,
7      'DD-MON-YY HH:MI:SS AM') > SYSDATE-(10/1440)
8 ORDER BY creation_time DESC;

REASON
-----
RESOLUT
-----
Metrics "Database Time Spent Waiting (%)" is at 21.12842 for
event class "Concurrency"
cleared

SQL>
SQL> EXIT;

$

```

Practice 5-1: Using Metrics (continued)

7. Check Alert History from EM.
 - a) In the Related Links section of the Database home page, click Alert History. Are the alerts that you generated in this practice visible?
 - b) Click the Database Time Spent Waiting (%) link.
Note: This action displays Database Time Spent Waiting (%) with the wait by class.
 - c) Click the Concurrency link.
Note: This action displays the Database Time Spent Waiting (%): Wait Class Concurrency page with the graph showing the waits and thresholds.
8. In a terminal window, change directory to `/home/oracle/workshops` and execute the `cleanup 5 1` script to clean up your environment.

```
$ ./cleanup 5 1

User dropped.

Tablespace dropped.

PL/SQL procedure successfully completed.

$
```

Practices for Lesson 6

The goal of this practice is to create a static baseline and a repeating baseline. The static baseline is used to create an adaptive threshold on a metric.

Practice 6-1: Using Baselines

In this practice, you create a static baseline over the snapshots that are present in the database, from 9:00 PM the day before this class started to 5:00 AM the following morning. This set of snapshots should be in the range of 1–100. (Your instructor may have different instructions). Name the baseline LAST_NIGHT.

1. Create the LAST_NIGHT baseline.

Step	Page	Action		
a	Database Instance home	Click the Server tab.		
b	Server	Click Automatic Workload Repository in the Statistics Management section.		
c	Automatic Workload Repository	Click the Snapshots link.		
d	Snapshots	Write the time and date of snapshots 2 and 100 (or the most recent, whichever is the least).	Date	Time
		2		
		100 (most recent)		
		Click Automatic Workload Repository in the navigation path at the top of the page to return to the Automatic Workload Repository page.		
e	Automatic Workload Repository	Click the Baselines link.		
f	AWR Baselines	Click Create.		
g	Create Baseline:	Select Single.		

Practice 6-1: Using Baselines (continued)

	Baseline Interval Type	Click Continue.		
h	Create Baseline: Single Baseline		Field	Value
		Complete the baseline name field.	Baseline Name	LAST_NIGHT
		Select Snapshot Range. Move the chart view range until you can see the set of snapshots 2 through 100.		
		Select the start and end snapshots.	Period Start Time	Time of Snapshot 2 Or as directed by the instructor
			Period End Time	Time of snapshot 100 Or as directed by the instructor
		Click Finish.		

2. On the AWR Snapshots page, calculate the statistics for the LAST_NIGHT baseline.

Step	Page	Action
a.	AWR Baselines	Select LAST_NIGHT.
		From the Actions drop-down list, select Schedule Statistics Computation.
		Click Go.
b	Compute Threshold Statistics: LAST_NIGHT	Select Immediately.
		Click Submit.

3. Set the Performance page to show the LAST_NIGHT baseline statistics.
- a) Click Settings on the Performance page.

Practice 6-1: Using Baselines (continued)

- b) On the Performance Page Settings page, in the Baseline display section, set the baseline to LAST_NIGHT.

Select “Show the 99th percentile line using a static baseline with computed statistics.” Click OK.

Step	Page	Action
a	AWR Baselines	Click the Database tab.
b	Database Home	Click the Performance tab.
c	Performance	Click Settings.
d	Performance Page Settings	From the Baseline Name drop-down list, select LAST_NIGHT.
		Select “Show the 99th percentile line using a static baseline with computed statistics.”
		Click OK.

4. Create a template that creates a baseline for the Monday maintenance window for the next month. Keep the baseline for 60 days. The weekday maintenance window start and end times can be found on the Windows page Home > Server > Windows (in the Scheduler section).

Step	Page	Action
a.	Performance	Click the Server tab.
b	Server	Click Windows in the Oracle Scheduler section.
c	Scheduler Windows	Record the start time of MONDAY_WINDOW. This is shown as Repeat Time (should be 10:00 PM).
		Record the duration of MONDAY_WINDOW.
		Click the Database tab.
d	Database home	Click the Server tab.
e	Server	Click AWR Baselines.

Practice 6-1: Using Baselines (continued)

f	AWR Baselines	Click Create.		
g	Create Baseline: Baseline Interval Type	Click Repeating.		
		Click Continue.		
h	Create Baseline: Repeating Baseline Template (See the example in the following screenshot.)	Set values on the page:	Field	Value
			Baseline Name Prefix	MONDAY
			Baseline Time Period: Start Time	10 PM
			Baseline Time Period: Duration	4
			Frequency	Weekly and Monday
			Interval of Baseline Creation: Start Time	Today 9:50 PM
			Interval of Baseline Creation: End Time	One Month from today 2:20 AM
			Retention Time (days)	60
		Click Finish.		
i	AWR Baselines	Click AWR Baseline Templates in the Related Links section.		
j	AWR Baseline Templates	Verify that the MONDAY template was created as expected.		

Practice 6-1: Using Baselines (continued)

h)

Create Baseline: Repeating Baseline Template

Cancel

Back

Finish

The repeating type of baseline has a time interval that repeats over a time period. For example, every Monday from 10:00 AM to 12:00 PM for the year 2007.

* Baseline Name Prefix

Baseline Time Period

Start Time

10

PM

Duration (Hours)

4

Frequency

☐ Daily
 ☒ Weekly
 ☐ Monday
 ☐ Tuesday
 ☐ Wednesday
 ☐ Thursday
 ☐ Friday
 ☐ Saturday
 ☐ Sunday

Interval of Baseline Creation

Start Time

Feb 5, 2010

9

50

AM

End Time

Mar 5, 2010

2

20

AM

Purge Policy

Retention Time (Days)

60

TIP A baseline template with the same name as the baseline name prefix will be created.

- Define an alert threshold for transactions per second based on the LAST_NIGHT baseline. Set the critical threshold to 120% of maximum and the warning to 110% of maximum.

Step	Page	Action		
a	AWR Baseline Templates	Click AWR Baselines in the Related Links section.		
b	AWR Baselines	Click Baseline Metric Thresholds in the Related Links section.		
c	Baseline Metric Thresholds	Change the view to Basic Metrics.		
		Edit the Number of Transaction (per second) metric. (Click the pencil icon on the right.)		
d	Edit Thresholds: Number of Transactions (per second)	Set values:	Field	Value
			Name	LAST_NIGHT
			Threshold Type	Set to Percentage of Maximum
			Critical	120

Practice 6-1: Using Baselines (continued)

			Warning	110
			Occurrences	1
		Click Preview. To move the active window slider (in the upper left) over the period of time that covers the LAST_NIGHT baseline, click the day that has the majority of the baseline statistics in it.		
		Write down the estimated number of transactions per second required to produce a critical alert:_____		
		Click Apply Thresholds.		

6. Run a workload by executing the `./workgen 6 1` script. The script will ask for the number of transactions per second that is required to produce a critical alert. If the script is running correctly, you can run the `ps` command and see the `baseline_wkld.sh` process.

```
$ cd /home/oracle/workshops
$ ./workgen 6 1
Enter the number of transactions needed to produce a critical
alert:12
$ ps
  PID TTY          TIME CMD
 29422 pts/2    00:00:00 bash
 31484 pts/2    00:00:00 baseline_wkld.sh
 31511 pts/2    00:00:00 sleep
 31512 pts/2    00:00:00 ps
```

7. On Database Home page, set View Data to Automatically (15 sec). After a short time, an alert appears in the Alerts section of the Database home page of EM. This alert is named “Number of Transactions (per second)” and the message indicates the value of the metric.

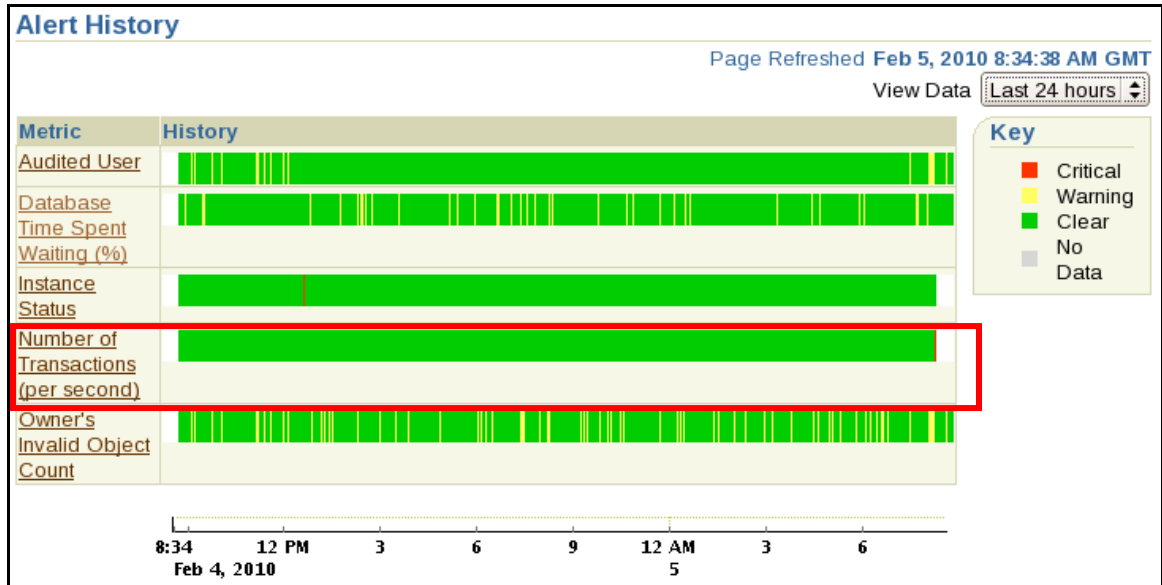
Alerts					
Category		All	Go	Critical ✖ 1	Warning ⚠ 2
Severity	Category	Name	Impact	Message	Alert Triggered
✖	Throughput	Number of Transactions (per second)		Metrics "User Transaction Per Sec" is at 34.80939	Feb 5, 2010 8:31:44 AM
⚠	User Audit	Audited User		User SYS logged on from edrsr10p1.us.oracle.com.	Feb 5, 2010 7:21:24 AM
⚠	Invalid Objects by Schema	Owner's Invalid Object Count		6 object(s) are invalid in the OE schema.	Jan 27, 2010 11:04:11 AM

Stop the workload by deleting the runload file from the `/home/oracle/workshops` directory.

Practice 6-1: Using Baselines (continued)

```
$ cd /home/oracle/workshops
$ rm runload
```

8. A short while later, the alert is cleared and no longer appears on the home page. The alert history shows that this metric had one or more critical alerts. View the alert history for this metric.
- a) Click Alert History in the Related Links section of the home page.



9. Clean up this practice by removing the adaptive threshold.

Step	Page	Action
a	Database Home page	Click Baseline Metric Thresholds in the Related Links section.
b	Baseline Metric Thresholds	Click “Number of Transactions (per second).”
c	Edit Thresholds: Number of Transactions (per second) page	Click Clear Thresholds.

Practice 7-1: Using AWR-Based Tools

In this practice, you use AWR-based tools to find and tune a database problem. Note that during this practice, you explicitly capture AWR snapshots. However, by default, the system automatically captures AWR snapshots every hour.

1. Execute the `./prepare 7 1` script to set up this practice. You can find this script in your `$HOME/workshops` directory. The script is displayed here to show the type of segment management that is specified for the tablespace.

```
$ cd $HOME/workshops
$ ./prepare 7 1

SQL> Connected.
SQL> Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area  836976640 bytes
Fixed Size                  1339740 bytes
Variable Size              805310116 bytes
Database Buffers           25165824 bytes
Redo Buffers                5160960 bytes
Database mounted.
Database opened.
SQL> drop tablespace tbsspc including contents and datafiles;

Tablespace dropped.

SQL>
SQL> CREATE SMALLFILE TABLESPACE "TBSSPC"
  2  DATAFILE 'tbsspc1.dbf' SIZE 50M
  3  AUTOEXTEND ON NEXT 10M MAXSIZE 200M
  4  LOGGING
  5  EXTENT MANAGEMENT LOCAL
  6  SEGMENT SPACE MANAGEMENT MANUAL;

Tablespace created.

SQL>
SQL> execute
dbms_workload_repository.modify_snapshot_settings(interval =>
1440);

PL/SQL procedure successfully completed.

SQL>
SQL> exec
dbms_advisor.set_default_task_parameter('ADDM','DB_ACTIVITY_MI
N',30);

PL/SQL procedure successfully completed.
```


Practice 7-1: Using AWR-Based Tools (continued)

```

SQL>
SQL> drop user spc cascade;

User dropped.

SQL>
SQL> create user spc identified by spc
      2 default tablespace tbsspc
      3 temporary tablespace temp;

User created.

SQL>
SQL> grant connect, resource, dba to spc;

Grant succeeded.

SQL>
SQL> connect spc/spc
Connected.
SQL>
SQL> drop table spct purge;
drop table spct purge
      *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> create table spct(id number, name varchar2(2000));

Table created.

SQL>
SQL> exec DBMS_STATS.GATHER_TABLE_STATS(-
> ownname=>'SPC', tabname=>'SPCT',-
> estimate_percent=>DBMS_STATS.AUTO_SAMPLE_SIZE);

PL/SQL procedure successfully completed.

SQL>
SQL> exit;
$

```

2. Take the first AWR snapshot. You can use either Enterprise Manager or the DBMS_WORKLOAD_REPOSITORY package.

Step	Page	Action
a	EM Database home	Click the Server tab.

Practice 7-1: Using AWR-Based Tools (continued)

b	Server	Click Automatic Workload Repository.	
c	Automatic Workload Repository	Click the Snapshots link.	
d	Snapshots	Click Create.	
e	Confirmation	Click Yes.	
f	Processing	Wait until the snapshot is created.	
g	Snapshots	Record the new snapshot number:	_____

3. Execute the `./workgen 7 1` script to generate the workload for this practice. The script is located in your `$HOME/workshops` directory. Wait until the “Load is finished” message appears, and then continue.

```
$ ./workgen 7 1
```

```
PL/SQL procedure successfully completed
```

```
...
```

```
Load is finished
```

```
$
```

4. Take a second AWR snapshot. You can use the `sol_07_01_02.sh` script located in your `$HOME/solutions` directory to take an AWR snapshot by using the `DBMS_WORKLOAD_REPOSITORY` package, or you can use Enterprise Manager to take a snapshot. Repeat the steps in step 2. Record the snapshot number: _____
5. Generate the AWR report. You can use the `awrrpt.sql` script located in the `$ORACLE_HOME/rdbms/admin` directory to create an AWR report. The solution in this document uses Enterprise Manager.

Step	Page	Action
a	Snapshots	Select the first snapshot number from step 2.
		Select View Report from Actions.
		Click Go.
b	View Report	Select the second snapshot number from step 4.
		Click OK.

Practice 7-1: Using AWR-Based Tools (continued)

c	Processing:View Report	Wait until the report is created.
d	Snapshot Details	

6. Review the AWR Report.

- a) What is the DB time? _____
 What is the elapsed time? _____
 Why is there a difference? _____

(DB Time includes the sum of all sessions in database calls, so DB Time could be as high as the product of the active sessions and the elapsed time.)

Snapshot Details

View A

Details

Report

58

WORKLOAD REPOSITORY report for

DB Name	DB Id	Instance	Inst num	Startup Time	Release	RAC
ORCL	1237161768	orcl	1	05-Feb-10 09:02	11.2.0.1.0	NO

Host Name	Platform	CPUs	Cores	Sockets	Memory (GB)
edrsr10p1.us.oracle.com	Linux IA (32-bit)	1	1	1	1.97

	Snap Id	Snap Time	Sessions	Cursors/Session
Begin Snap:	277	05-Feb-10 09:05:45	34	2.1
End Snap:	278	05-Feb-10 09:10:17	37	2.8
Elapsed:		4.54 (mins)		
DB Time:		13.46 (mins)		

- b) Scroll down to Top 5 Timed Foreground Events. (Or, use the Find command in your browser.)
 What is the top event? _____
 ("buffer busy waits" is expected.)

Practice 7-1: Using AWR-Based Tools (continued)**Top 5 Timed Foreground Events**

Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
buffer busy waits	877	398	454	49.24	Concurrency
log file sync	170	83	489	10.29	Commit
db file sequential read	4,835	52	11	6.43	User I/O
DB CPU		44		5.44	
free buffer waits	1,961	35	18	4.38	Configuration

- c) Scroll down to Time Model Statistics. (Or, select Main Report > Wait Event Statistics > Time Model Statistics.)

In which category is most of the DB time being collected?

(“sql execute elapsed time” is expected.)

Time Model Statistics

- Total time in database user-calls (DB Time): 807.9s
- Statistics including the word "background" measure background time statistic
- Ordered by % or DB time desc, Statistic name

Statistic Name	Time (s)	% of DB Time
sql execute elapsed time	711.88	88.12
parse time elapsed	66.64	8.25
hard parse elapsed time	66.04	8.17
DB CPU	43.95	5.44
connection management call elapsed time	31.81	3.94
PL/SQL compilation elapsed time	30.81	3.81
PL/SQL execution elapsed time	12.72	1.57
hard parse (sharing criteria) elapsed time	0.15	0.02
repeated bind elapsed time	0.06	0.01
hard parse (bind mismatch) elapsed time	0.00	0.00
sequence load elapsed time	0.00	0.00
DB time	807.89	
background elapsed time	139.93	
background cpu time	0.82	

Practice 7-1: Using AWR-Based Tools (continued)

- d) Scroll down to the Buffer Wait Statistics section. (Or, select Back to Top > Wait Statistics > Buffer Wait Statistics.)

What is the class of waits? _____ .
 (“Waits on data blocks” is the expected answer.)

Buffer Wait Statistics			
<ul style="list-style-type: none"> ordered by wait time desc, waits desc 			
Class	Waits	Total Wait Time (s)	Avg Time (ms)
data block	1,374	396	288
undo header	9	3	364
segment header	19	1	77

- e) Scroll down to “Segments by Buffer Busy Waits.” (Or, select Back to Top > Segment Statistics > Segments by Buffer Busy Waits.)

Which segment has the most buffer busy waits? _____ (SPCT)

Segments by Buffer Busy Waits						
<ul style="list-style-type: none"> % of Capture shows % of Buffer Busy Waits for each top segment compared with total Buffer Busy Waits for all segments captured by the Snapshot 						
Owner	Tablespace Name	Object Name	Subobject Name	Obj. Type	Buffer Busy Waits	% of Capture
SPC	TBSSPC	SPCT		TABLE	1,384	99.35
SYS	SYSTEM	AUD\$		TABLE	9	0.65

- f) At this point, you have determined that there are a large number of buffer busy waits on a single segment.

7. View the ADDM report that was generated when you created the snapshot in step 4.

Step	Page	Action		
a	Database home	Click Advisor Central in the Related Links section.		
b	Advisor Central	Enter search values:	Field	Value
			Advisory Type	ADDM

Practice 7-1: Using AWR-Based Tools (continued)

			Advisor Runs	Last 24 Hours
		Click Go.		

- c) Find and select the ADDM run that covers the snapshots you recorded in steps 2 and 4 by checking the description field of each ADDM report. The description will be in the “ADDM auto run: snapshots [177, 178], instance 1, database id 1159023676” format. Substitute your start and end snapshot numbers. (The latest ADDM report should be the one that meets the criteria.) Click View Result.
 - d) On the Automatic Database Diagnostic Monitor (ADDM) page, click the “Buffer Busy” finding. If more than one finding with this name appears on the page, check all of them.
 - e) On the Performance Finding Details page, you should see that ADDM recommends reorganizing the SPCT table to use Automatic Segment Space Management. Although you are able to infer this by looking at either a Statspack report or an AWR report, ADDM automatically tells you what to do. This recommendation may not appear. ADDM considers many factors before it gives a recommendation. A delay in creating the snapshot after the workload finishes is one reason that the recommendation may not appear.
8. View the ADDM findings from the Database home page.
 - a) Click the Database tab to return to the Database Instance home page. The latest ADDM report may not be visible until a few automatic refreshes take place. After a few automatic refreshes, you can see the latest performance analysis reported on the Database home page.
 - b) In the Performance Analysis section, click the “Buffer busy” finding. This shows the same page as in step 7c.
 9. Implement the recommendation.
 - a) To implement the suggested recommendation, execute the lab_07_01_09.sql script. This script moves the table used in the workload to a tablespace by using Automatic Segment Space Management. You can locate the lab_07_01_09.sql script in your \$HOME/labs directory.

```

$
$ sqlplus / as sysdba
SQL> @lab_07_01_09.sql
SQL>
SQL> drop tablespace tbsspc
      2 including contents and datafiles;

```

Practice 7-1: Using AWR-Based Tools (continued)

Tablespace dropped.

SQL>

```
SQL> CREATE SMALLFILE TABLESPACE "TBSSPC"
  2 DATAFILE 'tbsspc1.dbf' SIZE 50M
  3 AUTOEXTEND ON NEXT 10M MAXSIZE 200M
  4 LOGGING
  5 EXTENT MANAGEMENT LOCAL
  6 SEGMENT SPACE MANAGEMENT AUTO;
```

Tablespace created.

SQL>

SQL> connect spc/spc

Connected.

SQL>

SQL> drop table spct purge;

drop table spct purge

*

ERROR at line 1:

ORA-00942: table or view does not exist

SQL> create table spct(id number, name varchar2(2000))

2 tablespace tbsspc;

Table created.

SQL>

SQL> exec DBMS_STATS.GATHER_TABLE_STATS(-

> ownname=>'SPC', tabname=>'SPCT',-

> estimate_percent=>DBMS_STATS.AUTO_SAMPLE_SIZE);

PL/SQL procedure successfully completed.

SQL>

SQL> exit;

Disconnected from Oracle Database 11g Enterprise Edition

Release 11.1.0.6.0 - Production

With the Partitioning, Oracle Label Security, OLAP, Data

Mining

and Real Application Testing options

\$

10. Generate the workload again to observe the differences that the recommendation made. Execute the **./workgen 7 1** script to start generating the workload for this practice again.

\$ cd \$HOME/workshops

\$./workgen 7 1

PL/SQL procedure successfully completed

Practice 7-1: Using AWR-Based Tools (continued)

```
PL/SQL procedure successfully completed
PL/SQL procedure successfully completed
...
Load is finished
$
```

11. When the script finishes, take a third AWR snapshot. Record the snapshot number: _____
12. Generate a Compare Periods Report comparing the periods between the three snapshots.
 - a) You can use the `awrddrpt.sql` script located in your `$ORACLE_HOME/rdbms/admin` directory to generate the diff-diff report. The solution in this document uses Enterprise Manager. On the Snapshots page, select Compare Periods from the Actions drop-down list. Make sure that your first AWR snapshot is selected. (This snapshot number was recorded in step 2g.) Then, click Go.
 - b) On the Compare Periods: First Period End page, select your second AWR snapshot (the snapshot number recorded in step 4) and click Next.
 - c) On the Compare Periods: Second Period Start page, select your second AWR snapshot (the snapshot number recorded in step 4) and click Next.
 - d) On the Compare Periods: Second Period End page, select your last AWR snapshot (the snapshot number recorded in step 11) and click Next.
 - e) On the Compare Periods: Review page, click Finish.
13. Review the Compare Periods Report.
 - a) On the Compare Periods: Results page, select Per Transaction from the View Data drop-down list on the General tabbed page.
 - b) You can see the performance differences between the two analyzed periods. Click the Report tab.
 - c) The Processing page appears. Wait until the report is generated.
 - d) On the Report tabbed page, you can now see the Workload Repository Compare Period Report.
 - e) Scroll down to the Top Timed Events section.
What is the %DB time for buffer busy waits in the first period? _____
What is the %DB time for buffer busy waits in the second period?

 - f) In the Report Details section, click Wait Stats.
 - g) In the Wait Stats section, click Buffer Wait Statistics.

Practice 7-1: Using AWR-Based Tools (continued)

- h) In the Buffer Wait Statistics report:
What is the wait time % of DB time for data blocks for the first period?

What is the wait time % of DB time for data blocks for the second period?

- i) Click Back to Top to return to the Report Details section.
- j) In the Report Details section, click Segment Statistics.
- k) In the Segment Statistics section, click Top 5 Segments by Buffer Busy Waits.
- l) Examine the Top 5 Segments by Buffer Busy Waits report.
What is the number of buffer busy waits in the first period?

What is the number of buffer busy waits in the second period?

14. Generate the ASH Report from the Top Activity page for the period with the greatest load. You can use the `ashrpt.sql` script located in the `$ORACLE_HOME/rdbms/admin` directory to generate the corresponding ASH Report. The solution in this document uses Enterprise Manager.

Step	Page	Action
a	Database home	Click the Performance tab.
b	Performance	Note the time of the first peak in the Active Sessions Graph. Pick a 10-minute interval that brackets the peak.
		Click Run ASH Report.
c	Run ASH Report	Enter the Start time and End time from the times that you noted in the previous step. Click Generate Report.
d	Processing	Wait until the report is generated.

15. View the ASH report.
- a) From the first section, you can get various general details from the report.
- b) Scroll down to the Top User Events section.
What is the top event? _____ (buffer busy waits is in the top events.)
- c) Scroll down to the Top SQL Command Types section.
What is the type of the top SQL statement? _____ (INSERT)

Practice 7-1: Using AWR-Based Tools (continued)

- d) Scroll down to the Top SQL with Top Events section. What is the top SQL statement and what was the primary wait event for that statement?

What is the SQL ID? _____ (3csh3g3mjhmzh)

- e) Scroll down to the Activity Over Time section.
How many periods are reported? _____ (11, maybe less, some machines show only 1 period, when the activity otherwise is very low.)
What is the elapsed time of each period? _____ (This varies, usually 1 minute, except that the first and last periods may be shorter to allow periods to fall on minute boundaries.)
Which periods show buffer busy waits? _____ (possibly none if buffer busy waits are not the top event)
These answers may vary greatly.

16. Clean up this practice by running the `./cleanup 7 1` script.

```
$ ./cleanup 7 1

Tablespace dropped.

User dropped.

PL/SQL procedure successfully completed.
```

Practices for Lesson 8

In this practice, you monitor applications with services.

Practice 8-1: Using Services with a Single-Instance Oracle Database

The first step in creating a service configuration is to plan services for each application or application function that uses your database.

In this practice, you create the following configuration in the `orcl` database:

Service Name	Usage	Response Time (sec)– Warning/Critical
SERV1	Client service	0.4, 1.0

- 1) Use the `DBMS_SERVICE` package to create a service called `SERV1`. Then, make sure that you add your service name to your `tnsnames.ora` file.
 - a) The recommended method of adding a service name to the `tnsnames.ora` file is to use Net Manager. For this exercise, execute the `lab_08_01_01.sh` script. Review the `tnsnames.ora` file at `$ORACLE_HOME/network/admin` to confirm that the following lines are included. Substitute the output of the `hostname` command for `<hostname>` in the following screenshot:

```
SERV1 =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP)
      (HOST = <hostname>.us.oracle.com)(PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = SERV1.us.oracle.com)
    )
  )
```

```
$ cd /home/oracle/labs
$ ./lab_08_01_01.sh
edrsr3p1
$
```

- 2) Use the `DBMS_SERVICE.CREATE_SERVICE` procedure to create a service.

```
$ sqlplus / as sysdba

SQL> EXEC -
DBMS_SERVICE.CREATE_SERVICE('SERV1','SERV1.us.oracle.com')

PL/SQL procedure successfully completed.

SQL> exit;
```

Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

- 3) After you have created your services, try connecting to your database by using your service name. What happens and why? _____
(You cannot connect by using your service because, although it is defined, it is not started on your instance. You can verify this by looking at the SERVICE_NAMES initialization parameter, and by looking at the services known to the listener.)

```
$ sqlplus / as sysdba

SQL> show parameter service

NAME                                TYPE        VALUE
-----
service_names                       string      orcl.oracle.com

SQL> connect system@SERV1
Enter password: oracle_4U /* password is not displayed */
ERROR:
ORA-12514: TNS:listener does not currently know of service
requested in connect
descriptor

Warning: You are no longer connected to ORACLE.
SQL> exit

$ lsnrctl services

Connecting to
(DESCRIPTION=(ADDRESS=(PROTOCOL=IPC)(KEY=EXTPROC1521)))
Services Summary...
Service "orcl" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1465 refused:0 state:ready
        LOCAL SERVER
Service "orcl.us.oracle.com" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1465 refused:0 state:ready
        LOCAL SERVER
Service "orclXDB" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "D000" established:0 refused:0 current:0 max:1022
state:ready
```

Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

```
DISPATCHER <machine: edrsr9p0.us.oracle.com, pid:
3998>

(ADDRESS=(PROTOCOL=tcp)(HOST=edrsr9p0.us.oracle.com)(PORT=3277
4))
Service "orcl_XPT" has 1 instance(s).
Instance "orcl", status READY, has 1 handler(s) for this
service...
Handler(s):
"DEDICATED" established:1465 refused:0 state:ready
LOCAL SERVER
The command completed successfully
$
```

- 4) How would you make sure that you can connect using your service? Attempt your solution, and connect to your instance using your service.
- a) **Answer:** You must start the `SERV1` service on your instance.
- Note:** If the service does not appear in the `lsnrctl services` listing immediately, run the `ALTER SYSTEM REGISTER;` command, and then retry the connection by running the `connect system/oracle@SERV1` command.

```
$ sqlplus / as sysdba

SQL*Plus: Release 11.1.0.6.0 - Production on Thu Mar 27
09:10:31 2008

Copyright (c) 1982, 2007, Oracle. All rights reserved.

Connected to:
Oracle Database 11g Enterprise Edition Release 11.1.0.6.0 -
Production
With the Partitioning, Oracle Label Security, OLAP, Data
Mining
and Real Application Testing options

SQL> show parameter service

NAME                                TYPE                                VALUE
-----
service_names                       string                             orcl.us.oracle.com

SQL> EXEC DBMS_SERVICE.START_SERVICE('SERV1');

PL/SQL procedure successfully completed.

SQL> show parameter service

NAME                                TYPE                                VALUE
-----
```

Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

```
-----
service_names      string      orcl.us.oracle.com, SERV1.us.o
                        racle.com

SQL> host lsnrctl services

LSNRCTL for Linux: Version 11.1.0.6.0 - Production on 27-MAR-
2008 09:13:39

Copyright (c) 1991, 2007, Oracle. All rights reserved.

Connecting to
(DESCRIPTION=(ADDRESS=(PROTOCOL=IPC)(KEY=EXTPROC1521)))
Services Summary...
Service "SERV1.us.oracle.com" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1 refused:0 state:ready
      LOCAL SERVER
Service "orcl" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1 refused:0 state:ready
      LOCAL SERVER
Service "orcl.us.oracle.com" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1 refused:0 state:ready
      LOCAL SERVER
Service "orclXDB" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "D000" established:0 refused:0 current:0 max:1022
state:ready
      DISPATCHER <machine: edrsr9p0.us.oracle.com, pid:
3998>

(ADDRESS=(PROTOCOL=tcp)(HOST=edrsr9p0.us.oracle.com)(PORT=3277
4))
Service "orcl_XPT" has 1 instance(s).
  Instance "orcl", status READY, has 1 handler(s) for this
service...
    Handler(s):
      "DEDICATED" established:1 refused:0 state:ready
      LOCAL SERVER
The command completed successfully
```

Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

```
$ sqlplus /nolog
SQL> connect SYSTEM@SERV1
Enter password: oracle_4U /* Password is not displayed */
Connected.
SQL> exit
Disconnected from Oracle Database 11g Enterprise Edition
Release 11.1.0.6.0 - Production
With the Partitioning, Oracle Label Security, OLAP, Data
Mining
and Real Application Testing options
```

- 5) From the /home/oracle/workshops directory, execute the ./workgen 8 1 script. This workload will run until stopped; continue to the next step. This script starts three workloads, each connecting with a different service name: SERV1, ORCL, and SYS\$.

```
$ cd /home/oracle/workshops
$ ./workgen 8 1
$ ps
  PID TTY          TIME CMD
 2386 pts/4        00:00:00 bash
 23752 pts/4        00:00:00 workgen
 23753 pts/4        00:00:00 workgen
 23754 pts/4        00:00:00 wksh_w81_orcl.s
 23756 pts/4        00:00:00 sqlplus
 23781 pts/4        00:00:00 insert_orders.s
 23794 pts/4        00:00:00 sqlplus
 23797 pts/4        00:00:00 update_orders.s
 23805 pts/4        00:00:00 delete_orders.s
 23811 pts/4        00:00:00 insert_orders.s
 23816 pts/4        00:00:00 sleep
 23818 pts/4        00:00:00 update_orders.s
 23823 pts/4        00:00:00 sleep
 23824 pts/4        00:00:00 delete_orders.s
 23829 pts/4        00:00:00 sleep
 23834 pts/4        00:00:00 sleep
 23835 pts/4        00:00:00 insert_orders.s
 23843 pts/4        00:00:00 sleep
 23844 pts/4        00:00:00 sleep
 23846 pts/4        00:00:00 update_orders.s
 23847 pts/4        00:00:00 sleep
 23851 pts/4        00:00:00 sleep
 23852 pts/4        00:00:00 sqlplus
 23855 pts/4        00:00:00 ps
$
```


Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

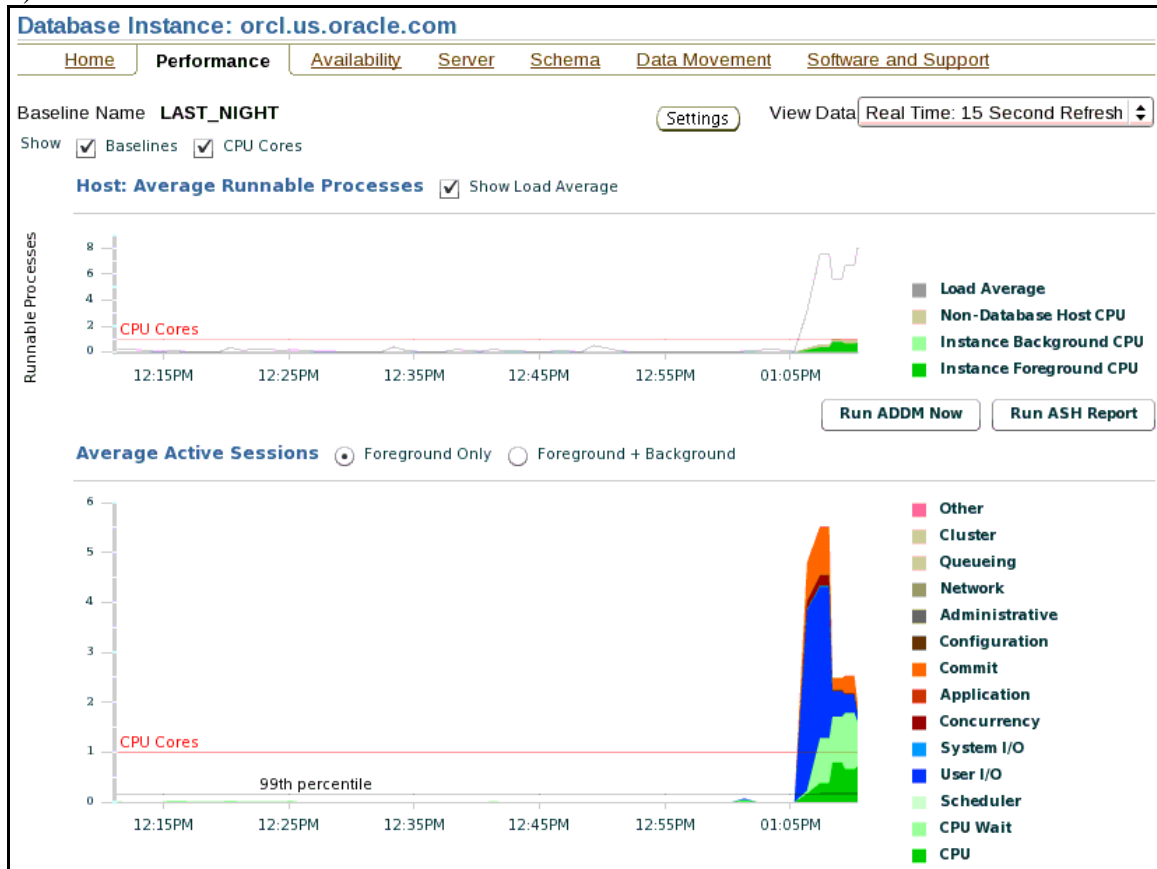
- 6) After the execution starts, access the EM Top Consumers page from the Performance tabbed page, and observe that SERV1 is using some resources. Also, check the statistics on your service with V\$SERVICE_STATS from a SQL*Plus session connected as SYSDBA.

Note: Screenshots of the steps follow the table.

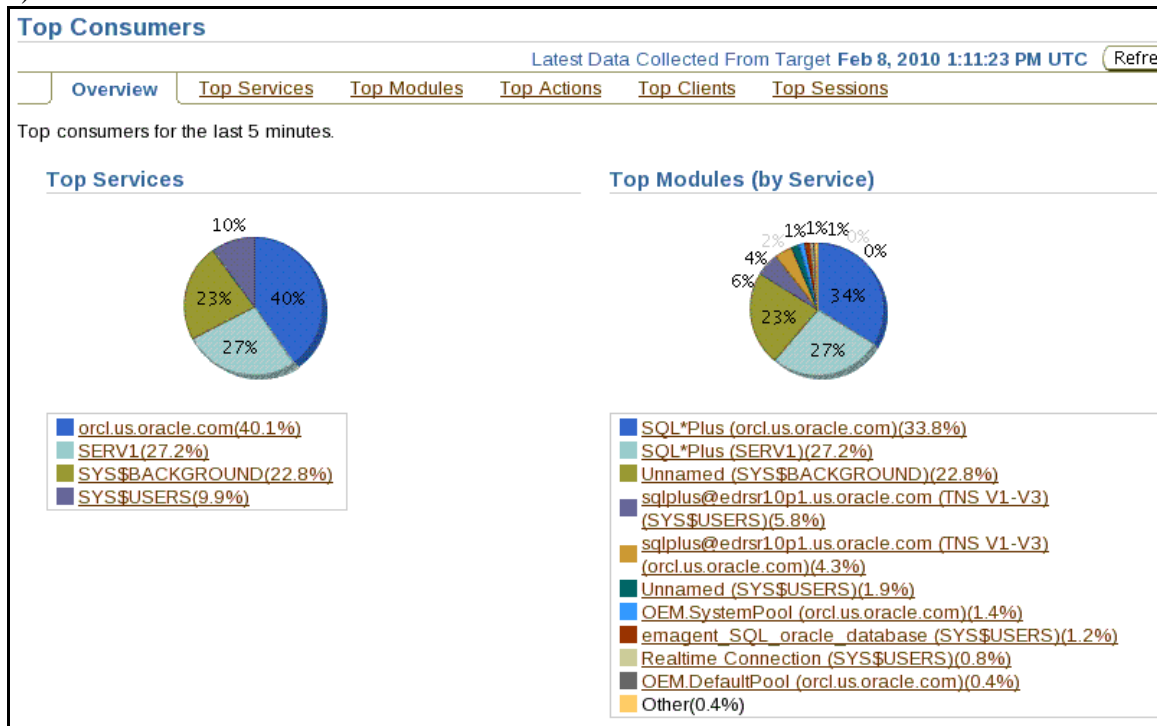
Step	Page	Action	Observation
a	Database Instance home	Click the Performance tab.	
b	Performance (See the example screenshot that follows the table.)		On the Performance page, wait at this point until the Average Active Sessions graph shows a significant rise.
		Click Top Consumers in the Additional Monitoring Links section.	
c	Top Consumers Overview (See the example screenshot, which follows the table.)	Click Refresh.	The names and number of services listed in the Top Services graph depend on the number and type of connections to the database. The network service name of each connection is recorded as a separate service. As a result, all the connections made without a service name will be aggregated, as will all the connections made as SERV1.
		Click the Top Services tab.	
d	Top Services	Click the SERV1 link.	
e	Service:SERV1 Modules tab	Click the Statistics tab.	
f	Service:SERV1 Statistics tab		Notice that the statistics are the same as the AWR report, but aggregated by service.

Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

b)



c)

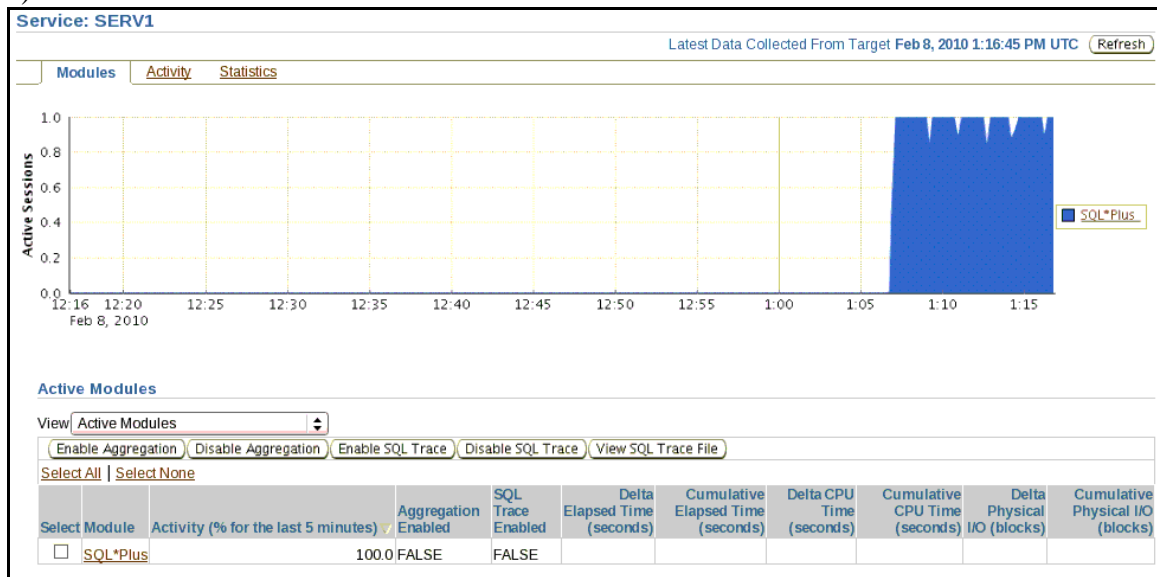


Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

d)

Top Consumers									
Latest Data Collected From Target Feb 8, 2010 1:13:23 PM UTC									
Overview Top Services Top Modules Top Actions Top Clients Top Sessions									
View <input type="text" value="Active Services"/>									
<input type="button" value="Enable SQL Trace"/> <input type="button" value="Disable SQL Trace"/> <input type="button" value="View SQL Trace File"/>									
Select All Select None									
Select	Service	Activity (% for the last 5 minutes)	SQL Trace Enabled	Delta Elapsed Time (seconds)	Cumulative Elapsed Time (seconds)	Delta CPU Time (seconds)	Cumulative CPU Time (seconds)	Delta Physical I/O (blocks)	Cumulative Physical I/O (blocks)
<input type="checkbox"/>	SERV1	42.6	FALSE	16	388	8	189	32	
<input type="checkbox"/>	orcl.us.oracle.com	32.1	FALSE	11	1341	3	430	352	
<input type="checkbox"/>	SYS\$BACKGROUND	20.6	FALSE	0	0	0	0	560	
<input type="checkbox"/>	SYS\$USERS	4.7	FALSE	0	3760	0	1720	24	

e)



Oracle University and En-Sof Informatica E Treinamento Ltda use only

Database Instance: [ord.us.oracle.com](#) > [Top Consumers](#) >

View Data

Real Time: 15 Second Refresh

Service: **SERV1**

Latest Data Collected From Target **Feb 8, 2010 1:18:00 PM UTC**

Refresh

Modules

Activity

Statistics

Previous

1-25 of 28

Next 3

Name	Delta Value	Cumulative Value
logons cumulative	0	9
user calls	0	170
DB time	13801585	656699136
DB CPU	7770819	330949696
parse count (total)	0	197
parse time elapsed	0	3502068
execute count	0	226
sql execute elapsed time	13801086	653518912
opened cursors cumulative	0	230
session logical reads	1426	43289
physical reads	1135	38112
physical writes	0	0
redo size	0	46684
user commits	0	0

- ```
$ rm $HOME/workshops/runload
```

- | Step | Page                       | Action                                                                                                                 |
|------|----------------------------|------------------------------------------------------------------------------------------------------------------------|
| a    | Database home              | Click the Metric and Policy Settings link in the Related Links section.                                                |
| b    | Metric and Policy Settings | Select All Metrics in View. (The page will refresh.)                                                                   |
|      |                            | Click the multiedit icon (the group of pencils) for the “Service Response Time (per user call) (microseconds)” metric. |

## Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

|   |                                                                                            |                                  |                       |           |
|---|--------------------------------------------------------------------------------------------|----------------------------------|-----------------------|-----------|
| c | Edit Advanced Settings:<br>Service Response Time (per<br>user call) (microseconds)         | Click Add.                       |                       |           |
| d | Specify Multiple<br>Thresholds: Service<br>Response Time (per user<br>call) (microseconds) | Specify the<br>threshold values: | Fields                | Values    |
|   |                                                                                            |                                  | Service Name          | SERV1     |
|   |                                                                                            |                                  | Warning<br>Threshold  | 40000000  |
|   |                                                                                            |                                  | Critical<br>Threshold | 100000000 |
|   |                                                                                            | Click Continue.                  |                       |           |
| e | Metric and Policy Settings                                                                 | Click OK.                        |                       |           |
| f | Confirmation                                                                               | Click OK.                        |                       |           |

9) Verify the metric settings.

| Step | Page                                                                     | Action                                                        |
|------|--------------------------------------------------------------------------|---------------------------------------------------------------|
| a    | Database Control home                                                    | Click the All Metrics link in the Related Links section.      |
| b    | All Metrics                                                              | Expand Database Services.                                     |
|      |                                                                          | Click “Service Response Time (per user call) (microseconds).” |
| c    | Service Response Time (per user call) (microseconds)                     | Click the SERV1 link in the Service Name column.              |
| d    | Service Response Time (per user call) (microseconds): Service Name SERV1 | Select Real Time: 30 Second Refresh from the View Data.       |

10) From your terminal emulator session, execute the  
`/home/oracle/labs/lab_08_01_10.sh` script. This script executes a workload under the SERV1 service. Observe the Metric Value graph on the Service Response Time (per user call) (microseconds): Service Name SERV1 page.

- Execute the script. It takes a while to complete. Do not wait until it completes. Look at the corresponding threshold history by using Database Control.




## Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)

```
$ cd $HOME/labs
$./lab_08_01_10.sh
```

11) While the workload is executing, observe the Service Response Time graph for the SERV1 service. You may have to wait a couple of minutes for the graph to respond.

| Step | Page                                                                     | Action                                                                     | Answer                                                                                                              |
|------|--------------------------------------------------------------------------|----------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------|
| a    | Service Response Time (per user call) (microseconds): Service Name SERV1 | Monitor the Metric Value graph.                                            | The graph shows the metric for the SERV1 service.                                                                   |
|      |                                                                          | What conclusion can you draw?                                              | Services can be monitored individually.                                                                             |
|      |                                                                          | When an Alert is shown, navigate to the home page. Click the Database tab. |                                                                                                                     |
| b    | Database Instance home                                                   | View Alerts.                                                               | An alert for the Service Response Time for the SERV1 service is present. See the screenshot that follows the table. |

b)

| Alerts                                                                              |                           |                                                      |        |                                                                         |                          |
|-------------------------------------------------------------------------------------|---------------------------|------------------------------------------------------|--------|-------------------------------------------------------------------------|--------------------------|
| Category                                                                            |                           | All                                                  | Go     | Critical 0                                                              | Warning 3                |
| Severity                                                                            | Category                  | Name                                                 | Impact | Message                                                                 | Alert Triggered          |
|  | User Audit                | Audited User                                         |        | User SYS logged on from edrsr10p1.us.oracle.com.                        | Feb 8, 2010 6:36:10 AM   |
|  | Database Services         | Service Response Time (per user call) (microseconds) |        | Metrics "Elapsed Time Per User Call" is at 59289022 for service "SERV1" | Feb 8, 2010 1:33:12 PM   |
|  | Invalid Objects by Schema | Owner's Invalid Object Count                         |        | 6 object(s) are invalid in the OE schema.                               | Jan 27, 2010 11:04:11 AM |

## ***Practice 8-1: Using Services with a Single-Instance Oracle Database (continued)***

12) Use Database Control to remove the thresholds that you specified during this practice.

| Step | Page                                                                         | Action                                                                                          |
|------|------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| a    | Database Control home                                                        | Click Metric and Policy Settings in the Related Links section.                                  |
| b    | Metrics and Policies                                                         | Click the multiedit icon on the right for Service Response Time (per user call) (microseconds). |
| c    | Edit Advanced Settings: Service Response Time (per user call) (microseconds) | Select SERV1.                                                                                   |
|      |                                                                              | Click Remove.                                                                                   |
|      |                                                                              | Click Continue.                                                                                 |
| d    | Metrics and Policies Settings                                                | Click OK.                                                                                       |
| e    | Confirmation                                                                 | Click OK.                                                                                       |

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment

Unless specified differently, log in as SYSDBA in your Database Control Console.

- 1) Execute the `lab_08_02_01.sh` script to remove all the trace files already generated in your `$ORACLE_BASE/diag/rdbms/orcl/orcl/trace` directory.

```
$ cd $HOME/labs
$./lab_08_02_01.sh

trace_dir=$ORACLE_BASE/diag/rdbms/orcl/orcl/trace

mv $trace_dir $trace_dir.old
mkdir $trace_dir
mv $trace_dir.old/alert_orcl.log $trace_dir/
rm -rf $ORACLE_BASE/diag/rdbms/orcl/orcl/trace.old

$
```

- 2) Execute the `start_servwork.sh` script to start four sessions by using the `SERV1` service.

```
$./start_servwork.sh
Started stream with pid=7904
Started stream with pid=7905
Started stream with pid=7906
Started stream with pid=7907
$
PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

$
```



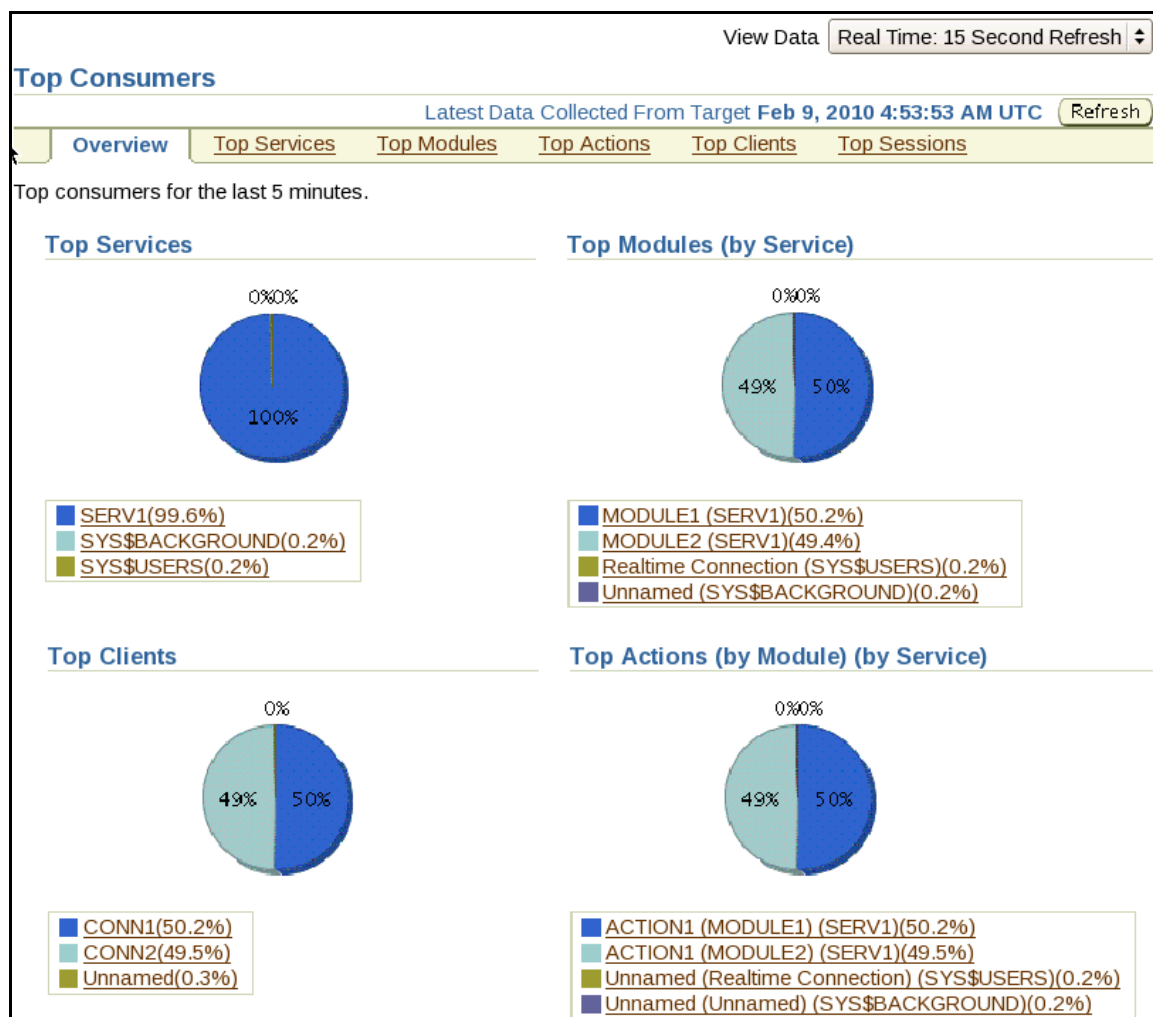
## ***Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)***

- 3) Using Database Control, determine the list of services, modules, and actions that the workload is using.

| Step | Page                                                                    | Action/Observe                                                                                                                                                                                               |
|------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a    | Database Control home                                                   | Click the Performance tab.                                                                                                                                                                                   |
| b    | Performance                                                             | Click the Top Consumers link in the Additional Monitoring Links section.                                                                                                                                     |
| c    | Top Consumers<br>(See the example screenshot, which follows the table.) | You should see that the current workload uses only the SERV1 service, and that half of the sessions are using the SERV1/MODULE1/ACTION1 action and the other half is using the SERV1/MODULE2/ACTION1 action. |

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)

The following screenshot displays the Top Consumers page:



- 4) Using Database Control, enable statistics aggregation for both SERV1/MODULE1/ACTION1 and SERV1/MODULE2/ACTION1.

| Step | Page          | Action/Observation                         |
|------|---------------|--------------------------------------------|
| a    | Top Consumers | Click the Top Actions tab.                 |
| b    | Top Actions   | Select the SERV1 /MODULE1 /ACTION1 action. |
|      |               | Click Enable Aggregation.                  |
| c    | Top Actions   | Select the SERV1 /MODULE2 /ACTION1 action. |
|      |               | Click Enable Aggregation.                  |

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)

|   |             |                                                                                                                                                                 |
|---|-------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| d | Top Actions | You should see that additional statistics at the right of the window are now tracked for both actions (may require a refresh to see the additional statistics). |
|---|-------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|

d)

| Top Consumers                                                                                                                                                                 |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|---------|---------|-------------------------------------|---------------------|-------------------|------------------------------|-----------------------------------|--------------------------|-------------------------------|
| Latest Data Collected From Target Feb 9, 2010 4:57:38 AM UTC <span>Refresh</span>                                                                                             |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
| <a href="#">Overview</a> <a href="#">Top Services</a> <a href="#">Top Modules</a> <a href="#">Top Actions</a> <a href="#">Top Clients</a> <a href="#">Top Sessions</a>        |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
| View: <span>Active Actions</span>                                                                                                                                             |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
| <a href="#">Enable Aggregation</a> <a href="#">Disable Aggregation</a> <a href="#">Enable SQL Trace</a> <a href="#">Disable SQL Trace</a> <a href="#">View SQL Trace File</a> |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
| <a href="#">Select All</a>   <a href="#">Select None</a>                                                                                                                      |                 |         |         |                                     |                     |                   |                              |                                   |                          |                               |
| Select                                                                                                                                                                        | Service         | Module  | Action  | Activity (% for the last 5 minutes) | Aggregation Enabled | SQL Trace Enabled | Delta Elapsed Time (seconds) | Cumulative Elapsed Time (seconds) | Delta CPU Time (seconds) | Cumulative CPU Time (seconds) |
| <input checked="" type="checkbox"/>                                                                                                                                           | SERV1           | MODULE1 | ACTION1 | 50.0                                | TRUE                | FALSE             | 0                            | 0                                 | 0                        | 0                             |
| <input checked="" type="checkbox"/>                                                                                                                                           | SERV1           | MODULE2 | ACTION1 | 50.0                                | TRUE                | FALSE             | 0                            | 0                                 | 0                        | 0                             |
| <input type="checkbox"/>                                                                                                                                                      | SYS\$BACKGROUND | Unnamed | Unnamed | .1                                  | FALSE               | FALSE             |                              |                                   |                          |                               |

- 5) Using Database Control, enable tracing for both SERV1/MODULE1/ACTION1 and SERV1/MODULE2/ACTION1. Let the system generate the trace files for one minute, and then disable tracing for both SERV1/MODULE1/ACTION1 and SERV1/MODULE2/ACTION1. After this is done, determine the list of generated trace files.

| Step | Page             | Action                                            |
|------|------------------|---------------------------------------------------|
| a    | Top Actions      | Select Actions with Aggregation Enable from View. |
|      |                  | Select the SERV1/MODULE1/ACTION1 action.          |
|      |                  | Click Enable SQL Trace.                           |
| b    | Enable SQL Trace | Click OK.                                         |
| c    | Top Actions      | Select Actions with Aggregation Enable from View. |
|      |                  | Select the SERV1/MODULE2/ACTION1 action.          |
|      |                  | Click Enable SQL Trace.                           |
| d    | Enable SQL Trace | Click OK.                                         |

- 6) Find the trace files. Use the following command:

```
$ ls -lt $ORACLE_BASE/diag/rdbms/orcl/orcl/trace/*.trc
```

```
$ ls -lt $ORACLE_BASE/diag/rdbms/orcl/orcl/trace/*.trc
-rw-r----- 1 oracle oinstall 10617510 Mar 30 07:13 orcl_ora_3952.trc
-rw-r----- 1 oracle oinstall 10827320 Mar 30 07:13 orcl_ora_3954.trc
```

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)

```
-rw-r----- 1 oracle oinstall 1624 Mar 30 07:12 orcl_dbrm_1633.trc
-rw-r----- 1 oracle oinstall 271200169 Mar 30 07:12 orcl_ora_3950.trc
-rw-r----- 1 oracle oinstall 253132737 Mar 30 07:11 orcl_ora_3959.trc
-rw-r----- 1 oracle oinstall 910 Mar 30 07:01 orcl_m000_3272.trc
$
```

- 7) Stop your workload by executing the `stop_servwork.sh` script, and then start it again by executing the `start_servwork.sh` script. Then check whether the statistics for your enabled aggregation actions are increasing.

```
$./stop_servwork.sh
Killing stream with pid=3067
Killing stream with pid=3064
Killing stream with pid=3063
Killing stream with pid=3062
$
$./start_servwork.sh
Started stream with pid=8535
Started stream with pid=8536
Started stream with pid=8539
Started stream with pid=8540
$
PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

$
```

- 8) Look again at your Top Actions page. You should see statistics refreshing for your two actions.
- 9) Using Database Control, disable statistic aggregations for both the actions.

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)

- a) From the Top Actions page, select both the actions, and click the Disable Aggregation button.

10) Stop the workload by using the `stop_servwork.sh` script.

```
$./stop_servwork.sh
Killing stream with pid=3944
Killing stream with pid=3942
Killing stream with pid=3941
Killing stream with pid=3940
$
```

11) Use the trace session utility and `tkprof` to interpret the generated trace files for the `SERV1` service and the `MODULE1` module.

```
$ cd $ORACLE_BASE/diag/rdbms/orcl/orcl/trace
$ trcsess output=serv1_module1.trc service=SERV1 *ora*.trc
$ tkprof serv1_module1.trc results.trc

TKPROF: Release 11.2.0.1.0 - Development on Tue Feb 9 05:08:10
2010

Copyright (c) 1982, 2009, Oracle and/or its affiliates. All
rights reserved.
```

12) View the output of `tkprof` with your favorite viewer or editor; the viewer `less` is shown.

```
$ less results.trc

TKPROF: Release 11.2.0.1.0 - Development on Tue Feb 9 05:08:10 2010

Copyright (c) 1982, 2009, Oracle and/or its affiliates. All rights
reserved.

Trace file: serv1_module1.trc
Sort options: default

count = number of times OCI procedure was executed
cpu = cpu time in seconds executing
elapsed = elapsed time in seconds executing
disk = number of physical reads of buffers from disk
query = number of buffers gotten for consistent read
current = number of buffers gotten in current mode (usually for update)
rows = number of rows processed by the fetch or execute call

SQL ID: 7r0kgzntdn7sq
Plan Hash: 2945320129
SELECT COUNT(*)
FROM DBA_OBJECTS
```

## Practice 8-2: Tracing Services in a Single-Instance Oracle Environment (continued)

| call    | count | cpu    | elapsed | disk    | query   | current | rows |
|---------|-------|--------|---------|---------|---------|---------|------|
| Parse   | 2     | 0.01   | 0.02    | 0       | 0       | 0       | 0    |
| Execute | 1240  | 0.05   | 0.05    | 0       | 0       | 0       | 0    |
| Fetch   | 1235  | 289.61 | 1601.80 | 1123850 | 2274870 | 0       | 1235 |
| total   | 2477  | 289.68 | 1601.88 | 1123850 | 2274870 | 0       | 1235 |

Misses in library cache during parse: 1  
Optimizer mode: ALL\_ROWS  
Parsing user id: 122 (recursive depth: 1)

Elapsed times include waiting on following events:

| Event waited on             | Times<br>Waited | Max. Wait | Total Waited |
|-----------------------------|-----------------|-----------|--------------|
| direct path read            | 34              | 1.12      | 12.85        |
| kfk: async disk IO          | 141171          | 0.51      | 3.53         |
| latch: cache buffers chains | 17              | 0.51      | 5.13         |
| asynch descriptor resize    | 690             | 0.00      | 0.00         |
| latch free                  | 5               | 0.50      | 1.80         |

\*\*\*\*\*

SQL ID: 1hb6a7lhd3d53  
Plan Hash: 0  
BEGIN dbms\_application\_info.set\_module('MODULE1','ACTION1'); END;

| call    | count | cpu  | elapsed | disk | query | current | rows |
|---------|-------|------|---------|------|-------|---------|------|
| Parse   | 0     | 0.00 | 0.00    | 0    | 0     | 0       | 0    |
| Execute | 2     | 0.00 | 0.00    | 0    | 0     | 0       | 2    |
| Fetch   | 0     | 0.00 | 0.00    | 0    | 0     | 0       | 0    |
| total   | 2     | 0.00 | 0.00    | 0    | 0     | 0       | 2    |

Misses in library cache during parse: 0  
Optimizer mode: ALL\_ROWS  
Parsing user id: 122  
...  
(type q to exit)

13) Clean up the database environment. Run the lab\_08\_02\_13.sh script.

```
$ cd $HOME/labs
$./lab_08_02_13.sh

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
$
```

## Practices for Lesson 9

In this practice, you examine and compare the information provided by some built-in diagnostic tools. You start with a poorly performing SQL statement and examine the output of the various tools to try to determine the reasons for poor performance.

Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN

Practice 9-2: Using the SQL TRACE facility

### **Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN**

In this practice, you examine a bad SQL statement and gather information about it by using the SQL\*Plus AUTOTRACE command.

1. Create the PLUSTRACE role. To use the SQL\*Plus AUTOTRACE command, the PLUSTRACE role must be granted to the user that wishes to use AUTOTRACE.

```
$ sqlplus / as sysdba

SQL> @?/sqlplus/admin/plustrce.sql
SQL>
SQL> drop role plustrace;
drop role plustrace
 *
ERROR at line 1:
ORA-01919: role 'PLUSTRACE' does not exist

SQL> create role plustrace;

Role created.

SQL>
SQL> grant select on v_$sesstat to plustrace;

Grant succeeded.

SQL> grant select on v_$statname to plustrace;

Grant succeeded.

SQL> grant select on v_$mystat to plustrace;

Grant succeeded.

SQL> grant plustrace to dba with admin option;

Grant succeeded.

SQL>
SQL> set echo off
SQL>
```

2. Grant the PLUSTRACE role to the SH user.

```
SQL> grant plustrace to sh;

Grant succeeded.
```



**Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN (continued)**

```
SQL> exit
```

3. The SQL statement has been isolated, and you can view it in the `bad_sql.sql` file. The script is in the `$HOME/labs/sqltune` directory. What can be determined by inspecting this statement? Each observation produces more specific questions.

```
$ cd $HOME/labs/sqltune
$ cat bad_sql.sql
select time_id, QUANTITY_SOLD, AMOUNT_SOLD
 from sales s, customers c
 where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
 order by time_id;
```

- a) This statement has a join between two tables: SALES and CUSTOMERS.

**Information:** These are large tables in the Sales History (SH) schema.

**Question:** Are the statistics up-to-date on these tables?

- b) The join column is `cust_id` in both tables.

**Questions:** Are there indexes on these columns?

Do the indexes have up-to-date statistics?

Are the indexes built properly?

- c) The selection predicate is `CUST_FIRST_NAME='Dina'`.

**Questions:** Does `CUST_FIRST_NAME` have an index?

- d) The ORDER BY clause may require a sort.

**Question:** Is the sort area in memory sized properly?

4. Determine the information provided by AUTOTRACE about the execution of this statement.

**Note:** The Explain Plan output has been reformatted for readability.

- a) Run the `bad_sql.sql` script with AUTOTRACE ON.

```
$ sqlplus /nolog

SQL> connect sh/sh
Connected.
SQL> set autotrace on
SQL> @bad_sql.sql

...
TIME_ID QUANTITY_SOLD AMOUNT_SOLD

24-DEC-01 1 32.27
24-DEC-01 1 33.25
```

## Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN (continued)

```
28-DEC-01 1 56.96
28-DEC-01 1 56.96
```

895 rows selected.

Execution Plan

-----  
Plan hash value: 1897253553

```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU)| Time | Pstart | Pstop |

0	SELECT STATEMENT		5557	179K		948 (2)	00:00:12		
1	SORT ORDER BY		5557	179K	272K	948 (2)	00:00:12		
* 2	HASH JOIN		5557	179K		897 (2)	00:00:11		
* 3	TABLE ACCESS FULL	CUSTOMERS	43	516		405 (1)	00:00:05		
4	PARTITION RANGE ALL		918K	18M		498 (2)	00:00:06	1	28
5	TABLE ACCESS FULL	SALES	918K	18M		498 (2)	00:00:06	1	28

```

Predicate Information (identified by operation id):

```

2 - access("C"."CUST_ID"="S"."CUST_ID")
3 - filter("CUST_FIRST_NAME"='Dina')
```

Statistics

```

13701 recursive calls
0 db block gets
6083 consistent gets
3213 physical reads
0 redo size
23039 bytes sent via SQL*Net to client
1068 bytes received via SQL*Net from client
61 SQL*Net roundtrips to/from client
103 sorts (memory)
0 sorts (disk)
895 rows processed
```

SQL> **set autotrace off**

b) What objects are accessed? \_\_\_\_\_

(The SALES and CUSTOMERS tables)

**Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN (continued)**

- c) What object is accessed first? \_\_\_\_\_  
(The CUSTOMERS table is read into a hash table.)
- d) Which object is partitioned? \_\_\_\_\_  
(The SALES table is partitioned.)
- e) How many partitions are accessed? \_\_\_\_\_  
(All 28 partitions)
- f) How many rows are returned (Cardinality)? \_\_\_\_\_  
(895 rows)
- g) How many rows does the execution plan expect? \_\_\_\_\_  
(5557 rows)
- h) What is the cost? \_\_\_\_\_  
(948, may vary due to specific hardware)
- i) Are any indexes used? \_\_\_\_\_  
(No)
- j) How much sort space is expected? \_\_\_\_\_  
(179 KB)
- k) How many consistent reads (logical reads) were done? \_\_\_\_\_  
(6083)
- l) How many physical reads were done?  
\_\_\_\_\_  
(6120, may vary due to the preceding activity)
- m) Did the sort spill to disk? \_\_\_\_\_  
(No, sorts, disk was 0)
5. Use DBMS\_XPLAN.DISPLAY to just get the explain plan for the SQL statement. The command requires that the statement be explained with the EXPLAIN PLAN command; then, the plan table is displayed. The terminal window must be resized and set for at least a 108-character width to display properly.
- a) Generate an explain plan.

```
SQL> EXPLAIN PLAN FOR
2> select time_id, QUANTITY_SOLD, AMOUNT_SOLD
3> from sales s, customers c
4> where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
5> order by time_id;

Explained
```

## Practice 9-1: Using AUTOTRACE and DBMS\_XPLAN (continued)

- b) Display the plan table. Note that this is the same explain plan that you received from the AUTOTRACE command.

```
SQL> SET LINESIZE 130
SQL> SET PAGESIZE 0
SQL> SELECT * FROM table(DBMS_XPLAN.DISPLAY);
Plan hash value: 1897253553
```

| Id  | Operation           | Name      | Rows | Bytes | TempSpc | Cost (%CPU) | Time     | Pstart | Pstop |
|-----|---------------------|-----------|------|-------|---------|-------------|----------|--------|-------|
| 0   | SELECT STATEMENT    |           | 5557 | 179K  |         | 948 (2)     | 00:00:12 |        |       |
| 1   | SORT ORDER BY       |           | 5557 | 179K  | 272K    | 948 (2)     | 00:00:12 |        |       |
| * 2 | HASH JOIN           |           | 5557 | 179K  |         | 897 (2)     | 00:00:11 |        |       |
| * 3 | TABLE ACCESS FULL   | CUSTOMERS | 43   | 516   |         | 405 (1)     | 00:00:06 |        |       |
| 4   | PARTITION RANGE ALL |           | 918K | 18M   |         | 489 (2)     | 00:00:06 | 1      | 28    |
| 5   | TABLE ACCESS FULL   | SALES     | 918K | 18M   |         | 489 (2)     | 00:00:06 | 1      | 28    |

```
Predicate Information (identified by operation id):

```

```
2 - access("C"."CUST_ID"="S"."CUST_ID")
3 - filter("CUST_FIRST_NAME"='Dina')
```

```
18 rows selected.
```

## Practice 9-2: Using the SQL TRACE Facility

In this practice, you trace a bad SQL statement and format the trace file with the tkprof utility.

1. Start a session to trace the SQL statement in the `bad_sql.sql` script as the SH user. Set the `TRACEFILE_IDENTIFIER` parameter for your session. Add the "badsql" string to the trace file name.

```
ALTER SESSION SET TRACEFILE_IDENTIFIER="badsql";
```

```
$ sqlplus /nolog

SQL> connect sh/sh
Connected.
SQL>

SQL> ALTER SESSION SET TRACEFILE_IDENTIFIER="badsql";
```

2. Enable tracing in your session with the command, and then run the `bad_sql.sql` script from the `$HOME/labs/sqltune` directory.

```
EXECUTE DBMS_SESSION.SET_SQL_TRACE(TRUE);
```

```
SQL> EXECUTE DBMS_SESSION.SET_SQL_TRACE(TRUE);

PL/SQL procedure successfully completed.

SQL> @$HOME/labs/sqltune/bad_sql.sql

...
TIME_ID QUANTITY_SOLD AMOUNT_SOLD

24-DEC-01 1 32.27
24-DEC-01 1 33.25
28-DEC-01 1 56.96
28-DEC-01 1 56.96

895 rows selected.

SQL> exit
```

3. Find and format the trace file. The trace file will be in the `$ORACLE_BASE/diag/rdbms/orcl/orcl/trace` directory with the `badsql` string embedded in the name. The `ls *badsql*.trc` command returns the name of the trace file. Use that name of the trace file in the `tkprof` command.

```
$ cd $ORACLE_BASE/diag/rdbms/orcl/orcl/trace
$ ls *badsql*.trc
orcl_ora_14804_badsql.trc
$ tkprof orcl_ora_14804_badsql.trc badsql.txt EXPLAIN=sh/sh
SYS=NO
```

## Practice 9-2: Using the SQL TRACE Facility (continued)

```
TKPROF: Release 11.2.0.1.0 - Development on Tue Feb 9 06:55:44
2010
```

```
Copyright (c) 1982, 2009, Oracle and/or its affiliates. All
rights reserved.
$
```

4. View the badsql.txt file. Locate the SQL statement.

```
select time_id, QUANTITY_SOLD, AMOUNT_SOLD
 from sales s, customers c
 where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
 order by time_id
```

Find the trace summary below the statement as shown in the following example.

The values shown are examples. The values you see may differ.

How much CPU time was used to fetch the records? \_\_\_\_\_(0.81)

How much elapsed time was used for the entire query? \_\_\_\_\_(0.89)

What was the number of physical reads (disk)? \_\_\_\_\_(3102)

What was the number of logical reads (query + current)? \_\_\_\_\_(3160)

```
$ less $ORACLE_BASE/diag/rdbms/orcl/orcl/trace/badsql.txt

TKPROF: Release 11.2.0.1.0 - Development on Tue Feb 9 06:55:44 2010

Copyright (c) 1982, 2009, Oracle and/or its affiliates. All rights reserved.

Trace file: orcl_ora_14804_badsql.trc
Sort options: default

count = number of times OCI procedure was executed
cpu = cpu time in seconds executing
elapsed = elapsed time in seconds executing
disk = number of physical reads of buffers from disk
query = number of buffers gotten for consistent read
current = number of buffers gotten in current mode (usually for update)
rows = number of rows processed by the fetch or execute call

...

select time_id, QUANTITY_SOLD, AMOUNT_SOLD
 from sales s, customers c
 where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
 order by time_id

call count cpu elapsed disk query current rows

Parse 1 0.00 0.00 0 0 0 0
Execute 1 0.00 0.00 0 0 0 0
Fetch 61 0.81 0.99 3088 3175 0 895

total 63 0.82 1.00 3088 3175 0 895

Misses in library cache during parse: 1
```

**Practice 9-2: Using the SQL TRACE Facility (continued)**

```

Optimizer mode: ALL_ROWS
Parsing user id: 85 (SH)

Rows Row Source Operation

 895 SORT ORDER BY (cr=3175 pr=3088 pw=0 time=1341 us cost=948 size=183381
card=5557)
 895 HASH JOIN (cr=3175 pr=3088 pw=0 time=132695 us cost=897 size=183381
card=5557)
 65 TABLE ACCESS FULL CUSTOMERS (cr=1457 pr=1454 pw=0 time=64 us cost=405
size=516 card=43)
 918843 PARTITION RANGE ALL PARTITION: 1 28 (cr=1718 pr=1634 pw=0 time=2431046
us cost=489 size=19295703 ca
rd=918843)
 918843 TABLE ACCESS FULL SALES PARTITION: 1 28 (cr=1718 pr=1634 pw=0
time=998014 us cost=489 size=1929570
3 card=918843)

Rows Execution Plan

 0 SELECT STATEMENT MODE: ALL_ROWS
 895 SORT (ORDER BY)
 895 HASH JOIN
 65 TABLE ACCESS MODE: ANALYZED (FULL) OF 'CUSTOMERS' (TABLE)
 918843 PARTITION RANGE (ALL) PARTITION: START=1 STOP=28
 918843 TABLE ACCESS MODE: ANALYZED (FULL) OF 'SALES' (TABLE)
 PARTITION: START=1 STOP=28

...

OVERALL TOTALS FOR ALL NON-RECURSIVE STATEMENTS

call count cpu elapsed disk query current rows

Parse 1 0.00 0.00 0 0 0 0
Execute 2 0.00 0.01 0 0 0 1
Fetch 61 0.81 0.99 3088 3175 0 895

total 64 0.82 1.01 3088 3175 0 896

Misses in library cache during parse: 1
Misses in library cache during execute: 1

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call count cpu elapsed disk query current rows

Parse 18 0.02 0.04 0 0 0 0
Execute 36 0.02 0.04 0 1 2 1
Fetch 133 0.00 0.04 8 110 0 348

total 187 0.06 0.13 8 111 2 349

Misses in library cache during parse: 18
Misses in library cache during execute: 19

 3 user SQL statements in session.
 35 internal SQL statements in session.
 38 SQL statements in session.
 1 statement EXPLAINED in this session.

```

## Practice 9-2: Using the SQL TRACE Facility (continued)

```

Trace file: orcl_ora_9029_badsql.trc
Trace file compatibility: 11.1.0.7
Sort options: default

 1 session in tracefile.
 3 user SQL statements in trace file.
 35 internal SQL statements in trace file.
 38 SQL statements in trace file.
 21 unique SQL statements in trace file.
 1 SQL statements EXPLAINED using schema:
 SH.prof$plan_table
 Default table was used.
 Table was created.
 Table was dropped.
 536 lines in trace file.
 14 elapsed seconds in trace file.

(type q to exit)
```



## Practices for Lesson 10

In the following practices, you demonstrate various ways to influence the optimizer by changing the SQL:

Practice 10-1: Using Extended Statistics

Practice 10-2: Stale Statistics

For all practices in this lesson, the environment must be changed by executing the `./prepare 10 1` script from the `$HOME/workshops` directory. The script needs to be run only once.

## Practice 10-1: Capturing Extended Statistics

In this practice, you create extended statistics over a multicolumn set and an expression.

1. Run the `prepare 10 1` script from the `$HOME/workshops` directory if you have not already done so.

```
$ $HOME/workshops/prepare 10 1

System altered.
```

2. In the `SH.CUSTOMERS` table, the `CUST_CITY`, `CUST_STATE_PROVINCE`, and `COUNTRY_ID` columns are frequently used together as predicates. Because these columns are related, the combined selectivity is very different from the selectivity that the optimizer would calculate. The optimizer assumes that the columns are independent. To capture the statistics over this set of columns, there are two methods: create a concatenated index over these columns and collect the statistics over the index, or create a column group.

Create a column group to collect statistics over this column group without creating an index. Examine the `lab_10_01_01.sql` script in the `$HOME/labs/sqltune` directory, and then execute it.

```
$ cd $HOME/labs/sqltune
$ sqlplus /nolog

SQL> connect sh/sh
Connected.
SQL> set echo on
SQL> @lab_10_01_01.sql
SQL> SET SERVEROUTPUT ON
SQL>
SQL> exec dbms_stats.drop_extended_stats(-
2 'sh','customers',-
3 '(CUST_CITY,CUST_STATE_PROVINCE,COUNTRY_ID)');

*
ERROR at line 1:
ORA-20000: extension
"(CUST_CITY,CUST_STATE_PROVINCE,COUNTRY_ID)" does not
exist
ORA-06512: at "SYS.DBMS_STATS", line 6893
ORA-06512: at "SYS.DBMS_STATS", line 28235
ORA-06512: at line 1

SQL>
SQL> DECLARE
2 cg_name VARCHAR2(30);
3 BEGIN
4 cg_name := dbms_stats.create_extended_stats(
5 'SH','CUSTOMERS',
6 '(CUST_CITY,CUST_STATE_PROVINCE,COUNTRY_ID)');
```

**Practice 10-1: Capturing Extended Statistics (continued)**

```

7
8 dbms_output.put_line('column group name is: ' || cg_name);
9
10 dbms_stats.gather_table_stats('SH','CUSTOMERS',
11 method_opt=>'for all columns size 1,for columns
12 (CUST_CITY,CUST_STATE_PROVINCE,COUNTRY_ID) size 3');
13 END;
14 /
column group name is: SYS_STUMZ$C3AIHLPBROI#SKA58H_N
PL/SQL procedure successfully completed.
SQL>

```

3. The SELECT statement in lab\_10\_01\_02.sql gives distinct values of each column in the group.

```

SQL> @lab_10_01_02.sql
SQL> select count(distinct(cust_city)) cities,
2 count(distinct(CUST_STATE_PROVINCE)) states,
3 count(distinct(country_id)) countries
4 from customers;

```

| CITIES | STATES | COUNTRIES |
|--------|--------|-----------|
| 620    | 145    | 19        |

4. The SELECT statement in lab\_10\_01\_03.sql shows the selectivity that the default algorithm in the optimizer would find.

```

SQL> @lab_10_01_03.sql
SQL> select 1/(count(distinct(cust_city))) *
2 1/(count(distinct(CUST_STATE_PROVINCE))) *
3 1/(count(distinct(country_id)))
default_selectivity
4 from customers
5 /

```

| DEFAULT_SELECTIVITY |
|---------------------|
| 5.8545E-07          |

5. The PL/SQL block in lab\_10\_01\_04.sql retrieves the column group name.

```

SQL> @lab_10_01_04.sql
SQL> declare
2 cg_name VARCHAR2(30);
3 BEGIN
4 cg_name := dbms_stats.show_extended_stats_name
5 ('SH','CUSTOMERS',
6 '(CUST_CITY,CUST_STATE_PROVINCE,COUNTRY_ID)');
7
8 dbms_output.put_line('column group name is: ' || cg_name);

```

**Practice 10-1: Capturing Extended Statistics (continued)**

```

9
10 END;
11 /
column group name is:SYS_STUMZ$C3AIHLPBROI#SKA58H_N

PL/SQL procedure successfully completed.

```

6. The lab\_10\_01\_05.sql script retrieves the statistics on the column group. NUM\_DISTINCT shows the number of distinct values for the three columns. The selectivity for the column group is 1/620 or 1.61E-3, rather than the 5.8E-7 that the default calculation would have given. The extended statistics help the optimizer make better choices, with more accurate selectivity when these columns are used together in a predicate.

```

SQL> @lab_10_01_05.sql
SQL> select e.extension col_group, t.num_distinct, t.histogram
2 from user_stat_extensions e, user_tab_col_statistics t
3 where e.extension_name=t.column_name
4 and t.table_name='CUSTOMERS';

COL_GROUP

NUM_DISTINCT HISTOGRAM

("CUST_CITY", "CUST_STATE_PROVINCE", "COUNTRY_ID")
620 HEIGHT BALANCED

```

7. Tables often have columns that are not populated that are called “virtual columns.” They are calculated as part of the SELECT statement. When these virtual columns are used in a predicate, ORDER BY, or join conditions, the optimizer needs statistics to make fully informed decisions. The lab\_10\_01\_07.sql script creates a table with four columns and populates it. The id column contains 1,000 distinct numbers: 1 to 999. The second column contains only an X. The Cid column contains id as a character string. The Nid column contains (id\*100)+6, so that an id of 27 becomes the Nid of 2706.

```

SQL> @lab_10_01_07.sql
SQL>
SQL> DROP TABLE EXP_TEST PURGE;

Table dropped.

SQL>
SQL> Create table exp_test (id Number(6,0),
2 pad varchar2 (20),
3 Cid VARCHAR(20),
4 Nid NUMBER(6,0));

```

**Practice 10-1: Capturing Extended Statistics (continued)**

Table created.

```
SQL>
SQL> Begin
 2 FOR i in 1..1000 LOOP
 3
 4 INSERT INTO exp_test values(
 5 i, 'X', TO_CHAR(i), i*100+6);
 6 END LOOP;
 7 END;
 8 /
```

PL/SQL procedure successfully completed.

8. When a column named in the predicate has a function applied to it, the index on the column cannot be used. A function-based index on the expression can be created to allow an index access to be used. The lab\_10\_01\_08 script creates a function-based index over the `(mod(id,10))` expression.

```
SQL> @lab_10_01_08.sql
SQL> set echo on
SQL>
SQL> CREATE index exp_mod on exp_test(mod(id,10));
```

Index created.

SQL>

9. With the lab\_10\_01\_09.sql script, create a virtual column over the `(mod(nid,10))` expression.

```
SQL> @lab_10_01_09.sql
SQL> set echo on
SQL>
SQL> select dbms_stats.create_extended_stats(
 2 'sh','exp_test','(mod(Nid,10))') from dual;

DBMS_STATS.CREATE_EXTENDED_STATS('SH','EXP_TEST','(MOD(NID,10))')

SYS_STUFYQ8TH8YVE2JG8NLWCN9XTR
```

10. Gather statistics on the EXP\_TEST table. The virtual column named in step 9 is explicitly named in the method\_opt parameter.

```
SQL> @lab_10_01_10.sql
SQL> set echo on
SQL>
SQL> exec dbms_stats.gather_table_stats('sh','exp_test',-
```

## Practice 10-1: Capturing Extended Statistics (continued)

```
> method_opt=>'for all columns size 1 for columns
(mod(Nid,10)) size 3');

PL/SQL procedure successfully completed.

SQL>
```

11. The extended statistics have been calculated for the EXP\_TEST table. The extended statistics have been defined for (MOD(nid,10)). What are the values of the extended statistics?

```
SQL> @lab_10_01_11.sql
SQL> select e.extension col_group, t.num_distinct, t.histogram
2 from user_stat_extensions e, user_tab_col_statistics t
3 where e.extension_name=t.column_name
4 and t.table_name='EXP_TEST';

COL_GROUP

NUM_DISTINCT HISTOGRAM

(MOD("ID",10))
 10 NONE

(MOD("NID",10))
 1 FREQUENCY

SQL> exit
```

The extended statistics were calculated over the EXP\_MOD functional index that was created in step 8, as well as over the virtual column defined in step 9.

## Practice 10-2: Stale Statistics

In this practice, you see that statistics can influence the optimizer performance.

1. Create a table and run statistics on it and the associated indexes as the SH user.  
Use the `cr_badstats.sql` script in the `$HOME/labs/sqltune` directory to create the CTEST table and the associated indexes, and to set up the scenario. Statistics are gathered, and then enough rows are inserted into the CTEST table to render the statistics stale.

```
$ cd $HOME/labs/sqltune
$ sqlplus /nolog

SQL> connect sh/sh
Connected
SQL> @cr_badstats.sql
SQL> -- Create a table with indexes that will yield a
different
SQL> -- execution plan when statistics are gathered
SQL>
SQL> -- Create a table, indexes and gather statistics
SQL>
SQL> DROP SEQUENCE ctest_pk_seq;
DROP SEQUENCE ctest_pk_seq
 *
ERROR at line 1:
ORA-02289: sequence does not exist

SQL>
SQL> DROP TABLE CTEST PURGE;
DROP TABLE CTEST PURGE
 *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
SQL> CREATE TABLE CTEST AS SELECT * FROM CUSTOMERS;

Table created.

SQL>
SQL> ALTER TABLE CTEST MODIFY (CUST_ID PRIMARY KEY);

Table altered.

SQL>
SQL> CREATE INDEX CTEST_LASTNAME ON CTEST (CUST_LAST_NAME);

Index created.

SQL>
```

**Practice 10-2: Stale Statistics (continued)**

```

SQL> CREATE INDEX CTEST_COUNTRY ON CTEST (COUNTRY_ID);
Index created.

SQL>
SQL> EXECUTE DBMS_STATS.GATHER_TABLE_STATS('SH','CTEST');
PL/SQL procedure successfully completed.

SQL>
SQL> -- Modify the table to create skewed data on Country_id
SQL>
SQL> CREATE SEQUENCE ctest_pk_seq START WITH 200000
NOMAXVALUE;

Sequence created.

SQL>
SQL> INSERT INTO CTEST
 2 SELECT ctest_pk_seq.nextval,CUST_FIRST_NAME,
 3 CUST_LAST_NAME, CUST_GENDER,CUST_YEAR_OF_BIRTH,
 4 CUST_MARITAL_STATUS, CUST_STREET_ADDRESS,
 5 CUST_POSTAL_CODE,CUST_CITY, CUST_CITY_ID,
 6 CUST_STATE_PROVINCE,CUST_STATE_PROVINCE_ID,COUNTRY_ID,
 7 CUST_MAIN_PHONE_NUMBER, CUST_INCOME_LEVEL,
 8 CUST_CREDIT_LIMIT,CUST_EMAIL,CUST_TOTAL,CUST_TOTAL_ID,
 9 CUST_SRC_ID,CUST_EFF_FROM, CUST_EFF_TO,CUST_VALID
 10 FROM CTEST
 11 WHERE country_id = 52790;

18520 rows created.

SQL>
SQL> INSERT INTO CTEST
 2 SELECT ctest_pk_seq.nextval, CUST_FIRST_NAME,
 3 CUST_LAST_NAME, CUST_GENDER,CUST_YEAR_OF_BIRTH,
 4 CUST_MARITAL_STATUS, CUST_STREET_ADDRESS,
 5 CUST_POSTAL_CODE,CUST_CITY, CUST_CITY_ID,
 6 CUST_STATE_PROVINCE,CUST_STATE_PROVINCE_ID,COUNTRY_ID,
 7 CUST_MAIN_PHONE_NUMBER,CUST_INCOME_LEVEL,
 8 CUST_CREDIT_LIMIT,CUST_EMAIL,CUST_TOTAL,CUST_TOTAL_ID,
 9 CUST_SRC_ID,CUST_EFF_FROM, CUST_EFF_TO,CUST_VALID
 10 FROM CTEST
 11 WHERE country_id = 52790;

37040 rows created.

SQL>

```

2. Autotrace the SQL statement in the `bad_stats.sql` script and observe the execution plan and the differences between the estimated cost and the actual cost. The script is in the `$HOME/labs/sqltune` directory.



**Practice 10-2: Stale Statistics (continued)**

```

SQL> @bad_stats.sql
SQL>
SQL> -- Show the choice of FTS with old statistics
SQL> SET ECHO ON
SQL> SET AUTOTRACE TRACEONL
SQL>
SQL> SELECT CUST_ID, CUST_LAST_NAME, CUST_TOTAL
 2 FROM CTEST
 3 WHERE COUNTRY_ID = 52779;

3833 rows selected.

```

## Execution Plan

Plan hash value: 3591594933

| Id  | Operation         | Name  | Rows | Bytes | Cost (%CPU) | Time     |
|-----|-------------------|-------|------|-------|-------------|----------|
| 0   | SELECT STATEMENT  |       | 2921 | 96393 | 406 (1)     | 00:00:05 |
| * 1 | TABLE ACCESS FULL | CTEST | 2921 | 96393 | 406 (1)     | 00:00:05 |

## Predicate Information (identified by operation id):

1 - filter("COUNTRY\_ID"=52779)

## Statistics

```

2 recursive calls
1 db block gets
3162 consistent gets
2883 physical reads
176 redo size
86727 bytes sent via SQL*Net to client
3224 bytes received via SQL*Net from client
257 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
3833 rows processed

```

```

SQL>
SQL> SET AUTOTRACE OFF
SQL> SET ECHO OFF

```



**Practice 10-2: Stale Statistics (continued)**

```
Statistics
```

```

 0 recursive calls
 0 db block gets
 3132 consistent gets
 2874 physical reads
 0 redo size
 86727 bytes sent via SQL*Net to client
 3224 bytes received via SQL*Net from client
 257 SQL*Net roundtrips to/from client
 0 sorts (memory)
 0 sorts (disk)
 3833 rows processed
```

```
SQL>
```

6. What was the number of rows estimated? \_\_\_\_\_ (4247, may vary)  
 What was the number of rows fetched? \_\_\_\_\_ (3833)  
 What was the number of logical reads? \_\_\_\_\_ (varies)  
 The number of logical reads is a rough measure of the cost of the statement. There should be a small reduction in the number of logical reads between this execution and the previous execution of this statement.
7. Clean up the objects created for this scenario. Use the `cleanup_bad_stats.sql` script.

```
SQL> @cleanup_bad_stats.sql
```

```
SQL> DROP SEQUENCE ctest_pk_seq;
```

```
Sequence dropped.
```

```
SQL>
```

```
SQL> DROP TABLE CTEST PURGE;
```

```
Table dropped.
```

```
SQL>
```

```
SQL> EXIT
```

```
Disconnected from Oracle Database 11g Enterprise Edition
Release 11.1.0.6.0 - Production
With the Partitioning, Oracle Label Security, OLAP, Data
Mining
and Real Application Testing options
$
```

8. In this exercise, stale statistics cause the optimizer to choose a full-table scan (FTS) rather than an index access. This example shows a small cost reduction. The differences in cost can be magnified by repeated execution of the same statements and by larger differences in the statistics.

## Practices for Lesson 11

In this practice, you reduce the cost of queries by reducing the number of blocks that must be retrieved.

## Practice 11-1: Excess Blocks

In this practice, you see the differences in the explain plan and the costs of using a table that has more blocks than necessary. This condition can be caused in a number of ways: direct loads followed by deletes, inserts and deletes with a high PCTFREE setting, and batch deletes. The immediate solution is to shrink or reorganize the table; the long-term solution is to modify the processing or PCTUSED so that blocks are reused for inserts.

1. Create a table that uses more blocks than necessary. Use the `crctest.sql` script as the SH user. This script creates a table with the same data as the SH.CUSTOMERS table, but with a PCTFREE=80, so that it uses many more blocks than needed.

```
$ cd $HOME/labs/sqltune
$ sqlplus /nolog

SQL> connect sh/sh
Connected.

SQL> @crctest.sql
SQL>
SQL> -- create a new revision of the Customers table with a
SQL> -- large number of blocks.
SQL> set echo on
SQL>
SQL> DROP TABLE CTEST PURGE;
DROP TABLE CTEST PURGE
 *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
SQL> CREATE TABLE CTEST
 2 PCTFREE 80 PCTUSED 10
 3 as SELECT * FROM customers;

Table created.

SQL>
SQL> EXEC DBMS_STATS.GATHER_TABLE_STATS('SH','CTEST');

PL/SQL procedure successfully completed.

SQL>
```

2. Run the ctest.sql script as the SH user with AUTOTRACE on.

```
SQL> @ctest.sql
SQL>
SQL> -- set autotrace on --
SQL> set echo on
SQL>
SQL> SET AUTOTRACE TRACEONLY
SQL>
SQL> select time_id, QUANTITY_SOLD, AMOUNT_SOLD
 2 from sales s, ctest c
 3 where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
 4 order by time_id;
```

895 rows selected.

#### Execution Plan

-----  
Plan hash value: 2054825740  
-----

| Id  | Operation           | Name  | Rows | Bytes | TempSpc | Cost (%CPU) | Time     | Pstart | Pstop |
|-----|---------------------|-------|------|-------|---------|-------------|----------|--------|-------|
| 0   | SELECT STATEMENT    |       | 5557 | 179K  |         | 2453 (1)    | 00:00:30 |        |       |
| 1   | SORT ORDER BY       |       | 5557 | 179K  | 272K    | 2453 (1)    | 00:00:30 |        |       |
| * 2 | HASH JOIN           |       | 5557 | 179K  |         | 2402 (1)    | 00:00:29 |        |       |
| * 3 | TABLE ACCESS FULL   | CTEST | 43   | 516   |         | 1912 (1)    | 00:00:23 |        |       |
| 4   | PARTITION RANGE ALL |       | 918K | 18M   |         | 498 (4)     | 00:00:06 | 1      | 28    |
| 5   | TABLE ACCESS FULL   | SALES | 918K | 18M   |         | 498 (4)     | 00:00:06 | 1      | 28    |

#### Predicate Information (identified by operation id):

-----  
2 - access("C"."CUST\_ID"="S"."CUST\_ID")  
3 - filter("CUST\_FIRST\_NAME"='Dina')

#### Statistics

-----  
13516 recursive calls  
0 db block gets  
11455 consistent gets  
8658 physical reads  
1242 redo size

```

 23039 bytes sent via SQL*Net to client
 1068 bytes received via SQL*Net from client
 61 SQL*Net roundtrips to/from client
 111 sorts (memory)
 0 sorts (disk)
 895 rows processed

SQL>
SQL> SET AUTOTRACE OFF
SQL>

```

- Note the cost, both in terms of the number of physical reads and logical reads (consistent gets + DB block gets), and the estimated cost.  
 Physical reads \_\_\_\_\_ (8658, may vary)  
 Logical reads \_\_\_\_\_ (11455, may vary)  
 Estimated cost \_\_\_\_\_ (2453)
- Reorganize the table to use fewer blocks. Run the `reorg_ctest.sql` script. This script sets `PCTFREE` to 10, and then moves the table. This script also checks the number of blocks used before and after the reorganization.

```

SQL> @reorg_ctest.sql
SQL> -- Rebuild the table with higher density
SQL> SET ECHO ON
SQL>
SQL> SELECT BLOCKS, EMPTY_BLOCKS FROM USER_TABLES
 2 WHERE TABLE_NAME = 'CTEST';

 BLOCKS EMPTY_BLOCKS

 7038 0

SQL>
SQL> ALTER TABLE CTEST
 2 PCTFREE 10;

Table altered.

SQL>
SQL> ALTER TABLE CTEST MOVE;

Table altered.

SQL>
SQL> EXEC DBMS_STATS.GATHER_TABLE_STATS('SH','CTEST');
PL/SQL procedure successfully completed.

SQL>
SQL> SELECT BLOCKS, EMPTY_BLOCKS FROM USER_TABLES
 2 WHERE TABLE_NAME = 'CTEST';

 BLOCKS EMPTY_BLOCKS

```

|      |   |
|------|---|
| 1485 | 0 |
|------|---|

SQL>

5. Run the ctest.sql script as the SH user with AUTOTRACE on.

```
SQL> @ctest.sql
SQL>
SQL> -- set autotrace on --
SQL> set echo on
SQL>
SQL> SET AUTOTRACE TRACEONLY
SQL>
SQL> select time_id, QUANTITY_SOLD, AMOUNT_SOLD
 2 from sales s, ctest c
 3 where c.cust_id = s.cust_id and CUST_FIRST_NAME='Dina'
 4 order by time_id;

895 rows selected.

Execution Plan

Plan hash value: 2054825740
```

| Id  | Operation           | Name  | Rows | Bytes | TempSpc | Cost (%CPU) | Time     | Pstart | Pstop |
|-----|---------------------|-------|------|-------|---------|-------------|----------|--------|-------|
| 0   | SELECT STATEMENT    |       | 5557 | 179K  |         | 948 (3)     | 00:00:12 |        |       |
| 1   | SORT ORDER BY       |       | 5557 | 179K  | 536K    | 948 (3)     | 00:00:12 |        |       |
| * 2 | HASH JOIN           |       | 5557 | 179K  |         | 897 (3)     | 00:00:11 |        |       |
| * 3 | TABLE ACCESS FULL   | CTEST | 43   | 516   |         | 405 (1)     | 00:00:05 |        |       |
| 4   | PARTITION RANGE ALL |       | 918K | 18M   |         | 489 (4)     | 00:00:06 | 1      | 28    |
| 5   | TABLE ACCESS FULL   | SALES | 918K | 18M   |         | 498 (4)     | 00:00:06 | 1      | 28    |

Predicate Information (identified by operation id):

```

2 - access("C"."CUST_ID"="S"."CUST_ID")
3 - filter("CUST_FIRST_NAME"='Dina')
```

Statistics

```

15301 recursive calls
0 db block gets
6306 consistent gets
```



```

 3217 physical reads
 0 redo size
 23039 bytes sent via SQL*Net to client
 1068 bytes received via SQL*Net from client
 61 SQL*Net roundtrips to/from client
 147 sorts (memory)
 0 sorts (disk)
 895 rows processed

SQL>
SQL> SET AUTOTRACE OFF

```

6. Note the various costs. Write down the number of physical reads and logical reads (consistent gets + DB gets), and the estimated cost.

Physical reads \_\_\_\_\_ (3217, may vary)

Logical reads \_\_\_\_\_ ( 6306, may vary)

Estimated cost \_\_\_\_\_ (948)

The actual numbers that you see may vary from those in the example, but you should see a reduction in the number of logical and physical reads and the estimated cost. These reductions are due to reducing the number of physical blocks used to hold the table rows. This effect is most pronounced with full-table scans. But even with index access methods, if the table uses more blocks, it will require more memory buffers and more block visits to access the same number of rows.

7. Clean up the objects for this exercise.

```

SQL> DROP TABLE CTEST PURGE;
SQL> EXIT

```

## Practice 11-2: Coalescing an Index

In this practice, you observe the effect of coalescing an index on a SQL statement and an execution plan.

1. The `coalesce_setup.sql` script creates the TESTLM tablespace, the CTEST table, and the CTEST\_PK index, with excess space in the index. CTEST is created as an empty version of the SH.CUSTOMERS table, and then populated with four times the rows from CUSTOMERS. To make the index sparse, every second row is deleted based on CUST\_ID, which is unique and is generated by a sequence.

```
$ cd $HOME/labs/sqltune
$ sqlplus /nolog

SQL> SET ECHO ON

SQL> @coalesce_setup.sql

SQL> CONNECT / as sysdba
Connected.
SQL>
SQL> DROP TABLESPACE TESTLM INCLUDING CONTENTS AND DATAFILES;
DROP TABLESPACE TESTLM INCLUDING CONTENTS AND DATAFILES
*
ERROR at line 1:
ORA-00959: tablespace 'TESTLM' does not exist

SQL>
SQL> CREATE TABLESPACE TESTLM
 2 DATAFILE '/u01/app/oracle/oradata/orcl/testlm1.dbf'
 3 SIZE 100M
 4 AUTOEXTEND ON NEXT 10M MAXSIZE 300M
 5 SEGMENT SPACE MANAGEMENT MANUAL;

Tablespace created.

SQL>
SQL> CONNECT sh/sh
Connected.

SQL>
SQL> -- Create a table with indexes that have large number of
SQL> -- partially empty blocks and will benefit from an index
SQL> -- coalesce
SQL>
SQL> -- Create a table, indexes and gather statistics
SQL> DROP SEQUENCE ctest_pk_seq;
DROP SEQUENCE ctest_pk_seq
*
ERROR at line 1:
```

## Practice 11-2: Coalescing an Index (continued)

ORA-02289: sequence does not exist

```
SQL> DROP TABLE CTEST PURGE;
DROP TABLE CTEST PURGE
 *
```

ERROR at line 1:

ORA-00942: table or view does not exist

SQL>

```
SQL> CREATE TABLE CTEST TABLESPACE TESTLM
 2 AS SELECT *
 3 FROM CUSTOMERS where 1=2;
```

Table created.

SQL>

```
SQL> ALTER TABLE CTEST ADD Constraint CTEST_PK
 2 Primary KEY (cust_id);
```

Table altered.

```
SQL> ALTER TABLE CTEST ADD (DUMMY VARCHAR2(300));
```

Table altered.

SQL>

SQL>

```
SQL> CREATE SEQUENCE ctest_pk_seq START WITH 200000
NOMAXVALUE;
```

Sequence created.

SQL>

```
SQL> INSERT INTO CTEST
 2 SELECT ctest_pk_seq.nextval, CUST_FIRST_NAME,
 3 CUST_LAST_NAME, CUST_GENDER, CUST_YEAR_OF_BIRTH,
 4 CUST_MARITAL_STATUS, CUST_STREET_ADDRESS, CUST_POSTAL_CODE,
 5 CUST_CITY, CUST_CITY_ID, CUST_STATE_PROVINCE,
 6 CUST_STATE_PROVINCE_ID, COUNTRY_ID, CUST_MAIN_PHONE_NUMBER,
 7 CUST_INCOME_LEVEL, CUST_CREDIT_LIMIT, CUST_EMAIL,
 8 CUST_TOTAL, CUST_TOTAL_ID, CUST_SRC_ID, CUST_EFF_FROM,
 9 CUST_EFF_TO, CUST_VALID,
 10 CUST_FIRST_NAME || ' ' || CUST_LAST_NAME
 11 FROM CUSTOMERS;
```

55500 rows created.

SQL>

SQL>

```
SQL> INSERT INTO CTEST
```

**Practice 11-2: Coalescing an Index (continued)**

```

2 SELECT ctest_pk_seq.nextval, CUST_FIRST_NAME,
3 CUST_LAST_NAME,CUST_GENDER, CUST_YEAR_OF_BIRTH,
4 CUST_MARITAL_STATUS,CUST_STREET_ADDRESS,CUST_POSTAL_CODE,
5 CUST_CITY,CUST_CITY_ID,CUST_STATE_PROVINCE,
6 CUST_STATE_PROVINCE_ID,COUNTRY_ID,CUST_MAIN_PHONE_NUMBER,
7 CUST_INCOME_LEVEL,CUST_CREDIT_LIMIT,CUST_EMAIL,
8 CUST_TOTAL,CUST_TOTAL_ID,CUST_SRC_ID,CUST_EFF_FROM,
9 CUST_EFF_TO,CUST_VALID,
10 CUST_FIRST_NAME||' '||CUST_LAST_NAME
11 FROM CUSTOMERS;

55500 rows created.

SQL>
SQL>
SQL> INSERT INTO CTEST
2 SELECT ctest_pk_seq.nextval, CUST_FIRST_NAME,
3 CUST_LAST_NAME,CUST_GENDER, CUST_YEAR_OF_BIRTH,
4 CUST_MARITAL_STATUS,CUST_STREET_ADDRESS,CUST_POSTAL_CODE,
5 CUST_CITY,CUST_CITY_ID,CUST_STATE_PROVINCE,
6 CUST_STATE_PROVINCE_ID,COUNTRY_ID,CUST_MAIN_PHONE_NUMBER,
7 CUST_INCOME_LEVEL,CUST_CREDIT_LIMIT,CUST_EMAIL,
8 CUST_TOTAL,CUST_TOTAL_ID,CUST_SRC_ID,CUST_EFF_FROM,
9 CUST_EFF_TO,CUST_VALID,
10 CUST_FIRST_NAME||' '||CUST_LAST_NAME
11 FROM CUSTOMERS;

55500 rows created.

SQL>
SQL>
SQL> INSERT INTO CTEST
2 SELECT ctest_pk_seq.nextval, CUST_FIRST_NAME,
3 CUST_LAST_NAME,CUST_GENDER, CUST_YEAR_OF_BIRTH,
4 CUST_MARITAL_STATUS,CUST_STREET_ADDRESS,CUST_POSTAL_CODE,
5 CUST_CITY,CUST_CITY_ID,CUST_STATE_PROVINCE,
6 CUST_STATE_PROVINCE_ID,COUNTRY_ID,CUST_MAIN_PHONE_NUMBER,
7 CUST_INCOME_LEVEL,CUST_CREDIT_LIMIT,CUST_EMAIL,
8 CUST_TOTAL,CUST_TOTAL_ID,CUST_SRC_ID,CUST_EFF_FROM,
9 CUST_EFF_TO,CUST_VALID,
10 CUST_FIRST_NAME||' '||CUST_LAST_NAME
11 FROM CUSTOMERS;

55500 rows created.

SQL>
SQL> DELETE FROM CTEST
2 WHERE mod(cust_id,2) = 1;

111000 rows deleted.

SQL>

```

**Practice 11-2: Coalescing an Index (continued)**

```
SQL> commit;

Commit complete.

SQL>
SQL> EXEC DBMS_STATS.GATHER_TABLE_STATS('SH','CTEST',-
> method_opt=>'FOR ALL INDEXED COLUMNS SIZE AUTO');

PL/SQL procedure successfully completed.

SQL>
```

2. Execute the SQL statement in the `coalesce_idx.sql` script. The script sets AUTOTRACE TRACEONLY.

What is the access method? \_\_\_\_\_ (Full-table scan)

What is the estimated cost? \_\_\_\_\_ (1693)

What is the number of consistent reads? \_\_\_\_\_ (8091, may vary)

```
SQL> @coalesce_idx_1.sql
SQL> -- Will yield different plans before and after index
coalesce
SQL>
SQL> Set AUTOTRACE TRACEONLY
SQL> select /*BEFORE*/ cust_id,cust_last_name, cust_first_name
 2 from ctest
 3 where cust_id < 257500
/

28749 rows selected.
```

## Execution Plan

Plan hash value: 3591594933

| Id  | Operation         | Name  | Rows  | Bytes | Cost (%CPU) | Time     |
|-----|-------------------|-------|-------|-------|-------------|----------|
| 0   | SELECT STATEMENT  |       | 28749 | 1094K | 1691 (1)    | 00:00:21 |
| * 1 | TABLE ACCESS FULL | CTEST | 28749 | 1094K | 1691 (1)    | 00:00:21 |

Predicate Information (identified by operation id):

1 - filter("CUST\_ID"<257500)

## Statistics

```
1 recursive calls
0 db block gets
8100 consistent gets
```

**Practice 11-2: Coalescing an Index (continued)**

```

 6228 physical reads
 0 redo size
 613586 bytes sent via SQL*Net to client
 21495 bytes received via SQL*Net from client
 1918 SQL*Net roundtrips to/from client
 0 sorts (memory)
 0 sorts (disk)
 28749 rows processed

SQL>

```

3. Execute the `coalesce_idx_size.sql` script to view index statistics.

```

SQL>
SQL> @coalesce_idx_size
SQL> -- Find the important index stats before and
SQL> -- after coalesce
SQL>
SQL> SELECT Blevel, LEAF_BLOCKS, DISTINCT_KEYS,
CLUSTERING_FACTOR
 2 FROM USER_INDEXES
 3 WHERE index_name = 'CTEST_PK';

 BLEVEL LEAF_BLOCKS DISTINCT_KEYS CLUSTERING_FACTOR

 1 417 111000 6224

SQL>

```

4. Coalesce the CTEST\_PK index.

```

SQL> @coalesce_fix
SQL> -- Fix the index (Coalesce)
SQL>
SQL> ALTER INDEX CTEST_PK COALESCE;

Index altered.

SQL>
SQL> -- GATHER STATS
SQL>
SQL> EXEC DBMS_STATS.GATHER_TABLE_STATS('SH','CTEST',-
 2 METHOD_OPT=>'FOR ALL INDEXED COLUMNS SIZE AUTO');

PL/SQL procedure successfully completed.

SQL>

```

## Practice 11-2: Coalescing an Index (continued)

5. View index statistics again. Notice the difference in the number of leaf blocks.

```
SQL> @coalesce_idx_size
SQL> -- Find the important index stats before and
SQL> -- after coalesce
SQL>
SQL> SELECT Blevel, LEAF_BLOCKS, DISTINCT_KEYS,
CLUSTERING_FACTOR
 2 FROM USER_INDEXES
 3 WHERE index_name = 'CTEST_PK';

 BLEVEL LEAF_BLOCKS DISTINCT_KEYS CLUSTERING_FACTOR

 1 232 111000 6224

SQL>
```

6. Execute the SQL statement in the coalesce\_idx.sql script again.

```
SQL> @coalesce_idx_2.sql
SQL> -- Will yield different plans before and after index
coalesce
SQL>
SQL> Set AUTOTRACE TRACEONLY
SQL> select /*AFTER*/ cust_id, cust_last_name, cust_first_name
 2 from ctest
 3 where cust_id < 257500
 4 /

28749 rows selected.
```

### Execution Plan

Plan hash value: 3972627342

| Id  | Operation                   | Name     | Rows  | Bytes | Cost (%CPU) | Time     |
|-----|-----------------------------|----------|-------|-------|-------------|----------|
| 0   | SELECT STATEMENT            |          | 28749 | 1094K | 1676 (1)    | 00:00:21 |
| 1   | TABLE ACCESS BY INDEX ROWID | CTEST    | 28749 | 1094K | 1676 (1)    | 00:00:21 |
| * 2 | INDEX RANGE SCAN            | CTEST_PK | 28749 |       | 62 (0)      | 00:00:01 |

Predicate Information (identified by operation id):

2 - access("CUST\_ID"<257500)

## Practice 11-2: Coalescing an Index (continued)

Statistics

```

 1 recursive calls
 0 db block gets
 5393 consistent gets
 1648 physical reads
 0 redo size
 613486 bytes sent via SQL*Net to client
 21495 bytes received via SQL*Net from client
 1918 SQL*Net roundtrips to/from client
 0 sorts (memory)
 0 sorts (disk)
 28749 rows processed
```

SQL>

SQL> SET AUTOTRACE OFF

SQL> **exit**

7. Observe the differences in the estimated cost and the number of logical reads (consistent gets).  
What is the access method? \_\_\_\_\_ (index-range scan)  
What is the estimated cost? \_\_\_\_\_ (1676)  
How many consistent reads were performed? \_\_\_\_\_ (5293, may vary)
8. Clean up this practice—execute the `./cleanup 10 4` script in the `$HOME/workshops` directory.

```
$ cd $HOME/workshops
```

```
$./cleanup OPT
```

```
Tablespace dropped.
```

```
$
```



## Practices for Lesson 12

The goal of this practice is to use SQL Performance Analyzer to identify the SQL statements that regress when the environment is upgraded.

## Practice 12-1: Using the SQL Performance Analyzer

In this practice, you simulate exporting a SQL Tuning Set (STS) from a 10g database and import it back into an 11g test environment. There, you access the performance of the SQL statements that you imported before upgrading the 10g database.

1. From a terminal window, referred to as the “first session,” execute the `setup_SPAbig10g.sh` script to set up your simulated 10g environment. In this simulated environment, the `OPTIMIZER_FEATURES_ENABLED` parameter is set to 10.2.0.2.

```
$ cd $HOME/labs/SPA
$ cat setup_SPAbig10g.sh

#!/bin/bash

rm /u01/app/oracle/admin/orcl/dpdump/apps.dmp
rm /u01/app/oracle/admin/orcl/dpdump/appsandstage.dmp

cp /home/oracle/labs/SPA/apps.dmp
/u01/app/oracle/admin/orcl/dpdump/apps.dmp

sqlplus /nolog <<FIN!
connect / as sysdba

@$HOME/workshops/spwkshSPA.sql

shutdown immediate

startup

set echo on

exec dbms_sqltune.drop_sqlset('STS_JFV','SYS');

drop user apps cascade;

host impdp system/oracle directory=DATA_PUMP_DIR
dumpfile=apps.dmp

select distinct last_analyzed from dba_tab_statistics where
owner='APPS';

EXECUTE DBMS_STATS.LOCK_SCHEMA_STATS('APPS');

ALTER DATABASE TEMPFILE
'/u01/app/oracle/oradata/orcl/temp01.dbf'
AUTOEXTEND ON NEXT 6400K MAXSIZE UNLIMITED;

FIN!
```

**Practice 12-1: Using the SQL Performance Analyzer (continued)**

```
$./setup_SPAbig10g.sh
```

2. **(Perform steps 2 and 3 at the same time.)** Generate a SQL Tuning Set (STS) called STS\_JFV that captures SQL statements from the cursor cache for approximately 12 minutes every five seconds. Make sure that you try to capture only those statements that come out of the SQL\_JFV module in the APPS schema. Also, this STS should belong to the SYS user. Use the capsts10g.sh script to perform this step.

```
$ cd /home/oracle/labs/SPA
$ cat capsts10g.sh
#!/bin/bash

sqlplus / as sysdba <<FIN!

SET ECHO ON;
SET TIMING ON;

begin
DBMS_SQLTUNE.CREATE_SQLSET (sqlset_name => 'STS_JFV');

dbms_sqltune.capture_cursor_cache_sqlset(
sqlset_name => 'STS_JFV' ,
basic_filter=> q'# module
like 'DWH_TEST%' and sql_text not like '%applicat%' and
parsing_schema_name in ('APPS') #' ,
time_limit => 8*60,
repeat_interval => 5);
end ;
/
show errors

FIN!

$./capsts10g.sh
```

3. **(Perform steps 2 and 3 at the same time.)** From a second terminal window, connected as the oracle user, execute your workload by using the wrkl10g\_jfv.sh script. This script runs a workload of 45 statements that are going to be captured in STS\_JFV automatically.

```
-- Second session
$ cd /home/oracle/labs/SPA
$./wrkl10g_jfv.sh

...
SQL> @6 Results in 1296 Ticks
SQL> SQL> SQL> SQL> @Statement 45
SQL> SQL> SQL> 2 3 4 5 6 7 8 9 10
11 12 13 14 15 16 17 18 19 20 21 22 23
```

**Practice 12-1: Using the SQL Performance Analyzer (continued)**

```

24 25 26 27 28 29 30 31 32 33 34 35 36
37 38 39 40 41 42 43 44 45 46 47 48 49
50 51 52 53 54 55 56 57 58 59 60 61 62
63 64 65 66 67 68 69 70 71 72 73 74 75
76 77 78 79 80 81 82 83 84 85 86 87 88
89 90 91 92 93 94 95 96 97 98 99 100 101
102 103 104
10 rows selected.

Elapsed: 00:00:30.24

Statistics

 92 recursive calls
 0 db block gets
6871559 consistent gets
10770 physical reads
 116 redo size
 2003 bytes sent via SQL*Net to client
 3147 bytes received via SQL*Net from client
 2 SQL*Net roundtrips to/from client
 2 sorts (memory)
 0 sorts (disk)
 10 rows processed

SQL> @10 Results in 1188 Ticks
SQL> SQL> SQL> SQL>
PL/SQL procedure successfully completed.

SQL> SQL> we are done
SQL> SQL> SQL> SQL> @End 2010:FEB-11:18:50:40

SQL> SQL> Disconnected ...

```

- After approximately eight minutes, both sessions should have finished. Check the content of STS\_JFV, and stage it in a table called APPS.STS\_JFV\_TAB. Use the stage\_sts.sh script to perform this step.

```

$ cd /home/oracle/labs/SPA
$ cat stage_sts.sh

#!/bin/bash
sqlplus /nolog <<FIN!

connect / as sysdba

set echo on
select name,statement_count from dba_sqlset;

drop table apps.sts_jfv_tab purge;

exec DBMS_SQLTUNE.CREATE_STGTAB_SQLSET('STS_JFV_TAB','APPS');

```

**Practice 12-1: Using the SQL Performance Analyzer (continued)**

```

exec
DBMS_SQLTUNE.PACK_STGTAB_SQLSET('STS_JFV','SYS','STS_JFV_TAB',
'APPS');
FIN!

$./stage_sts.sh

```

- Using Data Pump Export, export the APPS schema to the default Data Pump 10g directory. Use the `exportapps.sh` script to perform this step.

```

$ cd /home/oracle/labs/SPA
$ cat exportapps.sh

#!/bin/bash

sqlplus /nolog <<FIN!
connect apps/apps

drop table plan_table purge;

FIN!

rm /u01/app/oracle/admin/orcl/dpdump/appsandstage.dmp

expdp system/oracle_4U directory=DATA_PUMP_DIR
dumpfile=appsandstage schemas=apps

$./exportapps.sh

```

- In real 10g and 11g environments, you would copy the generated dump file from the default Data Pump 10g directory to the default Data Pump 11g directory. But in this simulated environment, they are the same directory. So nothing needs to be done.
- Now, restart your 11g environment to restore the database to an 11g environment. Use the `setup_SPAbig11g.sh` script to perform this step.

```

$./setup_SPAbig11g.sh

SQL> SQL> Connected.
SQL> SQL> Database closed.
Database dismounted.
ORACLE instance shut down.
SQL>
SQL> ORACLE instance started.

Total System Global Area 836976640 bytes
Fixed Size 1339740 bytes

```

## Practice 12-1: Using the SQL Performance Analyzer (continued)

```
Variable Size 796921508 bytes
Database Buffers 25165824 bytes
Redo Buffers 13549568 bytes
Database mounted.
Database opened.
SQL> SQL> Disconnected ...
```

- Using Data Pump Import, import the APPS schema into your 11g system. Use the `importapps.sh` script to perform this step.

```
$ cd /home/oracle/labs/SPA
$ cat ./importapps.sh

#!/bin/bash

sqlplus /nolog <<FIN!

Connect / as sysdba
drop user apps cascade;

host impdp system/oracle_4U directory=DATA_PUMP_DIR
dumpfile=appsandstage

FIN!

$./importapps.sh
```

- Unpack the previously imported staging table in the SYS schema. Use the `unpack_sts.sh` script to perform this step.

```
$ cd /home/oracle/labs/SPA
$ cat unpack_sts.sh

#!/bin/bash

sqlplus /nolog <<FIN!

Connect / as sysdba
set echo on

exec
DBMS_SQLTUNE.UNPACK_STGTAB_SQLSET('STS_JFV','SYS',TRUE,'STS_JF
V_TAB','APPS');

alter system flush buffer_cache;
alter system flush shared_pool;

FIN!

$./unpack_sts.sh
```

## Practice 12-1: Using the SQL Performance Analyzer (continued)

10. Use Enterprise Manager Database Control, as the SYS user, to test the behavior of STS\_JFV in the simulated 10g environment and compare it to the 11g environment. You will do this by changing the OPTIMIZER\_FEATURES\_ENABLED parameter. What are your conclusions?

**Note:** Screenshots of some of the steps described in the table are provided, as noted in the Page column, below the table. Refer to these screenshots as necessary.

| Step | Page                                                                    | Action                                                                  |                    |                                                                      |
|------|-------------------------------------------------------------------------|-------------------------------------------------------------------------|--------------------|----------------------------------------------------------------------|
| a    | Database home                                                           | Click the Software and Support tab.                                     |                    |                                                                      |
| b    | Software and Support page                                               | Click SQL Performance Analyzer in the Real Application Testing section. |                    |                                                                      |
| c    | SQL Performance Analyzer page                                           | Click Parameter Change.                                                 |                    |                                                                      |
| d    | Parameter Change (See the example screenshot, which follows the table.) | Set values.                                                             | Field              | Value                                                                |
|      |                                                                         |                                                                         | Task Name          | SPA_JFV1                                                             |
|      |                                                                         |                                                                         | SQL Tuning Set     | STS_JFV                                                              |
|      |                                                                         |                                                                         | Creation Method    | Execute SQLs                                                         |
|      |                                                                         |                                                                         | Per-SQL Time Limit | Unlimited                                                            |
|      |                                                                         |                                                                         | Parameter Name     | optimizer_features_enabled<br>(click the select, “flashlight”, icon) |
|      |                                                                         |                                                                         | Base Value         | 10.2.0.1                                                             |
|      |                                                                         |                                                                         | Changed Value      | 11.2.0.1                                                             |
|      |                                                                         |                                                                         | Comparis           | Elapsed Time                                                         |

## Practice 12-1: Using the SQL Performance Analyzer (continued)

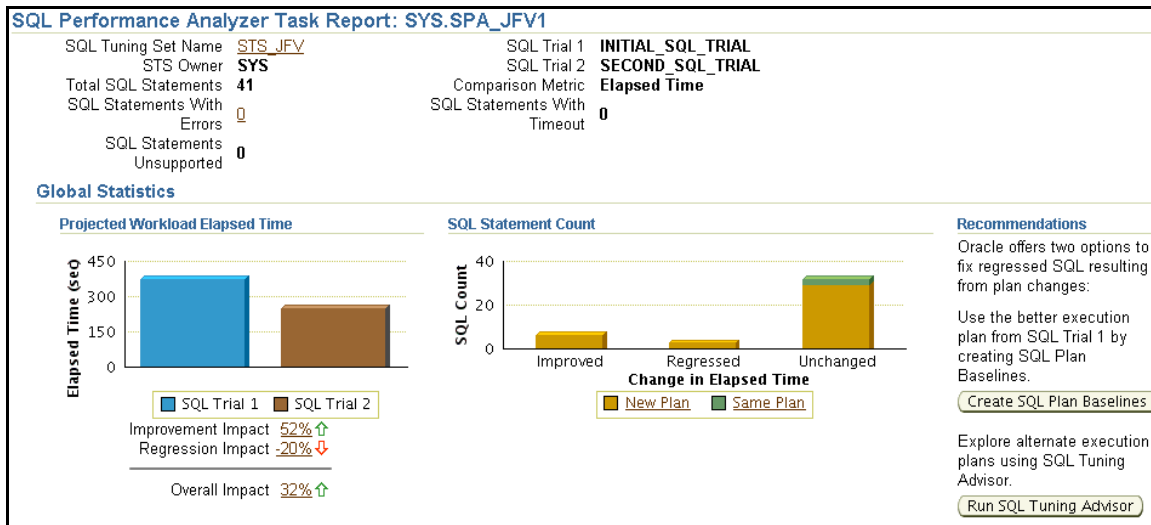
|   |                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                       |                  |             |
|---|------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|-------------|
|   |                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                       | on Metric        |             |
|   |                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                       | Schedule         | Immediately |
|   |                                                                                                                  | Click Submit.                                                                                                                                                                                                                                                                                                                                                         |                  |             |
| e | SQL Performance Analyzer                                                                                         | Let the screen refresh until the Status field shows Completed for your submitted task (SPA_JFV1). It may take up to 22 minutes to execute.                                                                                                                                                                                                                            |                  |             |
|   |                                                                                                                  | Click the SPA_JFV1 link.                                                                                                                                                                                                                                                                                                                                              |                  |             |
| f | SQL Performance Analyzer Task: SYS.SPA_JFV1                                                                      | This page shows the SQL Trials and the SQL Trial Comparisons. Click the eyeglass icon in the SQL Trial Comparisons table in the Comparison Report column. There should be only one entry in this table.                                                                                                                                                               |                  |             |
| g | SQL Performance Analyzer Task Report: SYS.SPA_JFV1<br><br>(See the example screenshots, which follow the table.) | You can tell that the second trial was faster than the first one. The second trial corresponds to the run that used the 11g optimizer features. The benefit depends on the statements that were captured and executed. You will see the number of the statements that use a new plan with 11g optimization. However, some statements regressed with 11g optimization. |                  |             |
|   |                                                                                                                  | From “Top 10 SQL Statements Based on Impact on Workload,” make a note of the SQL ID corresponding to the statements that have the highest negative impact on your workload. You use this information in Practice 13-1.                                                                                                                                                |                  |             |
|   |                                                                                                                  | Click the SQL ID link corresponding to the most regressing one.                                                                                                                                                                                                                                                                                                       |                  |             |
| h | SQL Details                                                                                                      | You can see the text of your statement and the different execution plans that were used during both trials.                                                                                                                                                                                                                                                           |                  |             |
|   |                                                                                                                  | Click the Back button on the browser.                                                                                                                                                                                                                                                                                                                                 |                  |             |
| i | SQL Performance Analyzer Task Result                                                                             | Click Run SQL Tuning Advisor in the Recommendations section.                                                                                                                                                                                                                                                                                                          |                  |             |
| j | Schedule SQL Tuning Task                                                                                         | Set values.                                                                                                                                                                                                                                                                                                                                                           | Field            | Value       |
|   |                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                       | Tuning Task Name | TUNE_JFV1   |
|   |                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                       | Schedule         | Immediately |



## Practice 12-1: Using the SQL Performance Analyzer (continued)

|   |                                                          |                                                                                                                                                                                                                                                                                                                                                    |
|---|----------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|   |                                                          | Click OK.                                                                                                                                                                                                                                                                                                                                          |
| k | SQL Performance Analyzer Task Result                     | Click the SQL Tune Report SYS.TUNE_JFV1 link in the Recommendations section.                                                                                                                                                                                                                                                                       |
|   |                                                          | Use the browser refresh button until you see a report that includes graphics. This may take a minute or longer.                                                                                                                                                                                                                                    |
| l | SQL Tuning Result Summary:SYS.TUNE_JFV1 (see screenshot) | <p>On this page, you can see that there are SQL Statements that can be improved by using profiles, indexes, or alternative SQL plans (using SQL Baselines). Using one of these recommendations, you can reduce or eliminate the regressions.</p> <p>Do not implement at this time; the following lessons explore these options in more detail.</p> |

g) The following screenshot displays the SQL Performance Analyzer Task Result: SYS.SPA\_JFV1 page:

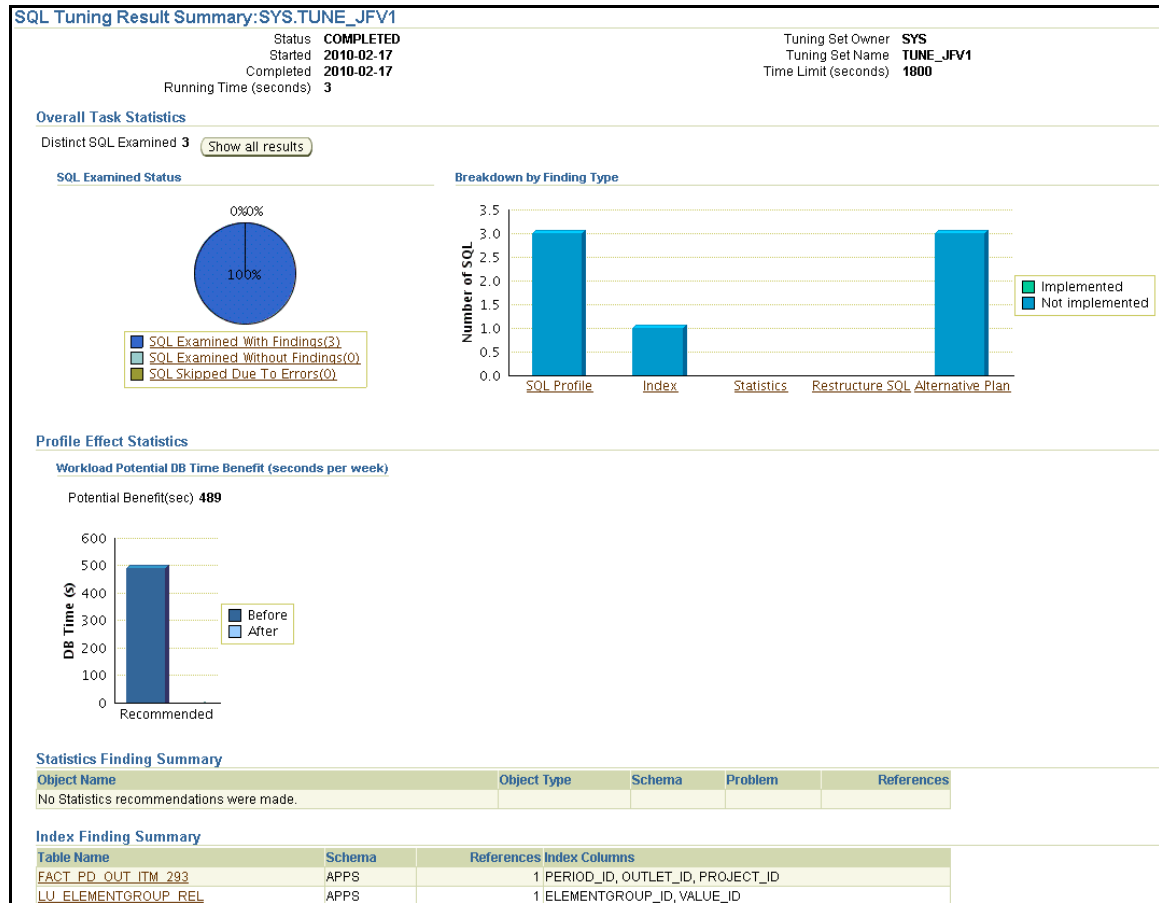


## Practice 12-1: Using the SQL Performance Analyzer (continued)

| Top 10 SQL Statements Based on Impact on Workload |                                |                            |                    |             |                       |          |
|---------------------------------------------------|--------------------------------|----------------------------|--------------------|-------------|-----------------------|----------|
|                                                   | SQL ID                         | Net Impact on Workload (%) | Elapsed Time (sec) |             | Net Impact on SQL (%) | New Plan |
|                                                   |                                |                            | SQL Trial 1        | SQL Trial 2 |                       |          |
| ↓                                                 | <a href="#">gkc6w4zhaysbk</a>  | -18.000                    | 103.999            | 171.012     | -64.440               | Y        |
| ↑                                                 | <a href="#">gfw9mbv2h44ns</a>  | 9.290                      | 34.936             | 0.368       | 98.950                | Y        |
| ↑                                                 | <a href="#">21t61c8b39njg</a>  | 9.200                      | 34.574             | 0.326       | 99.060                | Y        |
| ↑                                                 | <a href="#">94jmd58x6ch6d</a>  | 9.190                      | 34.611             | 0.407       | 98.820                | Y        |
| ↑                                                 | <a href="#">dx1c9zbr6w8h6</a>  | 8.870                      | 33.329             | 0.321       | 99.040                | Y        |
| ↑                                                 | <a href="#">2pg3srqh3qasz</a>  | 8.030                      | 30.244             | 0.351       | 98.840                | Y        |
| ↑                                                 | <a href="#">1wfywkvvcwm2sp</a> | 7.410                      | 73.676             | 46.098      | 37.430                | Y        |
| ↓                                                 | <a href="#">7866641pah2zg</a>  | -1.010                     | 3.457              | 7.235       | -109.300              | Y        |
| ↓                                                 | <a href="#">cpm5u0m6b68k5</a>  | -1.010                     | 3.543              | 7.317       | -106.560              | Y        |
| ↑                                                 | <a href="#">dvaxhpys7hpdv</a>  | 0.550                      | 2.923              | 0.860       | 70.560                | Y        |

## Practice 12-1: Using the SQL Performance Analyzer (continued)

1) The following screenshot displays the SQL Tuning Result Summary:SYS.TUNE\_JFV1.



## Practices for Lesson 13

SQL Plan Management (SPM) is an Oracle Database 11g feature that provides controlled execution plan evolution.

With SPM, the optimizer automatically manages execution plans and ensures that only known or verified plans are used.

When a new plan is found for a SQL statement, it will not be used until it has been verified to have comparable or better performance than the current plan.

SPM has three main components:

1. Plan Capture
2. Plan Selection
3. Plan Verification

In these practices, you see each of these components in action.

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor

In this practice, you use the result of your SQL Performance Analyzer comparison trial to seed SQL plan baselines for regressing SQL statements to avoid future plan regressions. The regressions could have come from a variety of sources. Regressions could be due to changes in statistics or changes in database parameters.

1. Use Enterprise Manager to check whether there are baselines for your most regressing statements. Then use Enterprise Manager to seed SQL Plan Baselines for your most regressing statements.

**Note:** Screenshots of some of the steps described in the table are provided, as noted in the Page column, below the table. Refer to these screenshots as necessary.

| Step | Page                                                                                                     | Action                                                                                                                                                                                                                                        |
|------|----------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a    | Current EM page                                                                                          | Click the Database tab.                                                                                                                                                                                                                       |
| b    | Database home                                                                                            | Click the Server tab.                                                                                                                                                                                                                         |
| c    | Server                                                                                                   | Click SQL Plan Control in the Query Optimizer section.                                                                                                                                                                                        |
| d    | SQL Plan Control                                                                                         | Click the SQL Plan Baseline tab.                                                                                                                                                                                                              |
|      |                                                                                                          | You should not see any baseline for your regressing statements.                                                                                                                                                                               |
|      |                                                                                                          | Click the Database tab.                                                                                                                                                                                                                       |
| e    | Database home                                                                                            | Click the Software and Support tab.                                                                                                                                                                                                           |
| f    | Software and Support                                                                                     | Click the SQL Performance Analyzer link in the Real Application Testing section.                                                                                                                                                              |
| g    | SQL Performance Analyzer                                                                                 | Click the SPA_JFV1 link in the SQL Performance Analyzer Tasks table.                                                                                                                                                                          |
| h    | SQL Performance Analyzer Task:<br>SYS.SPA_JFV1<br>(See the example screenshot, which follows the table.) | Click the eyeglass icon in the Comparison Report column for your first replay trial comparison.<br><br>Or you can click the View Latest Report link at the top of the page to view the same report, when there is only one Comparison Report. |
| i    | SQL Performance Analyzer Task Result:                                                                    | Click Create SQL Plan Baselines in the Recommendations section.                                                                                                                                                                               |

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

|   |                                                                                                                  |                                                                                                                                            |
|---|------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
|   | SYS.SPA_JFV1                                                                                                     |                                                                                                                                            |
| j | SQL Performance Analyzer Task Result: SYS.SPA_JFV1<br><br>(See the example screenshot, which follows the table.) | You should now see all regressing statements that meet the criteria. If the impact is before a certain threshold, they will not be listed. |
|   |                                                                                                                  | Set values                                                                                                                                 |
|   |                                                                                                                  | Field                                                                                                                                      |
|   |                                                                                                                  | Value                                                                                                                                      |
|   |                                                                                                                  | Job Name                                                                                                                                   |
|   |                                                                                                                  | baseline_jfv1                                                                                                                              |
|   |                                                                                                                  | Schedule                                                                                                                                   |
|   |                                                                                                                  | Immediately                                                                                                                                |
|   |                                                                                                                  | Click OK.                                                                                                                                  |
| k | SQL Performance Analyzer Task Result: SYS.SPA_JFV1                                                               | You should see an Information message, at the top of the page, saying that your job has been successfully submitted.                       |
|   |                                                                                                                  | Click the Database tab.                                                                                                                    |
| l | Database home                                                                                                    | Click the Server tab.                                                                                                                      |
| m | Server                                                                                                           | Click SQL Plan Control.                                                                                                                    |
| n | SQL Plan Control                                                                                                 | Click the SQL Plan Baseline tab.                                                                                                           |
|   |                                                                                                                  | You should now see the SQL plan baselines that were created. You should have one baseline for each regressing statement.                   |

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

h) The following screenshot displays the SQL Performance Analyzer Task: SYS.SPA\_JFV1 page:

**SQL Performance Analyzer Task: SYS.SPA\_JFV1**

[View Latest Report](#) Page Refreshed Feb 17, 2010 7:05:54 AM UTC [Refresh](#)

The SQL Performance Analyzer Task is a container for experimental results of executing a specific SQL Tuning Set under changed environmental conditions and assessing the impact of environmental changes on STS execution performance.

[SQL Tuning Set](#)

[SQL Trials](#)

A SQL Trial captures the execution performance of the SQL Tuning Set under specific environmental conditions. [Create SQL Trial](#)

| SQL Trial Name    | Description                                           | Created         | SQL Executed | Status    |
|-------------------|-------------------------------------------------------|-----------------|--------------|-----------|
| INITIAL_SQL_TRIAL | parameter optimizer_features_enable set to '10.2.0.1' | 2/17/10 6:07 AM | Yes          | COMPLETED |
| SECOND_SQL_TRIAL  | parameter optimizer_features_enable set to '11.2.0.1' | 2/17/10 6:17 AM | Yes          | COMPLETED |

[SQL Trial Comparisons](#)

Compare SQL Trials to assess change impact of environmental differences on SQL Tuning Set execution costs. [Run SQL Trial Comparison](#)

| Trial 1 Name      | Trial 2 Name     | Compare Metric | Created         | Status    | Comparison Report                 | SQL Tune Report                 |
|-------------------|------------------|----------------|-----------------|-----------|-----------------------------------|---------------------------------|
| INITIAL_SQL_TRIAL | SECOND_SQL_TRIAL | Elapsed Time   | 2/17/10 6:26 AM | COMPLETED | <a href="#">Comparison Report</a> | <a href="#">SQL Tune Report</a> |

i) The following screenshot displays the SQL Performance Analyzer Task Report: SYS.SPA\_JFV1 page:

**SQL Performance Analyzer Task Report: SYS.SPA\_JFV1**

SQL Tuning Set Name: [SYS\\_JFV](#) SQL Trial 1: [INITIAL\\_SQL\\_TRIAL](#)  
 STS Owner: **SYS** SQL Trial 2: [SECOND\\_SQL\\_TRIAL](#)  
 Total SQL Statements: **40** Comparison Metric: **Elapsed Time**  
 SQL Statements With Errors: **0** SQL Statements With Timeout: **0**  
 SQL Statements Unsupported: **0**

**Global Statistics**

**Projected Workload Elapsed Time**

■ SQL Trial 1 ■ SQL Trial 2

Improvement Impact: **56%** ↑  
 Regression Impact: **-25%** ↓  
 Overall Impact: **31%** ↑

**SQL Statement Count**

■ New Plan ■ Same Plan

**Recommendations**

Oracle offers two options to fix regressed SQL resulting from plan changes:

Use the better execution plan from SQL Trial 1 by creating SQL Plan Baselines.

[Create SQL Plan Baselines](#)

SQL Tune Report: [SYS.TUNE\\_JFV1](#)

**Top 10 SQL Statements Based on Impact on Workload**

| SQL ID                        | Net Impact on Workload (%) | Elapsed Time (sec) |             | Net Impact on SQL (%) | New Plan |
|-------------------------------|----------------------------|--------------------|-------------|-----------------------|----------|
|                               |                            | SQL Trial 1        | SQL Trial 2 |                       |          |
| <a href="#">gkc6w4zhaysbk</a> | -21.950                    | 99.152             | 165.800     | -67.220               | Y        |
| <a href="#">21t61c8b39njq</a> | 12.560                     | 38.582             | 0.456       | 98.820                | Y        |
| <a href="#">dx1c9zbr6w8h6</a> | 11.520                     | 35.386             | 0.405       | 98.860                | Y        |
| <a href="#">94imd58x6ch6d</a> | 11.380                     | 35.082             | 0.522       | 98.510                | Y        |
| <a href="#">qfw9mbv2h44ns</a> | 11.010                     | 33.881             | 0.458       | 98.650                | Y        |
| <a href="#">2pq3srqh3qasz</a> | 9.710                      | 29.902             | 0.416       | 98.610                | Y        |
| <a href="#">cpm5u0m6b68k5</a> | -1.570                     | 3.721              | 8.499       | -128.430              | Y        |
| <a href="#">7866641pah2zq</a> | -1.480                     | 4.091              | 8.582       | -109.750              | Y        |
| <a href="#">dvaxhpvs7hpdz</a> | 0.850                      | 3.173              | 0.585       | 81.550                | Y        |
| <a href="#">av7s4k9vz34hk</a> | 0.520                      | 1.797              | 0.233       | 87.050                | Y        |

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

j) The following screenshot displays the SQL Performance Analyzer Task Result: SYS.SPA\_JFV1 Create Baseline page:

**SQL Performance Analyzer Task Report: SYS.SPA\_JFV1** Cancel OK

**Create SQL Plan Baselines**

SQL Plan Baselines enable the optimizer to avoid performance regressions by requiring new plans to be at least as good as the better plans found in SQL trial 1.

**Regressed New Plan SQL Statements**

| SQL ID        | Net Impact on Workload (%) | Elapsed Time      |                  | Net Impact on SQL (%) | % of Workload     |                  |
|---------------|----------------------------|-------------------|------------------|-----------------------|-------------------|------------------|
|               |                            | INITIAL_SQL_TRIAL | SECOND_SQL_TRIAL |                       | INITIAL_SQL_TRIAL | SECOND_SQL_TRIAL |
| gkc6w4zhaysbk | -21.950                    | 99.152            | 165.800          | -67.220               |                   |                  |
| cqm5u0m6b68k5 | -1.570                     | 3.721             | 8.499            | -128.430              |                   |                  |
| 7866641pah2zg | -1.480                     | 4.091             | 8.582            | -109.750              |                   |                  |

**Job Parameters**

Job Name:

Description:

**Schedule**

☒ Immediately  
☐ Later

Time Zone:

Date:

(example: Feb 17, 2010)

Time:    ☒ AM ☐ PM

2. Using Enterprise Manager, set the USE\_SQL\_PLAN\_BASELINES parameter.

| Step | Page                  | Action                                                 |                                                  |
|------|-----------------------|--------------------------------------------------------|--------------------------------------------------|
| a    | Database home         | Click the Server tab.                                  |                                                  |
| b    | Server                | Click SQL Plan Control in the Query Optimizer section. |                                                  |
| c    | SQL Plan Control      | Click the SQL Plan Baseline tab.                       |                                                  |
| d    | SQL Baselines subpage | If                                                     | Use SQL Plan Baselines is set to TRUE.           |
|      |                       | Then                                                   | Go to step 3.                                    |
|      |                       | Else                                                   | Click the FALSE link for Use SQL Plan Baselines. |



### Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

|   |                           |                                               |
|---|---------------------------|-----------------------------------------------|
|   |                           | Click the Database tab.                       |
| e | Initialization Parameters | Set optimizer_use_sql_plan_baselines to TRUE. |
|   |                           | Click OK.                                     |

- Use Enterprise Manager to create a new SQL Performance Analyzer session to look at the difference between the previous 11g trial and a new one that takes into account the previously created baselines.

**Note:** Screenshots of some of the steps described in the table are provided, as noted in the Page column, below the table. Refer to these screenshots as necessary.

| Step | Page                                                                       | Action                                                                                    |                                         |                                         |
|------|----------------------------------------------------------------------------|-------------------------------------------------------------------------------------------|-----------------------------------------|-----------------------------------------|
| a    | Database home                                                              | Click the Software and Support tab.                                                       |                                         |                                         |
| b    | Software and Support                                                       | Click SQL Performance Analyzer.                                                           |                                         |                                         |
| c    | SQL Performance Analyzer                                                   | Click the SPA_JFV1 link.                                                                  |                                         |                                         |
| d    | SQL Performance Analyzer Task: SYS.SPA_JFV1                                | Click Create SQL Trial.                                                                   |                                         |                                         |
| e    | Create SQL Trial<br>(See the example screenshot, which follows the table.) | Set values.                                                                               | Field                                   | Value                                   |
|      |                                                                            |                                                                                           | SQLTrial Name                           | after_baselines                         |
|      |                                                                            |                                                                                           | Trial environment established check box | Select (on the lower right of the page) |
|      |                                                                            |                                                                                           | Per-SQL Time Limit                      | Unlimited                               |
|      |                                                                            |                                                                                           | Schedule                                | Immediately                             |
|      |                                                                            | Click Submit.                                                                             |                                         |                                         |
| f    | SQL Performance Analyzer Task: SYS.SPA_JFV1                                | Click Refresh until you see COMPLETED in the Status column for the AFTER_BASELINES trial. |                                         |                                         |
|      |                                                                            | Click Run SQL Trial Comparison.                                                           |                                         |                                         |

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

| g | Run SQL Trial Comparison<br><br>(See the example screenshot, which follows the table.) | Set values.                                                                                                                              | Field             | Value             |
|---|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|-------------------|-------------------|
|   |                                                                                        |                                                                                                                                          | Trial 1 Name      | INITIAL_SQL_TRIAL |
|   |                                                                                        |                                                                                                                                          | Trial 2 Name      | AFTER_BASELINES   |
|   |                                                                                        |                                                                                                                                          | Comparison Metric | Elapsed Time      |
|   |                                                                                        |                                                                                                                                          | Schedule          | Immediately       |
|   | Click Submit.                                                                          |                                                                                                                                          |                   |                   |
| h | SQL Performance Analyzer Task:<br>SYS.SPA_JFV1                                         | You should see a Confirmation message, and a new line in the SQL Trial Comparisons section of the page.                                  |                   |                   |
|   |                                                                                        | Click the corresponding eyeglass icon in the Comparison Report column.                                                                   |                   |                   |
| i | SQL Performance Analyzer Task Result:<br>SYS.SPA_JFV1                                  | You should no longer see any significant regressing statements, and all previously regressing statements should now be executing faster. |                   |                   |

### e) Create SQL Trial.

#### Create SQL Trial

SQL Trials capture execution performance of the SQL Tuning Set under a given optimizer environment.

SQL Performance Analyzer Task **SYS.SPA\_JFV1**

SQL Tuning Set **SYS.STS\_JFV**

\* SQL Trial Name

SQL Trial Description

Creation Method

Per-SQL Time Limit

☒ **TIP** Time limit is on elapsed time of test execution of SQL

#### Schedule

Time Zone

☒ Immediately

☐ Later

Date   
(example: Feb 17, 2010)

Time    ☒ AM ☐ PM

#### Trial environment determines results

The SQL Tuning Set remains constant under the SQL Performance Analyzer Task and its SQL is executed in isolation to create each SQL Trial. Performance differences between trials are thus attributed to environmental differences between trials.

Environmental changes affecting SQL optimization and performance may need to be made manually prior to execution of the Trial. These could include changing initialization parameters, gathering or setting optimizer statistics and creating indexes.

The Creation Method determines how the SQL Trial is created and what contents are generated, as follows:

- Executing SQLs generates both plans and statistics by actually running the SQL statements.
- Generating plans invokes the optimizer to create execution plans only without running the SQL statements.
- Remote execution and plan generation are done over a public database link on the remote system.

**NOTE: Be sure trial environment has been established prior to submitting.**

☒ Trial environment established

Cancel Submit

## Practice 13-1: Seeding SQL Plan Baselines from SQL Performance Advisor (continued)

g) Run SQL Trial Comparison.

### Run SQL Trial Comparison

Task Name

SYS.SPA\_JFV1

SQL Tuning Set

SYS.STS\_JFV

Trial 1 Name

INITIAL\_SQL\_TRIAL

Description

parameter optimizer\_features\_enable set to '11.2.0.1'

SQL Executed

Yes

Trial 2 Name

AFTER\_BASELINES

Description

SQL Executed

Yes

Comparison Metric

Elapsed Time

Compare trials to assess change impact

SQL Performance Analyzer trial comparison allows you to assess the impact on SQL Tuning Set performance of changes made between two trials.

It is important to know the difference between Trial 1 and Trial 2 execution environments in order to properly assign impacts to the changes between trials. Tracking environmental changes between trials is currently a user responsibility.

The selected comparison metric is used as the basis for comparison, and defaults to execute elapsed time when both trials contain test execution statistics. When execution statistics are not available, a less accurate comparison can be made using optimizer cost.

Schedule

Time Zone

Pacific/Pago\_Pago

☒ Immediately
 ☐ Later

Date

Feb 17, 2010

(example: Feb 17, 2010)

Time

6

58

05

☒ AM
 ☐ PM

4. Use Enterprise Manager to delete the created baselines.

| Step | Page                      | Action                                                          |
|------|---------------------------|-----------------------------------------------------------------|
| a    | Database home             | Click the Server tab.                                           |
| b    | Server                    | Click the SQL Plan Control link in the Query Optimizer section. |
| c    | SQL Plan Control          | Click the SQL Plan Baseline tab.                                |
| d    | SQL Plan Baseline subpage | Select all baselines.                                           |
|      |                           | Click Drop.                                                     |
| e    | Confirmation              | Click Yes.                                                      |

5. In a terminal window, run the `./cleanup 13 1` script.

```
$ cd $HOME/workshops
$./cleanup 13 1
```

## Practice 13-2: SQL Plan Management (SPM)

In this practice, you see all phases of using SQL Plan Management.

1. Before you can start this practice, you must set up a new user. Execute the `spm_setup.sh` script in the `$HOME/labs/SPM` directory to set up the environment for this practice. This script creates the SPM user that you use throughout the practice.

```
$ cd /home/oracle/labs/SPM
$ cat ./spm_setup.sh

#!/bin/bash

sqlplus /nolog <<FIN!
connect / as sysdba

set echo on

drop user spm cascade;

create user spm identified by spm;

grant dba to spm;

alter system set SHARED_POOL_SIZE = 220M SCOPE = MEMORY;

alter system flush shared_pool;

FIN!

$./spm_setup.sh
```

### Automatic Plan Capture

2. The first component of SPM is Plan Capture. There are two main ways to capture plans: automatically (at run time) or bulk load. In this practice, you turn on automatic plan capture so that the SPM repository is automatically populated for any repeatable SQL statement. Turn on automatic plan capture by setting the `optimizer_capture_sql_plan_baselines` initialization parameter to `TRUE` in your SQL\*Plus session, connected as the SPM user. After you have connected in your session, do not disconnect.

```
$ cd /home/oracle/labs/SPM

$ sqlplus spm/spm
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> show parameter baseline
```

| NAME                                 | TYPE    | VALUE |
|--------------------------------------|---------|-------|
| optimizer_capture_sql_plan_baselines | boolean | FALSE |
| optimizer_use_sql_plan_baselines     | boolean | FALSE |

```
SQL> alter session set optimizer_capture_sql_plan_baselines = TRUE;
```

```
Session altered.
```

```
SQL> alter session set optimizer_use_sql_plan_baselines = TRUE;
```

```
Session altered.
```

```
SQL>
```

3. Execute the following query in your SQL\*Plus session (**Note:** There are no spaces in /\*LOAD...):

```
select /*LOAD_AUTO*/ * from sh.sales
where quantity_sold > 40 order by prod_id;
```

Use the query1.sql script to execute the query.

```
SQL> @query1.sql
```

```
SQL> select /*LOAD_AUTO*/ * from sh.sales
2> where quantity_sold > 40 order by prod_id;
```

```
no rows selected
```

```
SQL>
```

4. Because this is the first time that you have seen this SQL statement, it is not yet repeatable, so there is no plan baseline for it. To confirm this, check whether there are any plan baselines that exist for your statement. (Use the check\_baselines.sql script.)

```
-- Should not return any rows
```

```
SQL> @check_baselines.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*LOAD_AUTO*/%';
```

```
no rows selected
```

## Practice 13-2: SQL Plan Management (SPM) (continued)

5. Execute the query that you ran in step 3.

```
SQL> @query1.sql
SQL> set echo on
SQL>
SQL> select /*LOAD_AUTO*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;

no rows selected

SQL>
```

6. The SQL statement is now known to be repeatable and a plan baseline is automatically captured. Check whether the plan baseline was loaded for the previous statement. What do you observe?

You can see from the output that a baseline has been created and enabled for this SQL statement. You can also tell that this plan was captured automatically by checking the values of the ORIGIN column. (Use the `check_baselines.sql` script.)

```
-- You should see one accepted entry

SQL> @check_baselines.sql
SQL> set echo on
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*LOAD_AUTO*/%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

8.0622E+18 SYS_SQL_6fe28d438dfc352f
select /*LOAD_AUTO*/ * from sh.sales where quantity_sold > 40
order by prod_id
SYS_SQL_PLAN_8dfc352f54bc8843 AUTO-CAPTURE YES YES NO YES

SQL>
```

7. Change or alter the optimizer mode to use FIRST\_ROWS optimization and execute your statement. Describe what happens.

## Practice 13-2: SQL Plan Management (SPM) (continued)

- a) This triggers the SQL statement to execute with a different plan.

```
SQL> alter session set optimizer_mode = first_rows;

Session altered.

SQL> @query1.sql
SQL> set echo on
SQL>
SQL> select /*LOAD_AUTO*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;

no rows selected

SQL>
```

- b) Because the SQL statement will have a new plan, another plan baseline is automatically captured. You can confirm this by checking the plan baseline again.

```
SQL> @check_baselines.sql
SQL> set echo on
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*LOAD_AUTO*/%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

8.0622E+18 SYS_SQL_6fe28d438dfc352f
select /*LOAD_AUTO*/ * from sh.sales where quantity_sold > 40
order by prod_id
SYS_SQL_PLAN_8dfc352f11df68d0 AUTO-CAPTURE YES NO NO YES

8.0622E+18 SYS_SQL_6fe28d438dfc352f
select /*LOAD_AUTO*/ * from sh.sales where quantity_sold > 40
order by prod_id
SYS_SQL_PLAN_8dfc352f54bc8843 AUTO-CAPTURE YES YES NO YES
```

- c) Now you see two plan baselines for your query, but notice that the new plan has not been accepted. This new plan will have to be validated before it is acceptable as a good plan to use.
8. Reset the optimizer mode to the default values and disable autocapture of plan baselines.

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> alter session set optimizer_mode = all_rows;

Session altered.

SQL> alter session set optimizer_capture_sql_plan_baselines =
FALSE;

Session altered.
```

9. Purge the plan baselines and confirm that the SQL plan baseline is empty. Use `purge_auto_baseline.sql`.

```
SQL> @purge_auto_baseline.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> exec :cnt :=
dbms_spm.drop_sql_plan_baseline('SYS_SQL_6fe28d438dfc352f');

PL/SQL procedure successfully completed.

SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*LOAD_AUTO*/%';

no rows selected

SQL>
```

**Loading Plans from SQL Tuning Sets**

10. Still connected to your SQL\*Plus session, check the execution plan for the following SQL statement, and then execute it (use `explain_query2.sql` and `query2.sql`).

```
select /*LOAD_STS*/ * from sh.sales
where quantity_sold > 40 order by prod_id;
```

```
SQL> @explain_query2.sql

-- You should see a Full Table Scan

SQL> set echo on
SQL>
SQL> explain plan for
2 select /*LOAD_STS*/ * from sh.sales
3 where quantity_sold > 40 order by prod_id;

Explained.
```



**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

SQL>
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));

PLAN_TABLE_OUTPUT

Plan hash value: 3803407550

| Id | Operation | Name |

0	SELECT STATEMENT	
1	SORT ORDER BY	
2	PARTITION RANGE ALL	
3	TABLE ACCESS FULL	SALES

10 rows selected.

SQL> @query2.sql
SQL> set echo on
SQL>
SQL> select /*LOAD_STS*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;

no rows selected

```

11. Alter the optimizer mode to use FIRST\_ROWS optimization, check the execution plan, and execute the query.

```

SQL> alter session set optimizer_mode = first_rows;

Session altered.

-- You should see a different plan
-- (Table Access By Local Index)

SQL> @explain_query2.sql
SQL> set echo on
SQL>
SQL> explain plan for
 2 select /*LOAD_STS*/ * from sh.sales
 3 where quantity_sold > 40 order by prod_id;

Explained.

SQL>
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));

PLAN_TABLE_OUTPUT

```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

Plan hash value: 899219946

```

| Id | Operation                         | Name            |
|----|-----------------------------------|-----------------|
| 0  | SELECT STATEMENT                  |                 |
| 1  | SORT ORDER BY                     |                 |
| 2  | PARTITION RANGE ALL               |                 |
| 3  | TABLE ACCESS BY LOCAL INDEX ROWID | SALES           |
| 4  | BITMAP CONVERSION TO ROWIDS       |                 |
| 5  | BITMAP INDEX FULL SCAN            | SALES_PROMO_BIX |

```

PLAN_TABLE_OUTPUT

```

```
12 rows selected.
```

```
SQL>
```

```
SQL> @query2.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> select /*LOAD_STS*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;
```

```
no rows selected
```

```
SQL>
```

12. Reset the optimizer mode to ALL\_ROWS optimization.

```
SQL> alter session set optimizer_mode = all_rows;
```

```
Session altered.
```

13. Verify that there are no baseline plans for your statement (use the check\_baselines2.sql script).

```
SQL> @check_baselines2.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*LOAD_STS*/%';
```

```
no rows selected
```

```
SQL>
```

14. Create a SQL tuning set that captures the SELECT statements that contain the LOAD\_STS hint. These statements are in the cursor cache. The STS should be called SPM\_STS and owned by the SPM user. Use the catchup\_sts.sql script to capture these statements.

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

SQL> @catchup_sts.sql
SQL> set echo on
SQL>
SQL> exec sys.dbms_sqltune.create_sqlset(-
> sqlset_name => 'SPM_STS', sqlset_owner => 'SPM');

PL/SQL procedure successfully completed.

SQL>
SQL> DECLARE
 2 stscur dbms_sqltune.sqlset_cursor;
 3 BEGIN
 4 OPEN stscur FOR
 5 SELECT VALUE(P)
 6 FROM TABLE(dbms_sqltune.select_cursor_cache(
 7 'sql_text like ''select /*LOAD_STS*/%'',
 8 null, null, null, null, null, null, 'ALL')) P;
 9
 10 -- populate the sqlset
 11 dbms_sqltune.load_sqlset(sqlset_name => 'SPM_STS',
 12 populate_cursor => stscur,
 13 sqlset_owner => 'SPM');
 14 END;
 15 /

PL/SQL procedure successfully completed.

SQL>

```

15. Verify which SQL statements are in SPM\_STS. (Use the check\_sts.sql script.)

SPM\_STS has your SQL statement with two different plans.

```

SQL> @check_sts.sql
SQL> set echo on
SQL>
SQL> select sql_text from dba_sqlset_statements
 2 where sqlset_name='SPM_STS';

SQL_TEXT

select /*LOAD_STS*/ * from sh.sales
where quantity_sold > 40 order by prod_id

select /*LOAD_STS*/ * from sh.sales
where quantity_sold > 40 order by prod_id

SQL>

```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

16. Populate the plan baseline repository with the plans found in SPM\_STS. Use the `populate_baseline.sql` script for that:

```
SQL> @populate_baseline.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> exec :cnt := dbms_spm.load_plans_from_sqlset(-
> sqlset_name => 'SPM_STS', -
> basic_filter => 'sql_text like ''select
/*LOAD_STS*/%'');

PL/SQL procedure successfully completed.
SQL>
```

17. Confirm that the plan baselines are loaded and note the value in the origin column. What do you observe? (Use `check_baselines2.sql`.)

You should see `MANUAL-LOAD` because you manually loaded these plans. Also note that this time, both plans are accepted.

```
SQL> @check_baselines2.sql
SQL> set echo on
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*LOAD_STS*/%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

1.2134E+19 SYS_SQL_a8632bd857a4a25e
select /*LOAD_STS*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_57a4a25e11df68d0 MANUAL-LOAD YES YES NO YES
1.2134E+19 SYS_SQL_a8632bd857a4a25e
select /*LOAD_STS*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_57a4a25e54bc8843 MANUAL-LOAD YES YES NO YES

SQL>
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

18. Purge the plan baselines and drop SPM\_STS. Use the `purge_sts_baseline.sql` script:

```
SQL> @purge_sts_baseline.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> exec :cnt := dbms_spm.drop_sql_plan_baseline(-
> 'SYS_SQL_a8632bd857a4a25e');

PL/SQL procedure successfully completed.

SQL>
SQL> print cnt;

 CNT

 2

SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*LOAD_STS*/%';

no rows selected

SQL>
SQL> exec sys.dbms_sqltune.drop_sqlset(-
> sqlset_name => 'SPM_STS', -
> sqlset_owner => 'SPM');

PL/SQL procedure successfully completed.

SQL>
```

**Loading Plans from the Cursor Cache**

19. Now, you see how to directly load plan baselines from the cursor cache. Before you begin, you need some SQL statements. Still connected to your SQL\*Plus session, check the execution plan for the following SQL statement, and then execute it. (Use the `explain_query3.sql` and `query3.sql` scripts.)

```
select /*LOAD_CC*/ * from sh.sales
where quantity_sold > 40 order by prod_id;
```

```
SQL> @explain_query3.sql
SQL> set echo on
SQL>
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> explain plan for
 2 select /*LOAD_CC*/ * from sh.sales
 3 where quantity_sold > 40 order by prod_id;
```

Explained.

```
SQL>
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));
```

PLAN\_TABLE\_OUTPUT

-----

Plan hash value: 3803407550

| Id | Operation           | Name  |
|----|---------------------|-------|
| 0  | SELECT STATEMENT    |       |
| 1  | SORT ORDER BY       |       |
| 2  | PARTITION RANGE ALL |       |
| 3  | TABLE ACCESS FULL   | SALES |

-----

10 rows selected.

```
SQL>
SQL> @query3.sql
SQL> set echo on
SQL>
SQL> select /*LOAD_CC*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;
```

no rows selected

SQL>

20. Now change the optimizer mode to use FIRST\_ROWS optimization and execute the same script as in the previous step. What do you observe?

a) You should see a different execution plan.

```
SQL> alter session set optimizer_mode = first_rows;
```

Session altered.

```
SQL> @explain_query3.sql
SQL> set echo on
SQL>
SQL> explain plan for
 2 select /*LOAD_CC*/ * from sh.sales
 3 where quantity_sold > 40 order by prod_id;
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

Explained.

SQL>

```
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));
```

PLAN\_TABLE\_OUTPUT

-----

Plan hash value: 899219946

-----

| Id | Operation                         | Name            |
|----|-----------------------------------|-----------------|
| 0  | SELECT STATEMENT                  |                 |
| 1  | SORT ORDER BY                     |                 |
| 2  | PARTITION RANGE ALL               |                 |
| 3  | TABLE ACCESS BY LOCAL INDEX ROWID | SALES           |
| 4  | BITMAP CONVERSION TO ROWIDS       |                 |
| 5  | BITMAP INDEX FULL SCAN            | SALES_PROMO_BIX |

PLAN\_TABLE\_OUTPUT

-----

12 rows selected.

SQL>

```
SQL> @query3.sql
```

```
SQL> set echo on
```

SQL>

```
SQL> select /*LOAD_CC*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;
```

no rows selected

SQL>

21. Reset the optimizer mode to ALL\_ROWS.

```
SQL> alter session set optimizer_mode = all_rows;
```

Session altered.

22. Now that the cursor cache is populated, you must get the SQL ID for your SQL statement. Use the SQL ID to filter the content of the cursor cache and load the baselines with these two plans. Use the load\_cc\_baseline.sql script.

```
SQL> @load_cc_baseline.sql
```

```
SQL> set echo on
```

SQL>

```
SQL> variable cnt number;
```

SQL>

```
SQL> variable sqlid varchar2(20);
```

SQL>

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

SQL> begin
 2 select distinct sql_id into :sqlid from v$sql
 3 where sql_text like 'select /*LOAD_CC*/%';
 4 end;
 5 /

PL/SQL procedure successfully completed.

SQL>
SQL> print sqlid;

SQLID

6qc9wxhgzzz8g

SQL>
SQL> exec :cnt := dbms_spm.load_plans_from_cursor_cache(-
> sql_id => :sqlid);

PL/SQL procedure successfully completed.

SQL>

```

23. Confirm that the baselines were loaded. (Use check\_baselines3.sql.)

```

SQL> @check_baselines3.sql
SQL> set echo on
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*LOAD_CC*/%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

1.7783E+19 SYS_SQL_f6cb7f742ef93547
select /*LOAD_CC*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_2ef9354711df68d0 MANUAL-LOAD YES YES NO YES

1.7783E+19 SYS_SQL_f6cb7f742ef93547
select /*LOAD_CC*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_2ef9354754bc8843 MANUAL-LOAD YES YES NO YES

-- You should see two accepted baselines

```



**Practice 13-2: SQL Plan Management (SPM) (continued)**

24. Purge the plan baselines.

```
SQL> @purge_cc_baseline.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> exec :cnt := dbms_spm.drop_sql_plan_baseline(-
> 'SYS_SQL_f6cb7f742ef93547');

PL/SQL procedure successfully completed.

SQL>
SQL> print cnt;

 CNT

 2

SQL>
SQL> REM Check that plan baselines were purged:
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*LOAD_CC*/%';

no rows selected

SQL>
```

**Optimizer Plan Selection**

25. Now that you know how to capture plans, look at how the optimizer selects which plan baselines to use. First, you create two baselines for a statement and show them being used. Then you disable one of the baselines and see the second baseline being used. Finally, you disable both baselines and show that now the optimizer falls back to the default behavior of a cost-based plan. Start by executing the same query twice, with different plans. Determine the execution of the following statement, and then execute it. (Use the `explain_query4.sql` and `query4.sql` scripts.)

```
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id
```

```
SQL> @explain_query4.sql
SQL> set echo on
SQL>
SQL> explain plan for
2 select /*SPM_USE*/ * from sh.sales
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
3 where quantity_sold > 40 order by prod_id;
```

Explained.

```
SQL>
```

```
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));
```

```
PLAN_TABLE_OUTPUT
```

```

Plan hash value: 3803407550
```

```

| Id | Operation | Name |

0	SELECT STATEMENT	
1	SORT ORDER BY	
2	PARTITION RANGE ALL	
3	TABLE ACCESS FULL	SALES

```

```
10 rows selected.
```

```
SQL>
```

```
SQL> @query4.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> select /*SPM_USE*/ * from sh.sales
2 where quantity_sold > 40 order by prod_id;
```

```
no rows selected
```

```
SQL>
```

26. Change the optimizer mode to use FIRST\_ROWS optimization and execute the same script as in the previous step. What do you observe?

You should see a different execution plan.

```
SQL> alter session set optimizer_mode = first_rows;
```

```
Session altered.
```

```
SQL> @explain_query4.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> explain plan for
2 select /*SPM_USE*/ * from sh.sales
3 where quantity_sold > 40 order by prod_id;
```

```
Explained.
```

```
SQL>
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> select * from
table(dbms_xplan.display(null,null,'basic'));

PLAN_TABLE_OUTPUT

Plan hash value: 899219946

Id	Operation	Name
0	SELECT STATEMENT	
1	SORT ORDER BY	
2	PARTITION RANGE ALL	
3	TABLE ACCESS BY LOCAL INDEX ROWID	SALES
4	BITMAP CONVERSION TO ROWIDS	
5	BITMAP INDEX FULL SCAN	SALES_PROMO_BIX

PLAN_TABLE_OUTPUT

12 rows selected.

SQL>
SQL> @query4.sql
SQL> set echo on
SQL>
SQL> select /*SPM_USE*/ * from sh.sales
2 where quantity_sold > 40 order by prod_id;

no rows selected

SQL>
```

27. Reset the optimizer mode to ALL\_ROWS.

```
SQL> alter session set optimizer_mode = all_rows;

Session altered.
```

28. Populate the baseline with the two plans for your statement directly from the cursor cache. Use the load\_use\_baseline.sql script for that. Then verify that baselines were loaded. What do you observe?

You should see both plan baselines loaded. Note that both plans have been marked acceptable. This is because both plans were present in the cursor cache at the time of the load, and because these plans have been manually loaded, it is assumed that both plans have acceptable performance.

```
SQL> @load_use_baseline.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

SQL>
SQL> variable sqlid varchar2(20);
SQL>
SQL> begin
 2 select distinct sql_id into :sqlid from v$sql
 3 where sql_text like 'select /*SPM_USE*/%';
 4 end;
 5 /

PL/SQL procedure successfully completed.

SQL>
SQL> print sqlid;

SQLID

2pma6tcaczdc8

SQL>
SQL> exec :cnt := dbms_spm.load_plans_from_cursor_cache(-
> sql_id => :sqlid);

PL/SQL procedure successfully completed.

SQL>
SQL> print cnt;

 CNT

 2

SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*SPM_USE*/%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

7.6492E+17 SYS_SQL_0a9d872600ece455
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_00ece45511df68d0 MANUAL-LOAD YES YES NO YES

7.6492E+17 SYS_SQL_0a9d872600ece455
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_00ece45554bc8843 MANUAL-LOAD YES YES NO YES

SQL>

```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

29. Determine the baseline and the execution plan used to execute the following query:

```
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id.
```

What do you observe?

The note at the end of the explain output tells you that the system is using a baseline. From the execution plan, you can see that you are using the first baseline, a full-table scan. Use the `explain_query4_note.sql` script.

```
SQL> @explain_query4_note.sql
SQL> set echo on
SQL>
SQL> explain plan for
 2 select /*SPM_USE*/ * from sh.sales
 3 where quantity_sold > 40 order by prod_id;

Explained.

SQL>
SQL> select * from table(dbms_xplan.display(null,null,'basic
+note'));
```

PLAN\_TABLE\_OUTPUT

---

Plan hash value: 3803407550

---

| Id | Operation           | Name  |
|----|---------------------|-------|
| 0  | SELECT STATEMENT    |       |
| 1  | SORT ORDER BY       |       |
| 2  | PARTITION RANGE ALL |       |
| 3  | TABLE ACCESS FULL   | SALES |

---

PLAN\_TABLE\_OUTPUT

---

Note

---

- SQL plan baseline "SYS\_SQL\_PLAN\_00ece45554bc8843" used for this statement

14 rows selected.

SQL>

**Practice 13-2: SQL Plan Management (SPM) (continued)**

30. Disable that plan baseline, and check whether the system uses the other plan baseline when executing the statement again. Use the `check_baseline_used.sql` script for that. What do you observe?
- a) Now from the execution plan, you see that you are using an index scan instead of a full-table scan. So this is the second baseline.

```
SQL> @check_baseline_used.sql
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> select sql_handle,plan_name
 2 from dba_sql_plan_baselines
 3 where sql_text like 'select /*SPM_USE*/%';
```

| SQL_HANDLE               | PLAN_NAME                     |
|--------------------------|-------------------------------|
| SYS_SQL_0a9d872600ece455 | SYS_SQL_PLAN_00ece45511df68d0 |
| SYS_SQL_0a9d872600ece455 | SYS_SQL_PLAN_00ece45554bc8843 |

```
SQL>
SQL>
SQL> exec :cnt := dbms_spm.alter_sql_plan_baseline(-
> sql_handle => 'SYS_SQL_0a9d872600ece455', -
> plan_name => 'SYS_SQL_PLAN_00ece45554bc8843', -
> attribute_name => 'ENABLED', -
> attribute_value => 'NO');
```

PL/SQL procedure successfully completed.

```
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*SPM_USE*/%';
```

| SIGNATURE  | SQL_HANDLE               | SQL_TEXT                                                                        | PLAN_NAME                     | ORIGIN      | ENA | ACC | FIX | AUT |
|------------|--------------------------|---------------------------------------------------------------------------------|-------------------------------|-------------|-----|-----|-----|-----|
| 7.6492E+17 | SYS_SQL_0a9d872600ece455 | select /*SPM_USE*/ * from sh.sales<br>where quantity_sold > 40 order by prod_id | SYS_SQL_PLAN_00ece45511df68d0 | MANUAL-LOAD | YES | YES | NO  | YES |
| 7.6492E+17 | SYS_SQL_0a9d872600ece455 | select /*SPM_USE*/ * from sh.sales                                              |                               |             |     |     |     |     |

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_00ece45554bc8843 MANUAL-LOAD NO YES NO YES

SQL>
SQL>
SQL> explain plan for select /*SPM_USE*/ * from sh.sales
 2 where quantity_sold > 40 order by prod_id;

Explained.

SQL>
SQL> select * from table(dbms_xplan.display(null, null, 'basic
+note'));

PLAN_TABLE_OUTPUT

Plan hash value: 899219946

| Id | Operation | Name |
|----|-----|
| 0 | SELECT STATEMENT | |
| 1 | SORT ORDER BY | |
| 2 | PARTITION RANGE ALL | |
| 3 | TABLE ACCESS BY LOCAL INDEX ROWID | SALES |
| 4 | BITMAP CONVERSION TO ROWIDS | |
| 5 | BITMAP INDEX FULL SCAN | SALES_PROMO_BIX |

PLAN_TABLE_OUTPUT

Note

- SQL plan baseline "SYS_SQL_PLAN_00ece45511df68d0" used
for this statement

16 rows selected.

SQL>

```

31. Disable this other plan baseline and check whether the system falls back to the cost-based approach when executing the explain plan for the statement. Use the `check_baseline_used2.sql` script.

You know that the optimizer has gone back to the default cost-based approach because there is no note at the end of the plan stating that a baseline was used.

```

SQL> @check_baseline_used2.sql
SQL> set echo on
SQL>

```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> variable cnt number;
SQL>
SQL> exec :cnt := dbms_spm.alter_sql_plan_baseline(-
> sql_handle => 'SYS_SQL_0a9d872600ece455', -
> plan_name =>
'SYS_SQL_PLAN_00ece45511df68d0', -
> attribute_name => 'ENABLED', -
> attribute_value => 'NO');
```

PL/SQL procedure successfully completed.

```
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*SPM_USE*/%';
```

```
SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

7.6492E+17 SYS_SQL_0a9d872600ece455
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_00ece45511df68d0 MANUAL-LOAD NO YES NO YES

7.6492E+17 SYS_SQL_0a9d872600ece455
select /*SPM_USE*/ * from sh.sales
where quantity_sold > 40 order by prod_id
SYS_SQL_PLAN_00ece45554bc8843 MANUAL-LOAD NO YES NO YES
```

```
SQL>
SQL>
SQL> explain plan for select /*SPM_USE*/ * from sh.sales
2 where quantity_sold > 40 order by prod_id;
```

Explained.

```
SQL>
SQL> select * from table(dbms_xplan.display(null, null, 'basic
+note'));
```

PLAN\_TABLE\_OUTPUT

Plan hash value: 3803407550

```

| Id | Operation | Name |

```



**Practice 13-2: SQL Plan Management (SPM) (continued)**

|   |                     |       |  |
|---|---------------------|-------|--|
| 0 | SELECT STATEMENT    |       |  |
| 1 | SORT ORDER BY       |       |  |
| 2 | PARTITION RANGE ALL |       |  |
| 3 | TABLE ACCESS FULL   | SALES |  |

---

10 rows selected.

SQL>

32. Drop the plan baselines and check whether they are purged. Use the `purge_use_baseline.sql` script.

```
SQL> @purge_use_baseline
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> exec :cnt := dbms_spm.drop_sql_plan_baseline(-
> 'SYS_SQL_0a9d872600ece455');

PL/SQL procedure successfully completed.

SQL>
SQL> print cnt;

 CNT

 2

SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like 'select /*SPM_USE*/%';

no rows selected

SQL>
```

33. One of the methods used to enable plan evolution (or plan verification) is Automatic SQL Tuning that is run as an automated task in a maintenance window. Automatic SQL Tuning targets only the high-load SQL statements; for them, it automatically implements actions such as making a successfully verified plan an accepted plan. Here, you manually trigger SQL tuning to find a better plan for a given SQL statement. First, determine the execution plan of the following statement:

```
select /*+ USE_NL(s c) FULL(s) FULL(c) */
 c.cust_id, sum(s.quantity_sold)
from sh.sales s, sh.customers c
```

## Practice 13-2: SQL Plan Management (SPM) (continued)

```
where s.cust_id = c.cust_id and c.cust_id < 2
group by c.cust_id
```

Some optimizer hints to the statements have been added to ensure that you get a less than optimal plan at first. Use the `check_evolve_plan.sql` script for that. What do you observe?

As you can see, the execution plan is being forced by the hints to perform two full-table scans, followed by a nest loop join.

```
SQL> @check_evolve_plan.sql
SQL> set echo on
SQL>
SQL> explain plan for
 2 select /*+ USE_NL(s c) FULL(s) FULL(c) */
 3 c.cust_id, sum(s.quantity_sold)
 4 from sh.sales s, sh.customers c
 5 where s.cust_id = c.cust_id and c.cust_id < 2
 6 group by c.cust_id;

Explained.

SQL>
SQL> select * from table(dbms_xplan.display(null, null));
```

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005616876

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time | Pstart | Pstop |
|----|-----------|------|------|-------|-------------|------|--------|-------|
|----|-----------|------|------|-------|-------------|------|--------|-------|

PLAN\_TABLE\_OUTPUT

|     |                     |           |   |    |         |          |   |    |
|-----|---------------------|-----------|---|----|---------|----------|---|----|
| 0   | SELECT STATEMENT    |           | 1 | 39 | 893 (2) | 00:00:11 |   |    |
| 1   | HASH GROUP BY       |           | 1 | 39 | 893 (2) | 00:00:11 |   |    |
| 2   | NESTED LOOPS        |           | 1 | 31 | 892 (2) | 00:00:11 |   |    |
| * 3 | TABLE ACCESS FULL   | CUSTOMERS | 1 | 5  | 405 (1) | 00:00:05 |   |    |
| 4   | PARTITION RANGE ALL |           | 1 | 8  | 488 (2) | 00:00:06 | 1 | 28 |
| * 5 | TABLE ACCESS FULL   | SALES     | 1 | 8  | 488 (2) | 00:00:06 | 1 | 28 |

PLAN\_TABLE\_OUTPUT

Predicate Information (identified by operation id):

```
3 - filter("C"."CUST_ID"<2)
5 - filter("S"."CUST_ID"<2 AND "S"."CUST_ID"="C"."CUST_ID")
```

18 rows selected.

SQL>

**Practice 13-2: SQL Plan Management (SPM) (continued)**

34. Now execute the statement so that you can get the plan in the cursor cache and load the corresponding plan baseline. Use the `load_evolve_baseline.sql` script for that. What do you observe?

You see that the current plan is both enabled and accepted, but not fixed.

```
SQL> @load_evolve_baseline
SQL> set echo on
SQL>
SQL> variable cnt number;
SQL>
SQL> variable sqlid varchar2(20);
SQL>
SQL> select /*+ USE_NL(s c) FULL(s) FULL(c) */
 2 c.cust_id, sum(s.quantity_sold)
 3 from sh.sales s, sh.customers c
 4 where s.cust_id = c.cust_id and c.cust_id < 2
 5 group by c.cust_id;

no rows selected

SQL>
SQL> begin
 2 select sql_id into :sqlid from v$sql
 3 where sql_text like 'select /*+ USE_NL(s c) FULL(s)
FULL(c) */%';
 4 end;
 5 /

PL/SQL procedure successfully completed.

SQL>
SQL> exec :cnt := dbms_spm.load_plans_from_cursor_cache(-
> sql_id => :sqlid);

PL/SQL procedure successfully completed.

SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like 'select /*+ USE_NL(s c) FULL(s)
FULL(c) */%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

1.7750E+18 SYS_SQL_18a1ef14c17f5b75
```

### ***Practice 13-2: SQL Plan Management (SPM) (continued)***

```
select /*+ USE_NL(s c) FULL(s) FULL(c) */ c.cust_id,
sum(s.quantity_sold)
from s
SYS_SQL_PLAN_c17f5b75dc5f94e5 MANUAL-LOAD YES YES NO YES

SQL>
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

35. Manually create and execute a SQL tuning task to tune your statement. Use the `tune_evolve_sql.sql` script.

```
SQL> @tune_evolve_sql.sql
SQL> set echo on
SQL>
SQL> variable sqltext varchar2(4000);
SQL>
SQL> BEGIN
 2 :sqltext := q'# select /*+ USE_NL(s c) FULL(s) FULL(c)*/
 3 c.cust_id, sum(s.quantity_sold)
 4 from sh.sales s, sh.customers c
 5 where s.cust_id = c.cust_id
 6 and c.cust_id < 2
 7 group by c.cust_id
 8 #';
 9 END;
10 /

PL/SQL procedure successfully completed.

SQL>
SQL> variable spmtune varchar2(30);
SQL>
SQL> exec :spmtune := dbms_sqltune.create_tuning_task(-
> sql_text => :sqltext);

PL/SQL procedure successfully completed.

SQL>
SQL> exec dbms_sqltune.execute_tuning_task(:spmtune);

PL/SQL procedure successfully completed.

SQL>
```

36. Now that the tuning task has been completed, run the report and see what recommendations have been made for your statement. What do you observe?

There are two recommendations: a SQL profile or a new index creation. Use the `report_evolve_tuning.sql` script.

```
SQL> @report_evolve_tuning.sql
SQL> set echo on
SQL>
SQL> set long 10000
SQL>
SQL> select dbms_sqltune.report_tuning_task(:spmtune, 'TEXT')
 2 from dual;
-- Report Follows --
```

## Practice 13-2: SQL Plan Management (SPM) (continued)

37. Accept the SQL profile proposed by the SQL Tuning Advisor. Use the `accept_evolve_baseline.sql` script to accept the recommended SQL profile. What happens?

Accepting the profile causes a new SQL profile and plan baseline to be created for your statement. Now you see two baselines for your statement. Both of them are enabled and accepted.

**Note:** One is MANUAL-LOAD and the other is MANUAL-SQLTUNE.

```
SQL> @accept_evolve_baseline.sql
SQL> set echo on
SQL>
SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like
 5 'select /*+ USE_NL(s c) FULL(s) FULL(c) */%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

1.7750E+18 SYS_SQL_18a1ef14c17f5b75
select /*+ USE_NL(s c) FULL(s) FULL(c) */ c.cust_id,
sum(s.quantity_sold)
from sh.sales s, sh.customers c
where s.cust_id = c.cust_id and c.cust_id < 2
group by c.cust_id
SYS_SQL_PLAN_c17f5b75dc5f94e5 MANUAL-LOAD YES YES NO YES

SQL>
SQL> exec dbms_sqltune.accept_sql_profile(-
> task_name => :spmtune,-
> name => 'SPM_SQL_PROF');

PL/SQL procedure successfully completed.

SQL>
SQL> select signature, category, name, sql_text
 2 from dba_sql_profiles where name like 'SPM%';

SIGNATURE CATEGORY NAME

SQL_TEXT

1.7750E+18 DEFAULT SPM_SQL_PROF
select /*+ USE_NL(s c) FULL(s) FULL(c) */ c.cust_id,
sum(s.quantity_sold)
```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

from sh.sales s, sh.customers c
where s.cust_id = c.cust_id and c.cust_id < 2
group by c.cust_id

SQL> select signature, sql_handle, sql_text, plan_name,
 2 origin, enabled, accepted, fixed, autopurge
 3 from dba_sql_plan_baselines
 4 where sql_text like
 5 'select /*+ USE_NL(s c) FULL(s) FULL(c) */%';

SIGNATURE SQL_HANDLE

SQL_TEXT

PLAN_NAME ORIGIN ENA ACC FIX AUT

1.7750E+18 SYS_SQL_18a1ef14c17f5b75
select /*+ USE_NL(s c) FULL(s) FULL(c) */ c.cust_id,
sum(s.quantity_sold)
from sh.sales s, sh.customers c
where s.cust_id = c.cust_id and c.cust_id < 2
group by c.cust_id
SYS_SQL_PLAN_c17f5b75a10c1dcf MANUAL-SQLTUNE YES YES NO YES

1.7750E+18 SYS_SQL_18a1ef14c17f5b75
select /*+ USE_NL(s c) FULL(s) FULL(c) */ c.cust_id,
sum(s.quantity_sold)
from sh.sales s, sh.customers c
where s.cust_id = c.cust_id and c.cust_id < 2
group by c.cust_id
SYS_SQL_PLAN_c17f5b75dc5f94e5 MANUAL-LOAD YES YES NO YES

SQL>

```

38. Determine the plan used for your statement when executing it. What do you observe?

The next time that you execute the query, it uses the plan baseline and the SQL profile. Use the `explain_query5.sql` script.

```

SQL> @explain_query5.sql
SQL> set echo on
SQL>
SQL> explain plan for
 2 select /*+ USE_NL(s c) FULL(s) FULL(c) */
 3 c.cust_id, sum(s.quantity_sold)
 4 from sh.sales s, sh.customers c
 5 where s.cust_id = c.cust_id and c.cust_id < 2
 6 group by c.cust_id;

```

**Practice 13-2: SQL Plan Management (SPM) (continued)**

```

Explained.

SQL>
SQL> select * from table(dbms_xplan.display(null,
2 null, 'basic +note'));

PLAN_TABLE_OUTPUT

Plan hash value: 3070788227

Id	Operation	Name
0	SELECT STATEMENT	
1	HASH GROUP BY	
2	NESTED LOOPS	
3	PARTITION RANGE ALL	
4	TABLE ACCESS BY LOCAL INDEX ROWID	SALES
5	BITMAP CONVERSION TO ROWIDS	

PLAN_TABLE_OUTPUT

| 6 | BITMAP INDEX RANGE SCAN | SALES_CUST_BIX |
| 7 | INDEX UNIQUE SCAN | CUSTOMERS_PK |

Note

- SQL profile "SPM_SQL_PROF" used for this statement

18 rows selected.

SQL>

```

39. Execute the `cleanup_spm.sql` script to purge your environment for this practice.

```

SQL> @cleanup_spm.sql
SQL> set echo on
SQL>
SQL> exec dbms_sqltune.drop_sql_profile(-
> name => 'SPM_SQL_PROF');

PL/SQL procedure successfully completed.

SQL>
SQL> exec :cnt := dbms_spm.drop_sql_plan_baseline(-
> 'SYS_SQL_18a1ef14c17f5b75');

PL/SQL procedure successfully completed.

SQL>

```



### **Practice 13-2: SQL Plan Management (SPM) (continued)**

```
SQL> select signature, sql_handle, sql_text, plan_name,
2 origin, enabled, accepted, fixed, autopurge
3 from dba_sql_plan_baselines
4 where sql_text like
5 'select /*+ USE_NL(s c) FULL(s) FULL(c) */%';

no rows selected

SQL>
SQL> select signature, category, name, sql_text
2 from dba_sql_profiles where name like 'SPM%';

no rows selected

SQL> alter system set optimizer_use_sql_plan_baselines=FALSE;

System altered.

SQL> exit
```

40. Prepare for the next practice. Use the `prepare` command with the `CAPTURE` option. This script is in the `$HOME/workshops` directory. Exit EM before running the script. This script takes about 9 minutes.

```
$ cd $HOME/workshops
$./prepare CAPTURE
```

## Practices for Lesson 14

In this practice, you use Database Replay.

## Practice 14-1: Using Database Replay

In this practice, you capture a workload, process the workload, and replay the workload with changes.

1. If you have not already done so, execute the `prepare` command with the `CAPTURE` option. This script performs a flashback database, restoring the database to a known state, and starts the database with the parameters set to simulate a 10g database with manual SGA tuning. This script is in the `$HOME/workshops` directory. Exit EM before running the script.


```
$ cd $HOME/workshops
$./prepare CAPTURE
```

2. Create a directory to capture the workload files named `$ORACLE_BASE/capture`.

```
$ mkdir $ORACLE_BASE/capture
```

3. Use Enterprise Manager to start a workload capture. Start EM and log in.

**Note:** Screenshots of some of the steps described in the table are provided, as noted in the Page column, below the table. Refer to these screenshots as necessary.

| Step | Page                               | Action                                                                                                                                                                   |       |       |
|------|------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------|-------|
| a    | Database home                      | Click the Software and Support tab.                                                                                                                                      |       |       |
| b    | Software and Support               | Click Database Replay.                                                                                                                                                   |       |       |
| c    | Database Replay                    | Click the Go to Task (  ) icon to the right of the Capture Workload task.             |       |       |
| d    | Capture Workload: Plan Environment | Acknowledge the prerequisites. (Select the check boxes.)                                                                                                                 |       |       |
|      |                                    | Click Next.                                                                                                                                                              |       |       |
| e    | Capture Workload: Options          | Select “Do not restart the database prior to the capture.” Normally, you would restart the database to be sure to capture all transactions and not partial transactions. |       |       |
|      |                                    | Review the filter mode and exclusions. Do not change the exclusions.                                                                                                     |       |       |
|      |                                    | Click Next.                                                                                                                                                              |       |       |
| f    | Workload Capture:                  | Set values.                                                                                                                                                              | Field | Value |

### Practice 14-1: Using Database Replay (continued)

|   |                                                                                      |                                                                     |                  |                         |
|---|--------------------------------------------------------------------------------------|---------------------------------------------------------------------|------------------|-------------------------|
|   | Parameters                                                                           |                                                                     | Capture Name     | CAPTURE_10gSim          |
|   |                                                                                      | Click Create Directory Object.                                      |                  |                         |
| g | Create Directory Object                                                              | Set values.                                                         | <b>Field</b>     | <b>Value</b>            |
|   |                                                                                      |                                                                     | Name             | CAPTURE_DIR             |
|   |                                                                                      |                                                                     | Path             | /u01/app/oracle/capture |
|   |                                                                                      | Click Test File System.                                             |                  |                         |
| h | Host Login                                                                           | Enter the host credentials:<br>Username: oracle<br>Password: oracle |                  |                         |
|   |                                                                                      | Click Login.                                                        |                  |                         |
| i | Confirmation                                                                         | Click Return.                                                       |                  |                         |
| j | Create Directory Object                                                              | Click OK.                                                           |                  |                         |
| k | Workload Capture: Parameters                                                         | Click Next.                                                         |                  |                         |
| l | Workload Capture: Schedule<br>(See the example screenshot, which follows the table.) | Set values.                                                         | <b>Field</b>     | <b>Value</b>            |
|   |                                                                                      | Set host credentials.                                               | Username         | oracle                  |
|   |                                                                                      |                                                                     | Password         | oracle                  |
|   |                                                                                      |                                                                     | Confirm password | oracle                  |

**Practice 14-1: Using Database Replay (continued)**

|   |                                             |               |
|---|---------------------------------------------|---------------|
|   |                                             | Click Next.   |
| m | Capture Workload:<br>Review                 | Click Submit. |
| n | Confirmation                                | Click Yes.    |
| o | View Workload<br>Capture:<br>CAPTURE_10gSim |               |

# **Practice 14-1: Using Database Replay (continued)**

1) The following screenshot displays the Capture Workload: Schedule page:

**Capture Workload: Schedule**

Database **orcl.us.oracle.com**
Cancel
Back
Step 4 of 5
Next

Logged In As **SYS**

---

**Job Parameters**

Job Name 
Description

---

**Job Schedule**

Choose a start time and a capture duration so that the workload you are interested in replaying at a later time can be captured.

**Start**

☒ Immediately
☐ Later

Date 
Time 
AM PM

**Capture Duration**

☒ Not Specified
☐ Duration

Hours  Minutes

---

**Job Credentials**

**Host Credentials**

Username 
Password 
Confirm Password

☐ Save as Preferred Credential

4. Start the charbench workload. Press the **Enter** key according to the instructions to start the workload. After about five minutes, stop the workload by pressing Enter.

```

$ cd /home/oracle/swingbench/bin
$./charbench
Author : Dominic Giles
Version : 2.2

Results will be written to results.xml.
Hit Return to Start & Terminate Run...

Users: 30 TPM: 417 Nested TPM: 0
Users: 0 TPM: 343 Nested TPM: 0
Completed Run.

```

5. Stop the capture and export the AWR data.

| Step | Page                                  | Action              |
|------|---------------------------------------|---------------------|
| a    | View Workload Capture: CAPTURE-10gSIM | Click Stop Capture. |
| b    | Confirmation                          | Click Yes.          |
| c    | Export AWR Data                       | Click Yes.          |

## Practice 14-1: Using Database Replay (continued)

|   |                                       |                                                                                                  |
|---|---------------------------------------|--------------------------------------------------------------------------------------------------|
| d | View Workload Capture: CAPTURE_10gsim | Click View Job.                                                                                  |
| e | Scheduler Jobs                        | Click Refresh until the job no longer appears in the running jobs pane. This may take 3 minutes. |

6. Process the captured workload.

| Step | Page                                           | Action                                                                           |                  |              |
|------|------------------------------------------------|----------------------------------------------------------------------------------|------------------|--------------|
| a    | View Workload Capture: CAPTURE-10gSIM          | Click the Database tab.                                                          |                  |              |
| b    | Database home                                  | Click the Software and Support tab.                                              |                  |              |
| c    | Software and Support                           | Click Database Replay in the Real Application Testing section.                   |                  |              |
| d    | Database Replay                                | Click the Go to Task icon to the right of the Preprocess Captured Workload task. |                  |              |
| e    | Preprocess Captured Workload                   | Set Directory Object to CAPTURE_DIR.                                             |                  |              |
|      |                                                | Click Preprocess Workload.                                                       |                  |              |
| f    | Preprocess Captured Workload: Database Version | Click Next.                                                                      |                  |              |
| g    | Preprocess Captured Workload: Schedule         | Set values.                                                                      | <b>Field</b>     | <b>Value</b> |
|      |                                                | Set host credentials.                                                            | Username         | oracle       |
|      |                                                |                                                                                  | Password         | oracle       |
|      |                                                |                                                                                  | Confirm password | oracle       |
|      |                                                | Click Next.                                                                      |                  |              |
| h    | Preprocess Captured Workload: Review           | Click Submit.                                                                    |                  |              |
| i    | Database Replay                                | In the Confirmation message, click View Job.                                     |                  |              |

## Practice 14-1: Using Database Replay (continued)

|   |                                  |                                                                        |
|---|----------------------------------|------------------------------------------------------------------------|
| j | Execution:<br>orcl.us.oracle.com | Refresh until the status shows succeeded. This takes about 45 seconds. |
|---|----------------------------------|------------------------------------------------------------------------|

- Adjust the database setup. Execute the flashback REPLAY script. This script restores the database to the same point (SCN, where the capture started) and changes the parameters to use 11g optimizer and Automatic Memory Management. Log out of EM before you run `./prepare REPLAY`.

```
$ cd $HOME/workshops
$./prepare REPLAY
```

- Log in to EM and replay the workload.

| Step | Page                    | Action                                                              |       |                         |
|------|-------------------------|---------------------------------------------------------------------|-------|-------------------------|
| a    | Current page            | Click the Database tab.                                             |       |                         |
| b    | Database home           | Click the Software and Support tab.                                 |       |                         |
| c    | Software and Support    | Click Database Replay in the Real Application Testing section.      |       |                         |
| d    | Database Replay         | Click the Go to Task icon to the right of the Replay Workload task. |       |                         |
| e    | Replay Workload         | Click Create Directory Object.                                      |       |                         |
| f    | Create Directory Object | Set values.                                                         | Field | Value                   |
|      |                         |                                                                     | Name  | CAPTURE_DIR             |
|      |                         |                                                                     | Path  | /u01/app/oracle/capture |
|      |                         | Click Test File System.                                             |       |                         |
| g    | Host Login              | Enter host credentials:<br>Username: oracle<br>Password: oracle     |       |                         |
|      |                         | Click Login.                                                        |       |                         |
| h    | Confirmation            | Click Return.                                                       |       |                         |
| i    | Create Directory Object | Click OK.                                                           |       |                         |



## Practice 14-1: Using Database Replay (continued)

|   |                                                 |                                                                                                                                                                                                                                 |
|---|-------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| j | Replay Workload                                 | Click Setup Replay.                                                                                                                                                                                                             |
| k | Replay Workload: Prerequisites                  | Notice the advice given for setup. The database has been restored, the database environment has been changed, and there are no external references in this database. The only thing left to do is to set up the replay clients. |
|   |                                                 | Click Continue.                                                                                                                                                                                                                 |
| l | Replay Workload: References to External Systems | Click Continue.                                                                                                                                                                                                                 |
| m | Replay Workload: Choose Initial Options         | Click Next.                                                                                                                                                                                                                     |
| n | Replay Workload: Customize Options              | Click Next.                                                                                                                                                                                                                     |
| o | Replay Workload: Prepare Replay Clients         | Open a terminal window and run the workload replay client in calibrate mode.                                                                                                                                                    |
|   |                                                 | <code>\$ cd /u01/app/oracle/capture</code><br><code>\$ wrc mode=calibrate</code>                                                                                                                                                |
|   |                                                 | Type in the command to start the workload replay client in replay mode, but <i>do not start it</i> .                                                                                                                            |
|   |                                                 | <code>\$ wrc system/oracle_4U@orcl mode=replay</code>                                                                                                                                                                           |
|   |                                                 | In EM, Click Next.                                                                                                                                                                                                              |
| p | Replay Workload: Wait for Client Connections    | Switch to the terminal window and start the workload client.                                                                                                                                                                    |
|   |                                                 | When the client connection appears on the connection list, click Next.                                                                                                                                                          |

## Practice 14-1: Using Database Replay (continued)

|   |                            |               |
|---|----------------------------|---------------|
| q | Replay Workload:<br>Review | Click Submit. |
|---|----------------------------|---------------|

As the replay progresses, the chart changes when the page is refreshed.

**Note:** The first update may take a few minutes.

When the replay completes, observe Divergence. The Divergence section should show zeros for Error Divergence and Data Divergence.

Expand the Detailed Comparison: Examine Detailed Comparison section. It shows the differences in several areas. The replay may be faster or slower. This section does not show whether the workload is actually performing better or worse.

9. Create a Compare Periods Report.

| Step | Page                                | Action                                                                                                                                                                                |
|------|-------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a    | View Replay Workload                | Click the Report tab after the workload is completed. (The Report tab is not available while the replay is running.)                                                                  |
| b    | View Replay Workload Report subpage | Click Run AWR Compare Period Report in the Compare Period section.                                                                                                                    |
| c    | Import AWR Data                     | Click Yes.                                                                                                                                                                            |
| d    | Select Schema                       | Enter SYSTEM.                                                                                                                                                                         |
|      |                                     | Click Import AWR Data.                                                                                                                                                                |
| e    | View Replay Workload, Report tab    | Click View Job.                                                                                                                                                                       |
| f    | Scheduler Jobs                      | Click Refresh until the Import job no longer appears on the Running tab. This job takes about a minute.                                                                               |
|      |                                     | Click the Back button on the browser.                                                                                                                                                 |
|      |                                     | Click the Reload button on the browser.                                                                                                                                               |
| g    | View Replay Workload, Report tab    | In the AWR Report section, test that you can produce an AWR report from the Capture data.<br><br>Set the Workload Capture or Replay field to Capture-10gSIM.<br><br>Click Run Report. |

**Practice 14-1: Using Database Replay (continued)**

|   |                                  |                                                                               |
|---|----------------------------------|-------------------------------------------------------------------------------|
| h | WORKLOAD REPOSITORY report for   | If the report appears, the import was successful.<br>Close the report window. |
| i | View Replay Workload, Report tab | Click Run AWR Compare Period Report in the Compare Period Report section.     |

10. When the report appears, review the report.

- a) Notice the difference between Elapsed time and DB Time.

|                              | Elapsed Time | DB Time |
|------------------------------|--------------|---------|
| <b>1<sup>st</sup> Period</b> |              |         |
| <b>2<sup>nd</sup> Period</b> |              |         |

- b) Are there some areas where the performance seems to be worse?  
Check Top Timed Events and Time Model. (Look for places where % of DB time increased.)
- c) What is the overall performance difference?  
(In Time Model, DB Time in seconds is much reduced.  
**Note:** The DB CPU in % DB Time is increased because the total DB Time is much smaller.)
- d) What can you determine about individual SQL statement performance?  
(Statistics on all SQL statements are not reported. In each of the categories—Elapsed Time, CPU Time, Buffer Gets, and others—only the top 10 statements are reported. You could get some comparisons, but DB Replay is intended for overall performance comparison. SQL Performance Analyzer is intended for comparison of an entire set of SQL statements.)

11. From the \$HOME/workshops directory, use the `./cleanup REPLAY` command to reset the database environment for the next practice.

```
$ cd $HOME/workshops
$./cleanup REPLAY
```

12. Perform the first two steps for Practice 15-1, and then stop.

## Practices for Lesson 15

In this practice, you tune the shared pool.

## Practice 15-1: Sizing the Shared Pool

In this practice, you use the ADDM reports to diagnose and fix the shared pool size.

- 1) Prepare the database instance for this scenario.
  - a) Log out of Enterprise Manager.
  - b) In a terminal window, run the `./prepare SP` command.

```
$ cd /home/oracle/workshops
$./prepare SP
Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area 171581440 bytes
Fixed Size 1335248 bytes
Variable Size 134217776 bytes
Database Buffers 33554432 bytes
Redo Buffers 2473984 bytes
Database mounted.
Database opened.
Finished Shared Pool setup
$
```

- 2) In a terminal window, determine the machine time and run the workload generator. The `workgen` script creates an AWR snapshot and a Statspack snapshot, and then runs the workload. Write the time returned by the `date` command: \_\_\_\_\_
 

```
date
./workgen SP
```

```
$ date
Tue Feb 23 12:24:25 UTC 2010
$./workgen SP

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
$
```

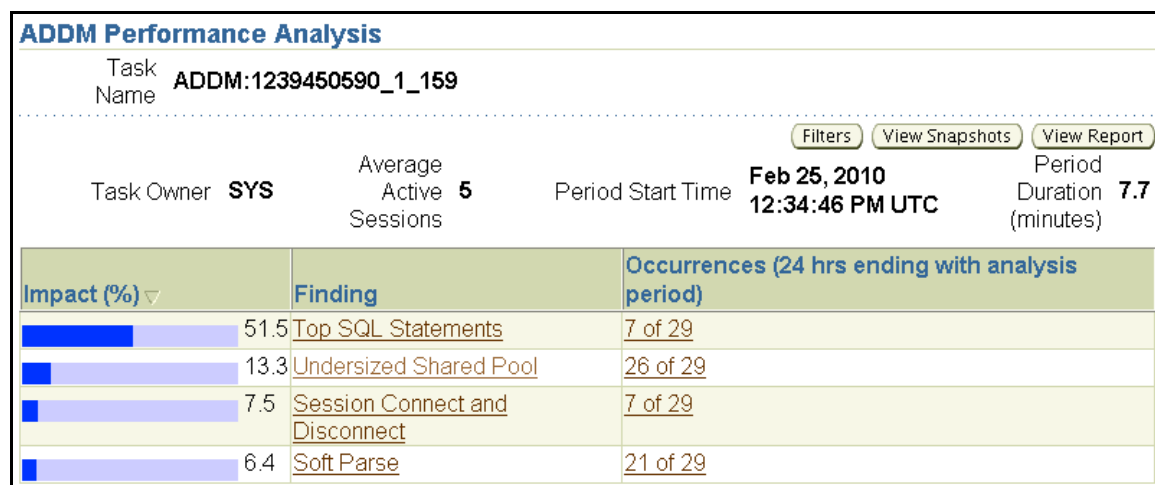
- 3) Check the date-time with the `date` command; if it has been more than 15 minutes since the time you recorded in step 2, continue, else wait until 15 minutes have elapsed. In a terminal window, stop the workload with the `rm runload` command.

```
$ rm runload
```

### Practice 15-1: Sizing the Shared Pool (continued)

- 4) Log in to Enterprise Manager. Click the Performance tab to see the Performance page. Wait until the Average Active Sessions graph has dropped to almost 0. Create an ADDM report. (Click Run ADDM Now, and then click Yes on the Confirmation page.) Notice any findings referring to the shared pool. These finding may include Hard Parse, Soft Parse, and Undersized Shared Pool.

The particular findings and the order of the findings may vary from the solution shown.



- 5) Create an AWR report between the last two snapshots. On the ADDM Report page, click View Snapshots to navigate to the snapshots page. Click the Report tab to generate the AWR report.
- 6) View the AWR report and note the values for the Top 5 Timed Foreground Events, Time Model Statistics, Instance Efficiencies, and the Library Cache Activity Statistics for the SQL AREA namespace.
- a) The Top 5 Timed Foreground Events may show time waiting on the library cache mutex, library cache load lock, shared pool latch, cursor pin waits, or row cache objects. Any of these indicate a possible problem in the shared pool.

| Top 5 Timed Foreground Events |       |         |               |           |             |
|-------------------------------|-------|---------|---------------|-----------|-------------|
| Event                         | Waits | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
| DB CPU                        |       | 249     |               | 10.86     |             |
| cursor: pin S wait on X       | 172   | 82      | 475           | 3.57      | Concurrency |
| latch: shared pool            | 77    | 49      | 633           | 2.13      | Concurrency |
| library cache load lock       | 59    | 45      | 755           | 1.94      | Concurrency |
| library cache: mutex X        | 12    | 34      | 2875          | 1.51      | Concurrency |

- b) Time Model Statistics, in this example, shows a significant % of DB Time being used for hard and soft parses (parse time elapsed).

## Practice 15-1: Sizing the Shared Pool (continued)

### Time Model Statistics

| Statistic Name                             | Time (s) | % of DB Time |
|--------------------------------------------|----------|--------------|
| sql execute elapsed time                   | 1,694.03 | 73.99        |
| parse time elapsed                         | 588.84   | 25.72        |
| hard parse elapsed time                    | 441.95   | 19.30        |
| DB CPU                                     | 248.58   | 10.86        |
| connection management call elapsed time    | 171.73   | 7.50         |
| hard parse (sharing criteria) elapsed time | 43.52    | 1.90         |
| hard parse (bind mismatch) elapsed time    | 30.16    | 1.32         |
| failed parse elapsed time                  | 11.90    | 0.52         |
| PL/SQL execution elapsed time              | 10.02    | 0.44         |
| PL/SQL compilation elapsed time            | 9.09     | 0.40         |
| repeated bind elapsed time                 | 0.72     | 0.03         |
| sequence load elapsed time                 | 0.51     | 0.02         |
| DB time                                    | 2,289.62 |              |
| background elapsed time                    | 25.27    |              |
| background cpu time                        | 2.42     |              |

- c) The latch and mutex waits taken with the hard parses shown in Time Model leads to a check of the Load Profile section. In this example, hard parses are a significant percentage of all parses.

**Practice 15-1: Sizing the Shared Pool (continued)**

| Load Profile      |            |                 |          |          |
|-------------------|------------|-----------------|----------|----------|
|                   | Per Second | Per Transaction | Per Exec | Per Call |
| DB Time(s):       | 5.0        | 12.6            | 0.01     | 0.01     |
| DB CPU(s):        | 0.5        | 1.4             | 0.00     | 0.00     |
| Redo size:        | 11,942.5   | 30,309.5        |          |          |
| Logical reads:    | 5,176.9    | 13,138.7        |          |          |
| Block changes:    | 55.7       | 141.3           |          |          |
| Physical reads:   | 1,647.2    | 4,180.4         |          |          |
| Physical writes:  | 3.9        | 9.8             |          |          |
| User calls:       | 732.7      | 1,859.7         |          |          |
| Parses:           | 288.0      | 731.0           |          |          |
| Hard parses:      | 61.1       | 155.1           |          |          |
| W/A MB processed: | 0.9        | 2.3             |          |          |
| Logons:           | 0.7        | 1.8             |          |          |
| Executes:         | 591.5      | 1,501.3         |          |          |
| Rollbacks:        | 0.0        | 0.0             |          |          |
| Transactions:     | 0.4        |                 |          |          |

- d) Navigate to Library Cache Statistics from the Main menu. Check the Reloads statistic of the SQL AREA namespace in the Library Cache Activity Statistics section. The hard parses can be caused by SQL that cannot be shared, or by cursors aging out of the shared pool before they can be reused. This example shows a high value for reloads and a high Pct Miss for Gets and Pins. This indicates that the cursor could have been shared, but the shared pool is too small. Reloads should be less than 1% of Pin Requests.



**Practice 15-1: Sizing the Shared Pool (continued)****Library Cache Activity**

| Namespace       | Get Requests | Pct Miss | Pin Requests | Pct Miss | Reloads | Invali- dations |
|-----------------|--------------|----------|--------------|----------|---------|-----------------|
| APP CONTEXT     | 1,764        | 1.70     | 1,764        | 1.70     | 0       | 0               |
| BODY            | 1,127        | 9.41     | 3,571        | 10.78    | 267     | 0               |
| CLUSTER         | 3,461        | 1.18     | 1,995        | 2.06     | 0       | 0               |
| DBLINK          | 593          | 13.66    | 0            |          | 0       | 0               |
| EDITION         | 304          | 27.63    | 606          | 27.89    | 0       | 0               |
| INDEX           | 624          | 51.60    | 624          | 52.08    | 1       | 0               |
| OBJECT ID       | 1            | 100.00   | 0            |          | 0       | 0               |
| QUEUE           | 948          | 2.00     | 1,063        | 5.55     | 34      | 0               |
| SCHEMA          | 18,550       | 0.06     | 0            |          | 0       | 0               |
| SQL AREA        | 189,895      | 56.89    | 390,715      | 16.35    | 20,167  | 2,505           |
| SUBSCRIPTION    | 1            | 100.00   | 1            | 100.00   | 0       | 0               |
| TABLE/PROCEDURE | 149,200      | 5.61     | 181,280      | 11.63    | 8,927   | 0               |
| TRIGGER         | 4,534        | 1.87     | 5,204        | 2.21     | 16      | 0               |

e) Further investigation of the dictionary cache shows a high number of misses.

**Dictionary Cache Stats**

| Cache                | Get Requests | Pct Miss | Scan Reqs | Pct Miss | Mod Reqs | Final Usage |
|----------------------|--------------|----------|-----------|----------|----------|-------------|
| dc_awr_control       | 12           | 66.67    | 0         |          | 2        | 1           |
| dc_files             | 17           | 100.00   | 0         |          | 0        | 0           |
| dc_global_oids       | 11,372       | 6.79     | 0         |          | 0        | 9           |
| dc_histogram_data    | 81,072       | 23.14    | 0         |          | 0        | 443         |
| dc_histogram_defs    | 159,785      | 27.07    | 0         |          | 2        | 924         |
| dc_object_grants     | 23,127       | 6.03     | 0         |          | 0        | 10          |
| dc_objects           | 346,928      | 5.04     | 0         |          | 3        | 587         |
| dc_profiles          | 326          | 13.50    | 0         |          | 0        | 0           |
| dc_rollback_segments | 65           | 0.00     | 0         |          | 0        | 22          |
| dc_segments          | 117,677      | 13.73    | 0         |          | 7        | 191         |
| dc_sequences         | 96           | 100.00   | 0         |          | 96       | 0           |
| dc_tablespace        | 81,028       | 0.06     | 0         |          | 0        | 4           |
| dc_users             | 170,746      | 0.66     | 0         |          | 0        | 64          |
| global database name | 1,022        | 2.25     | 0         |          | 0        | 1           |
| kqlsubheap_object    | 155          | 48.39    | 0         |          | 0        | 0           |

- Find an optimum shared pool size for this workload. The solution can be obtained either through the EM tools or by using the Shared Pool Advisory in the Statspack report.

## Practice 15-1: Sizing the Shared Pool (continued)

- a) Navigate to the Memory Advisors page under the Server tab.

**Memory Advisors**

Page Refreshed **March 30, 2010 9:13:47 AM UTC** [Refresh](#)

[Show SQL](#) [Revert](#) [Apply](#)

When Automatic Memory Management is enabled, the database will automatically set the optimal distribution of memory. The distribution of memory will change from time to time to accommodate changes in the workload.

Automatic Memory Management **Disabled** [Enable](#)


---

**SGA** [PGA](#)

The System Global Area (SGA) is a group of shared memory structures that contains data and control information for one Oracle database. The SGA is allocated in memory when an Oracle database instance is started.

Automatic Shared Memory Management **Disabled** [Enable](#)

|                |                                  |    |                           |
|----------------|----------------------------------|----|---------------------------|
| Shared Pool    | <input type="text" value="112"/> | MB | <a href="#">Advice</a>    |
| Buffer Cache   | <input type="text" value="32"/>  | MB | <a href="#">Advice</a>    |
| Large Pool     | <input type="text" value="4"/>   | MB |                           |
| Java Pool      | <input type="text" value="12"/>  | MB |                           |
| Other (MB)     | <input type="text" value="3"/>   |    |                           |
| Total SGA (MB) | <b>163</b>                       |    | <a href="#">Calculate</a> |

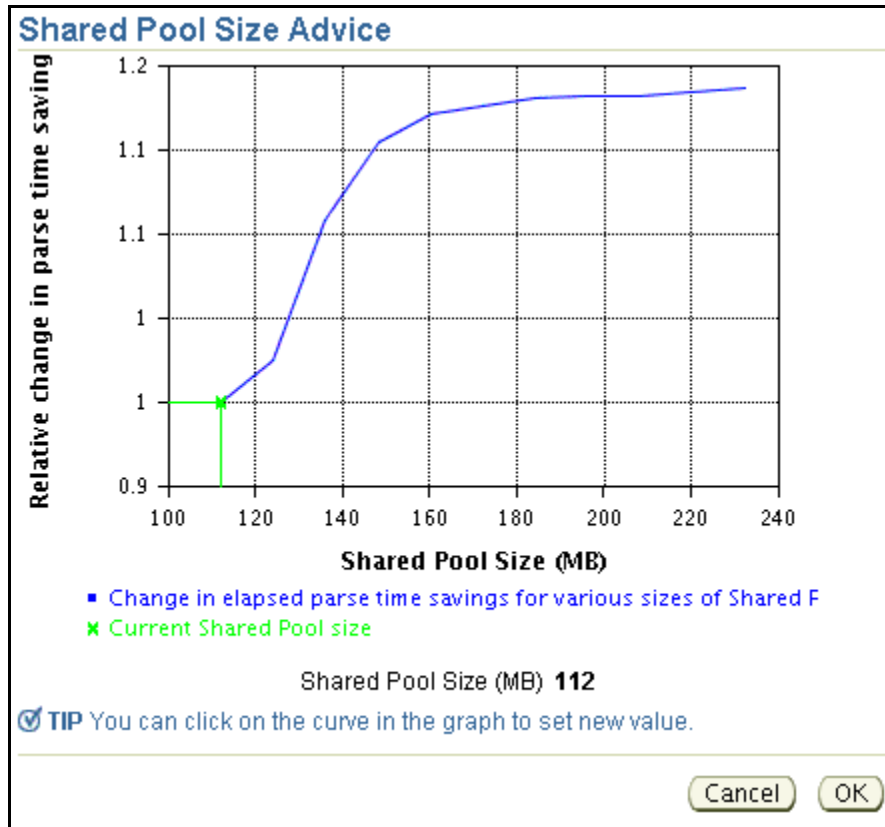


|                      |
|----------------------|
| Shared Pool (68.4%)  |
| Buffer Cache (19.6%) |
| Large Pool (2.4%)    |
| Java Pool (7.3%)     |
| Other (2.2%)         |

- b) Click Shared Pool Advice. The graph shows performance versus shared pool size. The point where the curve flattens is where there is no more improvement in performance as the shared pool size is increased. (The graph generated for your instance may vary.) Make a note of this point for your instance. Click Cancel.

**Note:** This graph displays estimates based on statistics collected during the previous workload. If the workload changes, so do the estimates.

## Practice 15-1: Sizing the Shared Pool (continued)



- c) On the Memory Advisors Page, change the Shared Pool value to 200 MB. Click Calculate.  
 What is the calculated value? \_\_\_\_\_  
 If the calculated value (Total SGA) is **less than** the Maximum SGA size, click Apply.  
 If the calculated value (Total SGA) is **greater than** the Maximum SGA size, an “Insufficient memory to grow pool” error message appears.
- Change the Maximum SGA size to the calculated value + 2 MB. Select the “Apply changes to SPFILE only” check box at the bottom of the page. Click Apply.
8. Run the `./prepare SP2` script from the `$HOME/workshops` directory. This script changes the `SHARED_POOL_SIZE` parameter to 200 MB; because, you are not using Automatic SGA Memory Management, the change requires the database to be restarted.

```
$./prepare SP2
Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area 263651328 bytes
Fixed Size 1335864 bytes
Variable Size 226495944 bytes
```

## Practice 15-1: Sizing the Shared Pool (continued)

```
Database Buffers 33554432 bytes
Redo Buffers 2265088 bytes
Database mounted.
Database opened.
Finished Shared Pool second setup
$
```

9. Record the current time: \_\_\_\_\_  
Run the workload generator again (. /workgen SP) for five minutes, and then stop the workload with the `rm runload` command.

a) Find the system time. Start the workload.

```
$ date
Thu Dec 1 08:10:03 PST 2005
$./workgen SP

PL/SQL procedure successfully completed.

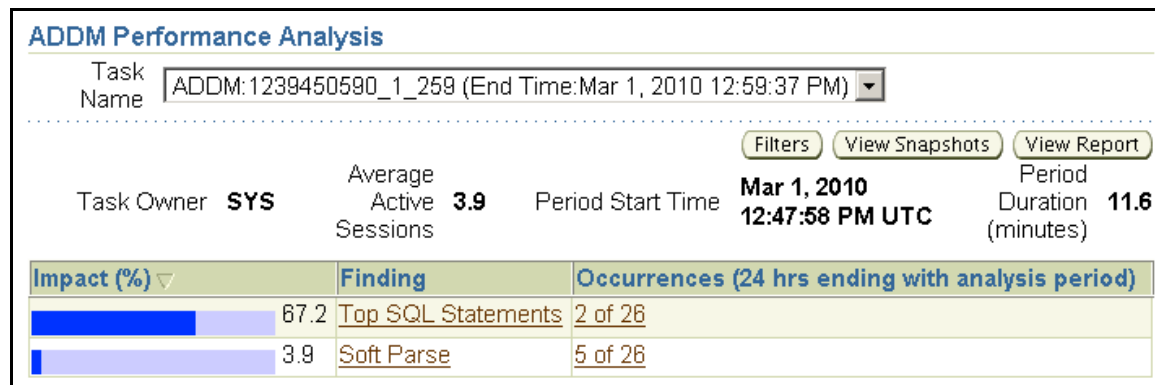
PL/SQL procedure successfully completed.
```

b) After five minutes, stop the workload.

```
$ rm runload
```

10. Log out of EM, log in to EM, navigate to the Performance page, and create another ADDM snapshot. What is the difference? What are the primary issues now?

a) Notice the findings. The hard-parse findings have been reduced or eliminated from the findings. There may be no findings.



11. From the ADDM page, click View Snapshots to navigate to the Snapshot Details page, and then click the Report tab to create an AWR report.

### Practice 15-1: Sizing the Shared Pool (continued)

12. View the AWR report and compare the values from the previous report.

**Note:** The actual values that you see may vary from those shown in the solutions.

- a) In the Top 5 Foreground Timed Events, the latch and mutex waits for the shared pool and cursors are reduced or eliminated.

| Top 5 Timed Foreground Events |         |         |               |           |            |
|-------------------------------|---------|---------|---------------|-----------|------------|
| Event                         | Waits   | Time(s) | Avg wait (ms) | % DB time | Wait Class |
| DB CPU                        |         | 328     |               | 12.17     |            |
| log file sync                 | 350     | 23      | 65            | 0.85      | Commit     |
| kfk: async disk IO            | 202,150 | 12      | 0             | 0.45      | System I/O |
| SQL*Net message to client     | 644,068 | 10      | 0             | 0.36      | Network    |
| latch free                    | 18      | 9       | 478           | 0.32      | Other      |

- b) Time Model (in the example) shows that hard parses are still contributing to parse elapsed time, but take a very small % of DB Time. Your values may vary.

**Practice 15-1: Sizing the Shared Pool (continued)****Time Model Statistics**

| Statistic Name                             | Time (s) | % of DB Time |
|--------------------------------------------|----------|--------------|
| sql execute elapsed time                   | 2,250.07 | 83.38        |
| DB CPU                                     | 328.40   | 12.17        |
| parse time elapsed                         | 126.85   | 4.70         |
| hard parse elapsed time                    | 21.24    | 0.79         |
| failed parse elapsed time                  | 13.79    | 0.51         |
| connection management call elapsed time    | 9.95     | 0.37         |
| PL/SQL compilation elapsed time            | 2.64     | 0.10         |
| PL/SQL execution elapsed time              | 1.47     | 0.05         |
| repeated bind elapsed time                 | 0.08     | 0.00         |
| hard parse (sharing criteria) elapsed time | 0.04     | 0.00         |
| sequence load elapsed time                 | 0.00     | 0.00         |
| DB time                                    | 2,698.62 |              |
| background elapsed time                    | 10.48    |              |
| background cpu time                        | 0.68     |              |

- c) The Load Profile section shows that hard parses are reduced (as are all parses).

| Load Profile      |            |                 |          |          |
|-------------------|------------|-----------------|----------|----------|
|                   | Per Second | Per Transaction | Per Exec | Per Call |
| DB Time(s):       | 3.9        | 15.6            | 0.02     | 0.00     |
| DB CPU(s):        | 0.5        | 1.9             | 0.00     | 0.00     |
| Redo size:        | 5,126.2    | 20,692.2        |          |          |
| Logical reads:    | 4,098.1    | 16,542.2        |          |          |
| Block changes:    | 20.9       | 84.2            |          |          |
| Physical reads:   | 2,251.5    | 9,088.4         |          |          |
| Physical writes:  | 2.2        | 8.7             |          |          |
| User calls:       | 1,065.0    | 4,299.0         |          |          |
| Parses:           | 176.9      | 714.1           |          |          |
| Hard parses:      | 9.5        | 38.3            |          |          |
| W/A MB processed: | 0.4        | 1.7             |          |          |
| Logons:           | 1.0        | 3.8             |          |          |
| Executes:         | 187.2      | 755.7           |          |          |
| Rollbacks:        | 0.0        | 0.0             |          |          |
| Transactions:     | 0.3        |                 |          |          |

## Practice 15-1: Sizing the Shared Pool (continued)

- d) In the Library Cache Activity section, reloads for the SQL AREA namespace are now close to an acceptable range. The % of Reloads versus Pin Requests is less than 4% in the example.

### Library Cache Activity

| Namespace       | Get Requests | Pct Miss | Pin Requests | Pct Miss | Reloads | Invali- dations |
|-----------------|--------------|----------|--------------|----------|---------|-----------------|
| APP CONTEXT     | 259          | 0.00     | 259          | 0.00     | 0       | 0               |
| BODY            | 2,112        | 2.04     | 3,906        | 1.28     | 0       | 0               |
| CLUSTER         | 145          | 0.00     | 74           | 0.00     | 0       | 0               |
| DBLINK          | 1,246        | 0.00     | 0            |          | 0       | 0               |
| EDITION         | 637          | 0.00     | 1,271        | 0.00     | 0       | 0               |
| INDEX           | 59           | 30.51    | 59           | 30.51    | 0       | 0               |
| OBJECT ID       | 1            | 100.00   | 0            |          | 0       | 0               |
| QUEUE           | 1,580        | 0.06     | 1,907        | 0.10     | 0       | 0               |
| SCHEMA          | 3,213        | 0.25     | 0            |          | 0       | 0               |
| SQL AREA        | 187,872      | 47.49    | 160,476      | 8.88     | 5,433   | 5,464           |
| SUBSCRIPTION    | 1            | 100.00   | 1            | 100.00   | 0       | 0               |
| TABLE/PROCEDURE | 22,339       | 4.27     | 50,937       | 4.63     | 2       | 0               |
| TRIGGER         | 5,632        | 0.05     | 5,654        | 0.05     | 0       | 0               |

- e) The Dictionary Cache Stats section shows a much reduced percentage of misses.

## Practice 15-1: Sizing the Shared Pool (continued)

### Dictionary Cache Stats

| Cache                | Get Requests | Pct Miss | Scan Reqs | Pct Miss | Mod Reqs | Final Usage |
|----------------------|--------------|----------|-----------|----------|----------|-------------|
| dc_awr_control       | 16           | 0.00     | 0         |          | 2        | 1           |
| dc_global_oids       | 15,587       | 0.29     | 0         |          | 0        | 80          |
| dc_histogram_data    | 11,056       | 9.35     | 0         |          | 0        | 2,126       |
| dc_histogram_defs    | 15,952       | 20.46    | 0         |          | 2        | 5,374       |
| dc_object_grants     | 26,506       | 0.29     | 0         |          | 0        | 109         |
| dc_objects           | 53,595       | 2.39     | 0         |          | 3        | 3,188       |
| dc_profiles          | 649          | 0.00     | 0         |          | 0        | 2           |
| dc_rollback_segments | 152          | 0.00     | 0         |          | 0        | 22          |
| dc_segments          | 6,825        | 12.19    | 0         |          | 8        | 1,357       |
| dc_sequences         | 3            | 100.00   | 0         |          | 3        | 3           |
| dc_tablespace_quotas | 4            | 50.00    | 0         |          | 0        | 2           |
| dc_tablespaces       | 30,048       | 0.00     | 0         |          | 0        | 19          |
| dc_users             | 81,144       | 0.02     | 0         |          | 0        | 97          |
| global database name | 1,101        | 0.00     | 0         |          | 0        | 1           |
| kqlsubheap_object    | 11           | 18.18    | 0         |          | 0        | 2           |
| outstanding_alerts   | 16           | 0.00     | 0         |          | 0        | 17          |



## Practice 15-2: Tuning a Hard-Parse Workload

In this section, you run a hard-parse workload, view the characteristic report events, and fix the problem.

1. Prepare the database for this scenario.
  - a) Log out of Enterprise Manager.
  - b) Set up the scenario by running the `./prepare HARDPARSE` script. Change the directory to the `/home/oracle/workshops` directory. Run the prepare script.

```
$ cd /home/oracle/workshops
$./prepare HARDPARSE
Database closed.
Database dismounted.
ORACLE instance shut down.

ORACLE instance started.

Total System Global Area 179949568 bytes
Fixed Size 1335304 bytes
Variable Size 142610424 bytes
Database Buffers 33554432 bytes
Redo Buffers 2449408 bytes
Database mounted.
Database opened.
Finished setup Hard Parse
$
```

2. Run the workload generator: `./workgen HARDPARSE`.  
The workload generator creates an AWR snapshot and a Statspack snapshot at the beginning. An additional Statspack snapshot is taken at the end.
  - a) Run the workload generator.

```
$./workgen HARDPARSE

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

System altered.

PL/SQL procedure successfully completed.

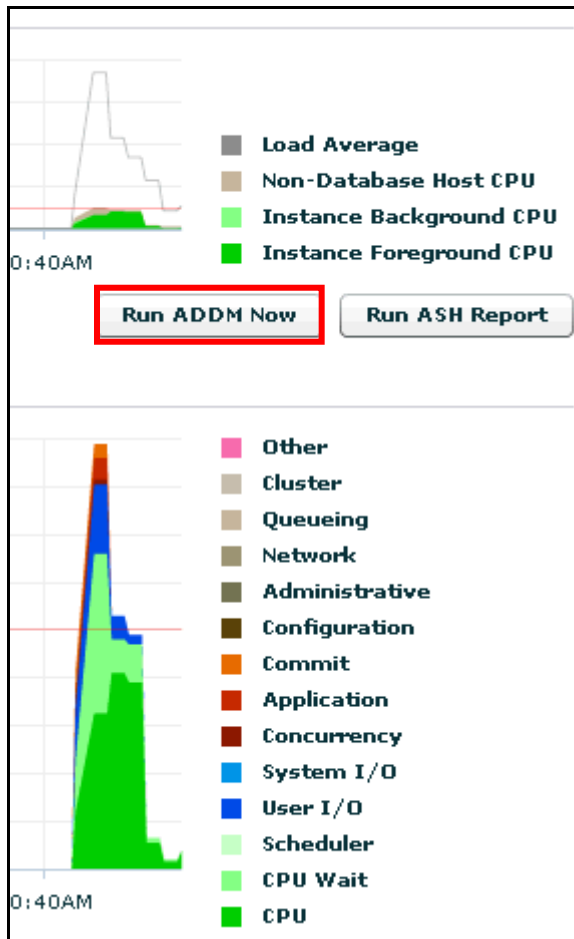
PL/SQL procedure successfully completed.
```

## Practice 15-2: Tuning a Hard-Parse Workload (continued)

PL/SQL procedure successfully completed.

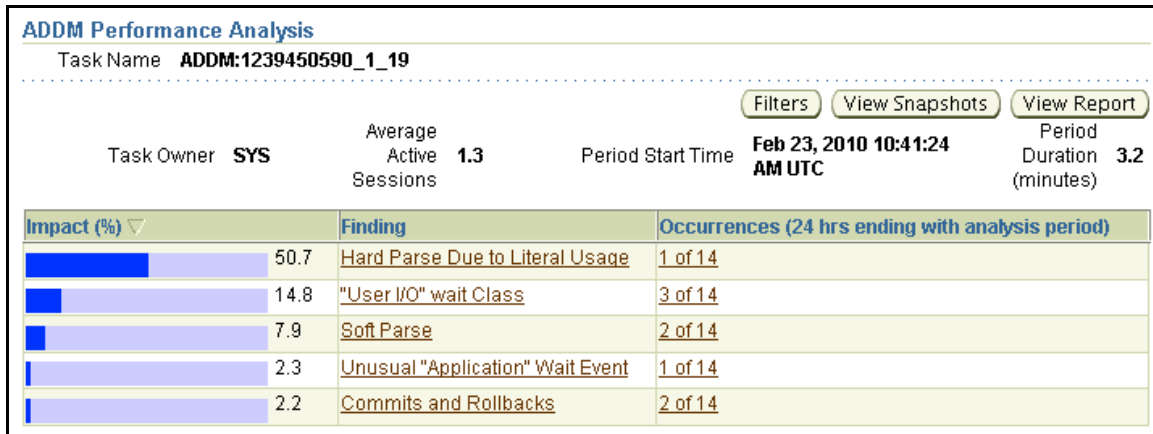
\$

- b) Log in to Enterprise Manager, navigate to the Database Performance page, wait until the workload begins to drop (about 5 minutes), and then click Run ADDM Now. Click Yes on the Confirmation page.



- c) View the ADDM report. Note that Benefit % is the same as % DB Time shown in the Top 5 Foreground Waits section of the AWR report. What is the value of Average Active Sessions? \_\_\_\_\_

## Practice 15-2: Tuning a Hard-Parse Workload (continued)



## Practice 15-2: Tuning a Hard-Parse Workload (continued)

- Drill down to the recommendations. Click the “Hard Parse Due to Literal Usage” finding.

**Recommendations**  
[Show All Details](#) | [Hide All Details](#)  

| Details                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | Category | Benefit (%) |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------|-------------|
| <a href="#">Hide</a> Application Analysis<br>Action <b>Alternatively, you may set the parameter "cursor_sharing" to "force".</b> <a href="#">Implement</a> <a href="#">Filters</a><br>Action <b>Investigate application logic for possible use of bind variables instead of literals.</b><br>Rationale <b>At least 7 SQL statements with FORCE_MATCHING_SIGNATURE 14812360487617340174 and PLAN_HASH_VALUE 4238351645 were found to be using literals. Look in V\$SQL for examples of such SQL statements.</b> |          | 50.7        |

**Findings Path**  
[Expand All](#) | [Collapse All](#)  

| Findings                                                                                                                                             | Percentage of Finding's Impact (%) | Additional Information |
|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------|------------------------|
| SQL statements were not shared due to the usage of literals. This resulted in additional hard parses which were consuming significant database time. | 50.7                               |                        |
| Hard parsing of SQL statements was consuming significant database time.                                                                              | 55.1                               |                        |

- Based on the ADDM report, what is the recommended action?

The recommendation is to “set the CURSOR\_SHARING parameter to FORCE.”

- Create an AWR or Statspack report from the last two snapshots. The following shows the solution using AWR and Enterprise Manager.

| Step | Page         | Action                            |
|------|--------------|-----------------------------------|
| a    | Current page | Navigate to the Snapshots page.   |
| b    | Snapshots    | Select the next to last snapshot. |
|      |              | Select View Report from Actions.  |
|      |              | Click Go.                         |
| c    | View Report  | Select the last snapshot.         |
|      |              | Click OK.                         |

- Review the AWR report for the characteristics of a hard-parse problem. The goal is to reduce hard parses because the cost of a hard parse is at least an order of magnitude greater than the cost of a soft parse.
  - The Top 5 Timed Events show that there are no significant wait events. CPU Time by itself does not indicate a problem.

## Practice 15-2: Tuning a Hard-Parse Workload (continued)

### Top 5 Timed Foreground Events

| Event                         | Waits  | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
|-------------------------------|--------|---------|---------------|-----------|-------------|
| DB CPU                        |        | 137     |               | 54.97     |             |
| db file sequential read       | 14,746 | 35      | 2             | 14.03     | User I/O    |
| SQL*Net break/reset to client | 1,862  | 6       | 3             | 2.26      | Application |
| log file sync                 | 118    | 5       | 46            | 2.18      | Commit      |
| buffer busy waits             | 2      | 3       | 1321          | 1.06      | Concurrency |

- b) Time Model gives the first indication of a problem. Parse time elapsed is a significant portion of DB Time, and hard parse elapsed time is the majority of the parse time elapsed.

### Time Model Statistics

- Total time in database user-calls (DB Time): 248.6s
- Statistics including the word "background" measure background process time, and so do not contribute to the DB time statistic
- Ordered by % or DB time desc, Statistic name

| Statistic Name                             | Time (s) | % of DB Time |
|--------------------------------------------|----------|--------------|
| sql execute elapsed time                   | 219.80   | 88.41        |
| parse time elapsed                         | 156.79   | 63.07        |
| hard parse elapsed time                    | 137.06   | 55.13        |
| DB CPU                                     | 136.66   | 54.97        |
| PL/SQL execution elapsed time              | 11.83    | 4.76         |
| PL/SQL compilation elapsed time            | 4.56     | 1.84         |
| hard parse (sharing criteria) elapsed time | 3.00     | 1.21         |
| connection management call elapsed time    | 1.88     | 0.75         |
| hard parse (bind mismatch) elapsed time    | 0.94     | 0.38         |
| repeated bind elapsed time                 | 0.09     | 0.04         |
| sequence load elapsed time                 | 0.06     | 0.03         |
| DB time                                    | 248.62   |              |
| background elapsed time                    | 18.59    |              |
| background cpu time                        | 1.13     |              |

- c) In the Load Profile section, compare Parses per second to Hard parses per second. A majority of the parses are hard parses.

**Practice 15-2: Tuning a Hard-Parse Workload (continued)****Load Profile**

|                   | Per Second | Per Transaction | Per Exec | Per Call |
|-------------------|------------|-----------------|----------|----------|
| DB Time(s):       | 1.3        | 1.7             | 0.00     | 0.04     |
| DB CPU(s):        | 0.7        | 1.0             | 0.00     | 0.02     |
| Redo size:        | 30,586.1   | 40,724.9        |          |          |
| Logical reads:    | 7,741.5    | 10,307.7        |          |          |
| Block changes:    | 128.5      | 171.1           |          |          |
| Physical reads:   | 97.8       | 130.3           |          |          |
| Physical writes:  | 16.4       | 21.9            |          |          |
| User calls:       | 32.2       | 42.9            |          |          |
| Parses:           | 448.1      | 596.6           |          |          |
| Hard parses:      | 343.2      | 456.9           |          |          |
| W/A MB processed: | 1.5        | 2.1             |          |          |
| Logons:           | 0.3        | 0.3             |          |          |
| Executes:         | 586.8      | 781.3           |          |          |
| Rollbacks:        | 0.0        | 0.0             |          |          |
| Transactions:     | 0.8        |                 |          |          |

- d) Further investigation into the Instance Efficiencies Percentages section shows that % Non-Parse CPU is low, indicating that much of the CPU time is spent parsing. The Soft Parse %: shows that very few statements are found in the cache.

**Instance Efficiency Percentages (Target 100%)**

|                               |        |                   |        |
|-------------------------------|--------|-------------------|--------|
| Buffer Nowait %:              | 100.00 | Redo NoWait %:    | 100.00 |
| Buffer Hit %:                 | 98.82  | In-memory Sort %: | 100.00 |
| Library Hit %:                | 70.64  | Soft Parse %:     | 23.42  |
| Execute to Parse %:           | 23.63  | Latch Hit %:      | 100.00 |
| Parse CPU to Parse Elapsed %: | 68.37  | % Non-Parse CPU:  | 31.94  |

7. Implement the ADDM recommendation: Either click the Implement button on the ADDM report or use the `ALTER SYSTEM SET CURSOR_SHARING = FORCE SCOPE=BOTH` command.
  - a) Using the SQL\*Plus interface, execute the `ALTER SYSTEM SET CURSOR_SHARING = FORCE SCOPE=BOTH` command.

## Practice 15-2: Tuning a Hard-Parse Workload (continued)

```
$ sqlplus / as sysdba

...

SQL> ALTER SYSTEM SET CURSOR_SHARING = FORCE SCOPE=BOTH;

System altered.

SQL> exit
Disconnected from Oracle Database 10g Enterprise Edition
Release 10.2.0.1.0 - Production
With the Partitioning, OLAP and Data Mining options
$
```

8. Run the workload generator again and observe the differences in the tool of your choice.
  - a) Run the ./workgen HARDPARSE script.
  - b) After about 3 minutes, run ADDM from the Performance page as in step 2. The load graph will be much reduced.
  - c) Using the ADDM tool, you see the impact of changing the CURSOR\_SHARING parameter. The performance penalty of hard-parsing has dropped significantly. The hard-parse finding may not appear in the ADDM report. It is possible that there will be no findings. Notice that the value for Average Active Sessions has dropped significantly. This indicates that fewer sessions are waiting.

| ADDM Performance Analysis                                                          |                                                    |                                                  |                                                       |                                      |
|------------------------------------------------------------------------------------|----------------------------------------------------|--------------------------------------------------|-------------------------------------------------------|--------------------------------------|
| Task Name <b>ADDM:1239450590_1_22</b>                                              |                                                    |                                                  |                                                       |                                      |
| Task Owner <b>SYS</b>                                                              |                                                    | Average Active Sessions <b>0.3</b>               | Period Start Time <b>Feb 23, 2010 11:00:50 AM UTC</b> | Period Duration <b>3.2</b> (minutes) |
| <a href="#">Filters</a> <a href="#">View Snapshots</a> <a href="#">View Report</a> |                                                    |                                                  |                                                       |                                      |
| Impact (%) ▾                                                                       | Finding                                            | Occurrences (24 hrs ending with analysis period) |                                                       |                                      |
| 19.8                                                                               | <a href="#">Undersized Shared Pool</a>             | 2 of 17                                          |                                                       |                                      |
| 19.7                                                                               | <a href="#">"User I/O" wait Class</a>              | 5 of 17                                          |                                                       |                                      |
| 9                                                                                  | <a href="#">Hard Parse Due to Sharing Criteria</a> | 3 of 17                                          |                                                       |                                      |
| 7.1                                                                                | <a href="#">Soft Parse</a>                         | 3 of 17                                          |                                                       |                                      |
| 4.7                                                                                | <a href="#">PL/SQL Compilation</a>                 | 2 of 17                                          |                                                       |                                      |
| 4.4                                                                                | <a href="#">Commits and Rollbacks</a>              | 3 of 17                                          |                                                       |                                      |
| 2.3                                                                                | <a href="#">Session Connect and Disconnect</a>     | 1 of 17                                          |                                                       |                                      |

- d) Create an AWR report. On the ADDM page, click View Snapshots. On the Snapshot detail page, click the Report tab.
- e) Using the AWR report, review the significant statistics. Notice that all significant statistics show improvement. Notice that hard parses are a much smaller percentage of the parses, Soft Parse % is much improved, and % Non-Parse CPU is much higher. You may have that noticed that the workload script ran more quickly as well.

## Practice 15-2: Tuning a Hard-Parse Workload (continued)

| Load Profile      |            |                 |          |          |
|-------------------|------------|-----------------|----------|----------|
|                   | Per Second | Per Transaction | Per Exec | Per Call |
| DB Time(s):       | 0.3        | 1.1             | 0.00     | 0.04     |
| DB CPU(s):        | 0.2        | 0.7             | 0.00     | 0.02     |
| Redo size:        | 15,565.7   | 67,046.8        |          |          |
| Logical reads:    | 6,253.0    | 26,933.9        |          |          |
| Block changes:    | 38.2       | 164.3           |          |          |
| Physical reads:   | 135.1      | 581.8           |          |          |
| Physical writes:  | 10.5       | 45.1            |          |          |
| User calls:       | 6.4        | 27.5            |          |          |
| Parses:           | 365.5      | 1,574.5         |          |          |
| Hard parses:      | 7.2        | 31.1            |          |          |
| W/A MB processed: | 0.6        | 2.4             |          |          |
| Logons:           | 0.1        | 0.4             |          |          |
| Executes:         | 441.0      | 1,899.7         |          |          |
| Rollbacks:        | 0.0        | 0.0             |          |          |
| Transactions:     | 0.2        |                 |          |          |



## Practice 15-2: Tuning a Hard-Parse Workload (continued)

### Time Model Statistics

- Total time in database user-calls (DB Time): 50.6s
- Statistics including the word "background" measure background process time, and so do not contribute to the DB time statistic
- Ordered by % of DB time desc, Statistic name

| Statistic Name                             | Time (s) | % of DB Time |
|--------------------------------------------|----------|--------------|
| sql execute elapsed time                   | 45.64    | 90.22        |
| DB CPU                                     | 30.62    | 60.53        |
| parse time elapsed                         | 20.86    | 41.23        |
| hard parse elapsed time                    | 17.28    | 34.15        |
| PL/SQL execution elapsed time              | 6.83     | 13.51        |
| hard parse (sharing criteria) elapsed time | 4.57     | 9.03         |
| PL/SQL compilation elapsed time            | 2.35     | 4.65         |
| connection management call elapsed time    | 1.15     | 2.27         |
| hard parse (bind mismatch) elapsed time    | 0.35     | 0.69         |
| repeated bind elapsed time                 | 0.09     | 0.18         |
| sequence load elapsed time                 | 0.00     | 0.00         |
| DB time                                    | 50.59    |              |
| background elapsed time                    | 4.62     |              |
| background cpu time                        | 0.31     |              |

### Instance Efficiency Percentages (Target 100%)

|                              |        |                   |        |
|------------------------------|--------|-------------------|--------|
| Buffer Nowait %:             | 100.00 | Redo NoWait %:    | 100.00 |
| Buffer Hit %:                | 97.84  | In-memory Sort %: | 100.00 |
| Library Hit %:               | 94.03  | Soft Parse %:     | 98.02  |
| Execute to Parse %:          | 17.12  | Latch Hit %:      | 100.00 |
| Parse CPU to Parse Elapsd %: | 78.67  | % Non-Parse CPU:  | 81.09  |

### Practice 15-3: Keeping Objects in the Shared Pool

The performance of some applications is hindered by having to reload frequently used objects into the shared pool. If an object (such as a procedure, package, or sequence) is used often but not frequently enough to prevent aging out of the shared pool, keeping it will reduce the number of reloads. Keeping objects that are used very frequently is not necessary, but it does not hurt performance because those objects will stay in the shared pool without being kept.

- 1) Start the workload generator: `./workgen KEEP`.
- 2) After the workload has run for a few minutes, choose a candidate object to KEEP in the shared pool. Write and save a query to find a candidate object. Write the query to test whether the object has been KEPT. A suggested query is:

```
SELECT name, type, sharable_mem, kept
FROM v$db_object_cache
WHERE sharable_mem > 4000
AND EXECUTIONS > 5
AND (type='FUNCTION' OR type='PROCEDURE')
/
```

This query is available in the `$HOME/workshops` directory as `sp_objects.sql`. The results of the query will vary. The `DEPLETE_INV` procedure should appear in all queries.

```
$ cd $HOME/workshops
$ sqlplus / as sysdba

SQL> @sp_objects.sql
SQL> set echo on
SQL>
SQL> set pagesize 100
SQL>
SQL> SELECT name, type, sharable_mem, kept
2 FROM v$db_object_cache
3 WHERE sharable_mem > 4000
4 AND EXECUTIONS > 5
5 AND (type='FUNCTION' OR type='PROCEDURE')
6 /
```

NAME

| TYPE             | SHARABLE_MEM | KEP |
|------------------|--------------|-----|
| GETEMKEY         |              |     |
| FUNCTION         | 8576         | NO  |
| DEPLETE_INV      |              |     |
| PROCEDURE        | 16768        | YES |
| SETEMUSERCONTEXT |              |     |
| PROCEDURE        | 20880        | NO  |

### Practice 15-3: Keeping Objects in the Shared Pool (continued)

```
DECRYPT
FUNCTION 20864 NO

SQL>
```

3. Use the KEEP procedure to keep the DEPLETE\_INV PROCEDURE object in the shared pool.

```
SQL> EXECUTE dbms_shared_pool.keep('OE.DEPLETE_INV');
```

PL/SQL procedure successfully completed.

```
SQL> @sp_objects.sql
```

```
SQL> set echo on
```

```
SQL>
```

```
SQL> set pagesize 100
```

```
SQL>
```

```
SQL> SELECT name, type, sharable_mem, kept
2 FROM v$db_object_cache
3 WHERE sharable_mem > 4000
4 AND EXECUTIONS > 5
5 AND (type='FUNCTION' OR type='PROCEDURE')
6 /
```

NAME

| NAME             | TYPE      | SHARABLE_MEM | KEP        |
|------------------|-----------|--------------|------------|
| GETEMKEY         | FUNCTION  | 8576         | NO         |
| DEPLETE_INV      | PROCEDURE | 16768        | <b>YES</b> |
| SETEMUSERCONTEXT | PROCEDURE | 20880        | NO         |
| DECRYPT          | FUNCTION  | 20864        | NO         |

4. Flush the shared pool and observe the state of KEPT objects.

```
SQL> ALTER SYSTEM FLUSH SHARED_POOL;
```

System altered.

```
SQL> @sp_objects
```

**Practice 15-3: Keeping Objects in the Shared Pool (continued)**

```

SQL> set echo on
SQL>
SQL> set pagesize 100
SQL>
SQL> SELECT name, type, sharable_mem, kept
2 FROM v$db_object_cache
3 WHERE sharable_mem > 4000
4 AND EXECUTIONS > 5
5 AND (type='FUNCTION' OR type='PROCEDURE')
6 /

```

NAME

-----

TYPE

SHARABLE\_MEM KEPT

-----

DEPLETE\_INV

PROCEDURE

16768 YES

SQL> **exit**

5. Stop the workload generator by executing the `rm runload` command.

```
$ rm runload
```

6. Clean up the database by running the `./cleanup SP` script.

```

$./cleanup SP
Cleanup Finished

```

## Practices for Lesson 16

The goal of this practice is to observe the symptoms of, diagnose, and apply solutions for buffer cache problems.

## Practice 16-1: Using the DB Cache Advisor

In this practice, assume that you have an OLTP workload that needs some additional tuning and that the SQL statements are not available to tune, or have already been tuned using the SQL Tuning Advisor and SQL Access Advisor. You use the diagnostics in the AWR report and the DB Cache Advisor to choose an optimal size for the database buffer cache. The workload generator for this scenario creates an AWR snapshot, and then runs the workload.

1. Complete the setup for this scenario.
  - a) Log out of the Enterprise Manager.
  - b) Run `./prepare BC`.

```
$ cd $HOME/workshops
$./prepare BC
Database closed.
Database dismounted.
ORACLE instance shut down.

ORACLE instance started.

Total System Global Area 284585984 bytes
Fixed Size 1335980 bytes
Variable Size 268438868 bytes
Database Buffers 8388608 bytes
Redo Buffers 6422528 bytes
Database mounted.
Database opened.
Finished setup Buffer Cache
$
```

2. Verify that `DB_CACHE_SIZE` is set to 8M.

```
$ sqlplus / as sysdba

SQL> show parameter DB_CACHE_SIZE

NAME TYPE VALUE
----- - -
db_cache_size big integer 8M

SQL> exit
$
```

3. Start the workload generator with `./workgen BC`. The `workgen` script creates an AWR snapshot and a Statspack snapshot before starting the workload in the background.

```
$./workgen BC

PL/SQL procedure successfully completed.
```

## Practice 16-1: Using the DB Cache Advisor (continued)

```
PL/SQL procedure successfully completed.
```

```
$
```

4. Wait for about five minutes, and then in a terminal window, stop the workload with the `rm runload` command.

```
$ rm runload
```

5. Log in to Enterprise Manager. From the Performance page, click Run ADDM Now to generate an ADDM report. View Performance Analysis and check the performance findings. Examine the top finding that you, as a DBA, can address.
  - a) Review ADDM Performance Analysis. When you select the findings with the greatest impact, you may find several that you can do nothing about, such as “Top SQL Statements” shown in the following example. Assume that you have already tuned these statements as covered in the previous lessons and move down the list of findings. Other findings that you may see are: “CPU Usage,” which usually recommends SQL tuning and “User I/O wait Class,” which requires more analysis, sometimes leading to a change in the buffer cache size. The finding of interest in the example list is “Undersized Buffer Cache.” The finding “Free Buffer waits” is often associated with an undersized buffer cache.

ADDM Performance Analysis

Task Name

ADDM:1239450590\_1\_555

Filters

View Snapshots

View Report

Task Owner

SYS

Average Active Sessions

0.8

Period Start Time

Mar 10, 2010 7:26:26 AM UTC

Period Duration

8 (minutes)

| Impact (%) ▾               | Finding                                 | Occurrences (24 hrs ending with analysis period) |
|----------------------------|-----------------------------------------|--------------------------------------------------|
| <div><div></div>51.2</div> | <a href="#">Top SQL Statements</a>      | <a href="#">18 of 41</a>                         |
| <div><div></div>16</div>   | <a href="#">Undersized Buffer Cache</a> | <a href="#">17 of 41</a>                         |
| <div><div></div>3.8</div>  | <a href="#">Hard Parse</a>              | <a href="#">17 of 41</a>                         |

- b) Drill down on the “Undersized Buffer Cache” finding. Notice the recommendation to increase the buffer cache size to 12 MB. This value is limited by the advisor, which takes estimates only from 50% to 200% of the current setting.
6. Create an AWR report. Save the report as `awr_BC_1.html`.
  - a) On the ADDM report page, click View Snapshots. On the Snapshot Details page, click the Report tab.
  - b) Set the download preferences. Note that these instructions are for Firefox 3.0.15. In the Browser menu, click Edit > Preferences. In the Preferences dialog box, in the Downloads section, select “Always ask me where to save files.” Click Close.

### Practice 16-1: Using the DB Cache Advisor (continued)

- c) At the top right of the report, click “Save to File.” When prompted, specify the name of the report. Notice which directory the report is being created in so that you can find it later, and write the directory path here.

7. Review the AWR report and identify the primary performance issue.

- a) The Top 5 Timed Foreground Events section shows one or more of the events that point to the buffer cache size: db file sequential read, db file scattered read, buffer busy waits, free buffer waits, latch: cache buffer lru chain, or latch: cache buffers chains.

From the report, determine the possible impact of the buffer cache problem. If all waits associated with the physical reads are eliminated, what percent increase in performance could you expect? \_\_\_\_\_ (10 to 25 % is expected. The following example shows ~17% = db file sequential reads + latch: cache buffers chains.) Of course, removing all the waits for physical reads is not possible. Reducing the physical reads will often reduce the DB CPU.

**Note:** Your results will vary.

| Top 5 Timed Foreground Events |           |         |               |           |             |
|-------------------------------|-----------|---------|---------------|-----------|-------------|
| Event                         | Waits     | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
| DB CPU                        |           | 257     |               | 64.05     |             |
| db file sequential read       | 3,722,448 | 67      | 0             | 16.69     | User I/O    |
| latch: cache buffers chains   | 59        | 3       | 51            | 0.76      | Concurrency |
| direct path read              | 105,223   | 2       | 0             | 0.60      | User I/O    |
| direct path write temp        | 63        | 1       | 19            | 0.30      | User I/O    |

- b) What is Time Model showing as the greatest loads? \_\_\_\_\_  
(sql execute elapsed time)
- c) What is the buffer cache size in the report in the Cache Sizes section of Report Summary? \_\_\_\_\_ (8M)
- d) What does the Load Profile section show for physical reads per second and transaction? \_\_\_\_\_ per second \_\_\_\_\_ per transaction. These numbers will have meaning when you compare them to the next test.
- e) What is the buffer hit % (buffer cache hit ratio) in the Instance Efficiency section? \_\_\_\_\_ (For OLTP type activity, less than 90% is considered a low buffer cache hit ratio.)



## Practice 16-1: Using the DB Cache Advisor (continued)

- f) To confirm the symptoms, check the top misses in the Latch Sleep Breakdown section. Look for misses and sleeps for latches related to the buffer cache. The number of misses and sleeps may be small enough in comparison to the get requests that they would not be considered at all, except that because they are the top in the latch sleep breakdown, they are additional evidence pointing to the buffer cache. What latches are listed at or near the top related to the buffer cache?

---

(You may see “cache buffers lru chain” or “cache buffers chains.”)

8. Use the Buffer Pool Advisory in the AWR report. The navigation path is Back to Top > Advisory Statistics > Buffer Pool Advisory.

What is an appropriate size for the database buffer cache?

\_\_\_\_\_ (16 MB or more)

In the advisory, what is the value shown for the estimated number of physical reads for the recommendation of 16 MB? \_\_\_\_\_ (The change in Est Physical Read Factor should be in the range of 0.5 to 0.03 depending on your machine.)

**Note:** Due to the limitations of the advisor, the maximum estimated cache is only 200% of the current size. A larger buffer pool may provide greater benefit.

9. Change the buffer cache size. To reduce the number of trials, increase DB\_CACHE\_SIZE to 20 MB.
- Navigate to the Memory Advisors page.
    - Click the Database tab.
    - Click the Server tab.
    - Click Memory Advisors in the Database Configuration section.
  - Reduce the shared pool size by 20 MB to 220 MB.
  - Click Apply. This reduced the memory allocated by 20 MB. Now you can increase the buffer cache size.
  - Change the buffer cache to 20 MB, and then click Apply as shown in the following example.

From the previous advisor run, you can see that the optimal size is greater than 16 MB. By changing DB\_CACHE\_SIZE to 20 MB, the next trial should allow the advisor to show a better value.

## Practice 16-1: Using the DB Cache Advisor (continued)

**Memory Advisors**

Page Refreshed **March 30, 2010 10:42:13 AM UTC** [Refresh](#) [Show SQL](#) [Revert](#) [Apply](#)

When Automatic Memory Management is enabled, the database will automatically set the optimal distribution of memory. The distribution of memory will change from time to time to accommodate changes in the workload.

Automatic Memory Management **Disabled** [Enable](#)

**SGA** **PGA**

The System Global Area (SGA) is a group of shared memory structures that contains data and control information for one Oracle database. The SGA is allocated in memory when an Oracle database instance is started.

Automatic Shared Memory Management **Disabled** [Enable](#)

Shared Pool  MB [Advice](#)

Buffer Cache  MB [Advice](#)

Large Pool  MB

Java Pool  MB

Other (MB)

Total SGA (MB) **263** [Calculate](#)

| Component    | Percentage |
|--------------|------------|
| Shared Pool  | 83.5%      |
| Buffer Cache | 7.6%       |
| Large Pool   | 1.5%       |
| Java Pool    | 4.6%       |
| Other        | 2.8%       |

10. Start the workload generator again: `./workgen BC`.

```
$./workgen BC
```

11. After about five minutes, stop the workload with the `rm runload` command.

```
$ rm runload
```

12. Create an ADDM report. View the Performance Analysis. (This is the same procedure as shown in step 5.) Check the performance findings. Did the primary issues change?

Yes, the issues have changed. The major change shown on the example report is that Average Active Sessions has reduced from the previous report. Each session in this workload is taking less time.

**Note:** The report you get could be completely different from the example report depending on the classroom machine. Your report may seem to show an increase in issues and impact. A close inspection of the AWR report shows that performance has improved. The example report shows a reduction in the impact of the Top SQL Statements finding and the elimination of the Undersized Buffer Cache finding.

ADDM Performance Analysis

Task Name

ADDM:1239450590\_1\_557 (End Time:Mar 10, 2010 8:31:17 AM)

Filters

View Snapshots

View Report

Task Owner

SYS

Average Active Sessions

0.6

Period Start Time

Mar 10, 2010 8:25:30 AM UTC

Period Duration (minutes)

5.8

Impact (%)

Finding

Occurrences (24 hrs ending with analysis period)

29.4

Top SQL Statements

18 of 42

2.3

"User I/O" wait Class

11 of 42

### Practice 16-1: Using the DB Cache Advisor (continued)

13. On the ADDM page, click View Snapshots. On the Snapshots page, click the Report tab to create an AWR report. Save this report as awr\_BC\_2.html.
14. In another window or browser tab, display the awr\_BC\_1.html file, so that you can compare the two reports.
15. Review the AWR report and identify the primary performance changes.
  - a) In Top 5 Timed Events, the order of events and the amount of time waiting changed. For “db\_sequential\_reads,” the % DB time spent waiting dropped now that more blocks are being held in the cache. Waits in terms of % DB time for the other buffer cache–related items have dropped. Overall, more time is going to useful work (SQL execution) and less time to physical reads. In the following example, almost all the wait time is part of the DB CPU.
 

**Note:** When the elapsed time of the test period is not the same, compare only the metrics that are normalized. Per second, per transaction, and % DB time metrics are normalized metrics.

| Top 5 Timed Foreground Events |        |         |               |           |             |
|-------------------------------|--------|---------|---------------|-----------|-------------|
| Event                         | Waits  | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
| DB CPU                        |        | 143     |               | 71.31     |             |
| db file sequential read       | 89,156 | 2       | 0             | 1.23      | User I/O    |
| direct path read              | 82,810 | 1       | 0             | 0.63      | User I/O    |
| direct path write temp        | 49     | 1       | 16            | 0.39      | User I/O    |
| latch: cache buffers chains   | 19     | 1       | 40            | 0.38      | Concurrency |

- b) What is the number of physical reads per second shown in the Load Profile section? \_\_\_\_\_

Compare this number with the number of physical reads per second in the AWR\_BC\_1 report. (The number of physical reads per second should have dropped significantly because the buffer cache size is increased to 20 MB.)

- c) What is the buffer hit ratio in the Instance Efficiencies section?  
 \_\_\_\_\_ (OLTP applications expect a buffer hit ratio of 90% or higher.)  
 You should see an improvement from the earlier report.
- d) The Buffer Cache Advisory indicates an optimal buffer cache size. What size is indicated for your instance and workload?  
 \_\_\_\_\_ (A value of approximately 24-36 MB is expected.)

An optimal size is the point where additional memory reduces the Estimated Physical Reads by a small amount. The optimal value for your instance may vary from the expected value.

## Practice 16-2: Using the Keep Pool

In this scenario, you assume the same workload as in scenario 16-1, but you are constrained on memory resources. Adjust the buffer pool memory areas to use as little memory as possible with the most benefit. The `prepare KC` script sets `DB_CACHE_SIZE` to 16 MB and `DB_KEEP_CACHE_SIZE` to 4 MB. A script that assigns a set of frequently used small tables to the keep pool is provided. Will this perform as well or better than having only `DB_CACHE_SIZE` set to 20 MB, thus allowing you to reduce the optimal size of the database buffer cache?

1. Prepare the database for this scenario.
  - a) Log out of Enterprise Manager.
  - b) Run the `./prepare KC` script.

```
$ cd $HOME/workshops
$./prepare KC
Database closed.
Database dismounted.
ORACLE instance shut down.
ORACLE instance started.

Total System Global Area 272019456 bytes
Fixed Size 1335924 bytes
Variable Size 247467404 bytes
Database Buffers 20971520 bytes
Redo Buffers 2244608 bytes
Database mounted.
Database opened.
Finished setup KC
$
```

2. Examine the first AWR report from Practice 16-1, `awr_BC_1.html`. Are there tables that could be cached in a keep pool?
  - a) The Top 5 Timed Foreground Events show one or more of the events that point to the buffer cache size: db file sequential read, db file scattered read, free buffer waits, latch: cache buffer chains, or latch: cache buffer lru chain.

### Top 5 Timed Foreground Events

| Event                       | Waits     | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
|-----------------------------|-----------|---------|---------------|-----------|-------------|
| DB CPU                      |           | 257     |               | 64.05     |             |
| db file sequential read     | 3,722,448 | 67      | 0             | 16.69     | User I/O    |
| latch: cache buffers chains | 59        | 3       | 51            | 0.76      | Concurrency |
| direct path read            | 105,223   | 2       | 0             | 0.60      | User I/O    |
| direct path write temp      | 63        | 1       | 19            | 0.30      | User I/O    |

## Practice 16-2: Using the Keep Pool (continued)

- b) How many full-table scans are being performed? Check the instance activity statistics. **Note:** Table scans (direct read) do not go into the buffer cache; these reads go directly into the session PGA. You should find that the number of table scans (long) and table scans (direct read) are almost the same.

| Statistic                 | Per second | Per transaction |
|---------------------------|------------|-----------------|
| table scans (direct read) |            |                 |
| table scans (long)        |            |                 |
| table scans (short)       |            |                 |

- c) Find the SQL statements that have the highest number of physical reads.  
**Hint:** Back to Top > SQL Statistics > SQL Ordered by Reads. The SQL ID may be different on your machine than that shown in the example.

SQL\_ID \_\_\_\_\_

SQL\_ID \_\_\_\_\_

### SQL Ordered by Reads

| Physical Reads | Executions | Reads per Exec | %Total | Elapsed Time (s) | %CPU  | %IO   | SQL Id                        | SQL Module | SQL Text                                |
|----------------|------------|----------------|--------|------------------|-------|-------|-------------------------------|------------|-----------------------------------------|
| 2,871,409      | 6          | 478,568.17     | 94.24  | 139.07           | 68.44 | 32.28 | <a href="#">9zb56qt165f3z</a> | SQL*Plus   | DECLARE CURSOR C2 IS<br>SELECT CU...    |
| 2,838,619      | 60,000     | 47.31          | 93.17  | 121.94           | 68.13 | 36.58 | <a href="#">431mwkvt65jbb</a> | SQL*Plus   | SELECT<br>SUM(AMOUNT_SOLD)<br>TOTAL ... |
| 142,623        | 96         | 1,485.66       | 4.68   | 16.15            | 55.09 | 6.52  | <a href="#">cpsda56ma8b15</a> | SQL*Plus   | DECLARE CURSOR C1 IS<br>SELECT * ...    |
| 135,338        | 96         | 1,409.77       | 4.44   | 11.24            | 57.89 | 8.49  | <a href="#">9a9ab05m4mptw</a> | SQL*Plus   | SELECT * FROM<br>ALL_TABLES WHERE...    |
| 18,344         | 432        | 42.46          | 0.60   | 5.47             | 72.07 | 2.20  | <a href="#">6g215w0qfstc2</a> | SQL*Plus   | SELECT<br>SUM(UNIT_COST) TOTAL<br>FR... |
| 13,580         | 432        | 31.44          | 0.45   | 4.47             | 80.16 | 3.07  | <a href="#">0u9dmyxy318w0</a> | SQL*Plus   | SELECT<br>SUM(AMOUNT_SOLD)<br>TOTAL ... |
| 9,011          | 29         | 310.72         | 0.30   | 89.32            | 62.47 | 0.15  | <a href="#">14zitd9unpsq9</a> | SQL*Plus   | DECLARE CURSOR C2 IS<br>SELECT * ...    |
| 5,768          | 672        | 8.58           | 0.19   | 1.29             | 60.87 | 5.40  | <a href="#">70uy0ku756w7a</a> | SQL*Plus   | SELECT<br>COUNT(INDEX_NAME)<br>FROM ... |
| 5,027          | 9,251      | 0.54           | 0.16   | 1.33             | 72.97 | 6.42  | <a href="#">av6kf9w9m362a</a> | SQL*Plus   | SELECT COUNT(*) FROM<br>ORDER_ITE...    |
| 3,911          | 5          | 782.20         | 0.13   | 4.80             | 25.41 | 3.49  | <a href="#">6qvch1xu9ca3g</a> |            | DECLARE job<br>BINARY_INTEGER := ...    |

**Practice 16-2: Using the Keep Pool (continued)**

- d) Check some of the SQL IDs for the full text of the statements with the highest number of physical reads to determine the segments involved. Then, find the table names that are being accessed.

What tables are most accessed? \_\_\_\_\_ (SALES, PRODUCTS, and COSTS)

The SQL IDs shown in this example correspond to the SQL IDs in the previous step. Your SQL text will be the same but the SQL ID may differ.

| SQL ID        | SQL Text                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|---------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 9zb56qt165f3z | <pre> DECLARE CURSOR C2 IS SELECT CUST_ID FROM CUSTOMERS WHERE CUST_ID BETWEEN 0 and 10000; CURSOR C1 IS SELECT PROD_ID FROM PRODUCTS; TYPE emp_list IS TABLE OF VARCHAR2(30); parameters emp_list; j NUMBER(10); tot_amount NUMBER(20); BEGIN FOR C1_ROW in C1 LOOP SELECT SUM(unit_cost) total INTO TOT_AMOUNT FROM sh.costs WHERE prod_id = C1_ROW.prod_id; END LOOP; FOR C2_ROW in C2 LOOP SELECT SUM(AMOUNT_SOLD) total INTO TOT_AMOUNT FROM sh.sales WHERE cust_id = C2_ROW.cust_id; END LOOP; FOR C1_ROW in C1 LOOP SELECT SUM(AMOUNT_SOLD) total INTO TOT_AMOUNT FROM sh.sales WHERE prod_id = C1_ROW.prod_id; END LOOP; END; </pre> |
| 431mwkyt65jbb | <pre> SELECT SUM(AMOUNT_SOLD) TOTAL FROM SH.SALES WHERE CUST_ID = :B1 </pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |

- e) Check the Segments by Physical Reads in the Segment Statistics section.  
What segments have the highest physical reads?

Partitions of the SH.SALES table are expected to appear. This is a very large table. The small tables that should be placed in the keep pool do not appear. Knowledge of the application is needed to properly choose the tables and indexes to keep.

## Practice 16-2: Using the Keep Pool (continued)

- The developers have provided a script named `keep_pool.sh` that will create a small keep pool of 4 MB. Assign some small but very active tables and indexes to the keep pool, and load these tables into the pool. The `keep_pool.sh` script also assigns the most active partitions of the `SALES` table to the keep pool. Review this script, and then run it. The full-table scan caused by the `SELECT * FROM <table_name>` statement will cause the table data to be loaded into the keep pool before the test run. The `ORDER BY` clause is used to force the use of an index so that the index is loaded into the keep cache. These accesses mean that the physical reads due to the first access is not included in the test statistics.

```
$ cat keep_pool.sh
#!/bin/bash
#-- create a KEEP pool
#-- and set the table properties to make use of it
#-- Select statements populate the keep cache

sqlplus -S /nolog <<-EOF
connect / as sysdba
ALTER SYSTEM SET db_keep_cache_size = 4M;
exit
EOF

sqlplus -S /nolog <<-EOF
connect hr/hr

ALTER TABLE EMPLOYEES STORAGE(BUFFER_POOL KEEP);
ALTER TABLE COUNTRIES STORAGE(BUFFER_POOL KEEP);
ALTER TABLE DEPARTMENTS STORAGE(BUFFER_POOL KEEP);
ALTER TABLE LOCATIONS STORAGE(BUFFER_POOL KEEP);
ALTER TABLE REGIONS STORAGE(BUFFER_POOL KEEP);
ALTER TABLE JOB_HISTORY STORAGE(BUFFER_POOL KEEP);
ALTER INDEX HR.DEPT_ID_PK STORAGE(BUFFER_POOL KEEP);
ALTER INDEX HR.DEPT_LOCATION_IX STORAGE(BUFFER_POOL KEEP);
ALTER INDEX HR.JHIST_EMPLOYEE_IX STORAGE(BUFFER_POOL KEEP);

select * from employees;
select * from departments;
select * from countries;
select * from locations;
select * from regions;
select * from job_history;
exit
EOF

sqlplus -S /nolog <<-EOF
connect oe/oe

ALTER TABLE orders STORAGE(BUFFER_POOL KEEP);
ALTER TABLE order_items STORAGE(BUFFER_POOL KEEP);
ALTER TABLE inventories STORAGE(BUFFER_POOL KEEP);
ALTER TABLE customers STORAGE(BUFFER_POOL KEEP);
```

**Practice 16-2: Using the Keep Pool (continued)**

```

ALTER TABLE product_information STORAGE(BUFFER_POOL KEEP);
ALTER INDEX INV_PRODUCT_IX STORAGE(BUFFER_POOL KEEP);
ALTER INDEX PRODUCT_INFORMATION_PK STORAGE(BUFFER_POOL KEEP);
ALTER INDEX OE.ORDER_PK STORAGE(BUFFER_POOL KEEP);
ALTER INDEX OE.ORD_CUSTOMER_IX STORAGE(BUFFER_POOL KEEP);
ALTER INDEX OE.ITEM_ORDER_IX STORAGE(BUFFER_POOL KEEP);
ALTER INDEX OE.PRODUCT_INFORMATION_PK STORAGE(BUFFER_POOL
KEEP);

select * from orders;
select * from order_items;
select * from inventories;
select * from customers
order by customer_id;
select * from product_information
order by product_id;
exit
EOF

sqlplus -S /nolog > /dev/null <<-EOF
set termout off
set feedback off
connect sh/sh
ALTER TABLE SALES MODIFY PARTITION SALES_Q3_1999
STORAGE(BUFFER_POOL KEEP);
ALTER TABLE SALES MODIFY PARTITION SALES_Q4_2001
STORAGE(BUFFER_POOL KEEP);
select * from sales;
ALTER TABLE COSTS STORAGE(BUFFER_POOL KEEP);
select * from costs;
EOF

$./keep_pool.sh
$

```

4. Test this new configuration with the workload from scenario 15-1. Run the `./workgen BC` script.

```

$./workgen BC

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

$

```

5. Wait for about five minutes, and then end the workload by deleting the runload file. While you are waiting, examine the graphs on the Performance page. Notice the peak values and the I/O per second values.

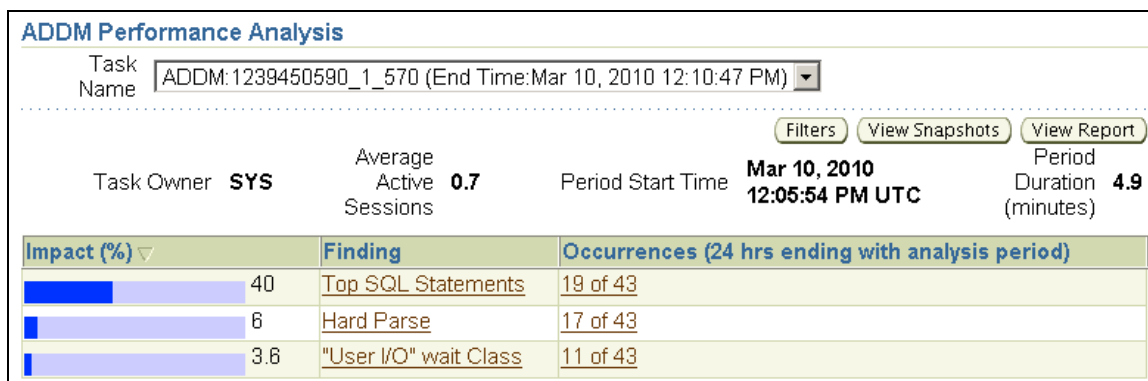


## Practice 16-2: Using the Keep Pool (continued)

6. Log in to Enterprise Manager. Create an ADDM report. Does the finding “Undersized Buffer Cache” appear? \_\_\_\_\_  
Has the impact for any of the findings changed? \_\_\_\_\_

(Yes. There may be a variety of findings shown, but Undersized Buffer Cache should not be one of them. The Keep Buffer cache may not improve performance and the Undersized Buffer Cache find may remain. The findings and impact numbers may not clearly indicate that performance has improved over the previous ADDM report associated with the second test in Practice 16-1, awr\_BC\_2. In the example, the clearest indicator on the ADDM report is Average Active Sessions. If the workload in two different periods is identical, you can deduce that the time required for each session is reduced, leading to a lower average. If the workload is not identical, this deduction may not hold true. In the following example, you can see that Average Active Sessions has increased somewhat over the second run in practice 16-1.

Recall that impact is measured in % DB Time. There are cases where the Total DB Time goes down, and the DB Time for a particular area does not go down proportionally, so the % DB Time for that area goes up.



7. Create an AWR report. Click View Snapshots. When the Snapshots page appears, click the Reports tab. Save this report as awr\_kc\_1.html.
8. Examine the AWR report and compare it with the previous report, awr\_BC\_2.html.
- Note:** If the elapsed time for awr\_BC\_2.html is very different from the elapsed time for awr\_kc\_1.html, the values of % DB Time and Physical Reads may not show the expected values because of the short durations of the workloads.
- Answer the following questions:
- a) What is the difference in Top 5 Timed Events? (Make the comparison based on % DB Time.) \_\_\_\_\_

The changes in % DB Time may be very small, but it may be enough to change the order of the top waits.

**Practice 16-2: Using the Keep Pool (continued)**

| Top 5 Timed Foreground Events    |         |         |               |           |             |
|----------------------------------|---------|---------|---------------|-----------|-------------|
| Event                            | Waits   | Time(s) | Avg wait (ms) | % DB time | Wait Class  |
| DB CPU                           |         | 123     |               | 62.85     |             |
| db file sequential read          | 361,613 | 6       | 0             | 2.95      | User I/O    |
| direct path read                 | 21,126  | 1       | 0             | 0.59      | User I/O    |
| enq: KO - fast object checkpoint | 2       | 1       | 393           | 0.40      | Application |
| latch: cache buffers chains      | 7       | 0       | 39            | 0.14      | Concurrency |

- b) In the Load Profile section, is there a difference in physical reads per second and physical reads per transactions for each of the two periods?

|                 | Statistic             | Per second | Per transaction |
|-----------------|-----------------------|------------|-----------------|
| <b>awr_BC_2</b> | <b>Physical reads</b> |            |                 |
| <b>awr_KC_1</b> | <b>Physical reads</b> |            |                 |

The number of physical reads in the KC\_1 report may be higher than the physical reads reported in BC\_2. This illustrates the efficiency of the buffer cache algorithms. If you recall, the physical reads for the first test were very high; this also illustrates that the keep pool can be an effective way of reducing physical reads when the application access patterns are known.

- c) In the Load Profile, compare DB Time per second for each of the two reports.

|                 | Statistic      | Per second |
|-----------------|----------------|------------|
| <b>awr_BC_2</b> | <b>DB Time</b> |            |
| <b>awr_KC_1</b> | <b>DB Time</b> |            |

The DB Time difference is the clearest indicator, in this practice, of a difference in performance.

- d) Is there a difference in Buffer Hit %? \_\_\_\_\_

The overall Buffer Hit % may change very little. The buffer cache hit ratio is a combination of the pools.

- e) Check the Buffer Pool Statistics. Is the keep pool being used?

\_\_\_\_\_

## Practice 16-2: Using the Keep Pool (continued)

Yes, the Buffer Pool Statistics section shows 100% hit ratio on the keep pool. The frequently used blocks are moved to the keep pool and a large number of hits are removed from the default pool calculation. You would expect the hit ratio of the default pool to decline. The statistics over such a small data set and short duration may not show any difference. A small number of physical reads may appear in the Keep Pool statistics; this indicates that the Pool has more blocks assigned to it than there is space available.

- f) How many full-table scans are being performed? Check the instance activity statistics for both reports: awr\_BC\_2 and awr\_KC\_1. (Table scans are unrelated to the physical reads when tables are in the keep buffer cache, but placing indexes in the keep pool can influence the optimizer to use an index rather than do a full-table scan) **Note:** All or nearly all of the full-table scans are direct read. This means that these reads are bypassing the buffer cache and going to PGA (server process memory).

| Statistic           | Per second<br>(awr_BC_2) | Per second<br>(awr_KC_1) |
|---------------------|--------------------------|--------------------------|
| table scans (long)  |                          |                          |
| table scans (short) |                          |                          |

9. Considering that this instance is memory constrained, what is the smallest DB\_CACHE\_SIZE that would significantly affect performance based on the Buffer Pool Advisor? \_\_\_\_\_

The advisor shows that there is a benefit by increasing the buffer cache to between 28 and 36 MB. The keep pool in this example did not show improvement in the amount of memory used. This practice did show the process to identify and cache tables in the keep cache.

10. Run the clean up script: From the /home/oracle/workshops directory, execute `./cleanup KC`.

```
$./cleanup KC
```

## Practices for Lesson 17

In this practice, you tune PGA Memory.

### **Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET**

You have a decision support system (DSS) that has a variety of queries being performed. Users are complaining that the queries are slow. PGA\_AGGREGATE\_TARGET is set at 20 MB. Determine the smallest setting that eliminates all multipass work areas.

1. Prepare the database for this scenario.
  - a) Log out of Enterprise Manager.
  - b) Execute the prepare script for this practice. Change directory to \$HOME/workshops. Execute ./prepare PGA.

```
$ cd $HOME/workshops
$./prepare PGA
Database closed.
Database dismounted.
ORACLE instance shut down.

ORACLE instance started.

Total System Global Area 192507904 bytes
Fixed Size 1335388 bytes
Variable Size 155193252 bytes
Database Buffers 33554432 bytes
Redo Buffers 2424832 bytes
Database mounted.
Database opened.

Tablespace created.

Database altered.

System altered.

Finished setup PGA
$
```

2. Start the workload generator with ./workgen PGA. This workload consists of several sessions performing a variety of queries that require different sizes of work areas.
3. Log in to Enterprise Manager. Wait a minute or two, and then, while the workload is running, use PGA Advisor and PGA Memory Usage Details to select a new PGA\_AGGREGATE\_TARGET.

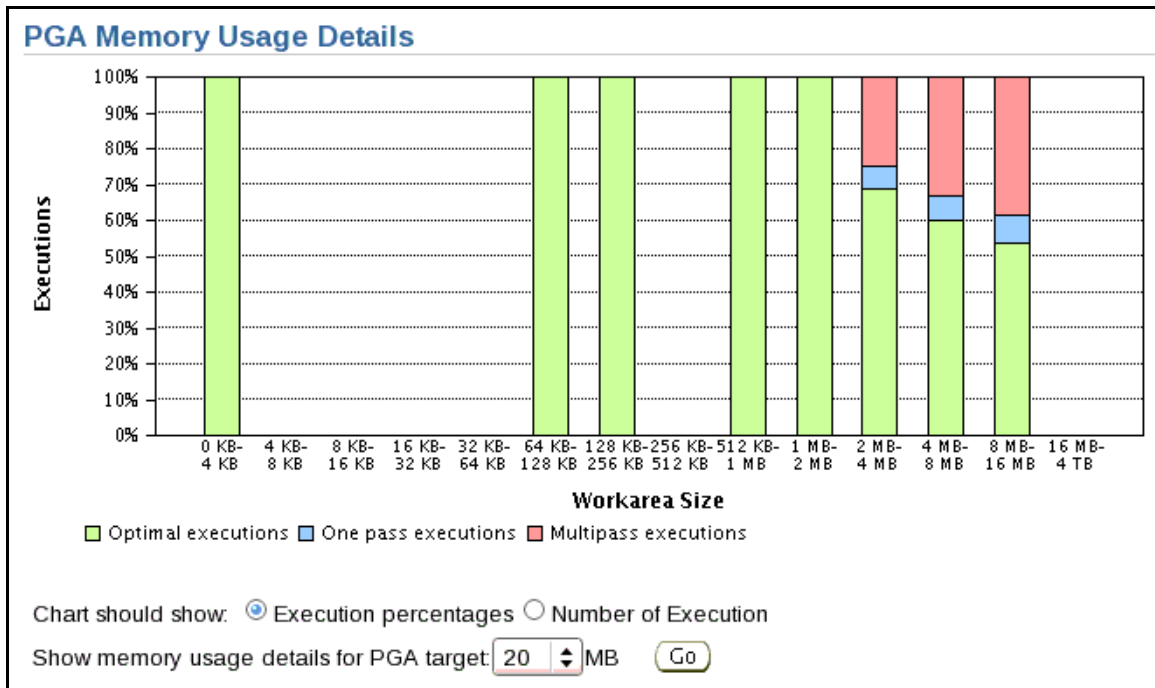
| Step | Page         | Action                  |
|------|--------------|-------------------------|
| a    | Current page | Click the Database tab. |

**Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)**

|   |                                                                                        |                                                                                                                                 |
|---|----------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| b | Database home                                                                          | Click the Server tab.                                                                                                           |
| c | Server page                                                                            | Click Memory Advisors in the Database Configuration section.                                                                    |
| d | Memory Advisors                                                                        | Click the PGA tab.                                                                                                              |
| e | Memory Advisor (PGA)                                                                   | Click PGA Memory Usage Details.                                                                                                 |
| f | PGA Memory Usage Details<br><br>(See the example screenshot, which follows the table.) | Set the value: Show memory usage details for the PGA target, 20 MB.                                                             |
|   |                                                                                        | Notice the sizes of the work areas that overflow to disk.                                                                       |
|   |                                                                                        | Notice the relative number of Optimal, One Pass, and Multipass executions.                                                      |
|   |                                                                                        | Click OK.                                                                                                                       |
| g | Memory Advisor (PGA)                                                                   | Click Advice.                                                                                                                   |
| h | PGA Aggregate Target Advice                                                            | Notice the overflow range (in red). This is the range where the required PGA for the number of sessions exceeds the PGA Target. |
|   |                                                                                        | Make a note of the PGA Target (MB) value where the red line changes to blue. (80 MB is expected.)                               |
|   |                                                                                        | Click Cancel.                                                                                                                   |
| i | Memory Advisor (PGA)                                                                   | Change Aggregate PGA Target to the value of the PGA Target that you noted in step 3h. (80 MB)                                   |
|   |                                                                                        | Click Apply.                                                                                                                    |

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

f) The following screenshot shows the PGA Memory Usage Details page:



4. Stop the workload with `rm runload`.

```
$ rm runload
```

5. Execute the workload with the `./workgen PGA` script.

```
$./workgen PGA
```

6. After the workload has been running for about five minutes, stop the workload with `rm runload`.

```
$ rm runload
```

7. Create an AWR snapshot and an AWR report for the last two snapshots.

| Step | Page                          | Action                               |
|------|-------------------------------|--------------------------------------|
| a    | Current page                  | Click the Database tab.              |
| b    | Database home                 | Click the Server tab.                |
| c    | Server tab page               | Click Automatic Workload Repository. |
| d    | Automatic Workload Repository | Click the Snapshot link.             |
| e    | Snapshots                     | Click Create.                        |

**Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)**

|   |                                     |                                                      |
|---|-------------------------------------|------------------------------------------------------|
| f | Confirmation                        | Click Yes.                                           |
| g | Snapshots                           | Select the beginning snapshot number.                |
|   |                                     | From the Actions drop-down list, select View Report. |
|   |                                     | Click Go.                                            |
| h | View Report: Select Ending Snapshot | Select the ending snapshot number.                   |
|   |                                     | Click OK.                                            |

8. Review the AWR report for symptoms and recommendations concerning PGA area sizing. Save the report as awr\_PGA\_1.html.
  - a) Check Top 5 Timed Foreground Events. Is there something that is taking a large amount of % DB Time? \_\_\_\_\_ (direct path read temp, direct path write temp)
  - b) Check Time Model Statistics. What areas are taking the greatest portion of DB Time? \_\_\_\_\_ (sql execute elapsed time, DB CPU). At this point, all that you know is that the SQL is slow, and the main diagnostics agree.
  - c) Determine the SQL with the greatest elapsed time (Back to Top > SQL Statistics > SQL ordered by Elapsed Time). There are several statements that are taking a large amount of elapsed time and CPU time, but very few executions. This combination of elapsed time and CPU time could be caused by excessive reads or sorting. Only one of the top six SQL statements is also in the SQL ordered by Reads, but five out of the top six are also in the SQL ordered by Gets. This indicates that many blocks are being accessed, but most of them are already in the buffer cache. Confirm this by checking the buffer cache hit ratio.



## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

### SQL ordered by Elapsed Time

| Elapsed Time (s) | Executions | Elapsed Time per Exec (s) | %Total | %CPU | %IO   | SQL Id        | SQL Module                  | SQL Text                                |
|------------------|------------|---------------------------|--------|------|-------|---------------|-----------------------------|-----------------------------------------|
| 557.00           | 2          | 278.50                    | 12.51  | 3.75 | 0.00  | cj4bx12rrw4p  | SQL*Plus                    | select s.prod_id,<br>SUM(s.amount...    |
| 423.26           | 2          | 211.63                    | 9.51   | 3.94 | 0.00  | 589mxw1dz9mhn | SQL*Plus                    | select s.prod_id,<br>sum(s.amount...    |
| 408.82           | 7          | 58.40                     | 9.18   | 2.51 | 0.00  | 8vjh4nbv1phu3 | SQL*Plus                    | select<br>sum(s.AMOUNT_sold) -<br>su... |
| 398.77           | 12         | 33.23                     | 8.96   | 3.43 | 0.00  | 2rmw64nrta33j | SQL*Plus                    | select s.prod_id,<br>sum(s.amount...    |
| 387.72           | 7          | 55.39                     | 8.71   | 2.95 | 0.00  | 6rz119f0r1wqv | SQL*Plus                    | select s.prod_id,<br>sum(s.amount...    |
| 334.84           | 18         | 18.60                     | 7.52   | 4.27 | 16.24 | 54dxhddg0mfdt | SQL*Plus                    | select * from (select *<br>from d...    |
| 309.48           | 14         | 22.11                     | 6.95   | 4.00 | 18.49 | fkmrxxv6fxqar | SQL*Plus                    | select * from (select *<br>from d...    |
| 295.78           | 12         | 24.65                     | 6.64   | 4.28 | 42.66 | bxwvcp4pjr43n | SQL*Plus                    | select * from (select *<br>from d...    |
| 283.28           | 17         | 16.66                     | 6.36   | 5.10 | 32.50 | cssuux75dw74c | SQL*Plus                    | select * from (select *<br>from d...    |
| 275.84           | 14         | 19.70                     | 6.20   | 5.85 | 44.03 | 3bg3pfux13pkr | SQL*Plus                    | select * from (select *<br>from d...    |
| 88.01            | 7          | 12.57                     | 1.98   | 2.91 | 0.36  | 3g36d1uum3662 | SQL*Plus                    | select PROD_ID,<br>SUM(amount_sol...    |
| 63.35            | 1          | 63.35                     | 1.42   | 2.64 | 0.04  | ab7ktcdksu27j | emagent_SQL_oracle_database | /* OracleOEM */ SELECT<br>/*+ IND...    |
| 60.30            | 7          | 8.61                      | 1.35   | 1.98 | 0.01  | 8rtsdaxm7apq7 | SQL*Plus                    | select<br>SUM(amount_sold),<br>PROD_... |

d) Navigate to the Instance Efficiencies Percentage report in the Report Summary section (Back to Top > Report Summary). What is the buffer cache hit ratio? \_\_\_\_\_ (99% expected)

e) Check Load Profile for logical reads and physical reads to confirm the buffer cache hit ratio.

Physical reads per second: \_\_\_\_\_

Logical reads per second: \_\_\_\_\_

With a buffer cache hit ratio of near 99%, the physical reads should be about 1% of the logical reads. Is this true? \_\_\_\_\_ (No, the physical reads are much higher.)

There are physical reads that are not counted in the buffer cache hit ratio because the blocks do not go through the buffer cache. These are direct reads.

f) Confirm that direct reads are a large part of the physical reads by checking the Instance Activity Statistics section for physical reads (Main Menu > Instance Activity Statistics). Of the physical reads in the following example, 99% are direct reads. Approximately a half of the direct reads, shown as “physical reads direct,” are to temporary tablespaces, shown as “physical reads direct temporary tablespace.” Physical reads on a temporary tablespace usually indicate a sorting operation and will have nearly matching physical writes. It appears that there may be a problem with temporary space usage, which usually means sort and join operations.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

### Instance Activity Statistics

|                                            |               |              |               |
|--------------------------------------------|---------------|--------------|---------------|
| physical read IO requests                  | 108,851       | 357.20       | 1,943.77      |
| physical read bytes                        | 1,753,186,304 | 5,753,131.58 | 31,306,898.29 |
| physical read total IO requests            | 110,959       | 364.12       | 1,981.41      |
| physical read total bytes                  | 1,787,584,512 | 5,866,010.29 | 31,921,152.00 |
| physical read total multi block requests   | 0             | 0.00         | 0.00          |
| physical reads                             | 214,012       | 702.29       | 3,821.64      |
| physical reads cache                       | 3,186         | 10.45        | 56.89         |
| physical reads cache prefetch              | 425           | 1.39         | 7.59          |
| physical reads direct                      | 210,826       | 691.83       | 3,764.75      |
| physical reads direct temporary tablespace | 103,804       | 340.64       | 1,853.64      |
| physical reads prefetch warmup             | 0             | 0.00         | 0.00          |

- g) Look at the SQL text of the top 5 SQL ordered by Elapsed Time (Back to Top > SQL Statistics > SQL ordered by Elapsed Time). All these statements have a group by or join, or both. These operations force sorts and additional reads if the PGA is not properly sized.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

| SQL ID        | Text                                                                                                                                                      |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| cj4bx12rruv4p | select s.prod_id, SUM(s.amount_sold) -<br>SUM(c.unit_cost) from costs c, sales s where<br>c.prod_id = s.prod_id and s.prod_id = 117<br>group by s.prod_id |
| 589mxwldz9mhn | select s.prod_id, sum(s.amount_sold) -<br>sum(c.unit_cost) from costs c, sales s where<br>c.prod_id = s.prod_id and s.prod_id = 114<br>group by s.prod_id |
| 8vjh4nbvlphu3 | select sum(s.AMOUNT_sold) - sum(c.unit_cost)<br>from costs c, sales s where c.prod_id =<br>s.prod_id and s.prod_id = 139                                  |
| 2rmw64nrra33j | select s.prod_id, sum(s.amount_sold) -<br>sum(c.unit_cost) from costs c, sales s where<br>c.prod_id = s.prod_id and s.prod_id = 144<br>group by s.prod_id |
| 6rz119f0rlwqy | select s.prod_id, sum(s.amount_sold) -<br>sum(c.unit_cost) from costs c, sales s where<br>c.prod_id = s.prod_id and s.prod_id = 142<br>group by s.prod_id |

- h) Determine the extra writes that are caused by the current PGA sizing. Check PGA Aggr Summary in the Advisory Statistics section (Back to Top > Advisory Statistics > PGA Aggr Summary).

What is the PGA Cache Hit ratio? \_\_\_\_\_ (25 in the example; your numbers may vary)

What is the Extra W/A (work area) MB Read/Written? \_\_\_\_\_ (3,600 MB in the example). The example indicates that 3,600 MB were read and written because the work area could not be completed in memory.

### PGA Aggr Summary

| PGA Cache Hit % | W/A MB Processed | Extra W/A MB Read/Written |
|-----------------|------------------|---------------------------|
| 25.46           | 1,242            | 3,634                     |

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

- i) Determine the recommended PGA sizing from the AWR report. In the Advisories section, find the PGA Memory Advisory (Back to Top > Advisory Statistics > PGA Memory Advisory). To find the minimum size with no multipass work areas, choose the smallest value in which Estd PGA Overalloc Count is 0. Using the example, the size is 128 MB. The estimated extra work area read from and written to disk does not drop significantly until the PGA is 128 MB. Your results may vary.

### PGA Memory Advisory

| PGA Target Est (MB) | Size Factr | W/A MB Processed | Estd Extra W/A MB Read/ Written to Disk | Estd PGA Cache Hit % | Estd PGA Overalloc Count | Estd Time  |
|---------------------|------------|------------------|-----------------------------------------|----------------------|--------------------------|------------|
| 10                  | 0.13       | 1,832.20         | 3,679.95                                | 33.00                | 81                       | 29,829,284 |
| 20                  | 0.25       | 1,832.20         | 3,679.95                                | 33.00                | 81                       | 29,829,284 |
| 40                  | 0.50       | 1,832.20         | 3,679.95                                | 33.00                | 81                       | 29,829,284 |
| 60                  | 0.75       | 1,832.20         | 3,679.95                                | 33.00                | 81                       | 29,829,284 |
| 80                  | 1.00       | 1,832.20         | 2,692.35                                | 40.00                | 80                       | 24,484,838 |
| 96                  | 1.20       | 1,832.20         | 2,689.44                                | 41.00                | 53                       | 24,469,095 |
| 112                 | 1.40       | 1,832.20         | 2,439.99                                | 43.00                | 13                       | 23,119,177 |
| 128                 | 1.60       | 1,832.20         | 1,148.66                                | 61.00                | 0                        | 16,131,062 |
| 144                 | 1.80       | 1,832.20         | 744.64                                  | 71.00                | 0                        | 13,944,711 |
| 160                 | 2.00       | 1,832.20         | 452.94                                  | 80.00                | 0                        | 12,366,113 |
| 240                 | 3.00       | 1,832.20         | 0.00                                    | 100.00               | 0                        | 9,915,029  |
| 320                 | 4.00       | 1,832.20         | 0.00                                    | 100.00               | 0                        | 9,915,029  |
| 480                 | 6.00       | 1,832.20         | 0.00                                    | 100.00               | 0                        | 9,915,029  |
| 640                 | 8.00       | 1,832.20         | 0.00                                    | 100.00               | 0                        | 9,915,029  |

- j) Observe the number of multipass and one-pass work areas that are being used. Find PGA Aggr Target Histogram in the Advisories section of the AWR report (Back to Top > Advisory Statistics > PGA Aggr Target Histogram). The work areas that cannot be contained in memory write to disk (temporary tablespace) and are recorded as one-pass or multipass work areas. In PGA Aggr Target Histogram, the column on the right shows the minimum requested size for a one-pass execution. The example shows a number of multipass executions and one-pass executions for sorts larger than 2 MB.

| Low Optimal | High Optimal | Total Execs | Optimal Execs | 1-Pass Execs | M-Pass Execs |
|-------------|--------------|-------------|---------------|--------------|--------------|
| 2K          | 4K           | 1,082       | 1,082         | 0            | 0            |
| 64K         | 128K         | 4           | 4             | 0            | 0            |
| 256K        | 512K         | 3           | 3             | 0            | 0            |
| 512K        | 1024K        | 330         | 330           | 0            | 0            |
| 1M          | 2M           | 2           | 2             | 0            | 0            |
| 2M          | 4M           | 38          | 0             | 8            | 30           |
| 4M          | 8M           | 62          | 0             | 4            | 58           |
| 8M          | 16M          | 56          | 2             | 0            | 54           |

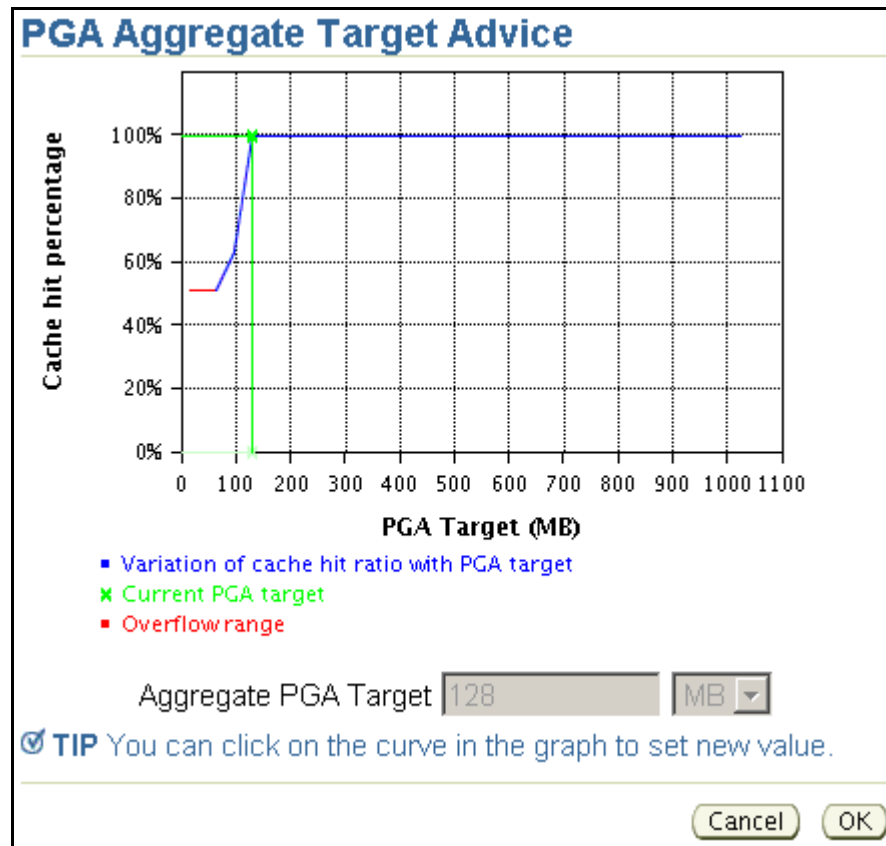
9. Make a new setting for PGA\_AGGREGATE\_TARGET on the basis of the value that you determined from the PGA Memory Advisory. The goal is to eliminate multipass executions. In a terminal window, connect as sysdba and change PGA\_AGGREGATE\_TARGET to 128 MB.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

```
$ sqlplus / as sysdba
SQL> ALTER SYSTEM SET PGA_AGGREGATE_TARGET = 128M;
SQL> exit;
```

10. Execute the workload with the . /workgen PGA script.
11. After a few minutes, view the Memory Advisors page, PGA tab (Server tab > Memory Advisors > PGA tab). Check the PGA Advisor and the PGA Memory Usage Details pages.
  - a) View the PGA Advice page. Have the recommended values changed?

Yes, the advice graph has changed; the following example indicates that a value of approximately 130 MB is optimal.



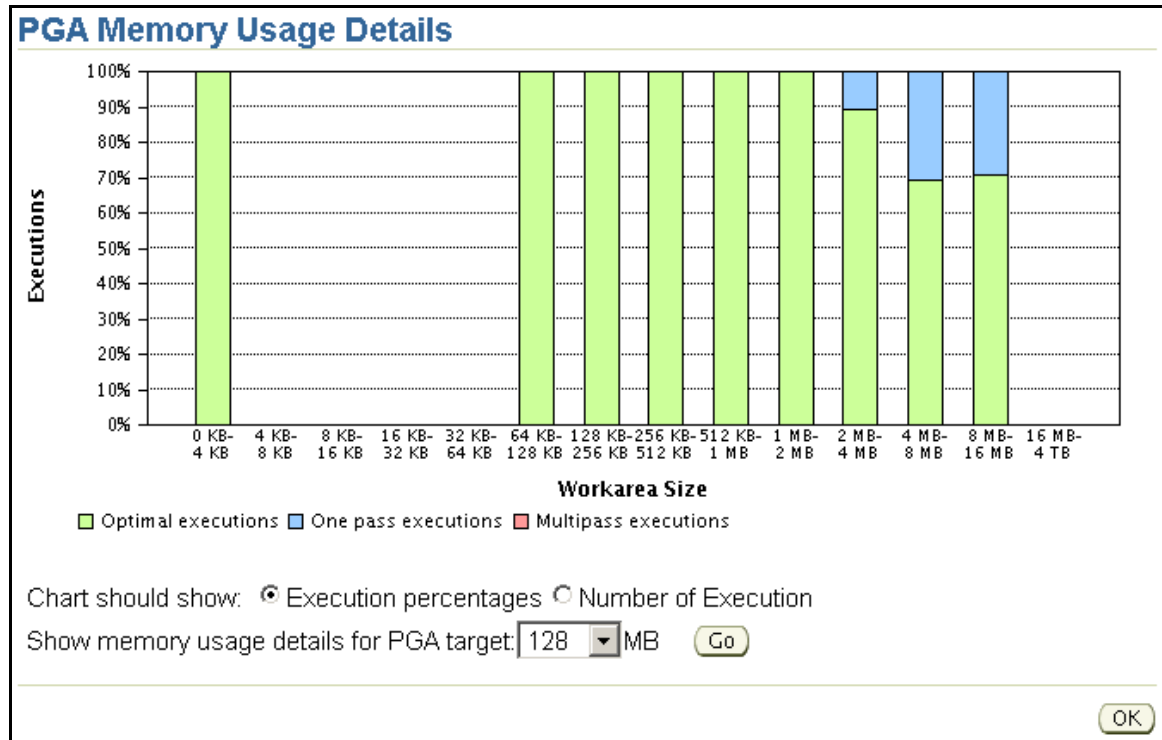
- b) View the PGA Memory Usage Details page.

When the PGA Memory Usage Details page appears, it starts with the lowest possible PGA Target value. Set the value to be displayed to your current value of PGA\_AGGREGATE\_TARGET.

Are there any multipass work areas?

**Answer:** No. The PGA Memory Usage Details histogram does not show any multipass work areas at 128 MB.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)



12. After the workload has been running for about five minutes, stop the workload with `rm runload`.

```
$ rm runload
```

13. Create an AWR snapshot and an AWR report for the last two snapshots. Use the same procedure as was used in step 6.
14. Review the AWR report for symptoms, recommendations concerning PGA area sizing, and the changes that show improved performance.
  - a) Save this report for later viewing: Click “Save to File.” Name the file `awr_PGA_2.html`.
  - b) Check Load Profile for physical reads and writes per second. The physical reads and writes per second have dropped significantly in the example compared to the previous report `awr_PGA_1.html`.
  - c) Have the SQL statements that have the highest elapsed time changed? Find the SQL statements identified in step 6c that have the greatest elapsed time or compare the two reports side by side. The order may have changed, but the same SQL IDs are in the Top SQL. The SQL ID is the hash value of the SQL statement and does not change between runs. Use the SQL ID to find the statements. Notice that the reads per execution have dropped significantly. Users would notice that the statements are responding more quickly based on this example.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

SQL ordered by Elapsed Time

| Elapsed Time (s) | Executions | Elapsed Time per Exec (s) | %Total | %CPU | %IO   | SQL Id                        | SQL Module | SQL Text                             |
|------------------|------------|---------------------------|--------|------|-------|-------------------------------|------------|--------------------------------------|
| 488.05           | 2          | 244.02                    | 13.29  | 3.52 | 0.00  | <a href="#">cj4bx12rruv4p</a> | SQL*Plus   | select s.prod_id,<br>SUM(s.amount... |
| 382.26           | 13         | 29.40                     | 10.41  | 3.84 | 0.00  | <a href="#">2rmw64nrra33j</a> | SQL*Plus   | select s.prod_id,<br>sum(s.amount... |
| 352.73           | 18         | 19.60                     | 9.61   | 3.89 | 2.98  | <a href="#">54dxhddg0mfdt</a> | SQL*Plus   | select * from (select * from d...    |
| 349.68           | 8          | 43.71                     | 9.52   | 3.32 | 0.00  | <a href="#">6vjh4nbv1phu3</a> | SQL*Plus   | select sum(s.AMOUNT_sold) -<br>su... |
| 348.84           | 7          | 49.85                     | 9.50   | 3.25 | 0.00  | <a href="#">6rz119f0r1wgy</a> | SQL*Plus   | select s.prod_id,<br>sum(s.amount... |
| 327.37           | 13         | 25.18                     | 8.92   | 3.45 | 30.75 | <a href="#">fkmrxw6fxqar</a>  | SQL*Plus   | select * from (select * from d...    |
| 318.64           | 16         | 19.92                     | 8.68   | 4.11 | 7.57  | <a href="#">cssuux75dw74c</a> | SQL*Plus   | select * from (select * from d...    |
| 301.11           | 12         | 25.09                     | 8.20   | 4.14 | 21.73 | <a href="#">3bq3pfux13pkr</a> | SQL*Plus   | select * from (select * from d...    |
| 276.71           | 14         | 19.77                     | 7.54   | 4.67 | 20.49 | <a href="#">bxwvcp4pjr43n</a> | SQL*Plus   | select * from (select * from d...    |
| 102.54           | 2          | 51.27                     | 2.79   | 6.07 | 0.00  | <a href="#">589mxw1dz9mhn</a> | SQL*Plus   | select s.prod_id,<br>sum(s.amount... |
| 62.99            | 7          | 9.00                      | 1.72   | 4.08 | 0.02  | <a href="#">3g36d1uum3662</a> | SQL*Plus   | select PROD_ID,<br>SUM(amount_sol... |
| 43.50            | 8          | 5.44                      | 1.18   | 3.12 | 3.65  | <a href="#">8rtsdaxm7apq7</a> | SQL*Plus   | select SUM(amount_sold),<br>PROD_... |

- d) Determine the improvement in the PGA cache hit percentage. How could the extra write statistic be deceiving? PGA Cache Hit % has increased, showing that more sorting was done in memory. Extra Work Area MB Read/Written could show an increase in I/O from the previous report because the statistic is a total over the report interval (especially if the durations of the reports are different).

### PGA Aggr Summary

| PGA Cache Hit % | W/A MB Processed | Extra W/A MB Read/Written |
|-----------------|------------------|---------------------------|
| 63.25           | 1,141            | 663                       |

- e) Determine the recommended PGA sizing from the AWR report PGA Memory Advisory. Following the process of finding the first value with the Estd PGA Overalloc Amount equal to 0, the recommended value should still be the same as in Step 8i.

## Practice 17-1: Tuning PGA\_AGGREGATE\_TARGET (continued)

### PGA Memory Advisory

| PGA Target Est (MB) | Size Factr | W/A MB Processed | Estd Extra W/A MB Read/ Written to Disk | Estd PGA Cache Hit % | Estd PGA Overalloc Count | Estd Time  |
|---------------------|------------|------------------|-----------------------------------------|----------------------|--------------------------|------------|
| 16                  | 0.13       | 1,148.33         | 2,265.97                                | 34.00                | 63                       | 20,715,288 |
| 32                  | 0.25       | 1,148.33         | 2,265.97                                | 34.00                | 63                       | 20,715,288 |
| 64                  | 0.50       | 1,148.33         | 2,265.97                                | 34.00                | 63                       | 20,715,288 |
| 96                  | 0.75       | 1,148.33         | 2,193.14                                | 34.00                | 46                       | 20,273,418 |
| 128                 | 1.00       | 1,148.33         | 331.44                                  | 78.00                | 0                        | 8,978,093  |
| 154                 | 1.20       | 1,148.33         | 331.44                                  | 78.00                | 0                        | 8,978,093  |
| 179                 | 1.40       | 1,148.33         | 238.44                                  | 83.00                | 0                        | 8,413,807  |
| 205                 | 1.60       | 1,148.33         | 79.08                                   | 94.00                | 0                        | 7,446,916  |
| 230                 | 1.80       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |
| 256                 | 2.00       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |
| 384                 | 3.00       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |
| 512                 | 4.00       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |
| 768                 | 6.00       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |
| 1,024               | 8.00       | 1,148.33         | 0.00                                    | 100.00               | 0                        | 6,967,150  |

f) Were there any multipass work areas in the reported period?

No multipass work areas are reported. The number of one-pass work areas cannot be directly compared due to the differences in time periods of the reports and the multipass sorts that are now one-pass. There is the additional complication that because the system is more efficient, more statements—and thus more sorts—can be performed in the same period.



## Practices for Lesson 18

The goal of this practice is to use the Automatic Memory Tuning capability of Oracle Database 11g.

## Practice 18-1: Enabling Automatic Shared Memory

In this practice, you enable Automatic Shared Memory Tuning. The initial configuration is simulating a memory-constrained system. The buffer cache is 16 MB, the keep pool is 4 MB, and the shared pool is 132 MB. The shared pool is loaded with kept objects as in a previous practice. For this practice, `SGA_MAX_SIZE` is set to 400 MB and `MEMORY_MAX_TARGET` is set to 800 MB so that you can easily determine the SGA optimal size. For this practice, you are using the Automatic Shared Memory Management feature of Oracle Database.

1. Prepare the instance for this practice.
  - a) Log out of Enterprise Manager.
  - b) Run the `./prepare ASMM` script.

```
$./prepare ASMM
Database closed.
Database dismounted.
ORACLE instance shut down.

ORACLE instance started.

Total System Global Area 313860096 bytes
Fixed Size 1336232 bytes
Variable Size 289410136 bytes
Database Buffers 16777216 bytes
Redo Buffers 6336512 bytes
Database mounted.
Database opened.
Finished setup ASMM
$
```

- c) Enable the Automatic Shared Memory Management feature in Enterprise Manager. Assume that SGA memory is constrained to 400 MB.

| Step | Page                                                                      | Action                                               |
|------|---------------------------------------------------------------------------|------------------------------------------------------|
| a    |                                                                           | Log in to Enterprise Manager.                        |
| b    | Database home                                                             | Click the Server tab.                                |
| c    | Server subpage                                                            | Click Memory Advisors.                               |
| d    | Memory Advisors<br>(See the example screenshot, which follows the table.) | Verify that Maximum SGA Size (MB) is set to 400 MB.  |
|      |                                                                           | Click Enable for Automatic Shared Memory Management. |

## Practice 18-1: Enabling Automatic Shared Memory (continued)

|   |                                           |                                                                                                                                                                                                                                                              |
|---|-------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| e | Enable Automatic Shared Memory Management | The memory size displayed is the sum of the current pool settings. Change this value to 168 MB. (The value must be evenly divisible by 4.) Note that smaller values of Total SGA Size (MB) will cause the Total SGA Size to default to the Maximum SGA Size. |
|   |                                           | Click OK.                                                                                                                                                                                                                                                    |
| f | Memory Advisors                           | A confirmation message is shown.                                                                                                                                                                                                                             |

d) The following screenshot shows the Memory Advisors page:

**Memory Advisors**

Page Refreshed May 2, 2010 7:31:57 PM UTC [Refresh](#) [Show SQL](#) [Revert](#) [Apply](#)

When Automatic Memory Management is enabled, the database will automatically set the optimal distribution of memory. The distribution of memory will change from time to time to accommodate changes in the workload.

Automatic Memory Management **Disabled** [Enable](#)

**SGA** [PGA](#)

The System Global Area (SGA) is a group of shared memory structures that contains data and control information for one Oracle database. The SGA is allocated in memory when an Oracle database instance is started.

Automatic Shared Memory Management **Disabled** [Enable](#)

Shared Pool  MB [Advice](#)

Buffer Cache  MB [Advice](#)

Large Pool  MB [Advice](#)

Java Pool  MB [Advice](#)

Other (MB)

Total SGA (MB) **159** [Calculate](#)

**Maximum SGA Size**

The Maximum SGA Size specifies the maximum memory that the database may allocate. If you specify the Maximum SGA Size, you can later dynamically change SGA component sizes (provided the total SGA size does not exceed the Maximum SGA Size).

Maximum SGA Size (MB)

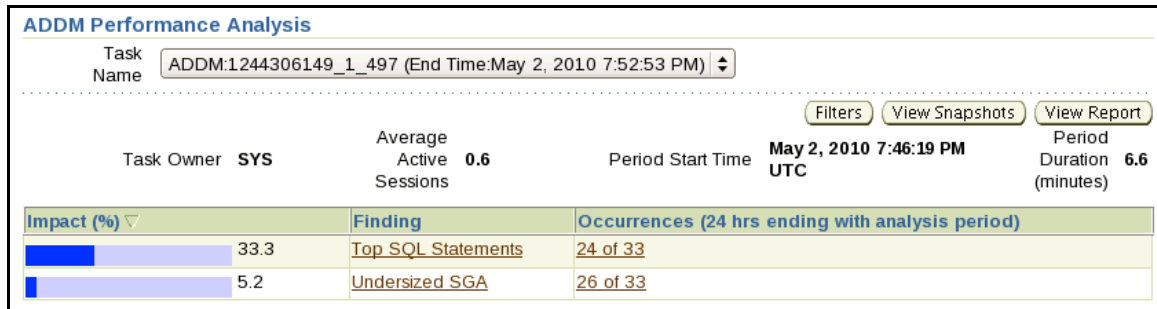
The database must be restarted before any changes to this value take effect.

2. Run the workload generator for Practice 16-1 again: `./workgen BC`.

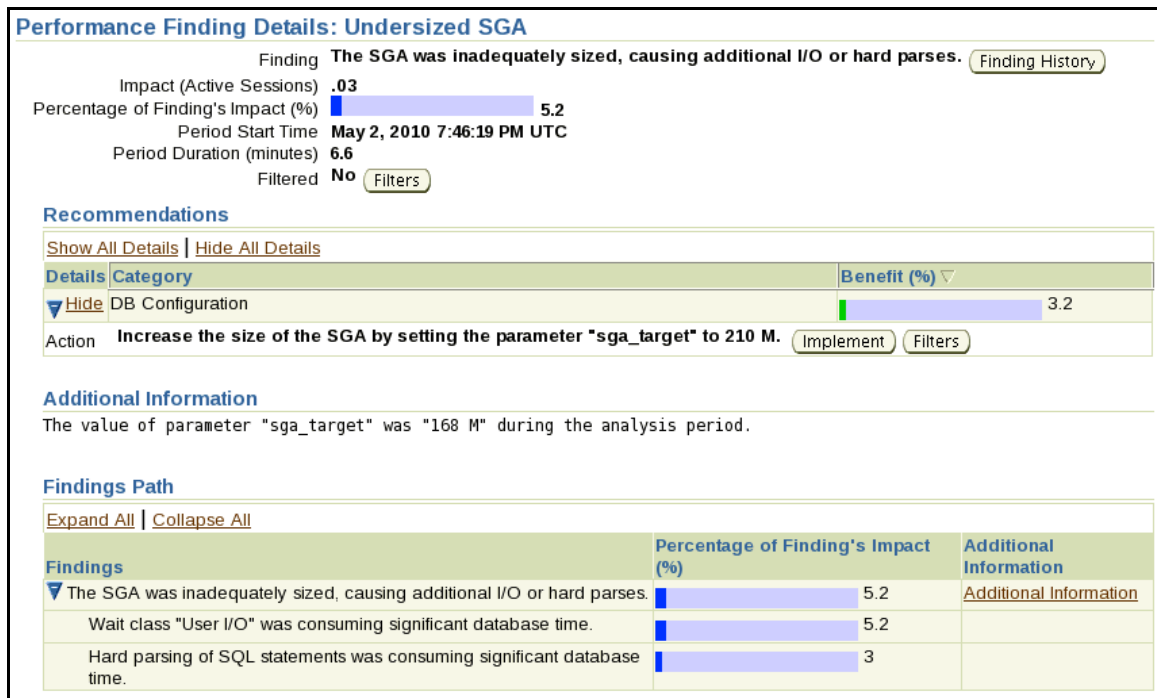
```
$./workgen BC
```

- Log in to EM and navigate to the Performance tabbed page. After about five minutes, run an ADDM report. For any findings related to memory sizing, drill down to the details.
  - View the ADDM report. There may be findings that are related to SQL statements and CPU usage, such as "CPU Usage" or "Top SQL by DB Time." There should be an "Undersized SGA" finding.

## Practice 18-1: Enabling Automatic Shared Memory (continued)



- b) View the finding details for “Undersized SGA” and view the recommendation. What is the recommended size for the SGA? \_\_\_\_\_  
(The recommendation will vary.)



- c) Return to the ADDM page.

**Practice 18-1: Enabling Automatic Shared Memory (continued)**

d) On the ADDM page, click View Snapshots, and then click the Report tab.

In the Report Summary section, examine the cache sizes.

What are the start and end sizes for the buffer cache and shared pool?

| Cache            | Begin | End |
|------------------|-------|-----|
| Buffer Cache     |       |     |
| Shared Pool Size |       |     |

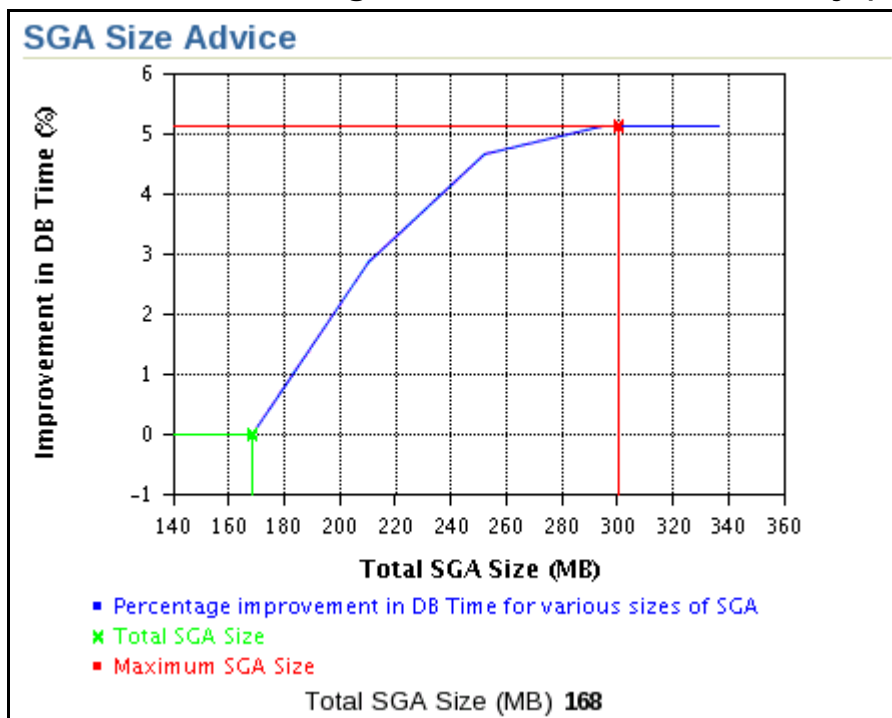
The values did not change because the sizes are minimum sizes specified by the initialization parameters for these pools.

4. View the Memory Advisors page (Server tab > Memory Advisors). What are the current memory parameter settings?
  - a) The memory parameter settings may vary depending on the machine that you are using.
    - Buffer Cache \_\_\_\_\_ (16 MB)
    - Shared Pool \_\_\_\_\_ (120 MB)
    - Large Pool \_\_\_\_\_ (4 MB)
    - Java Pool \_\_\_\_\_ (12 MB)
    - Other \_\_\_\_\_ (16 MB)

**Note:** The EM Memory Advisors page does not show the value of the keep pool separately. It is included in Other.
5. Check Total SGA Size Advisory for an optimal SGA size. Click Total SGA Size (MB) Advice on the Memory Advisors page. The advisory shows that increasing the SGA size gives some improvement in performance.

What is the recommended size? \_\_\_\_\_

In the following example, the greatest benefit for the memory added will be seen at approximately 210 MB, the first inflection point on the graph, but additional benefit will be gained until approximately 250 MB, the second inflection point, when there is little or no more benefit for adding memory.

**Practice 18-1: Enabling Automatic Shared Memory (continued)**

6. Stop the workload with the `rm runload` command.

```
$ rm runload
```

7. On the Memory Advisors page, change Total SGA Size to 260 MB. Click Apply.
8. Observe the initial settings for the SGA on the Memory Advisors page. (Click Refresh to view the new distribution of memory.)

What is the value of the buffer cache? \_\_\_\_\_

What is the value of the shared pool size? \_\_\_\_\_

The values for the shared pool and buffer cache may vary based on the machine you are using.

9. Start the workload again: `./workgen BC`.

```
$./workgen BC
```

10. On the Performance page, what are the differences between this and the previous run?

The Average Active Sessions graph and the Load average graphs may show small differences. The I/O graph shows a large difference in the number of buffer cache reads.

11. Observe the current settings for the SGA on the Memory Advisors page. (Click Refresh to view any changes in the distribution of memory.)

What is the value of the buffer cache? \_\_\_\_\_

What is the value of the shared pool size? \_\_\_\_\_

## Practice 18-1: Enabling Automatic Shared Memory (continued)

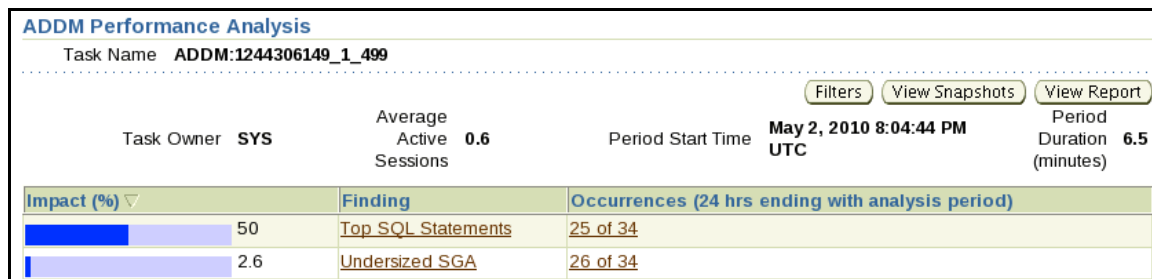
The values for the shared pool and buffer cache may vary based on the machine you are using. It is expected that the values will stay the same as the initial change seen in step 8.

12. After about five minutes, stop the workload with `rm runload`.

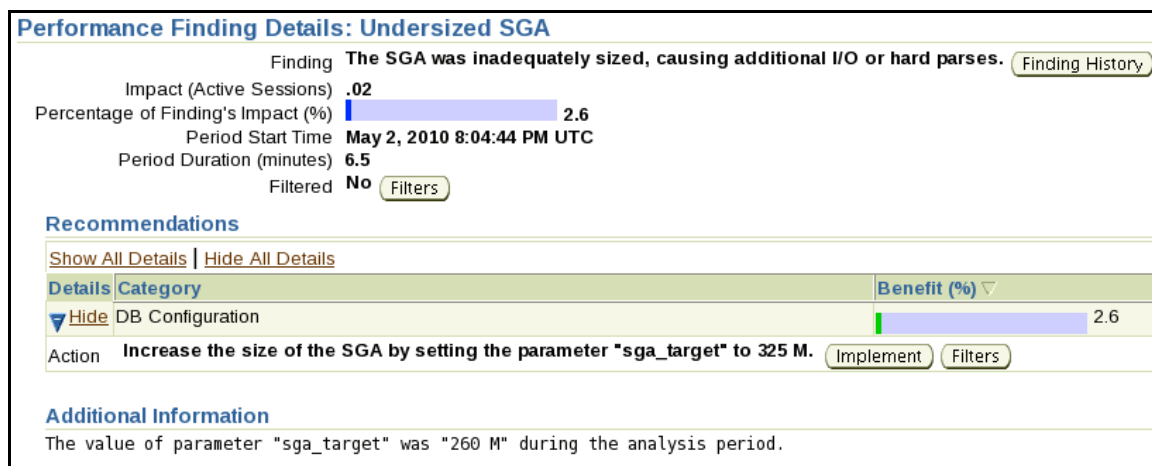
```
$ rm runload
```

13. Click Run ADDM Now on the Performance tabbed page to run an ADDM report. View the results and drill down to the detail pages.

- a) The number of findings and the impact of the issues may be misleading. It is common that correcting memory issues has a much greater impact than is estimated. It is possible that there are more findings and Impact % may be larger. Notice that the number of Average Active Sessions is reduced even though the load was identical. This indicates that the instance is running more efficiently.



- b) The finding “Undersized SGA” may still be reported. In the example, the recommendation is to increase the SGA to 325 MB. The impact of this issue as shown in the screenshot is only 2.6%. The recommended size may range to a much higher value.



- c) A careful comparison of the AWR reports from this and the previous runs will show an improvement. In the load profile, you may see a small improvement in DB Time per second, and a significant reduction in physical reads.

14. What would be a reasonable next tuning step based on the ADDM report?

## ***Practice 18-1: Enabling Automatic Shared Memory (continued)***

**Note:** The number and ranking of the issues shown in the ADDM reports of various machines will vary.

- a) **Hard Parse:** This is likely to be caused by an undersized shared pool. This is confirmed by the Undersized SGA finding in the preceding screenshot. Increase the SGA size if allowed by memory constraints.
- b) **CPU Usage:** The hard-parse activity uses CPU. Implementing a solution for the hard-parse finding will impact CPU usage.
- c) **Soft Parse:** This is primarily an application issue. Check the `SESSION_CURSOR_CACHE` value and set it to a higher value if it is already set.
- d) The sizing of the SGA has an impact on a variety of issues, which is usually larger than the impact rating assigned to it. Therefore, it is reasonable to assume that if the SGA is increased again, there will be further improvements in performance.



## Practice 18-2: Adjusting Memory as Workloads Change

On the basis of the recommendation of the ADDM report in the previous scenario, the SGA memory constraint was raised to 260 MB. You have been using an allocation of 200 MB of memory for the PGA work areas. The current application has seen an increase in activity, and now another application is being added to the database. With it comes an additional allocation of instance memory. Make use of Automatic Memory Management to allocate this additional memory for best performance.

1. You are given an absolute maximum of 800 MB of instance memory and are instructed to use as little as required to get reasonable performance. Set the maximum memory size to 800 MB. Enable Automatic Memory Management. Set the Memory target to 500 MB. Note by setting the value larger than needed during testing, you can find the optimal value, and then lower the setting. To set the parameters to enable AMM, you disable ASMM and remove the SGA\_TARGET value.
  - a) On the Memory Advisors page in the Current Allocation section, disable Automatic Shared Memory Management. Click Disable.
  - b) On the Disable Automatic Shared Memory Management page, change the Buffer Cache to 16 MB, so that the value will not impose a minimum.
  - c) Click Calculate; the Total SGA value should change.
  - d) Click OK.
  - e) On the Memory Advisors page, a Confirmation message saying that ASMM has been disabled appears.
2. On the Memory Advisors page, enable Automatic Memory Management. Click Enable.

The screenshot shows the 'Memory Advisors' page. At the top, it says 'Page Refreshed March 12, 2010 9:40:27 AM UTC' with buttons for 'Refresh', 'Show SQL', 'Revert', and 'Apply'. Below this, a message states: 'When Automatic Memory Management is enabled, the database will automatically set the optimal distribution of memory. The distribution of memory will change from time to time to accommodate changes in the workload.' A red box highlights the 'Automatic Memory Management Disabled' status with an 'Enable' button next to it. At the bottom, there are tabs for 'SGA' and 'PGA', with 'PGA' currently selected.

- a) On the Enable Automatic Memory Management page, set the maximum memory size to 800 MB and the Total Memory Size for Automatic Memory Management to 500 MB.

## Practice 18-2: Adjusting Memory as Workloads Change (continued)

**Enable Automatic Memory Management**

When Automatic Memory Management is enabled, the database will automatically set the optimal distribution of memory. The distribution of memory will change from time to time to accommodate changes in the workload. The Maximum Memory Size specifies the maximum memory that the database may allocate and must be set in order to use Automatic Memory Management.

Current Memory Usage (PGA+SGA) (MB) 460

Maximum Memory Size 800 MB

Changing the maximum memory size requires a restart of the database.

Total Memory Size for Automatic Memory Management 500 MB

CancelOK

- b) Navigate to the Memory Advisors page. (The allocation of the buffer cache and the shared pool memory will vary.)

The Total SGA Size is \_\_\_\_\_ MB

The buffer cache is \_\_\_\_\_ MB.

The shared pool is \_\_\_\_\_ MB.

- c) Click the PGA tab. What is the value of the current PGA allocated memory?  
\_\_\_\_\_ (Current PGA Allocated (KB))

3. Start the workload generator. The new workload is simulated by:  
./workgen AMM

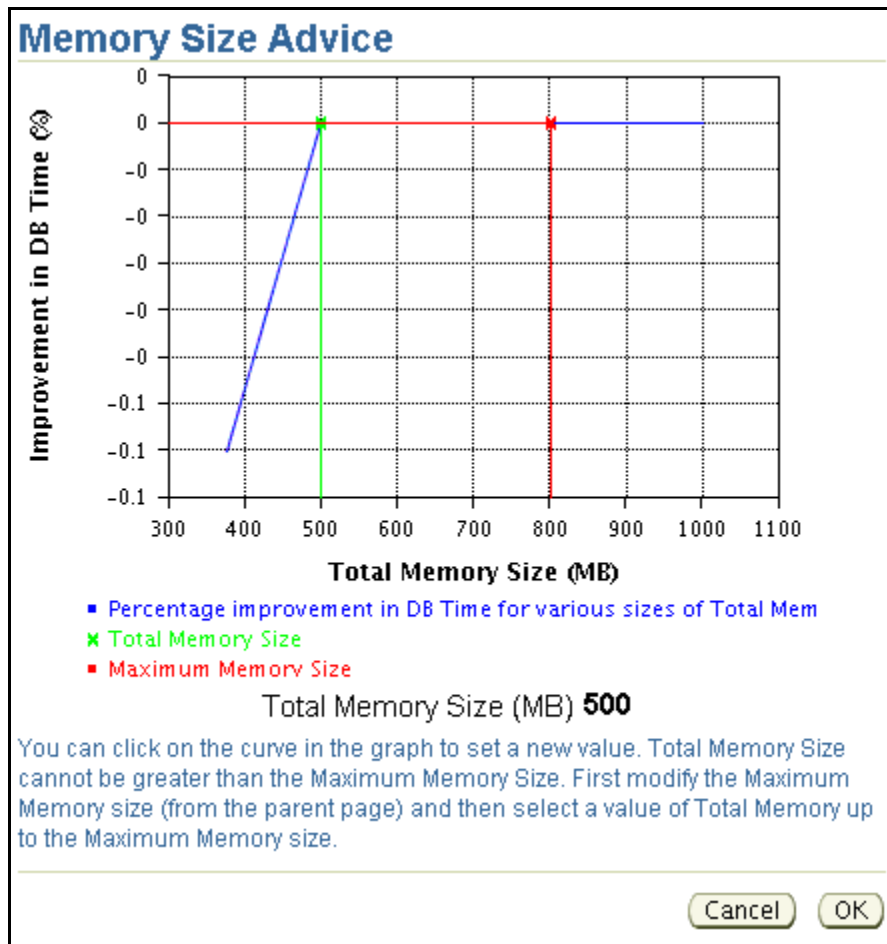
```
$./workgen AMM
```

4. Observe the Memory advisor graphs for about three minutes. Be sure to click Refresh.
- a) Did the memory settings change? \_\_\_\_\_
- Note:** These values change during the course of the workload.
- b) What is the value of the buffer cache? \_\_\_\_\_
- c) What is the value of the shared pool? \_\_\_\_\_
- d) What is the value of the PGA? \_\_\_\_\_ (Current PGA Allocated (KB) changes.)
- e) Click Advice. View the Memory Size Advice Graph. What would be a better value of Total Memory? \_\_\_\_\_

In the following example, the graph shows no improvement when memory is increased to more than 500 MB.

**Note:** The graph produced by your machine may vary.

## Practice 18-2: Adjusting Memory as Workloads Change (continued)



- On the Memory Advisors page, change the Total Memory Size to 750 MB and click Apply.
- Observe the Memory advisor graphs for about two or three minutes.
  - Did the memory settings change? \_\_\_\_\_

**Note:** These values can change during the course of the workload.

  - What is the value of the buffer cache? \_\_\_\_\_
  - What is the value of the shared pool? \_\_\_\_\_
  - What is the value of the PGA? \_\_\_\_\_
  - Click Advice. View the Memory Size Advice graph. What would be a better value of Total Memory? \_\_\_\_\_  
At this point, the graph should show that there is little or no advantage to increasing the total memory.
- Stop the workload with `rm runload`.

```
$ rm runload
```

## ***Practice 18-2: Adjusting Memory as Workloads Change (continued)***

8. In this practice, the SGA changed very little. It may not have changed at all on some machines. The PGA, however, was being used by the new workload and the distribution of memory between the PGA and SGA was affected.
9. How much tuning was required to find an optimum memory size for this workload?

The tuning did not require a shutdown after the maximum allowed memory value was set. The Advice graph gave a clear indication of the optimal size, so the number of iterations was reduced.

10. Run the `./cleanup AMM` script.

```
$./cleanup AMM
```

## Practices for Lesson 19

There are no practices for this lesson.

## Practices for Lesson 20

There are no practices for this lesson.

## Practices for Appendix B

This practice is optional and may be used with Appendix B.

## Practice for Appendix B-1: Installing Statspack

The goal of this practice is to install the Statspack repository and packages in the `orcl` database. For this exercise, install the Statspack repository in the `SYSAUX` tablespace. The Statspack repository may be installed in any permanent tablespace. Make sure that you use the `perfstat` password for the `PERFSTAT` user. The `sp*.sql` scripts are located in the `$ORACLE_HOME/rdbms/admin` directory.

1. To prepare the database for this practice, connect with `/ as SYSDBA` to drop the Statspack repository. Use the `spdrop.sql` script located in the `$ORACLE_HOME/rdbms/admin` directory.

```
$ sqlplus / as sysdba

SQL> @?/rdbms/admin/spdrop.sql

SQL>
```

2. Still connected as `SYSDBA` using `SQL*Plus`, execute the `spcreate.sql` script to create the Statspack repository and packages.

The script prompts you for the name and password of the owner of the Statspack object. Make the username `perfstat` with a password of `perfstat`. The default tablespace for the Statspack objects is `SYSAUX`. The default tablespace for the `perfstat` user is `TEMP`.

```
SQL> @?/rdbms/admin/spcreate.sql
```

3. Make sure that you did not have errors while installing Statspack. Check the log files that were generated during installation. The following is a command-line solution:

```
SQL> host ls *.lis
spcpkg.lis spctab.lis spcusr.lis spdtab.lis spdusr.lis

SQL> host grep error *.lis
spcpkg.lis:No errors.
spcpkg.lis:No errors.
spcpkg.lis:SPCPKG complete. Please check spcpkg.lis for any
errors.
spctab.lis:SPCTAB complete. Please check spctab.lis for any
errors.
spcusr.lis:SPCUSR complete. Please check spcusr.lis for any
errors.
spdtab.lis:SPDTAB complete. Please check spdtab.lis for any
errors.
spdusr.lis:SPDUSR complete. Please check spdusr.lis for any
errors.

SQL> exit
```



## Practice for Appendix B-2: Creating Snapshots

The goal of this practice is to run a workload on the `orcl` database and to capture Statspack snapshots for that workload.

1. Before you capture your first Statspack snapshot, execute the `lab_21_02_01.sh` script located in your `$HOME/labs` directory. This script creates a tablespace and a user as the setup for the rest of the practice.

```
$ cd $HOME/labs

$./lab_21_02_01.sh
...
SQL> Connected.
SQL> SQL> SQL> SQL> drop tablespace tbsspc including contents
and datafiles
*
ERROR at line 1:
ORA-00959: tablespace 'TBSSPC' does not exist

SQL> SQL> 2 3 4 5
Tablespace created.

SQL> SQL>
PL/SQL procedure successfully completed.

SQL> SQL> drop user spc cascade
*
ERROR at line 1:
ORA-01918: user 'SPC' does not exist

SQL> SQL> 2 3
User created.

SQL> SQL>
Grant succeeded.

SQL> SQL> Connected.
SQL> SQL> drop table spct purge
*
ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
Table created.

SQL> SQL> > >
PL/SQL procedure successfully completed.

SQL> SQL> Disconnected ...
```

## Practice for Appendix B-2: Creating Snapshots (continued)

```
$
```

2. Create a level 7 Statspack snapshot. Make sure that you retrieve the corresponding snapshot ID. Set the `i_snap_level` parameter of the `statspack.snap` procedure to 7. This script is `$HOME/solutions/sol_21_02_02.sh`.

```
$ sqlplus /nolog
SQL> CONNECT perfstat/perfstat

SQL> variable snap number;
SQL> begin
 2 :snap := statspack.snap(i_snap_level=>7);
 3 end;
 4 /

PL/SQL procedure successfully completed.

SQL> print snap

 SNAP

 1

SQL> exit
```

3. Execute the `start_21_02_03.sh` script from the `labs` directory. This script starts the workload on your `orcl` database. The script takes about one minute to execute.

```
$ cd $HOME/labs
$./start_21_02_03.sh
PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
...
PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.
$
```

## ***Practice for Appendix B-2: Creating Snapshots (continued)***

4. After the workload finishes running, capture a new level 7 Statspack snapshot. Make sure that you retrieve the corresponding snapshot ID. The number is system-assigned and is not controllable. The snap numbers are not necessarily consecutive.

```
$ cd $HOME/labs
$ sqlplus perfstat/perfstat
...

SQL> variable snap number;
SQL> begin
 2 :snap := statspack.snap(i_snap_level=>7);
 3 end;
 4 /

PL/SQL procedure successfully completed.

SQL> print snap
 SNAP

 2

SQL> exit
...
$
```

## Practice for Appendix B-3: Generating Statspack Reports

The goal of this practice is to generate a Statspack report, examine it, and fix any issues that you discover while interpreting the Statspack report.

1. Generate a Statspack report between the two previously captured snapshots. Use the `spreport.sql` script. Be sure to use the snap numbers that you found in the previous practice. When prompted for a file name, use `sp_1_2.lst`.

```
$ cd $HOME/labs
$ sqlplus perfstat/perfstat
...
SQL> set echo on

SQL> @?/rdbms/admin/spreport
...
Listing all Completed Snapshots
```

| Instance<br>Comment | DB Name | Snap Id | Snap Started      | Snap<br>Level |
|---------------------|---------|---------|-------------------|---------------|
| orcl                | ORCL    | 1       | 16 Mar 2010 09:54 | 7             |
|                     |         | 2       | 16 Mar 2010 09:57 | 7             |

```

Specify the Begin and End Snapshot Ids
~~~~~
Enter value for begin_snap: 1
Begin Snapshot Id specified: 1

Enter value for end_snap: 2
End Snapshot Id specified: 2

Specify the Report Name
~~~~~
The default report file name is sp_1_2. To use this name,
press <return> to continue, otherwise enter an alternative.

Enter value for report_name: sp_1_2.lst
...
sga_max_size 419430400
sga_target 0
shared_pool_size 0
spfile
/u01/app/oracle/product/11.2.0/dbhome_1/dbs/spfileorcl.ora
streams_pool_size 4194304
undo_tablespace UNDOTBS1
~~~~~

```

## Practice for Appendix B-3: Generating Statspack Reports (continued)

```
End of Report ( sp_1_2.lst )

SQL>exit;
```

2. Examine the report that you created. You can use the `less` utility, which uses `vi` navigation commands, or the `gedit` graphic text viewer. What are your conclusions?
  - a) Invoke a viewer. (`gedit` is shown.) Navigation in `gedit` is similar to other graphic-based editors. To find a section of the report, use Search > Find and the report section name.

```
$ cd $HOME/labs
$ gedit sp_1_2.lst
```

- b) Scroll to the Top 5 Timed Events section. “Buffer busy waits” is among Top 5 Timed Events.
- c) View Time Model System Stats. The top entry is “sql execute elapsed time.”
- d) Looking at the Buffer Wait Statistics section, you can see that most of the waits are due to the data block category.
- e) Looking at the “Segments by Buffer Busy Waits” section, SPCT is shown as the culprit segment in the TBSSPC tablespace.
- f) All this leads you to think that there is an issue with the SPCT table. Looking at the tablespace definition, you discover that the table was defined without Automatic Segment Space Management.

```
$ sqlplus spc/spc
...
SQL> SELECT segment_space_management
2     FROM dba_tablespaces
3     WHERE tablespace_name=(select tablespace_name
4                             from user_tables);

SEGMENT
-----
MANUAL

SQL> exit;
```

3. Fix the problem by implementing the recommendations from the previous step. You decide to re-create the table by using Automatic Segment Space Management. This script is available as `sol_21_03_03.sql` in the solutions directory.

```
$ sqlplus / as sysdba
```

## Practice for Appendix B-3: Generating Statspack Reports (continued)

```
...
SQL> @../solutions/sol_21_03_03
SQL>
SQL> drop tablespace tbsspc including contents and datafiles;

Tablespace dropped.

SQL>
SQL> CREATE SMALLFILE TABLESPACE "TBSSPC"
  2 DATAFILE 'tbsspc1.dbf' SIZE 50M
  3 LOGGING
  4 EXTENT MANAGEMENT LOCAL
  5 SEGMENT SPACE MANAGEMENT AUTO;

Tablespace created.

SQL>
SQL> connect spc/spc
Connected.
SQL>
SQL> create table spct(id number, name varchar2(2000))
tablespace tbsspc;

Table created.

SQL>
SQL> exec DBMS_STATS.GATHER_TABLE_STATS(-
> ownname=>'SPC', tabname=>'SPCT',-
> estimate_percent=>DBMS_STATS.AUTO_SAMPLE_SIZE);

PL/SQL procedure successfully completed.

SQL>
SQL> exit;
...
$
```

4. Regenerate the same workload, and verify that the problem disappears.
  - a) Take a new snapshot like you did in the previous practice. Run the same workload again and take another snapshot. Using the `$HOME/solutions/sol_21_03_04a.sh` script, you can create a snapshot, run the workload, and take a second snapshot script. Note the snap ID values. The script takes two minutes to complete.

```
$ $HOME/solutions/sol_21_03_04a.sh
```

- b) Generate the corresponding Statspack report using the snap ID values generated in the previous step.

## Practice for Appendix B-3: Generating Statspack Reports (continued)

```
$ sqlplus perfstat/perfstat

SQL> @?/rdbms/admin/spreport

...
Listing all Completed Snapshots
```

| Instance<br>Comment | DB Name | Snap Id | Snap Started      | Snap<br>Level |
|---------------------|---------|---------|-------------------|---------------|
| orcl                | ORCL    | 1       | 16 Mar 2010 09:54 | 7             |
|                     |         | 2       | 16 Mar 2010 09:57 | 7             |
|                     |         | 3       | 16 Mar 2010 10:10 | 7             |
|                     |         | 4       | 16 Mar 2010 10:12 | 7             |

```
Specify the Begin and End Snapshot Ids
~~~~~
Enter value for begin_snap: 3
Begin Snapshot Id specified: 3

Enter value for end_snap: 4
End Snapshot Id specified: 4

Specify the Report Name
~~~~~
The default report file name is sp_3_4. To use this name,
press <return> to continue, otherwise enter an alternative.

Enter value for report_name: sp_3_4.lst
```

- c) Look at the generated report by using your preferred text browser. You can see that the number of buffer busy waits has been reduced.





# B

## Using Statspack

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

# Objectives

After completing this lesson, you should be able to do the following:

- Install Statspack
- Create Statspack snapshots
- Generate Statspack reports
- Identify the major sections of the Statspack report

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

# Introduction to Statspack

- Statspack is a set of scripts that capture and report on performance data from within the Oracle database.
- With Statspack scripts, you can:
  - Capture instance data by taking a “snapshot”
  - Store snapshot data in the database in a separate schema
  - Create reports between two snapshots
  - Mark snapshots as baseline information
  - Use the reports as part of a performance tuning method

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Introduction to Statspack

Statspack is a set of scripts and PL/SQL packages. Statspack allows you to capture point-in-time performance data called a snapshot into a schema dedicated to Statspack. You can create a report between snapshots and view the database activity and performance for the time period covered by the snapshots. You can preserve sets of snapshots by marking them as baseline information. Marking the snapshots protects them from deletion by the purge procedure.

The reports generated by Statspack provide performance tuning information that can be used as part of a tuning methodology.

## Statspack Scripts

- Install Statspack by using `spcreate.sql`.
- Collect statistics with `statspack.snap`.
- Automate the collection of statistics with `spauto.sql`.
- Produce a report by using `spreport.sql` or `sprepsql.sql`.
- Purge Statspack data with `statspack.purge` or `sppurge.sql`.
- Truncate all Statspack tables with `sptrunc.sql`.
- Drop the Statspack repository with `spdrop.sql`.
- Export the Statspack repository with `spexp.par`.

**ORACLE**

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Statspack Scripts

### Installation of the Statspack Package

The installation of the Statspack utility creates the PERFSTAT user, which owns all PL/SQL code and database objects created (including the Statspack tables, the constraints, and the Statspack package). During the installation, you are prompted for the PERFSTAT user's default and temporary tablespaces.

The Statspack package has been available with the Oracle database since Oracle 8.1.6. When initially installed, approximately 80 MB of the PERFSTAT user's default tablespace is used. This may increase later when the tables contain snapshot information.

### Collecting Statistics

To take a snapshot of performance data, log on to SQL\*Plus as the PERFSTAT user, and then execute the `STATSPACK.SNAP` procedure. This procedure stores the current performance statistics in the Statspack tables, which can be used with another snapshot taken at a later time to produce a report or a base line.

### Automatically Collecting Statistics

To compare performance from one day, week, or year to the next, there must be multiple snapshots taken over a period of time. The best method to gather snapshots is to automate the collection at regular time intervals.

## Statspack Scripts (continued)

### Automatically Collecting Statistics (continued)

The `spauto.sql` script makes use of the `DBMS_JOB` package to provide an automated method for collecting statistics. The supplied script schedules a snapshot every hour, on the hour. This can be changed to suit the requirements of the system.

### Producing a Report

To examine the change in statistics between two time periods, execute the `spreport.sql` file while connected as the `PERFSTAT` user. You are shown a list of collection periods and can select a start and end period. The difference between the statistics at each end point is then calculated and put into a file with a name of your choice.

### Truncating Statspack Tables

If you want to truncate all performance data indiscriminately, it is possible to do this using `sptrunc.sql`. This script truncates all statistics data gathered, including snapshots marked as baselines.

### Dropping the Statspack Repository

To drop the Statspack Repository and remove the Statspack package, use the `spdrop.sql` script. If there are errors during the repository creation, then you must drop the repository before you attempt to re-create it.

### Exporting the Statspack Repository

If you want to share data with other sites, for example if Oracle Support requires the raw statistics, it is possible to export the `PERFSTAT` user. An export parameter file (`spuexp.par`) has been supplied for this purpose. To use this file, supply the export command with the `userid` parameter, along with the export parameter file name:

```
exp userid=perfstat/perfstat_password parfile=spuexp.par
```

This creates a file called `spuexp.dmp` and the `spuexp.log` log file.

**Note:** It is also possible to upgrade Statspack tables to the latest version of the Oracle database that you are using. Scripts are provided to accomplish this task.

# Installing Statspack

- Install Statspack:
  - Supplying parameters as requested interactively
  - Defining parameters in the session for batch mode

```
SQL> connect / as sysdba
SQL> define default_tablespace='SYSAUX'
SQL> define temporary_tablespace='TEMP'
SQL> define perfstat_password='erg8oiw'
SQL> @?/rdbms/admin/spcreate
SQL> undefine perfstat_password
```

- The `spcreate.sql` script creates the `PERFSTAT` user to own the Statspack repository and package.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Installing Statspack

- Installation scripts create a user called `PERFSTAT`, which owns all PL/SQL code and database objects created including the Statspack tables, constraints, and the `STATSPACK` package.
- There are two ways to install Statspack: interactively or in batch mode as shown in the slide. Batch mode is useful when you do not want to be prompted for the `PERFSTAT` user's password and the default and temporary tablespaces. To install in batch mode, you must assign values to the SQL\*Plus variables that specify the password and the default and temporary tablespaces before running `spcreate.sql`. Both methods require the default tablespace to be defined. You cannot use the `SYSTEM` tablespace as the default tablespace for the `PERFSTAT` user. `SYSAUX` is the recommended default tablespace.
- If you receive an error during the installation, execute `spdrop.sql` before you rerun `spcreate.sql`.

**Note:** To execute the installation script, you must use SQL\*Plus and connect as a user with the `SYSDBA` privilege.

## Capturing Statspack Snapshots

- Capture a snapshot:

```
SQL> connect perfstat/erg8oiw
SQL> execute statspack.snap;
```

- Use the function to determine the snapshot number:

```
SQL> variable snap number;
SQL> begin :snap := statspack.snap; end;/
SQL print snap
```

- Automate snapshot collection:

```
SQL> connect perfstat/erg8oiw
SQL> @spauto
```

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Capturing Statspack Snapshots

- The simplest, interactive way to take a snapshot is to log in to SQL\*Plus as the PERFSTAT user and execute the `statspack.snap` procedure as shown in the slide.
- This stores the current values for the performance statistics in the Statspack tables, and it can be used as a baseline snapshot for comparison with another snapshot taken at a later time.
- If you want to automate the gathering and reporting phases (such as during a benchmark), you may need to know the `snap_id` of the snapshot just taken. To take a snapshot and display the `snap_id`, call the `statspack.snap` function as shown in the slide.
- To use an Oracle-automated method for collecting statistics, you can use DBMS\_JOB. A sample script is supplied in `spauto.sql`, which schedules a snapshot every hour, on the hour. This script should be run as PERFSTAT. The `JOB_QUEUE_PROCESSES` initialization parameter should be set to a number greater than zero before automatic statistics gathering can run. The script also reports the corresponding job number. You can use the DBMS\_JOB package to change the interval of statistics collection, force the job to run immediately, or to remove automatic job collection.

**Note:** On UNIX systems, you could use utilities such as `cron` or `at` to automate the snapshot collection.

## Configuring Snapshot Data Capture

- Temporarily change a capture variable:

```
SQL> execute statspack.snap(i_snap_level=>6);
```

- Change defaults and take a snapshot:

```
SQL> execute statspack.snap -  
      (i_snap_level=>10, i_modify_parameter=>'true');
```

- Modify defaults without taking a snapshot:

```
SQL> execute statspack.modify_statspack_parameter -  
      (i_snap_level=>10, i_buffer_gets_th=>10000, -  
      i_disk_reads_th=>1000);
```

- Include session details in a snapshot:

```
SQL> execute statspack.snap(i_session_id=>3);
```



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Configuring Snapshot Data Capture

- Both the snapshot level and the thresholds specified affect the amount of data Statspack captures. You can control the amount of information gathered by the package by specifying a snapshot level. The higher the snapshot level, the more the data gathered. The default level set by the installation is level 5. Level details are shown in a later slide.
- There are threshold parameters that can be configured when collecting data on SQL statements. Data is captured on any SQL statements that breach the specified thresholds.
- To temporarily use a snapshot level or threshold that is different from default snapshot values, simply specify the required threshold or snapshot level when taking the snapshot. This value is used only for the immediate snapshot taken, as shown in the slide.
- You can save the new value as the instance default, by using `i_modify_parameter`, as shown in the slide. You can also change the defaults without taking a snapshot, using the `STATSPACK.MODIFY_STATSPACK_PARAMETER` procedure.
- To gather session statistics and wait events for a particular session, in addition to the instance statistics and wait events, specify the `session id` in the call to Statspack.



## Statspack Snapshot Levels

| Level => | Statistics captured          |
|----------|------------------------------|
| 0        | General performance          |
| 5        | SQL statements               |
| 6        | SQL plans and SQL plan usage |
| 7        | Segment-level statistics     |
| 10       | Parent and child latches     |

**AWR captures them all by default!**

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Statspack Snapshot Levels

All levels contain general performance statistics including the following: wait statistics, system events, system statistics, rollback segment data, row cache, SGA, background events, session events, lock statistics, buffer pool statistics, latch statistics, resource limit, and enqueue statistics. Statistics are included for each of the following, if enabled: Automatic Undo Management, buffer cache advisory data, automatic PGA memory management, and Cluster DB statistics.

Levels 5 and above gather the performance data on SQL statements that exceed one of six predefined threshold parameters for resource usage. The time required for the snapshot to complete depends on the value of `SHARED_POOL_SIZE` and the number of SQL statements in the shared pool at the time the snapshot is taken. You can change the predefined threshold values. The thresholds are:

- Number of executions of the SQL statement (default 100)
- Number of disk reads performed by the SQL statement (default 1,000)
- Number of parse calls performed by the SQL statement (default 1,000)
- Number of buffer gets performed by the SQL statement (default 10,000)
- Size of sharable memory used by the SQL statement (default 1 MB)
- Version count for the SQL statement (default 20)

## Statspack Snapshot Levels (continued)

Level 6 includes all statistics gathered at the lower levels, and gathers optimizer execution plans and plan usage data for each of the high-resource-usage SQL statements captured. A level 6 snapshot gathers information that is invaluable when determining whether the execution plan used for a SQL statement has changed. Therefore, level 6 snapshots should be used whenever there is the possibility that a plan may change, such as after large data loads, or after gathering new optimizer statistics. The plan for a SQL statement is captured only if the statement is in the shared pool at the time that the snapshot is taken, and exceeds one of the SQL thresholds. You can gather plans for all statements in the shared pool, if you temporarily specify the executions threshold (`i_executions_th`) to be zero (0) for those snapshots.

Level 7 includes all statistics gathered at the lower levels, and gathers the performance data on highly used segments. A level 7 snapshot captures segment-level statistics for segments that exceed one of seven threshold parameters that measure access or contention rates. You can change the default threshold values for:

- Number of logical reads on the segment (default 10,000)
- Number of physical reads on the segment (default 1,000)
- Number of buffer busy waits on the segment (default 100)
- Number of row lock waits on the segment (default 100)
- Number of ITL waits on the segment (default 100)
- Number of global cache Consistent Read blocks served (default 1,000)
- Number of global cache Current blocks served (default 1,000)

Levels  $\geq 10$  include all statistics gathered at the lower levels, and gathers parent and child latch information. Data gathered at this level can sometimes cause the snapshot to take longer to complete, and should only be used when advised by Oracle personnel.

**Note:** Refer to the `spdoc.txt` file for more information about the various levels and thresholds.

## Statspack Baselines and Purging

```
SQL> exec statspack.make_baseline -
      2 (i_begin_snap => 45, i_end_snap   => 50);
```

```
SQL> exec statspack.statspack.make_baseline(2,4,false);
```

```
SQL> exec statspack.make_baseline -
      2 (to_date('31-AUG-2009 09:00','DD-MON-YYYY HH24:MI'), -
      3 to_date('31-AUG-2009 12:00','DD-MON-YYYY HH24:MI'));
```

```
SQL> exec statspack.clear_baseline -
      2 (to_date('13-DEC-2009 23:00','DD-MON-YYYY HH24:MI'), -
      3 to_date('14-FEB-2010 02:00','DD-MON-YYYY HH24:MI'));
```

```
SQL> exec statspack.purge -
      2 (i_begin_date=>to_date('01-JAN-2009','DD-MON-YYYY'), -
      3 i_end_date   =>to_date('02-JAN-2009','DD-MON-YYYY'), -
      4 i_extended_purge=>TRUE);
```



Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Statspack Baselines and Purging

The PERFSTAT schema can grow very large with many snapshots. You can purge unnecessary data from the PERFSTAT schema by using `STATSPACK.PURGE`. By default, `PURGE` removes all snapshots, but you can identify and keep snapshot data by creating baselines. After you have determined which snapshot IDs or times of day most represent a particular workload whose performance data you would like to keep, you can mark the data as baselines. You can later decide to clear one or more baselines. Clearing the baseline does not remove the data, it just identifies the data as candidates for purging. You cannot roll back a purge operation.

The slide shows you examples of managing baselines of Statspack data:

- Making a baseline of snaps 45 and 50 including the range of snapshots in between
- Making a baseline of only snapshots 2 and 4 (excluding the snapshots in between)
- Making a baseline of snapshots taken between August 31, 2009 at 9:00 AM and August 31, 2009 at 12:00 noon
- Clearing an existing baseline including all the snapshots between December 13, 2009 at 11:00 PM and February 14, 2010 at 2:00 AM.
- Purging all snapshots that fall between January 1, 2009 and January 2, 2009; also performing an extended purge

## Statspack Baselines and Purging (continued)

In earlier releases of Statspack, Statspack identifier tables that contained SQL Text, SQL Execution plans, and Segment identifiers were not purged. Now you can purge unreferenced data in these tables by requesting that the extended purge be performed at the same time as the normal purge. Purging this data can be resource intensive, so you may choose to perform an extended purge less frequently than the normal purge.

# Reporting with Statspack

```
SQL> connect perfstat/perfstat_password
SQL> @?/rdbms/admin/spreport
```

```
SQL> connect perfstat/perfstat_password
SQL> define begin_snap=1
SQL> define end_snap=2
SQL> define report_name=batch_run
SQL> @?/rdbms/admin/spreport
```

```
SQL> connect perfstat/perfstat_password
SQL> define begin_snap=39
SQL> define end_snap=40
SQL> define hash_value=1988538571
SQL> define report_name=batch_sql_run
SQL> @sprepsql
```

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Reporting with Statspack

You can generate a performance report from two snapshots. There are two reports available:

- The instance report (`spreport.sql` and `sprepins.sql`) is a general instance health report, covering all aspects of instance performance. The instance report calculates and prints ratios, increases, and so on for all statistics between the two snapshot periods.
- The SQL report (`sprepsql.sql` and `sprsqins.sql`) is a report for a specific SQL statement. The SQL report is usually run after examining the high-load SQL sections of the instance health report. The SQL report displays SQL-specific statistics, the complete SQL text, and, if a level 6 snapshot or above has been taken, information about any SQL plan(s) associated with that statement.

Both reports prompt for the beginning snapshot ID, the ending snapshot ID, and the report name. The SQL report also requests a hash value for the SQL statement to be reported on. You can configure each report by using parameters such as the total number of rows of SQL output to display in each SQL section or the number of days of snapshots to list. Separating the phase of data gathering from producing a report provides the flexibility of basing a report on any data points selected.

**Note:** The instance must not have been shut down between the times that the beginning and ending snapshots were taken.

## Statspack Considerations

- Set `STATISTICS_LEVEL = TYPICAL`.
- Set `TIMED_STATISTICS = TRUE` to collect timing information.
- Avoid overlap between Statspack and AWR.
- Choose a sensible snapshot interval.
- Fix interval when “#####” appears in the report output.
- Gather optimizer statistics on the `PERFSTAT` schema.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Statspack Considerations

- In Oracle Database 11g, the `STATISTICS_LEVEL` parameter is set to `TYPICAL` by default. When `STATISTICS_LEVEL` is set to `TYPICAL` or `ALL`, Time Model data is populated by the server, and can be captured by Statspack. At a minimum, to have the most useful data available for tuning, you should set `TIMED_STATISTICS` to `TRUE` before the creation of a snapshot. Whenever `STATISTICS_LEVEL` is set to `TYPICAL` or `ALL`, `TIMED_STATISTICS` is automatically set to `TRUE`.
- If `STATISTICS_LEVEL` is `TYPICAL` or `ALL`, and you do not use the Database Diagnostics pack, then you should disable AWR data capture by setting the AWR snapshot schedule interval to zero. If you use both Statspack and AWR, schedule the Statspack snapshots to avoid the times when AWR is attempting to capture data. By default, AWR snapshots are taken every hour, on the hour.
- In general, snapshot data should be collected in intervals of an hour. In the unusual situation that finer-grained analysis is required, the snapshot interval can be changed for the duration of that problem analysis. Snapshot intervals of less than 15 minutes are seldom useful.

## Statspack Considerations (continued)

- When a value is too large for a Statspack field, it is represented by a series of pound signs, such as #####. The main reason for this overflow is an analysis interval that is too large. Reduce the length of time between snapshots selected for the report, or choose snapshots that are closer in time.
- For best performance when running the performance reports, optimizer statistics should be gathered on the Statspack schema. In Oracle Database 10g and later, the Oracle server automatically gathers optimizer statistics on database segments when the segments become stale. If you have disabled the automatic job, you should gather statistics manually using the following command:

```
execute dbms_stats.gather_schema_stats(ownname=>'PERFSTAT',cascade=>true);
```

## Statspack and AWR Reports

The Statspack and AWR reports are designed to be used top down.

- The first few pages contain the following:
  - Summary Information
  - Load Profile (useful with baseline)
  - Instance Efficiency (useful with baseline)
  - Top 5 Timed Events
  - Time Model
- The following pages show detailed data:
  - SQL
  - IO
  - Wait Statistics

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Statspack and AWR Reports

- The Statspack and AWR reports are designed to be used top down. The first few pages always indicate the biggest bottleneck to investigate. The Summary Information section shows the time interval being reported. Load Profile shows the level and type activity. This is very useful as a sanity check, to be sure that the workload has not changed significantly. The Instance Efficiency section gives the traditional ratios plus additional efficiency percentages that help identify and confirm the areas that need tuning.
- The Top 5 Timed Events section identifies which events are accounting for the most time by percentage. The Time Model section shows the time and the ratio of time consumed by a specific event to DB\_TIME.
- The Time Model entries that relate to the Top 5 Timed Events indicate the possible impact tuning each event could have.
- The subsequent pages of the Statspack and AWR reports give detailed information about a variety of database areas. These detailed sections are used to confirm diagnoses and plan remedies for the symptoms shown on the first pages.



# Reading a Statspack or AWR Report

- Start with the summary data at the top.

- **Top 5 Timed Events**

- Confirm {
- Wait Events and Background Wait Events
  - Wait Event Histogram (only in Statspack)
  - Load Profile (useful with baseline)
  - Instance Efficiency (useful with baseline)
  - Time Model

- Drill down to specific sections.
  - Indicated by top wait event



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Reading a Statspack or AWR Report

When diagnosing a possible database problem, start by checking the host hardware and OS, the middle tier or the network, and other resources outside the database. If these are not being completely consumed, it is time to look at the database.

### Start at the Top

Use the Top 5 Timed Events section to identify which events are accounting for the most time by percentage. After these are identified, look at the Wait Events and Background Wait Events sections for the average wait time for the highest-ranking events. This data can sometimes provide insight into the scale of the wait. Compare the average with data in the Wait Event Histogram section. Is the average indicative of the typical wait time, or is the average wait time severely skewed by a small number of atypical waits?

Look at the Load Profile, Instance Efficiency, and Time Model sections. Pay attention to any statistics or ratios that are related to the top wait events. Is there a single consistent picture? If not, note other potential issues to investigate while looking at the top events, but do not be distracted from the top wait events. Scan the other statistics. Are there any statistics in the Load Profile that are unusually high for this site, or any ratios in the Instance Efficiency section, which are atypical for this site (when compared to a baseline)? Next, drill down to the appropriate section in the Statspack report for additional data.

## Statspack/AWR Report Drilldown Sections

| Drilldown | Section                    |
|-----------|----------------------------|
| LC        | Shared Pool Statistics     |
| CPU       | Host CPU                   |
| CPU       | Instance CPU               |
| MEM       | Virtual Memory Paging      |
| CPU       | SQL ordered by CPU         |
| IO        | SQL ordered by Elapsed     |
| CPU       | SQL ordered by Gets        |
| IO        | SQL ordered by Reads       |
| LC        | SQL ordered by Parse Calls |

**CONT:** Block contention

**CPU:** CPU consumption

**ENQ:** Enqueue

**IO:** IO consumption

**LAT:** Latch contention

**LC:** Library cache

**MEM:** Memory consumption

**PGAM:** PGA memory consumption

**RBS:** Rollback segment

**UNDO:** Automatic undo

**SP:** Shared pool

**ORACLE**

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Statspack/AWR Report Drilldown Sections

- The relevant sections to examine are indicated by the top wait events.
- The slide shows you a partial list of the most important Statspack sections indicating which sections can be used for possible drilldown. Each wait event falls into a type of problem. These types are listed in the slide. Use the type to determine the section to examine for more detailed information.
- While drilling down, determine whether the data in these sections is consistent with the wait events? If not, identify related areas in the report for an alternative data source.
- Analyze the detailed data in the drilldown sections after identifying candidate problems and resources showing contention. Consider whether there is sufficient data to build a plausible theory for the cause and resolution of the problem.
- Use the *Oracle Database Performance Tuning Guide* as a resource to identify the cause and the resolution of the contention. The manual includes detailed descriptions of how to diagnose causes and solutions for wait events and how to optimally configure the instance to avoid problems.
- For more detail on using Statspack for tuning, see MetaLink notes; two are especially helpful: Note 228913.1: *Systemwide Tuning using STATSPACK Reports* and Note 942241.1: *FAQ Statspack Complete Reference*.

## Statspack/AWR Report Drilldown Sections (continued)

### Note

In Oracle Database, many of the commonly requested resources have their own wait events, such as common latches and all enqueuees. This makes it easier to identify contention. For example, library cache latch–related wait events would include “latch: library cache,” “latch: library cache lock,” and “latch: library cache pin.”

| <b>Drilldown</b> | <b>Section</b>                 |
|------------------|--------------------------------|
| LC               | SQL ordered by Sharable Memory |
| LC               | SQL ordered by Version Count   |
| IO               | IO Sections                    |
| IOMEM            | Buffer Cache                   |
| PGAMIO           | PGA Memory Stats               |
| PGAM             | Process Memory Stats           |
| ENQ              | Enqueue activity               |
| RBS UNDO         | Rollback and Undo              |
| LAT              | Latch Statistics               |
| CPU              | Segments by Logical Reads      |
| IO               | Segments by Physical Reads     |
| CONT             | Segments by Row Lock Waits     |
| CONT             | Segments by ITL Waits          |
| CONT             | Segments by Buffer Busy Waits  |
| LC               | Dictionary Cache               |
| LC               | Library Cache Activity         |
| STREA            | Streams                        |
| SP LC            | SGA breakdown difference       |
| SP LC            | SQL Memory Statistics          |

## Report Drilldown Examples

- If the top timed event is related to I/O waits, look at:
  - SQL ordered by Reads
  - SQL ordered by Elapsed
  - Tablespace IO Stats
  - File IO Stats
  - File IO Histogram
- If the top timed event is related to CPU usage, look at:
  - Load Profile
  - Time Model
  - SQL ordered by CPU
  - SQL ordered by Gets

The Oracle logo, consisting of the word "ORACLE" in a white, sans-serif font on a red rectangular background.

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Report Drilldown Examples

If the top events are I/O related (db file sequential read and db file scattered read), look at the “SQL ordered by Reads” or “SQL ordered by Elapsed” sections, and the Tablespace IO Stats, File IO Stats, and File IO Histogram sections.

If the top timed event is CPU time, look at the number of logical reads in Load Profile, and the top statistics in the Time Model data. Look at SQL by CPU or SQL by Gets. Correlate the output. Does the evidence point to SQL execution?

**Note:** On a well-performing system, the top events are likely to be CPU time, db file scattered read, and db file sequential read.

## Load Profile Section

- Allows characterization of the application
- Can point toward potential problems:
  - High hard parse rate
  - High I/O rate
  - High login rate

| Load Profile                | Per Second | Per Transaction         |
|-----------------------------|------------|-------------------------|
| Redo size:                  | 186,146.10 | 518.05                  |
| Logical reads:              | 1,670.36   | 4.65                    |
| Block changes:              | 1,451.40   | 4.04                    |
| Physical reads:             | 0.01       | 0.00                    |
| Physical writes:            | 0.66       | 0.00                    |
| User calls:                 | 8.22       | 0.02                    |
| Parses:                     | 11.20      | 0.03                    |
| Hard parses:                | 1.03       | 0.00                    |
| Sorts:                      | 9.01       | 0.03                    |
| Logons:                     | 0.20       | 0.00                    |
| Executes:                   | 382.20     | 1.06                    |
| Transactions:               | 359.32     |                         |
| % Blocks changed per Read:  | 86.89      | Recursive Call %: 99.39 |
| Rollback per transaction %: | 0.01       | Rows per Sort: 8.63     |

- Is more useful if you have a comparable baseline
- Answers “What has changed?”
  - Txn/sec change implies changed workload.
  - Redo size/txn implies changed transaction mix.
  - Physical reads/txn implies changed SQL or plan.

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Load Profile Section

The Load Profile section indicates the nature of the workload on the system. It includes the key V\$SYSSTAT statistics needed to characterize the system. From the Load Profile, you can determine the rates of:

- Logical and physical writes
- Logical and physical reads
- Hard and soft parses
- Session logons
- User activity

The values are shown per second (for absolute rates) and per transaction (for easier comparison with your application per-transaction rates to determine whether the application code, application use, or execution plans have changed).

This is where knowing what is typical on a system helps. If, for example, you know the typical number of logical reads/sec and redo size/sec when the system performs well, you can look at these values and see whether they are different when the system performs badly. To continue with the example, is the system doing twice the number of logical I/Os now than when the system performed well last week? Has the workload increased, or has the execution plan of key statements changed?

## Time Model Section

The Time Model section lists a set of statistics; it:

- Gives an overview of where time is spent inside the Oracle database server
- Is reported from the V\$SYS\_TIME\_MODEL view
- Gauges the performance impact of any entity of the database, and provides drill down.

| Time Model System Stats DB/Inst: ORCL/orcl Snaps: 1-2 |          |              |
|-------------------------------------------------------|----------|--------------|
| -> Ordered by % of DB time desc, Statistic name       |          |              |
| Statistic                                             | Time (s) | % of DB time |
| -----                                                 | -----    | -----        |
| sql execute elapsed time                              | 620.7    | 98.8         |
| DB CPU                                                | 78.1     | 12.4         |
| PL/SQL execution elapsed time                         | 10.7     | 1.7          |
| parse time elapsed                                    | 7.9      | 1.3          |
| hard parse elapsed time                               | 4.1      | .7           |
| connection management call elapsed                    | 2.5      | .4           |
| PL/SQL compilation elapsed time                       | 0.4      | .1           |
| hard parse (sharing criteria) elaps                   | 0.0      | .0           |
| repeated bind elapsed time                            | 0.0      | .0           |
| sequence load elapsed time                            | 0.0      | .0           |
| DB time                                               | 628.4    |              |
| background elapsed time                               | 62.1     |              |
| background cpu time                                   | 13.6     |              |
| -----                                                 | -----    | -----        |

ORACLE

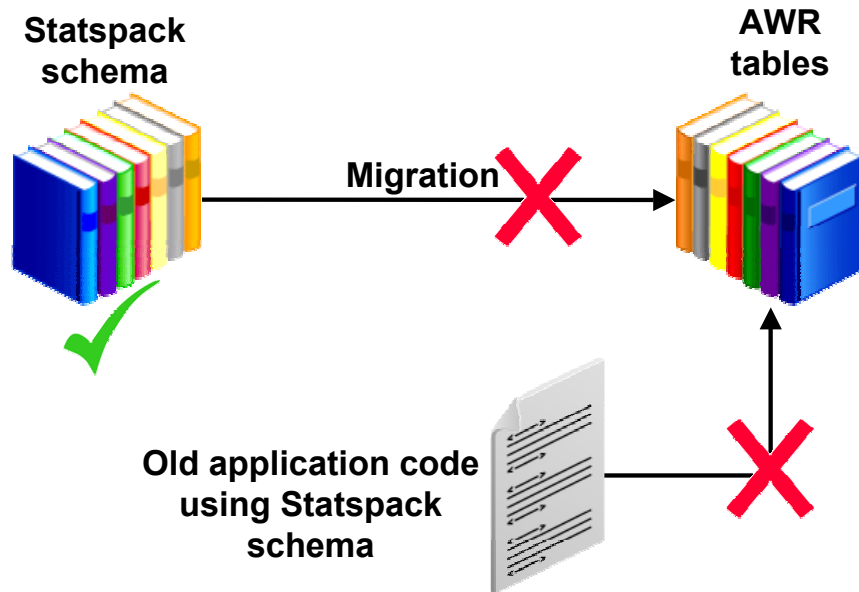
Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

### Time Model Section

The Time Model section is captured from a view called V\$SYS\_TIME\_MODEL. Time Model data is intended to be used to help identify which sections to drill down to. For example, in earlier releases, a strange set of wait events and latches would indicate that the application was performing an excess of connection management (logon/logoff activities). In prior releases, only by knowing which latches to look for was it possible to track down this inefficient activity. The Time Model tracks time by different types of operations in the database, thus making it simpler to identify the types of operations that the instance is spending its time performing.

Some of the statistics measure the CPU time used for an operation, and others measure the total elapsed time for that operation. The name of the statistic indicates which of these two it is measuring. Elapsed time includes wait time.

# Statspack and AWR



ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

## Statspack and AWR

- In the past, historical data could be obtained manually by using Statspack. Some DBAs wrote applications to report on Statspack data. These reports can still be used in Oracle Database 11g; however, if you want to use Automatic Workload Repository instead, then you need to change your application code.
- Statspack users should switch to Automatic Workload Repository in Oracle Database 11g Enterprise Edition. AWR is not available to Standard Edition users.
- There is no supported path to migrate Statspack data into AWR. Also, there is no view created on top of AWR to simulate the Statspack schema.

## Summary

In this lesson, you should have learned how to:

- Install Statspack
- Create Statspack snapshots
- Generate Statspack reports
- Identify the major sections of the Statspack report

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.



## Practice Overview: Use Statspack

This practice covers the following topics:

- Installing Statspack
- Creating Statspack snapshots
- Generating Statspack reports

ORACLE

Copyright © 2010, Oracle and/or its affiliates. All rights reserved.

